

Solution 4

1.

```

for (p = 0; p < 2; p++) {
    for (q = 0; q < 2; q++) {
        for (i = p*64; i < (p+1)*64; i++) {
            for (j = q*64; j < (q+1)*64; j++) {
                Y[i] = Y[i] + A[i][j]*X[j];
            }
        }
    }
}

```

When $p = 0$ we compute $Y[0:63]$ in 2 steps: first, we use $A[0:63][0:63]$ and $X[0:63]$ when $q = 0$; then, we use $A[0:63][64:127]$ and $X[64:127]$ when $q = 1$. When $p = 1$ we compute $Y[64:127]$ in 2 steps: first, we use $A[64:127][0:63]$ and $X[0:63]$ when $q = 0$; then, we use $A[64:127][64:127]$ and $X[64:127]$ when $q = 1$.

Storing one 64×64 block of 32-bit numbers (for matrix **A**) requires $64 \times 64 \times 4 = \mathbf{16KB}$ of memory, and storing two 128×1 blocks of 32-bit numbers (for vectors **X** and **Y**) requires $2 \times 128 \times 4 = \mathbf{1KB}$ of memory. Hence, the cache size should be at least **17KB**.

2. The total size of `int x[256][256]` is $256 \times 256 \times 4 = \mathbf{256KB}$, and each row requires $256 \times 4 = \mathbf{1KB}$. For every iteration of the outer loop (index i), we have 2 reads and 1 write per row element $x[i][j]$, i.e., each row i requires $(2+1) \times 256 = 768$ accesses to it.

If we have four **1KB**-pages, we get 4 page faults per 4 rows, or 4 page fault per 4×768 accesses (the size of each row is **1KB**, and there are 768 accesses to each row); therefore, the page fault rate is $4/(4 \times 768) = 0.13\%$. On the other hand, if we have one **4KB**-page, we get 1 page fault per 4×768 accesses; therefore, the page fault rate is $1/(4 \times 768) = 0.03\%$. In both cases the allocated memory amount is the same (**4KB** total), but the page fault rates are different.

3. If we allocate 512 **4KB**-pages for the `int a[1024][1024]` array, then the **Good** example will have 1 page fault per 1024 row element accesses ($p \approx \mathbf{0.1\%}$), and the **Bad** example will have 1 page fault per 1 row element access ($p = \mathbf{100\%}$), i.e., increasing the number of allocated **4KB**-pages from 1 to 512 does not provide any benefits for this particular application code. If we allocate one **4KB**-page for the `int a[512][512]` array (i.e., one page holds 2 rows), then the **Good** example will have 1 page fault per $512+512$ row element accesses ($p \approx \mathbf{0.1\%}$), and the **Bad** example will have 1 page fault per $1+1$ row element accesses ($p = \mathbf{50\%}$).

4. General formula: $T_{ave} = h_1C_1 + (1-h_1)(h_2C_2 + (1-h_2)((1-p)M + pD))$.

(a) The page table must have a dedicated entry for every possible VPN. As we have at most **4GB/1MB = 4K** pages (i.e., **4K** VPNs), the page table has **4K** entries.

(b) For **$h_1 = 0.95$** , **$h_2 = 0.90$** , and **$p = 0$** , we have:

$$T_{ave} = 0.95 \cdot 1\tau + 0.05(0.9 \cdot 4\tau + 0.1(1 \cdot 16\tau + 0 \cdot 10,000\tau)) = 1.21\tau.$$

(c) For **$h_1 = 0$** , **$h_2 = 0.90$** , and **$p = 0.001$** , we have:

$$T_{ave} = 0 \cdot 1\tau + 1(0.9 \cdot 4\tau + 0.1(0.9999 \cdot 16\tau + 0.0001 \cdot 10,000\tau)) = 5.3\tau.$$

(d) For **$h_1 = 0.95$** , **$h_2 = 0.90$** , and **$p = 0.001$** , we have:

$$T_{ave} = 0.95 \cdot 1\tau + 0.05(0.9 \cdot 4\tau + 0.1(0.9999 \cdot 16\tau + 0.0001 \cdot 10,000\tau)) = 1.215\tau.$$