

Solution 4

1.

```
for (p = 0; p < 2; p++) {
    for (q = 0; q < 2; q++) {
        for (i = p*64; i < (p+1)*64; i++) {
            for (j = q*64; j < (q+1)*64; j++) {
                Y[i] = Y[i] + A[i][j]*X[j];
            }
        }
    }
}
```

When $p = 0$ we compute $Y[0:63]$ in 2 steps: first, we use $A[0:63][0:63]$ and $X[0:63]$ when $q = 0$; then, we use $A[0:63][64:127]$ and $X[64:127]$ when $q = 1$. When $p = 1$ we compute $Y[64:127]$ in 2 steps: first, we use $A[64:127][0:63]$ and $X[0:63]$ when $q = 0$; then, we use $A[64:127][64:127]$ and $X[64:127]$ when $q = 1$.

Storing one 64×64 block of 32-bit numbers (for matrix **A**) requires $64 \times 64 \times 4 = \mathbf{16KB}$ of memory, and storing two 128×1 blocks of 32-bit numbers (for vectors **X** and **Y**) requires $2 \times 128 \times 4 = \mathbf{1KB}$ of memory. Hence, the cache size should be at least **17KB**.

2. The size of `double x[256][256]` is $256 \times 256 \times 8 = \mathbf{512KB}$, and each row of **x** requires $256 \times 8 = \mathbf{2KB}$. For every iteration of the outer loop (index *i*), we have 1 read per row element $x[i][j]$ in the first inner loop, plus 1 read and 1 write per row element $x[i][j]$ in the second inner loop, i.e., each row *i* requires $(2+1) \times 256 = 768$ accesses to it. The total number of accesses is $256 \times 768 = 196,608$.

If we have two **2-KB** pages, we get 1 page fault per row (per 768 accesses), or 256 page faults in total (per 196,608 accesses). Therefore, the page fault rate is $1/768$, or equivalently, $256/196,608 = 0.130\%$.

If we have one **4-KB** page, we get 1 page fault per 2 rows (per 2×768 accesses), or $256/2 = 128$ page faults in total (per 196,608 accesses). Therefore, the page fault rate is $1/(2 \times 768)$, or equivalently, $128/196,608 = 0.065\%$.

In both cases the allocated memory amount is the same (**4KB** total), but the page fault rates are different.

3. The size of `float x[256][256]` is $256 \times 256 \times 4 = \mathbf{256KB}$, where each row requires $256 \times 4 = \mathbf{1KB}$. For the first loop (computing `trace`), we have 1 read per row *i*, i.e., there are 256 accesses required. For the other two nested loops (normalizing $x[i][j]$), we have 1 read and 1 write per row element $x[i][j]$, i.e., each row *i* requires $(1+1) \times 256 = 512$ accesses to it. Hence, the total number of accesses is $256 + 256 \times 512 = 131,328$.

If we have four **1-KB** pages, we get 256 page faults due to the first loop, and 256 page faults due to the other two nested loops; therefore, the page fault rate is $(256+256)/131,328 = 0.3899\%$.

If we have one **4-KB** page, we get $256/4=64$ page faults due to the first loop, and $256/4=64$ page faults due to the other two nested loops; therefore, the page fault rate is $(64+64)/131,328 = 0.0975\%$.

In both cases the allocated memory amount is the same (**4KB** total), but the page fault rates are different.