

## Solution 5

1.

```
for (p = 0; p < 2; p++) {
    for (q = 0; q < 2; q++) {
        for (i = p*64; i < (p+1)*64; i++) {
            for (j = q*64; j < (q+1)*64; j++) {
                Y[i] = Y[i] + A[i][j]*X[j];
            }
        }
    }
}
```

When  $p = 0$  we compute  $Y[0:63]$  in 2 steps: first, we use  $A[0:63][0:63]$  and  $X[0:63]$  when  $q = 0$ ; then, we use  $A[0:63][64:127]$  and  $X[64:127]$  when  $q = 1$ . When  $p = 1$  we compute  $Y[64:127]$  in 2 steps: first, we use  $A[64:127][0:63]$  and  $X[0:63]$  when  $q = 0$ ; then, we use  $A[64:127][64:127]$  and  $X[64:127]$  when  $q = 1$ .

Storing one  $64 \times 64$  block of 32-bit numbers (for matrix **A**) requires  $64 \times 64 \times 4 = \mathbf{16KB}$  of memory, and storing two  $128 \times 1$  blocks of 32-bit numbers (for vectors **X** and **Y**) requires  $2 \times 128 \times 4 = \mathbf{1KB}$  of memory. Hence, the cache size should be at least **17KB**.

2. The size of `int x[256][256]` is  $256 \times 256 \times 4 = \mathbf{256KB}$ , and each row of **x** requires  $256 \times 4 = \mathbf{1KB}$ . For every iteration of the outer loop (index *i*), we have 2 reads and 1 write per row element  $x[i][j]$ , i.e., each row *i* requires  $(2+1) \times 256 = 768$  accesses to it. The total number of accesses is  $256 \times 768 = 196,608$ .

If we have four **1-KB** pages, we get 1 page fault per row (per 768 accesses), or 256 page faults in total (per 196,608 accesses). Therefore, the page fault rate is  $1/768$ , or equivalently,  $256/196,608 = 0.1302\%$ .

If we have one **4-KB** page, we get 1 page fault per 4 rows (per  $4 \times 768$  accesses), or  $256/4 = 64$  page faults in total (per 196,608 accesses). Therefore, the page fault rate is  $1/(4 \times 768)$ , or equivalently,  $64/196,608 = 0.0326\%$ .

In both cases the allocated memory amount is the same (**4KB** total), but the page fault rates are different.

3. The size of `float x[256][256]` is  $256 \times 256 \times 4 = \mathbf{256KB}$ , where each row requires  $256 \times 4 = \mathbf{1KB}$ . For the first loop (computing `trace`), we have 1 read per row *i*, i.e., there are 256 accesses required. For the other two nested loops (normalizing  $x[i][j]$ ), we have 1 read and 1 write per row element  $x[i][j]$ , i.e., each row *i* requires  $(1+1) \times 256 = 512$  accesses to it. Hence, the total number of accesses is  $256 + 256 \times 512 = 131,328$ .

If we have four **1-KB** pages, we get 256 page faults due to the first loop, and 256 page faults due to the other two nested loops; therefore, the page fault rate is  $(256+256)/131,328 = 0.3899\%$ .

If we have one **4-KB** page, we get  $256/4=64$  page faults due to the first loop, and  $256/4=64$  page faults due to the other two nested loops; therefore, the page fault rate is  $(64+64)/131,328 = 0.0975\%$ .

In both cases the allocated memory amount is the same (**4KB** total), but the page fault rates are different.