

# I/O

D.N.Rakhmatov

Adopted (with modifications) from:  
F. Kuhns (WU-St.Luis)  
R. Hoelzeman (Pittsburgh)  
D. Patterson (UC-Berkeley)  
J. Armstrong (Virginia Tech)  
G. Buttazzo (Scuola Superiore)

C. Hamacher et al, *Computer Organization*, 6/E, © 2011 McGraw-Hill  
W. Stallings, *Computer Organization and Architecture*, 8/E, © 2010 Pearson  
A. Silberschatz et al, *Operating System Concepts Essentials*, 8/E, © 2011 Wiley

Fall 2016

UVic - CENG 355 - I/O

1

## I/O Considerations

### ■ Timing

- Typically, CPU and I/O device will not be operating at the same clock frequency
  - We must have a means of synchronizing (at least momentarily) the two in order to effect the information transfer

### ■ Speed

- During I/O operations, objective is to keep both CPU and I/O device busy
  - Not easy because of the range of CPU's and I/O device's operating speeds

### ■ Coding

- Information in CPU is represented in binary
  - The data representation is most likely not in a form suitable for external use

Fall 2016

UVic - CENG 355 - I/O

2

## I/O Interface I

### ■ I/O interface consists of two parts:

- The **hardware** interface: the electrical connections and signal paths
- The **software** interface: provides a flexible means for manipulating the data

### ■ It can be viewed by the software programmer as a "system" of special registers

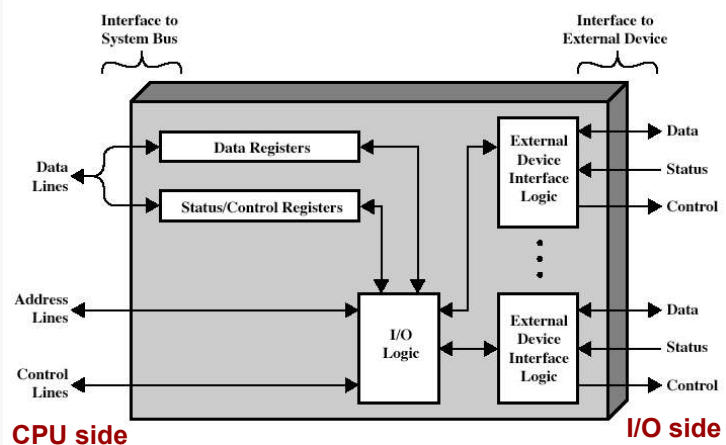
- **Control**: defines the operational characteristics of the interface
- **Status**: tracks the use of the interface (is it busy now?)
- **Data**: provides the actual data transfer mechanism

Fall 2016

UVic - CENG 355 - I/O

3

## I/O Interface II



Fall 2016

UVic - CENG 355 - I/O

4

## I/O Addressing

### ■ Isolated (standard) I/O

- I/O devices have their own unique address space
  - Additional signal on the bus indicates whether accessing the memory or an I/O device
- Individual devices selected based on:
  - Valid device address being placed on the address bus
  - Dedicated signal indicating I/O operation
  - Valid read or write pulse

### ■ Memory-mapped I/O

- I/O device address is a part of the memory address space
  - Certain memory addresses are reserved for I/O registers
  - More flexibility in accessing the device, but at a loss of real memory locations


Fall 2016

UVic - CENG 355 - I/O

5

## I/O Scenario

- CPU reads from **Status Register** in loop, waiting for I/O device to set **Ready** bit ( $0 \Rightarrow 1$ )
- CPU then reads from (input) or writes to (output) **Data Register**
  - Upon transferring data to/from **Data Register**, **Ready** bit of **Status Register** must be reset ( $1 \Rightarrow 0$ )
    - I/O interface circuit will typically do this automatically for you
- This is a **polling** scenario, **synchronous** with respect to current program execution



Fall 2016

UVic - CENG 355 - I/O

6

## Alternative to Polling

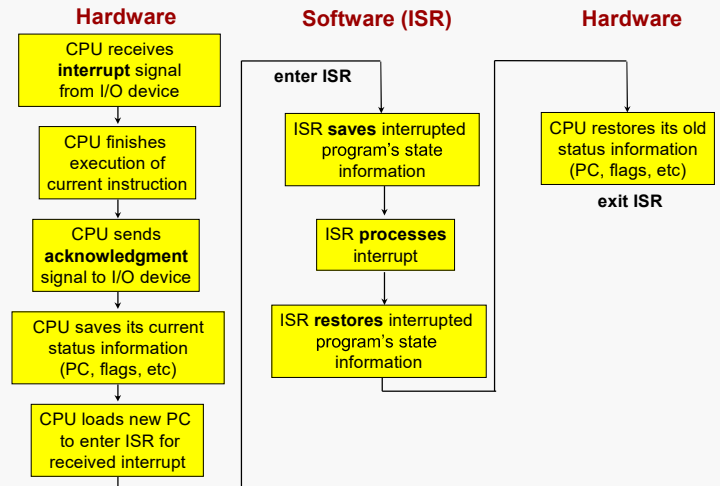
- We could have an unplanned procedure call that would be invoked only when I/O device is ready
  - Need exception in current program's control flow
    - **Interrupt** when I/O device is ready: CPU enters **ISR** (interrupt service routine)
    - **Return** when done with I/O data transfer: CPU exits **ISR**
- An I/O interrupt is **asynchronous** with respect to current program execution:
  - I/O interrupt is not associated with any instruction, but it can happen in the middle of any given instruction
  - I/O interrupt does not prevent any instruction from completion

Fall 2016

UVic - CENG 355 - I/O

7


## Simple Interrupt Processing



Fall 2016


UVic - CENG 355 - I/O

8



## Three I/O Control Mechanisms

- Programmed I/O (polling)
  - CPU initiates and controls all I/O operations
- Interrupt-driven I/O
  - CPU performs I/O operations upon receiving external interrupts from I/O devices ready to transfer data
- Direct memory access (DMA) controlled I/O
  - I/O operations are initiated and controlled by non-CPU hardware



Fall 2016

UVic - CENG 355 - I/O

10

## Programmed I/O

- CPU executes program segments that initiate, direct, and terminate (synchronous) I/O operations
  - Advantage: simple to implement, requiring little special hardware and software
  - Disadvantage: low CPU efficiency, as it is slowed down to the speed of I/O device
  - Data can be transferred using one of two methods:
    - Conditional transfers, taking place only after CPU determines that I/O device is ready
      - ✓ Guaranteed that I/O device won't be flooded by CPU or that CPU won't read the same data more than once
    - Unconditional transfers, taking place without checking that I/O device is actually ready to send/receive the data
      - ✓ Generally used to exchange data with a port that is known to be always "ready"

Fall 2016

UVic - CENG 355 - I/O

11

## Interrupt-Driven I/O and DMA

- Interrupt-driven I/O (asynchronous)
  - CPU is **not** required to poll I/O device
    - Interrupt asserted to notify CPU that I/O device is ready
    - ISR is "called" by CPU hardware itself, not by software
  - Need an extra pin or pins to accept interrupt signal(s) (e.g., **IRQ**)
- DMA (can be synchronous or asynchronous)
  - Used to transfer large blocks of data at high speed between an I/O device and the main memory directly
    - CPU sends the starting address, the number of data words in that block, and the direction of transfer to DMA controller
    - CPU grants DMA controller authority to control memory access, and DMA controller performs the data transfer (meanwhile, CPU is free to do other things)
    - Once the transfer is complete, the DMA controller sends an interrupt signal to inform CPU

Fall 2016

UVic - CENG 355 - I/O

12

## Example I

- Let's assume...
  - I/O device
    - Data transfer rate  $R_{I/O} = 8 \text{ MB/s} = 8 \times 2^{20} \text{ B/s} \approx 8.39 \times 10^6 \text{ B/s}$
    - 5% active (i.e., ready for transfer), not ready 95% of the time
    - Data transferred in chunks of  $d_{I/O} = 16 \text{ B}$  at a time (when ready)
  - CPU
    - Clock frequency  $f_{\text{clk}} = 500 \text{ MHz} = 500 \times 10^6 \text{ Hz}$
    - To perform a poll (i.e., call polling routine, access I/O device, return), it takes either  $N_{\text{poll-ready}} = 400$  cycles when I/O device is ready (transferring  $d_{I/O}$  of data), or  $N_{\text{poll-not-ready}} = 200$  cycles when I/O device is not ready (transferring no data)
    - To service an interrupt (i.e., enter ISR, access I/O device, exit), it takes  $N_{\text{int}} = 500$  cycles
    - **Note:** I/O interrupt processing is more expensive than polling for the same I/O access, i.e.,  $N_{\text{int}} > N_{\text{poll-ready}}$

Fall 2016

UVic - CENG 355 - I/O

13

## Example II

- CPU runs through  $500 \times 10^6$  cycles per sec ( $f_{\text{clk}}$ )
- I/O device can be accessed at the maximum rate of  $R_{I/O} / d_{I/O} \approx 0.52 \times 10^6$  times/s
  - Polling scenario:  $0.52 \times 10^6$  polls/s
  - Interrupt scenario:  $0.52 \times 10^6$  interrupts/s
- Cost of polling
  - $(5\% \times R_{I/O} / d_{I/O}) \times N_{\text{poll-ready}} + (95\% \times R_{I/O} / d_{I/O}) \times N_{\text{poll-not-ready}} \approx 109 \times 10^6$  cycles/s
  - CPU busy:  $109 \times 10^6 / 500 \times 10^6 \approx 22\%$  (too much!)
- Cost of interrupts
  - $(5\% \times R_{I/O} / d_{I/O}) \times N_{\text{int}} \approx 13 \times 10^6$  cycles/s
  - CPU busy:  $13 \times 10^6 / 500 \times 10^6 \approx 2.6\%$

Fall 2016

UVic - CENG 355 - I/O

14

## Example III

- Let's further assume...
  - I/O device
    - **DMA:** Data is transferred in chunks of  $d_{I/O-DMA} = 1 \text{ KB} = 1024 \text{ B}$  at a time
  - CPU
    - To initiate a DMA transfer, it takes  $N_{\text{DMA-start}} = 1000$  cycles
    - To complete a DMA transfer, it takes  $N_{\text{DMA-end}} = 500$  cycles
- I/O device can be accessed at the maximum rate of  $R_{I/O} / d_{I/O-DMA} \approx 8.2 \times 10^3$  times/s
  - DMA scenario:  $8.2 \times 10^3$  accesses/s
- Cost of DMA
  - $(5\% \times R_{I/O} / d_{I/O-DMA}) \times (N_{\text{DMA-start}} + N_{\text{DMA-end}}) \approx 0.6 \times 10^6$  cycles/s
  - CPU busy:  $0.6 \times 10^6 / 500 \times 10^6 \approx 0.12\%$

Fall 2016

UVic - CENG 355 - I/O

15

## Single Interrupt

- I/O device asserts its interrupt-request signal **IRQ** and keeps it asserted until the interrupt request is acknowledged by the processor
- How do we avoid successive interruptions while **IRQ** is asserted?
  - Use interrupt-disable instruction at the beginning of ISR and place interrupt-enable instruction(s) before return
    - Enabling/disabling interrupts can be done by setting/clearing the interrupt-enable bit (**IE**) in the processor's status register (**PSR**)
  - Design the processor hardware, so that **PSR[IE]** is cleared before entering ISR, and then set again upon return from ISR, automatically
  - Make **IRQ** circuit edge-sensitive

Fall 2016

UVic - CENG 355 - I/O

16

## Single Interrupt Scenario

- I/O device asserts interrupt request **IRQ**
- The processor interrupts the program currently being executed
- If needed, further interrupts are disabled by clearing **PSR[IE]**
- ISR handles the interrupt, while I/O device is informed that its request has been accepted
  - I/O device deasserts interrupt request **IRQ**
- Interrupts are enabled again and execution of the interrupted program is resumed

Fall 2016

UVic - CENG 355 - I/O

17

## Multiple Interrupts

- What if many I/O devices are connected to the processor, each capable of initiating interrupts?
  - How can the processor identify the interrupt source?
    - E.g., the processor can check appropriate status bits (interrupt flags) of all potential interrupt sources to figure out who requested an interrupt
  - How can the processor find the appropriate ISR for a given interrupt?
    - An interrupt source may identify itself by sending its **vector** code (ISR pointer) to the processor over the data bus
      - ✓ This is called **vectored interrupt scheme**
  - Is it allowed for another interrupt source to interrupt already running ISR?
  - How does the processor arbitrate multiple simultaneous interrupt requests?


Fall 2016

UVic - CENG 355 - I/O

18

## Nested Interrupts

- What if handling an interrupt is taking too long while another source is waiting to be serviced?
  - Need a **programmable** priority structure
  - An interrupt request from a higher-priority source should be accepted even when the processor is still handling the earlier request from a lower-priority source
- Use separate **IRQ** (request) and **IRQA** (acknowledge) pairs for each device with different priority levels




Fall 2016

UVic - CENG 355 - I/O

19

## Simultaneous Interrupts

- Easy with separate prioritized **IRQ** lines: pick the request with the highest priority
- What if the **IRQ** line is shared?
  - E.g., poll the status registers of I/O devices in the order of their priority and service the first interrupt source detected to request interrupt processing
  - Important:** when using a vectored interrupt scheme, must ensure that only one I/O device is selected to send an interrupt vector code over the shared bus



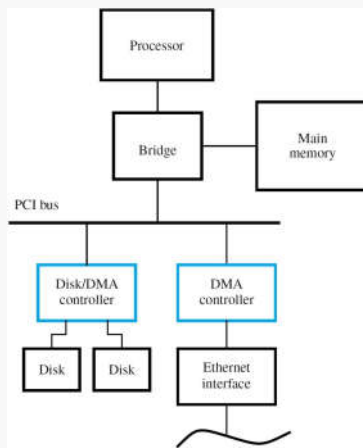
Fall 2016

UVic - CENG 355 - I/O

20

## DMA Data Transfer

- The CPU initializes DMA interface registers
- When ready, I/O device informs the DMA controller
- The DMA controller:
  - Obtains the bus by going through bus arbitration
  - Places starting memory address and control signals
  - Completes data transfer and releases the bus
  - Updates starting memory address and word count value
  - If more to transfer, loops back to repeat the process
- When done, the DMA controller notifies the CPU



Fall 2016


UVic - CENG 355 - I/O

21

## DMA Interface

- The OS puts the program that requested the transfer in the **suspended** state, initiates the DMA operation, and starts execution of some other program
- When the transfer is complete, the OS responds to the interrupt from the DMA controller by putting the suspended program into the **ready** state, so that it can be scheduled for execution

### DMA Interface Registers



Fall 2016

UVic - CENG 355 - I/O

22

## DMA Problems

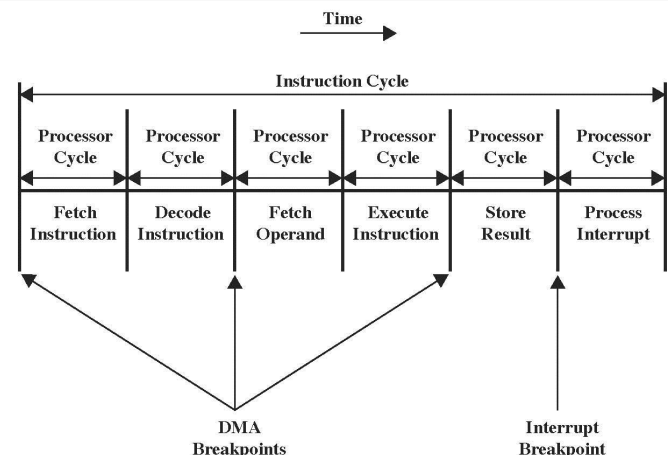
- Memory accesses by the processor and the DMA controllers must be **safe**
  - Memory locations accessed by the processor are protected against DMA, and vice versa
- DMA normally has higher priority than the CPU access
- DMA cycles can be:
  - Interwoven with the processor's access cycles on the system bus (cycle stealing)
  - Granted exclusive use of the bus without interruption (burst mode)
    - Often a DMA controller will have a data storage buffer that will be written or read in the burst mode
- What if multiple controllers try to use the bus at the same time to access the main memory?
  - Need an arbitration procedure to resolve such conflicts

Fall 2016

UVic - CENG 355 - I/O

23

## Instruction Cycle Breakpoints

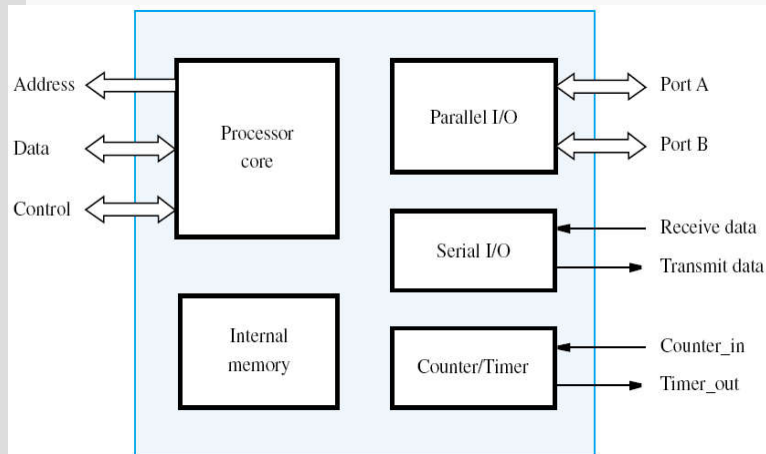


Fall 2016

UVic - CENG 355 - I/O

24

## Example Microcontroller

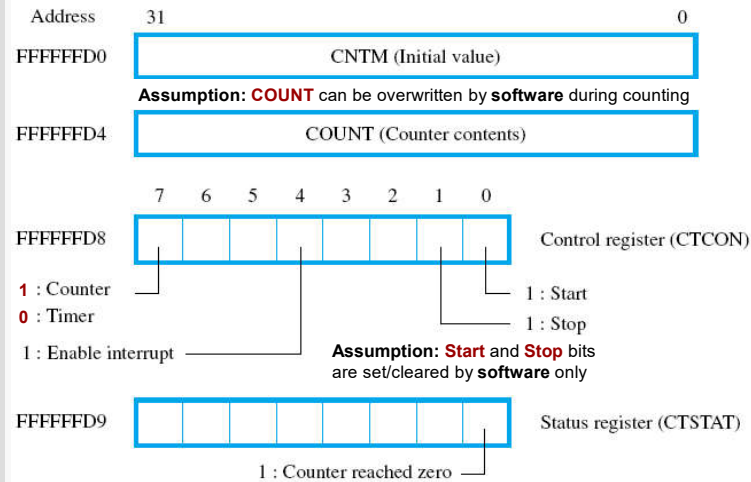


Fall 2016

UVic - CENG 355 - I/O

25

## Counter/Timer Registers

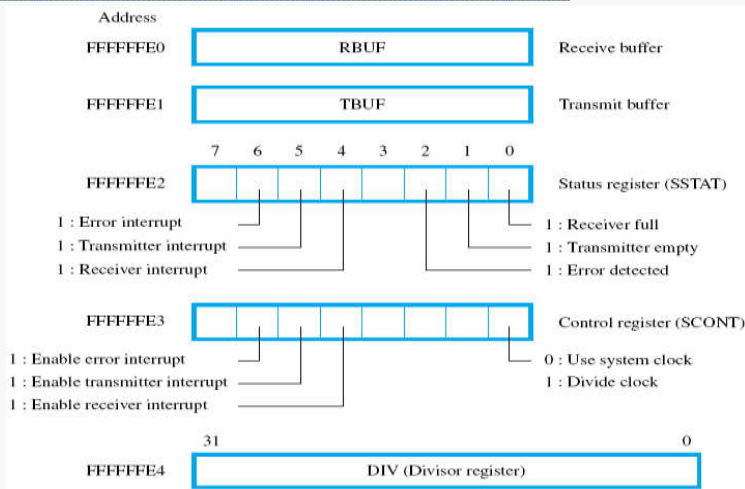


Fall 2016

UVic - CENG 355 - I/O

26

## Serial I/O Registers

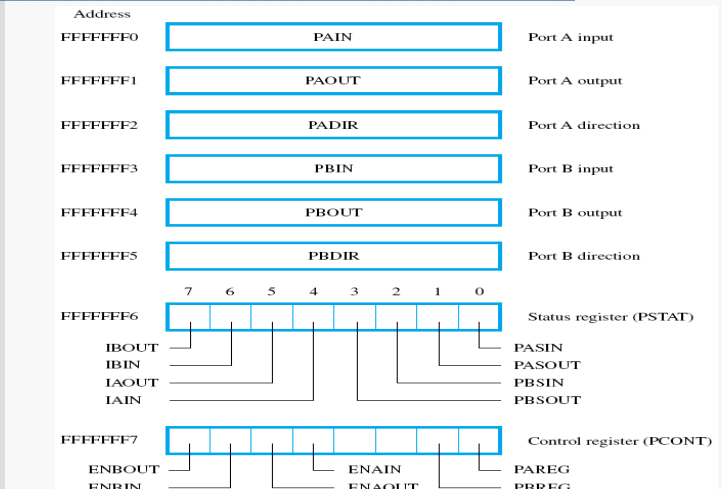


Fall 2016

UVic - CENG 355 - I/O

27

## Parallel I/O Registers



Fall 2016

UVic - CENG 355 - I/O

28

## Parallel Port Registers

- Direction registers **PADIR** and **PBDIR**
  - Bit  $PxDIR[k] = 1 \rightarrow$  pin  $Px[k]$  is output (otherwise, input)
- Control register **PCONT**
  - $ENxIN = 1 \rightarrow$  enable interrupts when input register  $PxIN$  is ready for reading
  - $ENxOUT = 1 \rightarrow$  enable interrupts when output register  $PxOUT$  is ready for writing
  - $PxREG = 1 \rightarrow$  store input data in  $PxIN$ ; otherwise, feed data directly from the pins
- Status register **PSTAT**
  - $PxSIN = 1 \rightarrow PxIN$  is ready for reading
  - $PxSOUT = 1 \rightarrow PxOUT$  is ready for writing
  - $IxIN = 1 \rightarrow PxSIN \cdot ENxIN = 1 \rightarrow$  raised interrupt
  - $IxOUT = 1 \rightarrow PxSOUT \cdot ENxOUT = 1 \rightarrow$  raised interrupt

Fall 2016

UVic - CENG 355 - I/O

29

## Interrupt Support

- Internal interrupt **IRQ**
  - Interrupt vector is at memory location **0x20**
  - CPU responds to **IRQ** interrupts only if  $PSR[6] = 1$  (i.e., bit 6 of the processor status register must be set to 1)
  - If multiple interrupts are enabled, all of them will share the same **IRQ** line
    - When **IRQ** becomes asserted, CPU must first determine the interrupt source by checking the status registers of the parallel port, serial port, and timer/counter
- External interrupt **XRQ**
  - Interrupt vector is at memory location **0x24**
  - CPU responds to **XRQ** interrupts only if  $PSR[7] = 1$  (i.e., bit 7 of the processor status register must be set to 1)
  - **XRQ** has higher priority than **IRQ**

Fall 2016

UVic - CENG 355 - I/O

30



# Subroutine Support

## ■ Processor register LR

- When entering a called subroutine, PC is first copied into LR (i.e., the current return address), and then loaded with the subroutine address

- When exiting a subroutine, the contents of LR are transferred back into PC

## ■ Processor registers IRA and IPSR

- When entering an ISR, PC is first copied into IRA, and then loaded with the ISR address (either 0x20, or 0x24); at the same time, PSR is copied into IPSR

- When exiting an ISR, the contents of IRA and IPSR are transferred back into PC and PSR, respectively

Fall 2016

UVic - CENG 355 - I/O

31

# Character Transfer (Polling)

```
/* Define register addresses */
#define RBUF (volatile unsigned char *) 0xFFFFFE0
#define SSTAT (volatile unsigned char *) 0xFFFFFE2
#define PAOUT (volatile unsigned char *) 0xFFFFF1
#define PADIR (volatile unsigned char *) 0xFFFFF2
/* Alternatively:
volatile unsigned char *RBUF = (unsigned char *) 0xFFFFFE0
volatile unsigned char *SSTAT = (unsigned char *) 0xFFFFFE2
volatile unsigned char *PAOUT = (unsigned char *) 0xFFFFF1
volatile unsigned char *PADIR = (unsigned char *) 0xFFFFF2
*/

int main() {
/* Initialize the parallel port */
    *PADIR = 0xFF;

/* Transfer characters */
    while (1) {
        while ((*SSTAT & 0x1) == 0);
        *PAOUT = *RBUF;
    }
    exit(0);
}
```

RBUF → PAOUT

Fall 2016

UVic - CENG 355 - I/O

32

# Character Transfer (Interrupts)

```
#define RBUF (volatile unsigned char *) 0xFFFFFE0
#define SCNT (volatile unsigned char *) 0xFFFFFE3
#define PAOUT (volatile unsigned char *) 0xFFFFF1
#define PADIR (volatile unsigned char *) 0xFFFFF2
#define IVECT (volatile unsigned int *) (0x20)

interrupt void intserv();

int main() {
/* Initialize the parallel port */
    *PADIR = 0xFF;

/* Initialize the interrupt mechanism */
    *IVECT = (unsigned int *) &intserv;
    asm(" MoveControl PSR, #0x40 ");
    *SCNT = 0x10;
    while (1);
    exit(0);
}

/* Interrupt service routine */
interrupt void intserv() {
    *PAOUT = *RBUF;
}
```

RBUF → PAOUT

Fall 2016

UVic - CENG 355 - I/O

33

# Timing Considerations


## ■ Our assumption so far:

- The **output** device connected to PAOUT is **faster** than the **input** device connected to RBUF
  - Characters can be sent directly from RBUF to PAOUT

## ■ What if the output device is slower than the input device in our example? – Need a memory **buffer**

- E.g., define a circular buffer (**mbuffer**) as a **64**-character array with two indices indicating the input (**fin**) and the output (**fout**) of the character queue
  - Characters are written into **mbuffer[fin]** and read from **mbuffer[fout]**

```
#define BSIZE 64
...
unsigned char mbuffer[BSIZE];
int fin = 0; int fout = 0;
```



Fall 2016

UVic - CENG 355 - I/O

34

# Polling with Buffering

```
int main() {
    unsigned char mbuffer[BSIZE];
    int fin = 0; int fout = 0;


    *PADIR = 0xFF;

    while (1) {
        while ((*SSTAT & 0x1) == 0) {
            if ((*PSTAT & 0x2) != 0) {
                *PAOUT = mbuffer[fout];
                if (fout < BSIZE-1) fout++;
                else fout = 0;
            }
        }
        mbuffer[fin] = *RBUF;
        if (fin < BSIZE-1) fin++;
        else fin = 0;
    }
    exit(0);
}
```


/\* Define circular buffer \*/  
/\* Initialize input/output indices \*/

/\* Configure Port A as output \*/

/\* While RBUF is not ready... \*/  
/\* If PAOUT is ready... \*/  
/\* Send character to PAOUT \*/  
/\* Update output index \*/



/\* Get character from RBUF \*/  
/\* Update input index \*/



Fall 2016

UVic - CENG 355 - I/O

35

# Interrupts with Buffering

```
unsigned char mbuffer[BSIZE];
int fin = 0; int fout = 0;

int main() {
    *PADIR = 0xFF;
    *IVECT = (unsigned int *) &intserv;
    asm(" MoveControl PSR, #0x40 ");
    *SCNT = 0x10;
    *PCNT = 0x20;
    while (1);
    exit(0);
}

interrupt void intserv() {
    if ((*SSTAT & 0x1) != 0) {
        mbuffer[fin] = *RBUF;
        if (fin < BSIZE-1) fin++;
        else fin = 0;
    }
    if ((*PSTAT & 0x2) != 0) {
        *PAOUT = mbuffer[fout];
        if (fout < BSIZE-1) fout++;
        else fout = 0;
    }
}
```

/\* Define circular buffer outside main() \*/  
/\* Initialize input/output indices \*/

/\* Configure Port A as output \*/  
/\* Set up interrupt vector \*/  
/\* CPU responds to IRQ \*/  
/\* Enable RBUF interrupts \*/  
/\* Enable PAOUT interrupts \*/  
/\* Empty loop, but can code other tasks here \*/

/\* Receiver interrupt flag set? \*/  
/\* Get character from RBUF \*/  
/\* Update input index \*/

/\* PAOUT flag set? \*/  
/\* Send character to PAOUT \*/  
/\* Update output index \*/

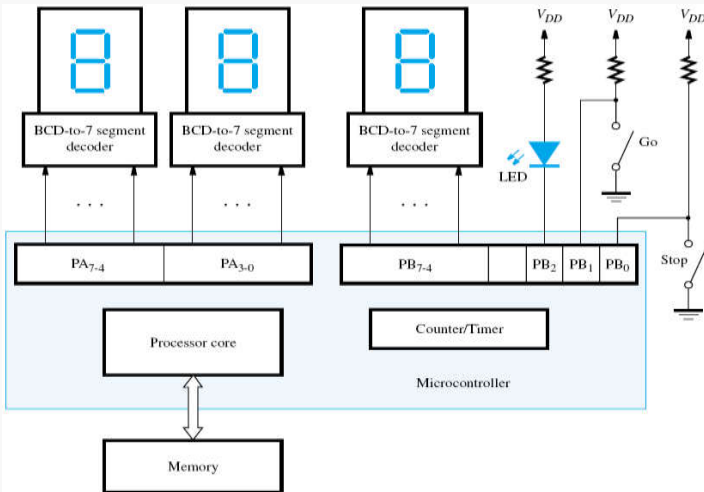
/\* Note: RBUF interrupts have higher priority → SSTAT is checked first \*/

Fall 2016

UVic - CENG 355 - I/O

36

## Reaction Timer Circuit



Fall 2016

UVic - CENG 355 - I/O

37

## Reaction Timer Operation

- The system is activated by pressing the **Go** switch
- Upon activation, the 3-digit display is set to **000** and the LED is turned off
- After a 3-second delay, the LED is turned **on** and the timing process begins
  - Timing clock is **100 MHz**
- When the **Stop** switch is pressed, the timing process is stopped, the LED is turned **off**, and the elapsed time is displayed on the 3-digit display
  - The elapsed time is calculated and displayed in hundredths of a seconds (assuming < 10 seconds)

Fall 2016

UVic - CENG 355 - I/O

38

## Implementation Ideas

- Use a **wait** loop to **poll** the state of the **Go** switch
- Once **Go** has been closed (**PB[1] = 0**), **wait** for 3 seconds, and turn the LED **on** (**PB[2] = 0**)
- Set the initial counter value to 0xFFFFFFFF, which is decremented each clock cycle during the timing process
- Start the timing process
- Use a **wait** loop to **poll** the state of the **Stop** switch
- Once **Stop** has been pressed (**PB[0] = 0**), stop the timer and turn the LED **off** (**PB[2] = 1**)
  - Delay = (0xFFFFFFFF – Counter Value)/1,000,000
- Convert the elapsed time into 3 digits and send them to the display

Fall 2016

UVic - CENG 355 - I/O

39

## C Program I

```
#define PAOUT (volatile unsigned char *) 0xFFFFFFFF1
#define PADIR (volatile unsigned char *) 0xFFFFFFFF2
#define PBIN (volatile unsigned char *) 0xFFFFFFFF3
#define PBOUT (volatile unsigned char *) 0xFFFFFFFF4
#define PBDIR (volatile unsigned char *) 0xFFFFFFFF5
#define CNTM (volatile unsigned int *) 0xFFFFFDD0
#define COUNT (volatile unsigned int *) 0xFFFFFDD4
#define CTCON (volatile unsigned char *) 0xFFFFFDD8
#define CTSTAT (volatile unsigned char *) 0xFFFFFDD9

int main() {
    unsigned int counter_value, total_count;
    unsigned int actual_time, seconds, tenths, hundredths = 0;
    *CTCON = 0x2;           /* Stop Timer (if running) */
    *PADIR = 0xFF;          /* Configure Port A */
    *PBDIR = 0xF4;          /* Configure Port B */
    *PAOUT = 0x0;           /* Display 0's */
    *PBOUT = 0x4;           /* Turn off LED, display 0 */
}
```

Fall 2016

UVic - CENG 355 - I/O

40

## C Program II

```
while (1) {
    while ((*PBIN & 0x2) != 0); /* Infinite loop */
    *CNTM = 300000000;          /* Wait for Go to be pressed */
    *CTSTAT = 0x0;              /* Counting for 3 seconds */
    *CTCON = 0x1;               /* Clear "Reached 0" flag */
    while ((*CTSTAT & 0x1) == 0); /* Start countdown */
    *CTCON = 0x2;               /* Wait until 0 is reached */
    *CTCON = 0x2;               /* Stop countdown */

    /* Start the timing process */
    *CNTM = 0xFFFFFFFF;         /* Initial counter value */
    *CTSTAT = 0x0;              /* Clear "Reached 0" flag */
    *PBOUT = 0x0;               /* Turn on LED */
    *CTCON = 0x1;               /* Start countdown */
    while ((*PBIN & 0x1) != 0); /* Wait for Stop to be pressed */

    /* Stop timing process */
    *CTCON = 0x2;               /* Stop countdown */
    *PBOUT = (char)((hundredths << 4 | 0x4)); /* Turn off LED */
}
```

Fall 2016

UVic - CENG 355 - I/O

41

## C Program III

```
/* Compute elapsed time */
counter_value = *COUNT; /* Read current counter value */
total_count = (0xFFFFFFFF - counter_value);
actual_time = total_count / 1000000; /* Units = second/100 */

seconds = actual_time / 100;
tenths = (actual_time - seconds*100) / 10;
hundredths = actual_time - (seconds*100 + tenths*10);

/* Display */
*PAOUT = (char)((seconds << 4 | tenths);
*PBOUT = (char)((hundredths << 4 | 0x4); /* Keep LED off */

}

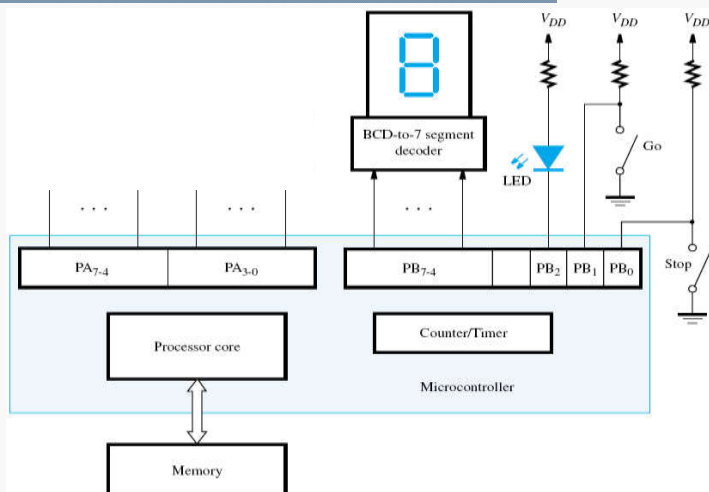
exit(0);
}
```

Fall 2016

UVic - CENG 355 - I/O

42

## Another Example I



Fall 2016

UVic - CENG 355 - I/O

43

## Another Example II

- The microcontroller is responsible for two tasks:
  - Measuring reaction time at Port B (main program)
  - Redirecting data from RBUF to PAOUT (interrupt service routine accessed via memory location **0x20**)
    - Port A is always ready to receive data
    - Processor's interrupt-enable (IE) bit is **PSR[6]**
- Specifications:
  - When **Go** is pressed (**PB[1] = 0**), display shows **0**, and LED is off (**PB[2] = 1**)
    - After **2 seconds** LED is turned on, and the timing process begins; during this timing process interrupts are not allowed
  - When **Stop** is pressed (**PB[0] = 0**), the timing process stops, LED is turned off, and the elapsed time is displayed

Fall 2016

UVic - CENG 355 - I/O

44

## Another Example III

```

/* Define register addresses */
#define RBUF (volatile unsigned char *) 0xFFFFFE0
#define SCNT (volatile unsigned char *) 0xFFFFFE3
#define PAOUT (volatile unsigned char *) 0xFFFFF1
#define PADIR (volatile unsigned char *) 0xFFFFF2
#define PBIN (volatile unsigned char *) 0xFFFFF3
#define PBOUT (volatile unsigned char *) 0xFFFFF4
#define PBDIR (volatile unsigned char *) 0xFFFFF5
#define CNTM (volatile unsigned int *) 0xFFFFFD0
#define COUNT (volatile unsigned int *) 0xFFFFFD4
#define CTCON (volatile unsigned char *) 0xFFFFFD8
#define CTSTAT (volatile unsigned char *) 0xFFFFFD9
#define IVECT (volatile unsigned int *) (0x20)

interrupt void intserv();
...
    
```

Fall 2016

UVic - CENG 355 - I/O

45

## Another Example IV

```

int main() {
    unsigned int count, time;

    *CTCON = 0x2;
    *PADIR = 0xFF;
    *PBDIR = 0xF4;
    *IVECT = (unsigned int *) &intserv;
    asm( " MoveControl PSR, #0x40 " );
    *SCNT = 0x10;
    *PBOUT = 0x4;
    while (1) {
        while ((*PBIN & 0x2) != 0);
        *CNTM = 200000000;
        *CTSTAT = 0x0;
        *CTCON = 0x1;
        while ((*CTSTAT & 0x1) == 0);
        *CTCON = 0x2;

        /* Elapsed/displayed time */
        /* Stop Timer (if running) */
        /* Configure Port A */
        /* Configure Port B */
        /* Set up interrupt vector */
        /* CPU responds to IRQ */
        /* Enable RBUF interrupts */
        /* Turn off LED, display 0 */
        /* Polling loop */
        /* Wait for Go */
        /* Counting for 2 seconds */
        /* Clear "Reached 0" flag */
        /* Start countdown */
        /* Wait until 0 is reached */
        /* Stop countdown */
    }
}
    
```

Fall 2016

UVic - CENG 355 - I/O

46

## Another Example V

```

*CNTM = 0xFFFFFFFF;
*CTSTAT = 0x0;
asm( " BitClear #6, PSR " );
*PBOUT = 0x0;
*CTCON = 0x1;
while ((*PBIN & 0x1) != 0);
*CTCON = 0x2;
*PBOUT = 0x4;
asm( " BitSet #6, PSR " );
count = (0xFFFFFFFF - *COUNT);
time = count/10000000;
*PBOUT = (char) ((count << 4 | 0x4);

/* Initial counter value */
/* Clear "Reached 0" flag */
/* Disable CPU interruption */
/* Turn on LED */
/* Start countdown */
/* Wait for Stop */
/* Stop countdown */
/* Turn off LED */
/* Enable CPU interruption */
/* Elapsed clock cycles */
/* Elapsed time, in 1/10 sec */
/* Update display, LED off */

}
exit(0);

interrupt void intserv() {
    *PAOUT = *RBUF;
}
    
```

Fall 2016

UVic - CENG 355 - I/O

47

## Operating System (OS)

- **OS** coordinates all activities in a computer system
  - OS manages processing, memory, I/O resources
  - OS interprets user commands, allocates storage, transfers information, handles I/O operations
  - OS uses a loader to execute application programs
- **Loader** is invoked when user types commands or clicks on icons in GUI (graphical user interface)
  - User's input identifies the object file that has information on the memory address and length of a program
  - Loader transfers the program from disk to memory and branches to its starting address
  - At program termination, loader recovers space in memory and awaits the next command

Fall 2016

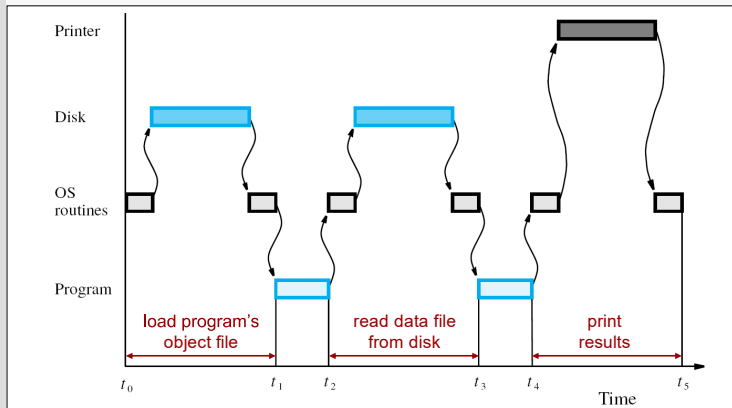
UVic - CENG 355 - I/O

48



## Single-Program Example

- Read a data file from the disk into the memory, perform some computations, and print the results



Fall 2016

UVic - CENG 355 - I/O

49

## Support for OS

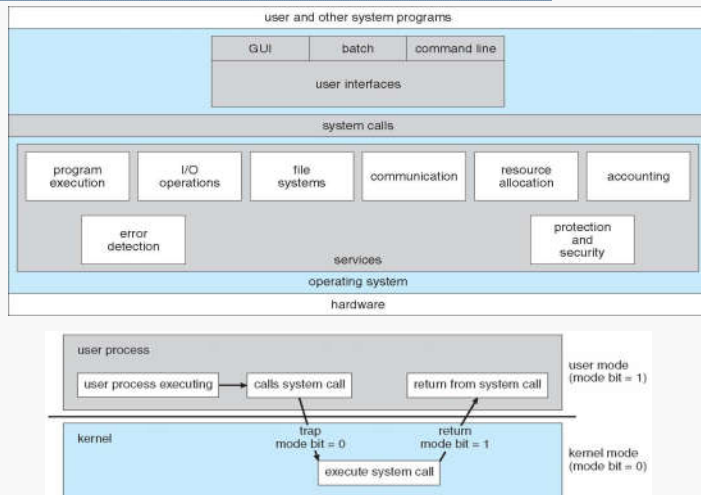
- Bootstrap program** is loaded and executed at power-up or reboot
  - Typically stored in ROM (read-only memory), generally known as **firmware**
  - Initializes all aspects of the system, loads the OS **kernel** into the main memory, and starts its execution
- OS operation is driven by OS service requests (traps), I/O interrupts, and other exceptions
- Dual-mode operation: **user mode** and **kernel mode**
  - Hardware provides for a **mode bit** to distinguish when system is running a user code or a kernel code
  - Some instructions can be designated as **privileged**, executable only in the kernel mode

Fall 2016

UVic - CENG 355 - I/O

50

## OS Services




Fall 2016

UVic - CENG 355 - I/O

51

## System Calls

- Each system call typically has a number associated with it
- System-call interface:
  - Maintains a parameter table indexed according to system call numbers
  - Invokes an intended system call in the OS kernel
  - Returns status of the system call and any return values



Fall 2016

UVic - CENG 355 - I/O

52

## Multitasking Example

- To allow for fair CPU usage when managing execution of multiple programs (tasks), OS uses hardware timer interrupts for time slicing
  - Each task gets its time slice
  - More important tasks can be given longer time slices
- Assume: Programs **A** and **B** have been initiated, and the CPU is currently executing **A**
  - When **A**'s time slice expires, the **SCHEDULER** program is entered, and **A**'s registers are saved
  - OS then selects **B**, restores **B**'s register values, and uses return-from-interrupt to resume **B**
- Process = program + program state**
  - Example **process states**: *Runnable*, *Running*, *Blocked*

Fall 2016

UVic - CENG 355 - I/O

53

OSINIT	Set interrupt vectors: Timer interrupt ← SCHEDULER Software interrupt ← OSSERVICES I/O interrupt ← IODATA
OSSERVICES	Examine stack or processor registers to determine requested operation. Call appropriate routine.
SCHEDULER	Save program state of current running process. Select another runnable process. Restore saved program state of new process. Return from interrupt.
(a) OS initialization, services, and scheduler	
IOINIT	Set requesting process state to Blocked. Initialize memory buffer address pointer and counter. Call device driver to initialize device and enable interrupts in the device interface. Return from subroutine.
IODATA	Poll devices to determine source of interrupt. Call appropriate driver. If END = 1, then set I/O-blocked process state to Runnable. Return from interrupt.

(b) I/O routines

Fall 2016

UVic - CENG 355 - I/O

54

# Process Concept

- OS is responsible for creating, suspending, resuming, deleting processes and providing mechanisms for process synchronization, communication, deadlock handling
- A process includes multiple parts:
  - **Text** section containing program code
  - **Data** section containing global variables and static data
  - **Stack** containing temporary data
    - Function parameters, return addresses, local variables
  - **Heap** containing dynamically allocated memory objects
  - Current activity information (**program state**), including PC and processor registers
    - Program state is different from process state (see next slide)


Fall 2016

UVic - CENG 355 - I/O

55

# Process State

- During its lifetime, a process changes its **state**
  - **new**: process is being created
  - **ready (runnable)**: process is ready to execute
  - **running**: process is being executed
  - **waiting (blocked)**: process is waiting for some event to occur
  - **terminated**: process has been finished or aborted
- Other states (e.g., **suspended**) may also be present

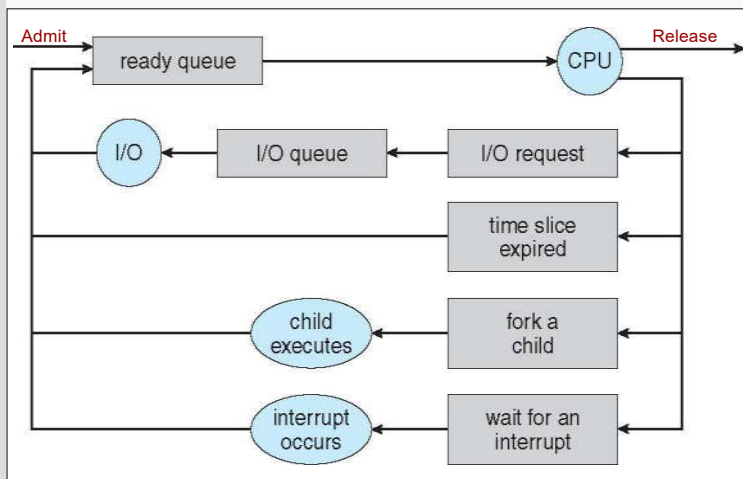


Fall 2016

UVic - CENG 355 - I/O

56

# Process Scheduling View

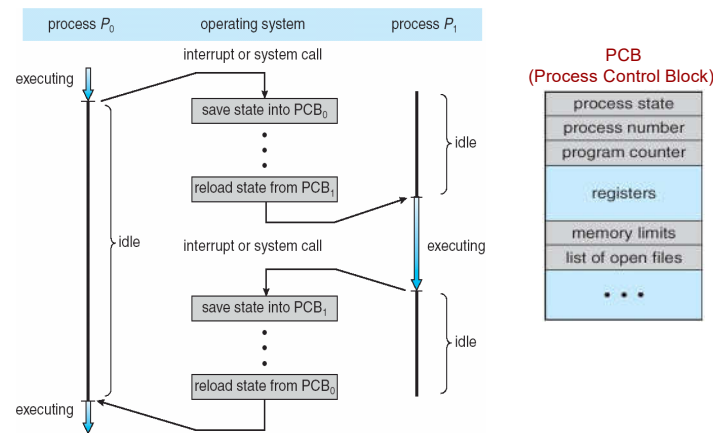


Fall 2016

UVic - CENG 355 - I/O

57

# Process (Context) Switching




Fall 2016

UVic - CENG 355 - I/O

58

# Real-Time Periodic Tasks



Task  $T_i = (C_i, D_i, P_i, \phi_i)$ , where:

$C_i$  – worst case execution time (WCET),  $\phi_i$  – initial delay,  
 $P_i$  – period,  $D_i$  – deadline (often same as  $P_i$ )

Fall 2016

UVic - CENG 355 - I/O

59

# Priority-Driven Scheduling I

- Single-processor scheduling problem:
  - Given task set  $\{T_1, T_2, \dots\}$ , for each task  $T_i$  and its every arrival (activation) at time  $\phi_i + kP_i$ , determine its start time  $s_{ik} \geq \phi_i + kP_i$ , such that  $s_{ik} + C_i \leq \phi_i + kP_i + D_i$  (i.e., its finish time meets task deadline)
- Single-processor scheduling algorithm:
  - When task  $T_i$  arrives at time  $\phi_i + kP_i$ , assign a certain priority value  $\tau_{ik}$  to it and place  $T_i$  into the prioritized queue of tasks ready for execution
    - Tasks in the ready queue are always ordered in the increasing order of their priorities
  - If  $\tau_{ik}$  is greater than the priority of a task currently being executed, suspend that task and start  $T_i$
  - Otherwise, once a current task is finished, say at time  $t$ , start the next highest-priority task in the ready queue

Fall 2016

UVic - CENG 355 - I/O

60

## Priority-Driven Scheduling II

- For a task to have any meaningful priority (before entering the priority queue), it must arrive and be ready for execution
- Examples of **static priority** assignment
  - Rate Monotonic (**RM**):  $\tau_{ik} = 1/P_i$  (independent of  $k$ )
  - Deadline Monotonic (**DM**):  $\tau_{ik} = 1/D_i$  (independent of  $k$ )
- Examples of **dynamic priority** assignment
  - Earliest Deadline First (**EDF**):  $\tau_{ik} = 1/(\phi_i + kP_i + D_i)$
  - Least Laxity First (**LLF**):  $\tau_{ik} = 1/(\phi_i + kP_i + D_i - t - \Delta C_i)$ , where  $\Delta C_i$  is the remaining execution time of  $T_i$ 
    - **Note:**  $\Delta C_i = C_i$  if  $T_i$  has not been suspended previously
- CPU utilization:  $C_1/P_1 + C_2/P_2 + \dots$

## RM Scheduling Example

- Set of three pre-emptive tasks  $T_i = (C_i, D_i, P_i, \phi_i)$ 
  - $\{T_1=(10,30,30,0), T_2=(17,30,40,0), T_3=(10,120,120,0)\}$
- RM scheduling prioritization:
  - $\tau_{1k} = 1/30$  (highest),  $\tau_{2k} = 1/40$ ,  $\tau_{3k} = 1/120$  (lowest)
- CPU utilization:
  - $10/30 + 17/40 + 10/120 = 84.2\%$

