

Transmission Control Protocol

Internet Architecture

Internet Architecture

- End-to-end semantics:
 - functions placed at low levels of a system may be redundant or of little value when compared with the cost of providing them at that low level

e.g., end-to-end caretaking

A --- B --- C --- D --- E

J.H. Saltzer, D.P. Reed and D.D. Clark. End-to-end design argument in system design, ACM TOCS, Vol 2, Number 4, November 1984, pp.277-288.

Internet Architecture (Cont'd)


- Intermediate node (e.g., routers): packet forwarding, FIFO (first in first out), discard packets when overflow, stateless, know almost nothing about end-to-end sessions
→ simple and robust
- End host: know almost nothing about network internals, manage all end-to-end session-related states
→ complex and intelligent

Transmission Control Protocol (TCP)

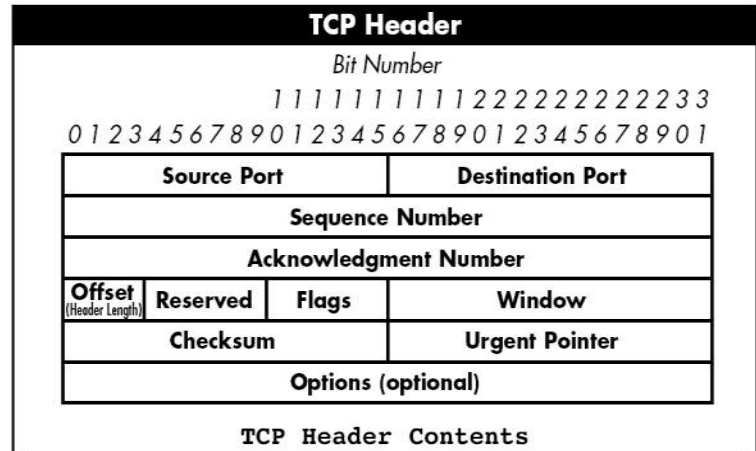
Transmission Control Protocol

- TCP
 - a *connection-oriented, end-to-end reliable*, and stream-like transport protocol over the connectionless, unreliable, and datagram-based IP service

Connection-oriented




- connection establishment: associate endpoints
 - 3-way handshakes open bi-directional data channels
 - data only transferred after connection established (with `bsd` socket interface)
 - 2-way handshakes close data channels individually (i.e., graceful close, other termination forms exist)



Error Control

- error control: detect and recover packet errors
 - lost packet: accumulative acknowledgment from receiver, timeout at sender, retransmit if timeout occurs
 - duplicated packet: discard according to sequence number (note: the relation between window space and sequence space) at receiver, no further recovery needed




Error Control (Cont'd)

- corrupted packet: discard according to checksum (also include a pseudo IP header) at receiver, as a lost packet
- reordered packet: reorder according to sequence number at receiver, no further recovery needed

Flow Control

- flow control: coordinate endpoints
 - slide window based flow control (in sequence space)
 - receiver's acknowledgment advances window and adjust window size
 - sender can not exceed window-bottom plus window-size in sequence space




Congestion Control

- congestion control: coordinate endpoints and network
 - was introduced in later 1980's
 - now standard in every TCP implementation
 - at the heart of TCP/IP research


Internet Congestion Collapse

- In the late 80s, the Internet suffered a series of congestion collapse




Congestion Control Objective

network power (throughput/delay)



- load-power curve
 - low-load: significant gain
 - mid-load: negligible gain
 - high-load: negative gain (collapse)

ack self-clocking



In equilibrium: a new packet isn't put into the network until an old packet leaves.


TCP Variants – TCP Tahoe

- cwnd* (congestion window)
- initial operation
 - $cwnd = 1 \text{ mss}$ (maximum segment size)
 - ssthresh* (slow-start threshold)
 - max. outstanding data
 - $swnd = \min\{cwnd, rwnd, \text{sender buffer size}\}$

Van Jacobson and Mike Karels. Congestion avoidance and control. ACM Computer Communication Review, 18(4):314–329, August 1990. Revised version of his SIGCOMM '88 paper.


Slow-start

- slow-start: when $cwnd < ssthresh$
 - $cwnd += 1 \text{ mss}$ on a new ack
 - for ack-every-segment receiver
 $cwnd_{i+1} = cwnd_i * 2$
 - for delayed-ack receiver
 $cwnd_{i+1} \approx cwnd_i * 1.5$



Congestion Avoidance

- congestion avoidance: when $cwnd \geq ssthresh$
 - new ack, $cwnd += \frac{mss}{cwnd} \text{ mss}$
 - for ack-every-segment receiver,
 $cwnd_{i+1} = cwnd_i + 1 \text{ mss}$
 - for delayed-ack receiver $cwnd_{i+1} \approx cwnd_i + \frac{mss}{2}$




Timeout

- packet is assumed lost when timeout occurs
- packet loss is assumed due to network congestion
- retransmit lost packet

$$ssthresh = \max\left\{2 \text{ mss}, \frac{\min\{cwnd, \text{amount-of-in-flight-data}\}}{2}\right\}$$

$$cwnd = 1 \text{ mss}$$


- followed by slow start



Fast Retransmit

- retransmission after timeout is slow
- receiver returns duplicated acknowledgment on non-in-order packets (due to reordered or lost packets)
- sender assumes the most unacknowledged packet is lost if number of *dupacks* exceeds a certain threshold (3 in most popular TCP implementations)
- retransmit the most unacknowledged packet, adjust *cwnd* and *ssthresh* as a timeout occurs
- followed by slow-start procedure

TCP Variants – TCP Reno



- Fast recovery: for triple-dupack, retransmit most unacknowledged packets, set *ssthresh* as fast retransmit, set *cwnd* = *ssthresh*

TCP New Reno

- partial acknowledgment: recover multiple packet losses
 - when entering fast recovery, record highest *seqno* ever sent as (*record*)
 - a *newack* not covers *record* is evaluated as a *pack*, retransmit the most unacknowledged packet
 - when *newack* covers *record*, exit fast recovery

TCP SACK

- selective acknowledgment: recover multiple packet losses
 - receiver returns *sack* option and the first field specifies a *hole* in queue when a packet arrives
 - duplicates the first two *sack* fields in previous *sack* option
 - if a packet is reported in a *hole* three times, it is assumed lost and is arranged for retransmission

TCP Throughput


- For saturated sender, TCP throughput (bytes/sec) can be approximated as the sender window size (bytes) over round-trip time (seconds).
- Assuming $rwnd \gg cwnd$, the sender window size equals $cwnd$, which is determined according to the loss events.

25

TCP Throughput

- triple-*dupack* model

$$\frac{3w^2}{8} = \frac{1}{p}, w: cwnd, p: \text{loss rate}$$



- limitation: small packet loss rate, periodically loss
- other models: consider *delack* and assume drop-window-tail loss

26

TCP Throughput

- model with timeout
 - BSD timeout backoff: 1, 2, 4, 8, 16, 32, 64, 64, ...
 - correlated losses in rtt , independent between rtt packet
 - $\frac{1}{rtt\sqrt{\frac{2bp}{3}} + T_0 \min(1, 3\sqrt{\frac{3bp}{8}p(1+32p^2)})}$, b : *ack* pattern
 - TCP-Friendly Rate Control (TFRC) protocol

27

28