

Building up a heap in linear
time

Building up a heap

Building up a heap

- n standard insert-operations for a heap result in $O(n \log(n))$ time.

Building up a heap

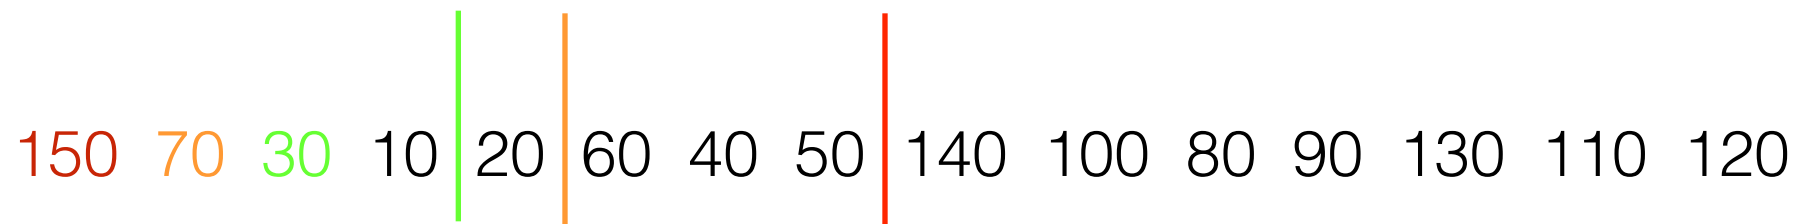
- n standard insert-operations for a heap result in $O(n \log(n))$ time.
- Can we build up a heap for n given elements faster? Is $O(n)$ possible?

150 70 30 10 20 60 40 50 | 140 100 80 90 130 110 120


150 70 30 10 20 60 40 50 | 140 100 80 90 130 110 120

150 70 30 10 20 | 60 40 50 | 140 100 80 90 130 110 120


150 70 30 10 20 60 40 50 140 100 80 90 130 110 120



150 70 30 10 | 20 | 60 40 50 | 140 100 80 90 130 110 120




150 70 30 10 20 60 40 50 140 100 80 90 130 110 120




150 70 30 10 20 60 40 50 140 100 80 90 130 110 120


Vertical lines are placed between 10 and 20 (green), 20 and 60 (orange), and 50 and 140 (red).




150 70 30 10 20 60 40 50 140 100 80 90 130 110 120




150 70 30 10 20 60 40 50 | 140 100 80 90 130 110 120




150 70 30 10 20 60 40 50 140 100 80 90 130 110 120




150 70 30 10 20 60 40 50 140 100 80 90 130 110 120




150 70 30 10 20 60 40 50 140 100 80 90 130 110 120




150 70 30 10 20 60 40 50 140 100 80 90 130 110 120







150 70 30 10 20 60 40 50 140 100 80 90 130 110 120







150 70 30 10 20 60 40 50 140 100 80 90 130 110 120
































150 70 30 10 20 60 40 50 140 100 80 90 130 110 120



































































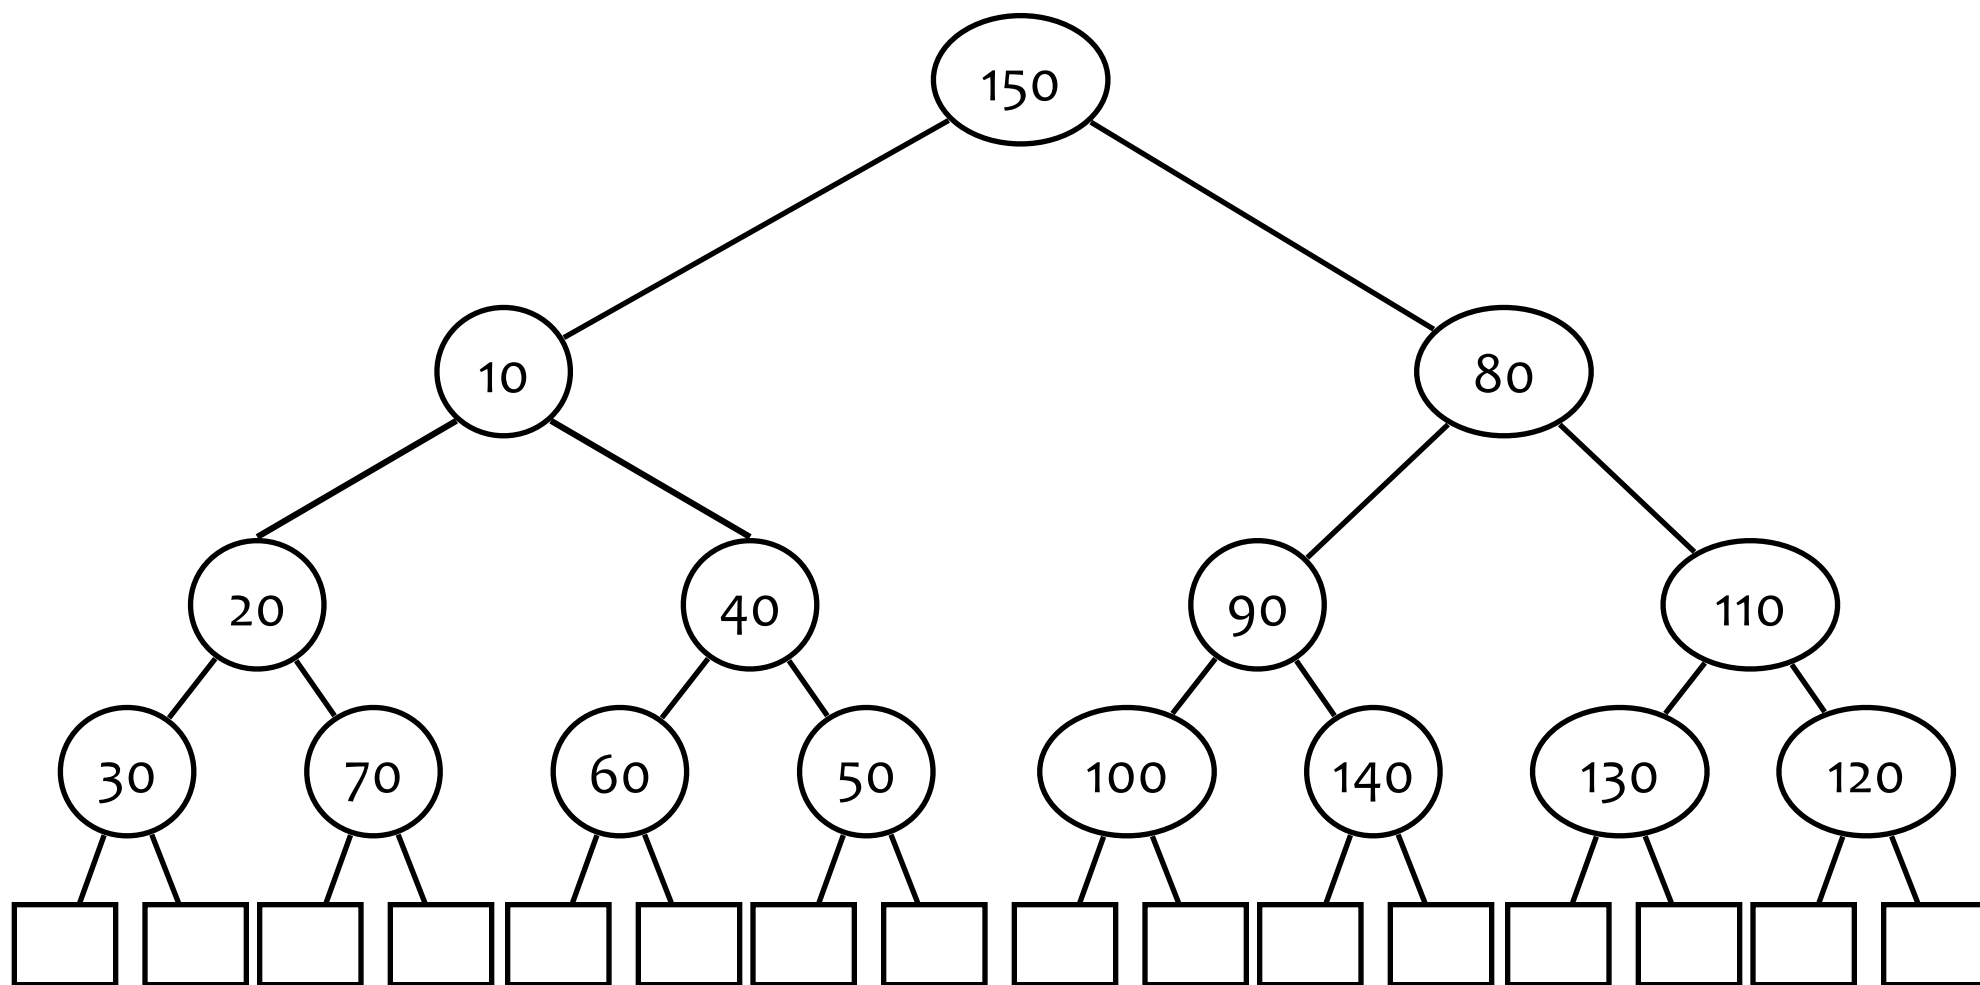




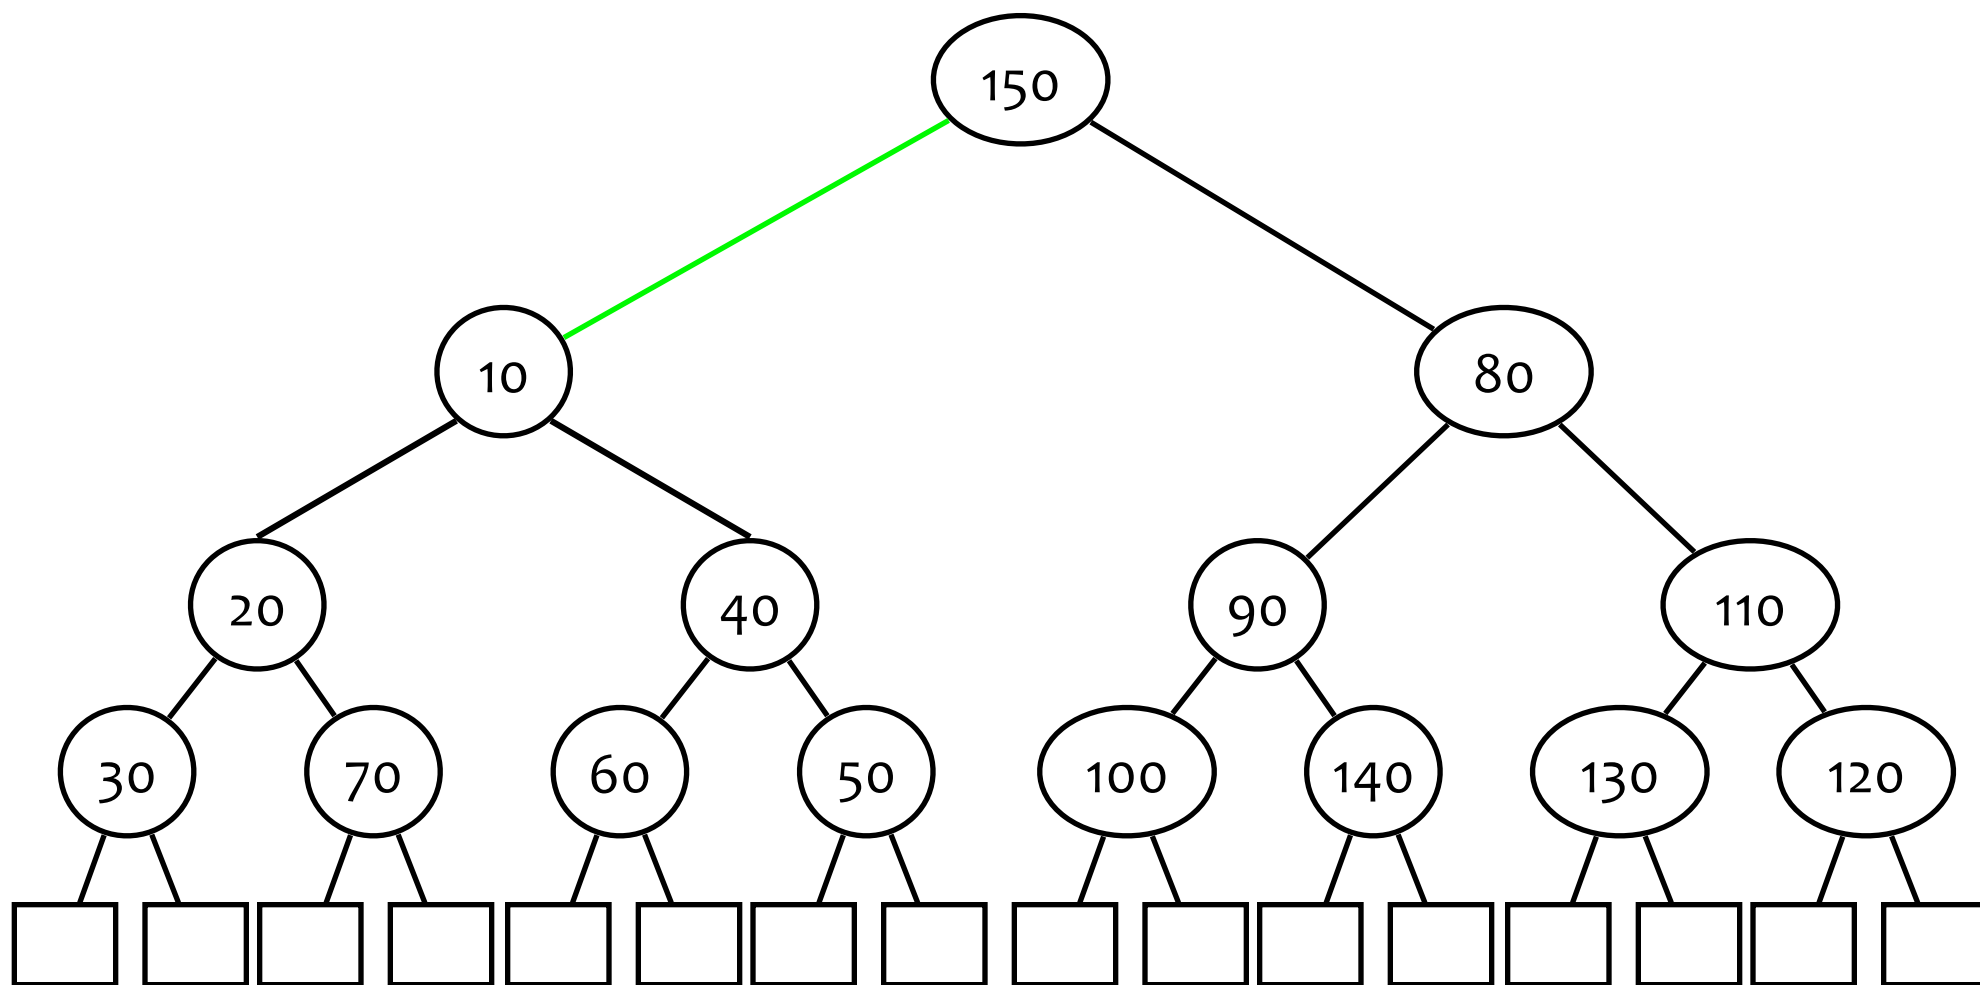
150	70	30	10	20	60	40	50	140	100	80	90	130	110	120
-----	----	----	----	----	----	----	----	-----	-----	----	----	-----	-----	-----



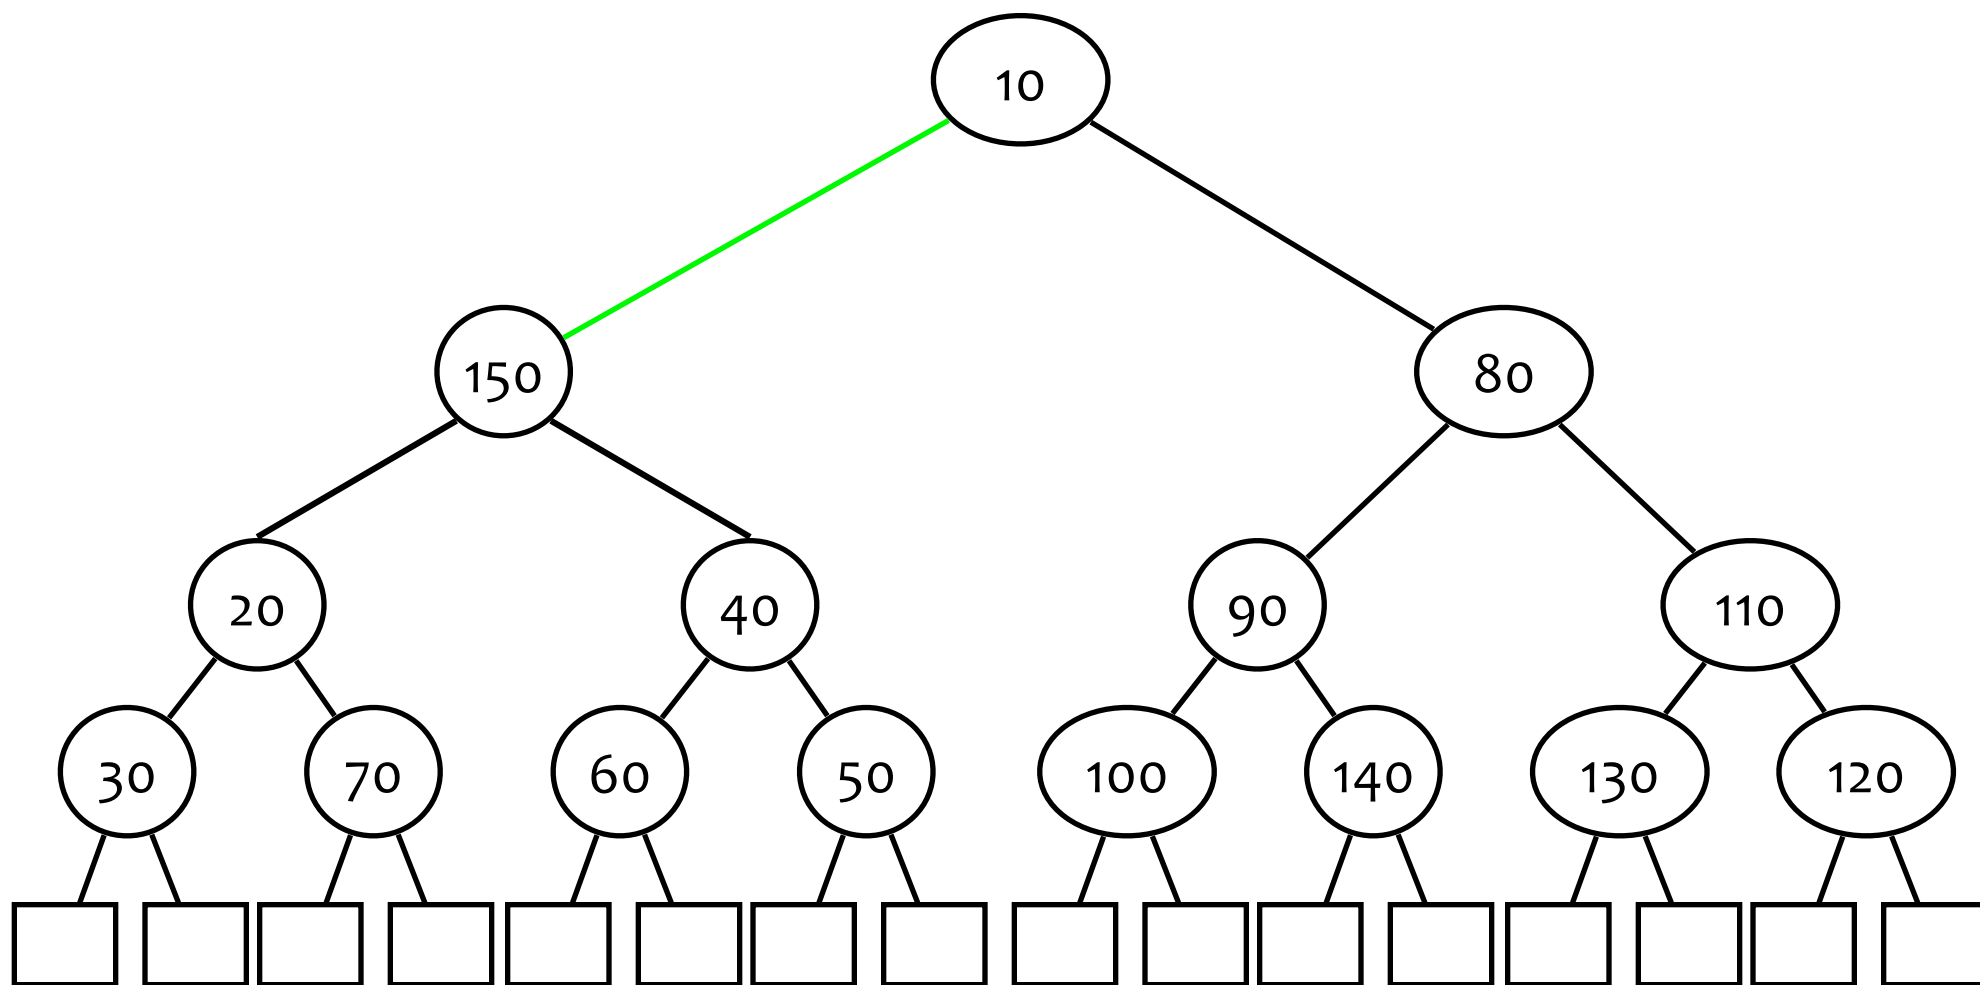
150	70	30	10	20	60	40	50	140	100	80	90	130	110	120
-----	----	----	----	----	----	----	----	-----	-----	----	----	-----	-----	-----



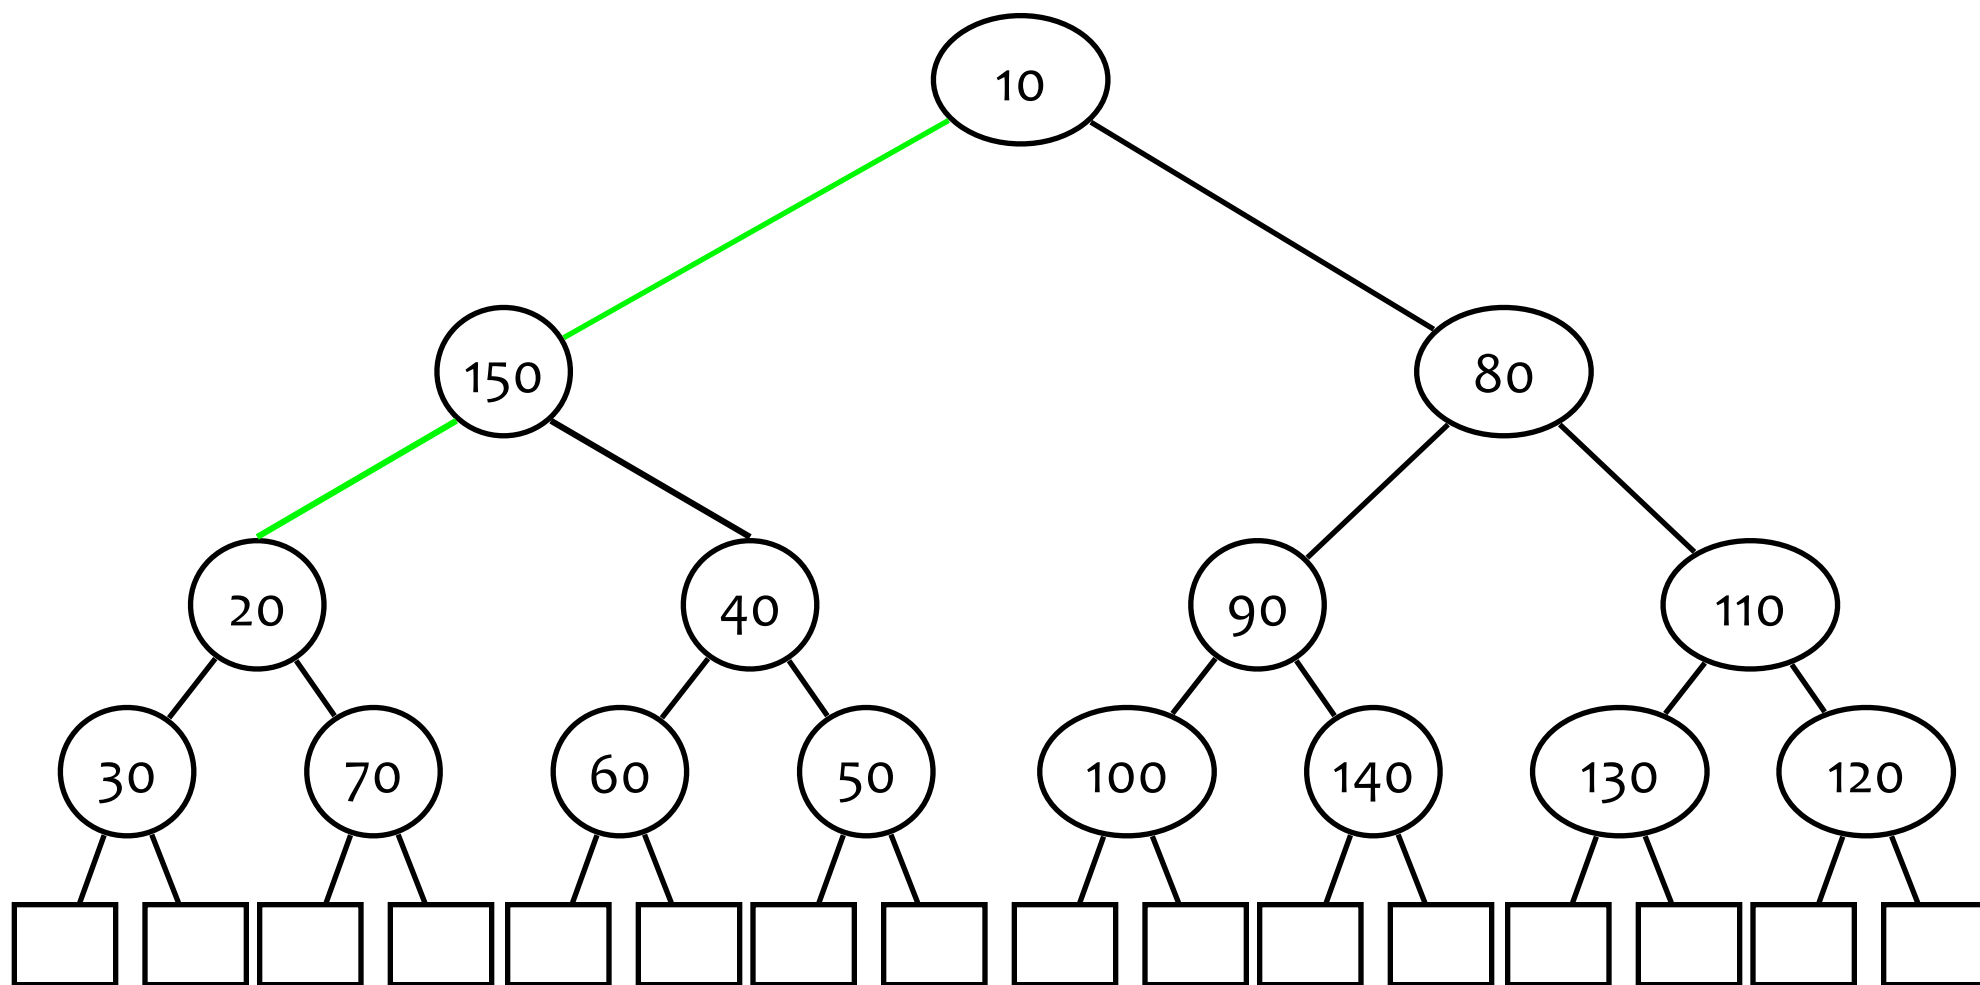
150	70	30	10	20	60	40	50	140	100	80	90	130	110	120
-----	----	----	----	----	----	----	----	-----	-----	----	----	-----	-----	-----



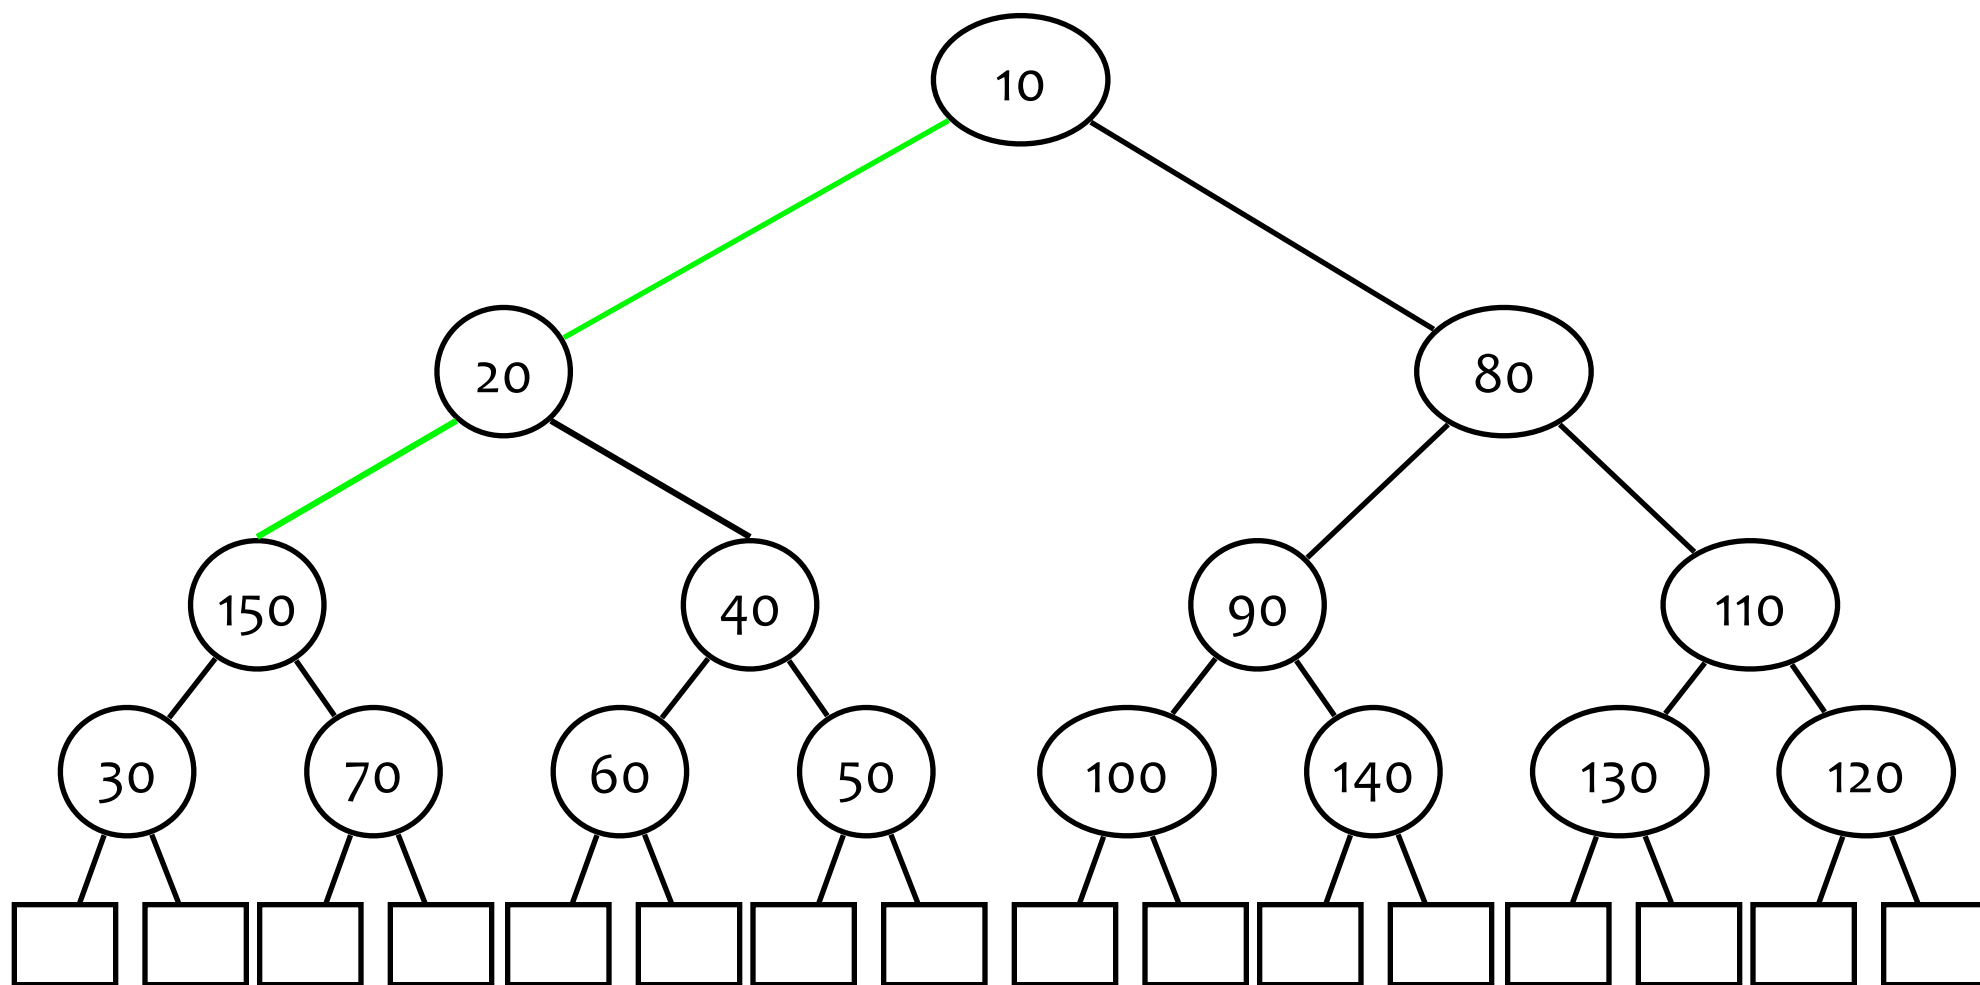
150	70	30	10	20	60	40	50	140	100	80	90	130	110	120
-----	----	----	----	----	----	----	----	-----	-----	----	----	-----	-----	-----



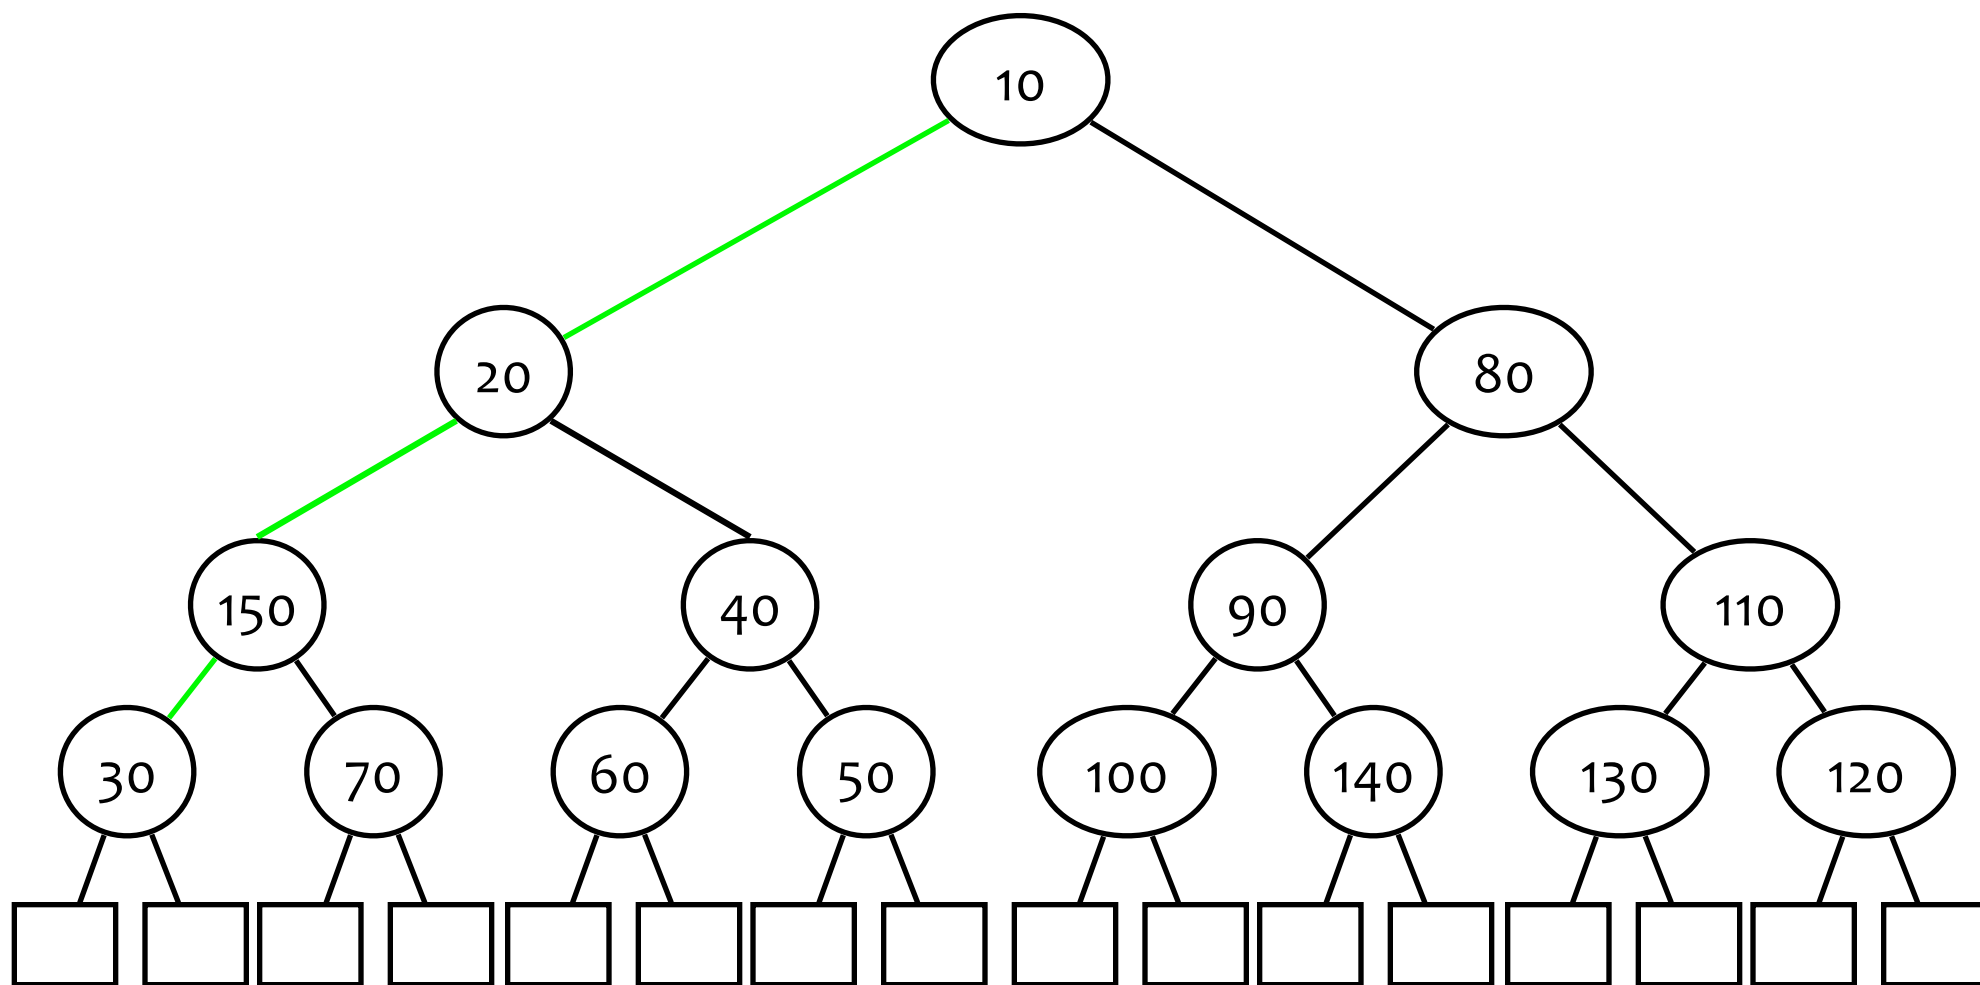
150	70	30	10	20	60	40	50	140	100	80	90	130	110	120
-----	----	----	----	----	----	----	----	-----	-----	----	----	-----	-----	-----



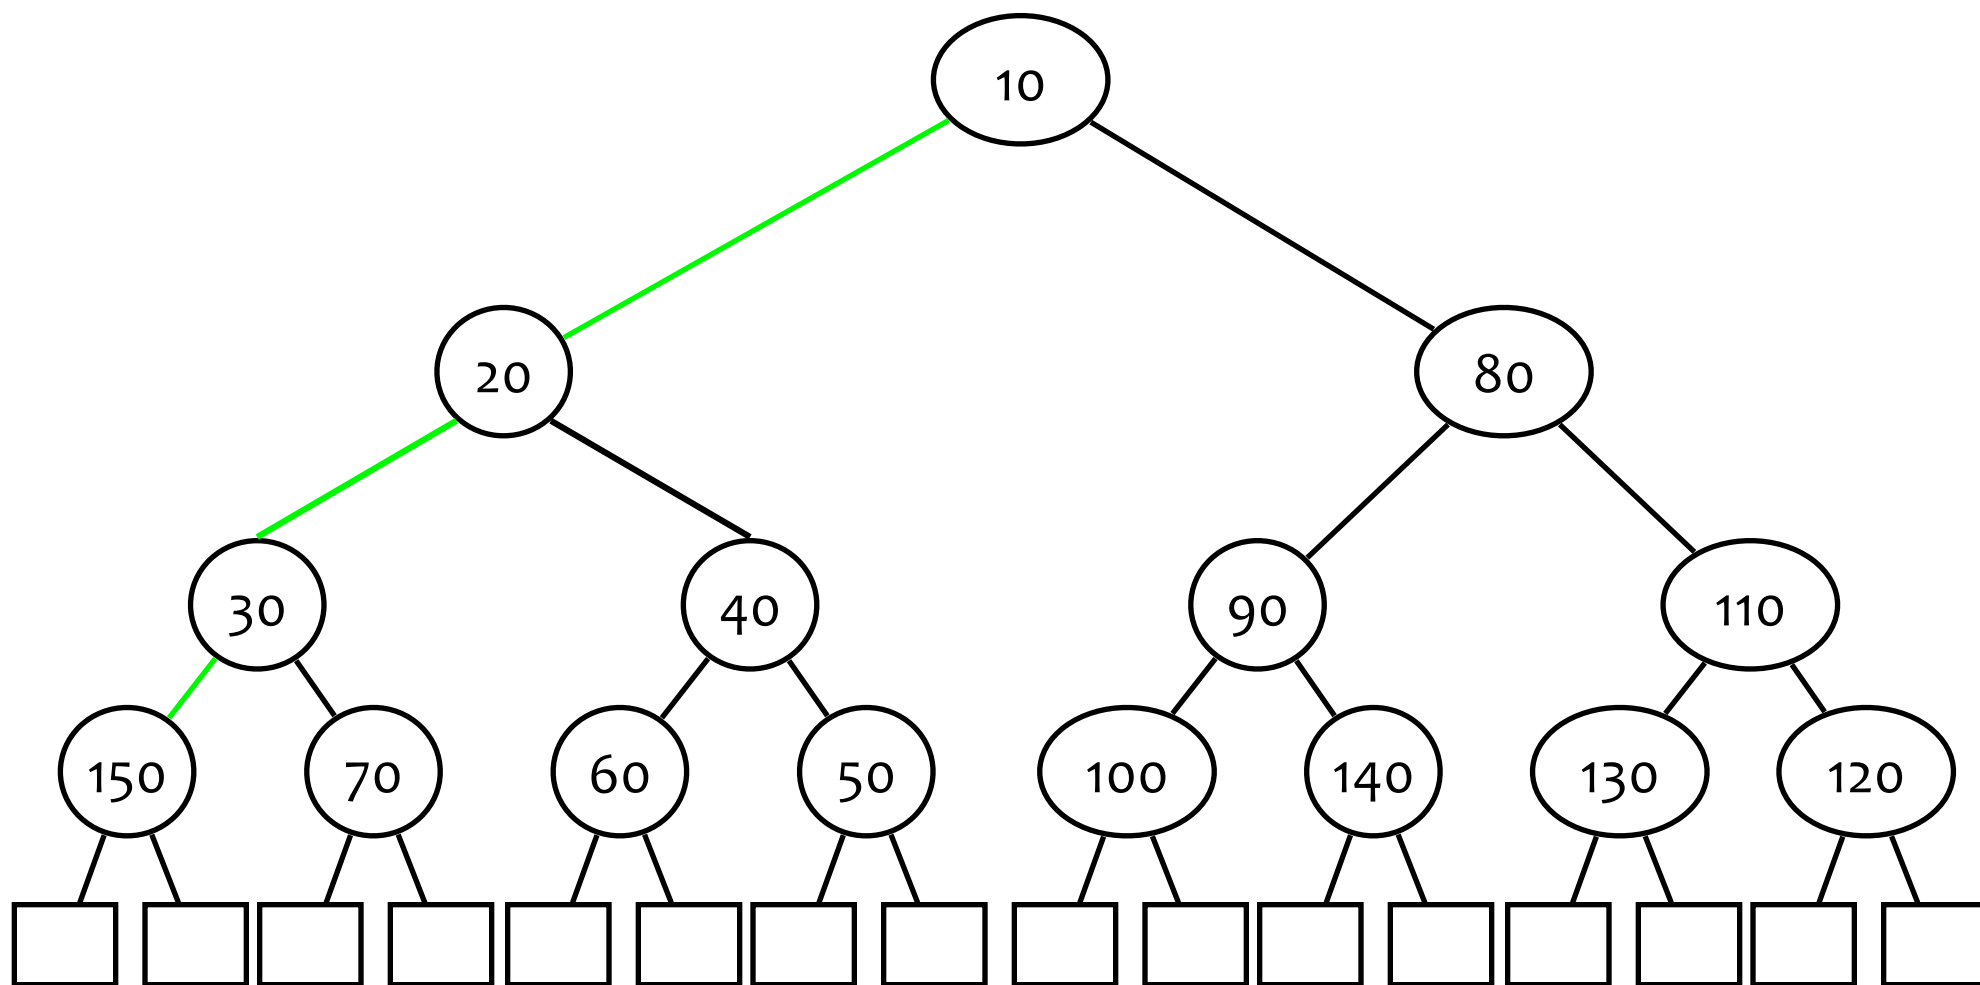
150	70	30	10	20	60	40	50	140	100	80	90	130	110	120
-----	----	----	----	----	----	----	----	-----	-----	----	----	-----	-----	-----



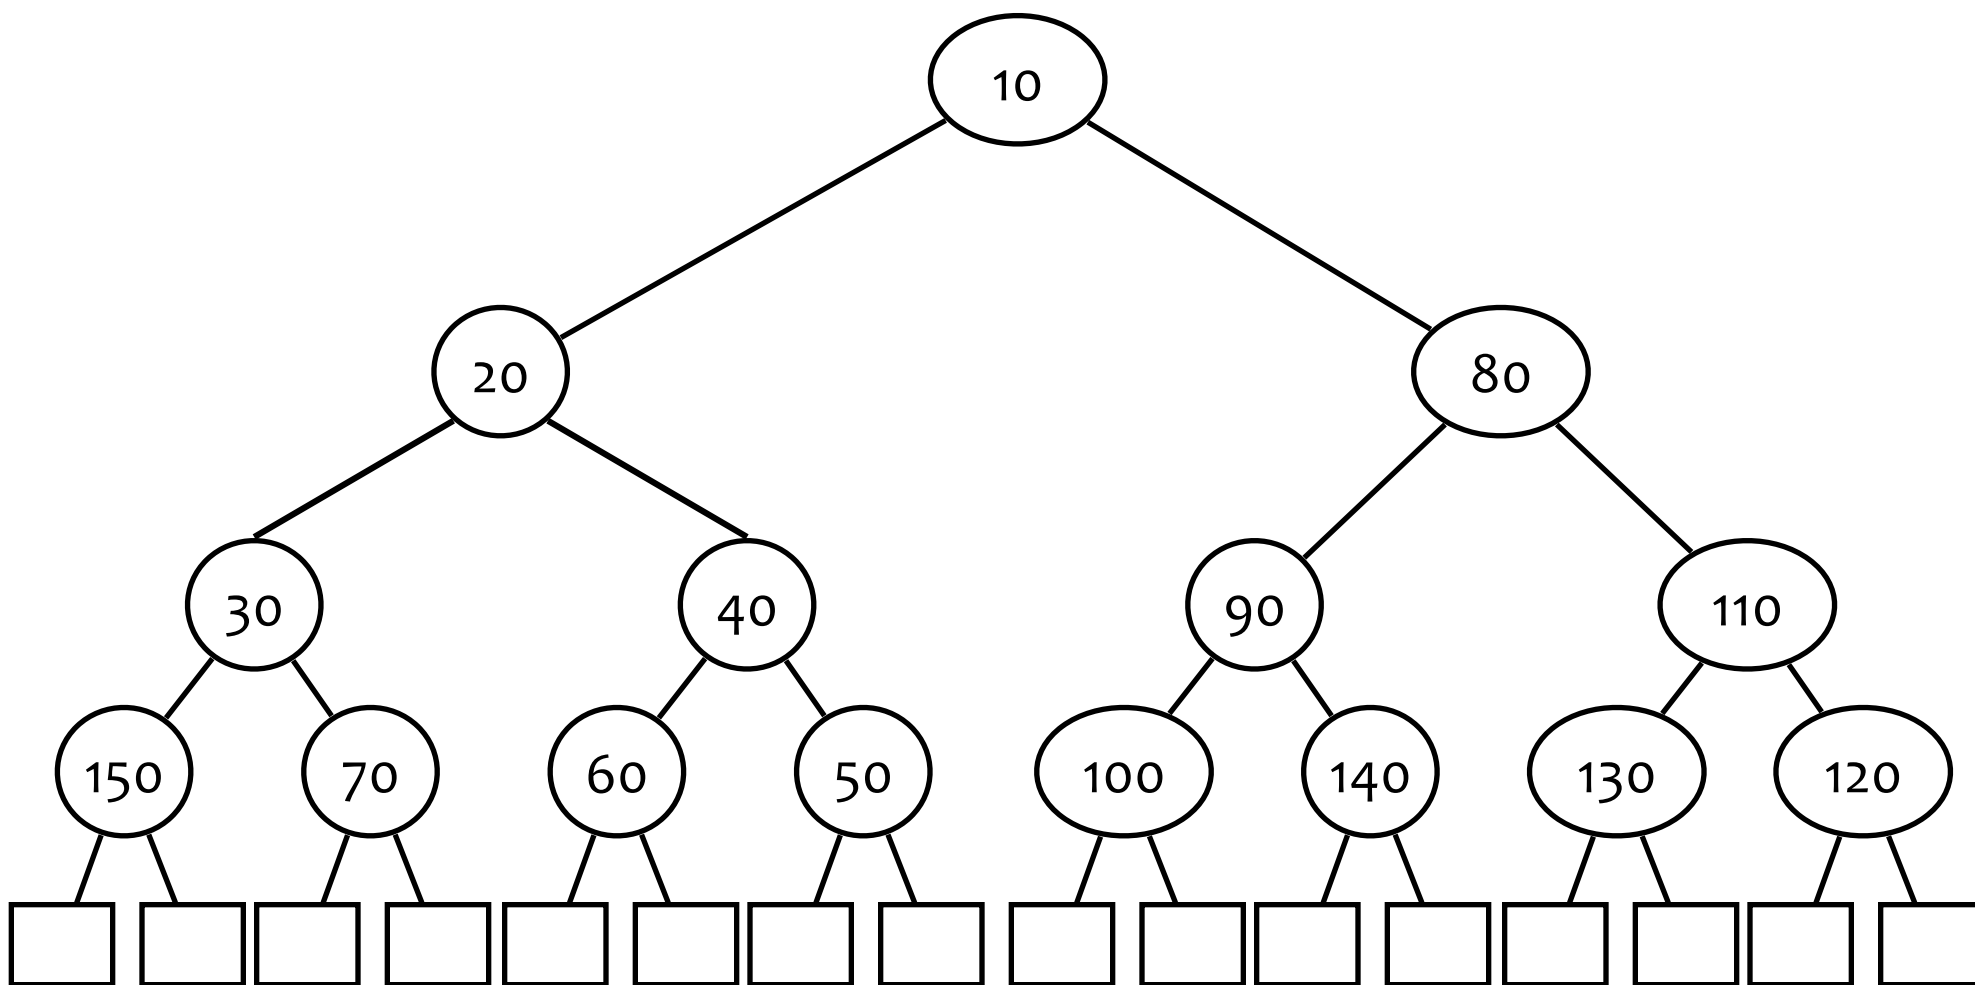
150	70	30	10	20	60	40	50	140	100	80	90	130	110	120
-----	----	----	----	----	----	----	----	-----	-----	----	----	-----	-----	-----




150	70	30	10	20	60	40	50	140	100	80	90	130	110	120
-----	----	----	----	----	----	----	----	-----	-----	----	----	-----	-----	-----




150	70	30	10	20	60	40	50	140	100	80	90	130	110	120
-----	----	----	----	----	----	----	----	-----	-----	----	----	-----	-----	-----



Did we really insert all n
elements in $O(n)$ time??




We show: $\max \# \text{ bubble down ops} \leq \# \text{ heap edges}$



Proof idea


For each new node joining two heaps: mark path of maximum number of bubble-down operations

For marking: first go right,
then left until bottom




For each new node joining two heaps:
mark path with maximum number of
bubble-down operations


Heaps of height 1:
leftmost branch is
not marked



For each new node joining two heaps:
mark path with maximum number of
bubble-down operations




For each new node joining two heaps:
mark path with maximum number of
bubble-down operations



For each height-2 heap, leftmost branch not marked


For each new node joining two heaps:
mark path with maximum number of
bubble-down operations

Heaps of
height 3



For each new node joining two heaps:
mark path with maximum number of
bubble-down operations


Heaps of
height 3



For each height-3 heap, leftmost branch not marked


For each new node joining two heaps:
mark path with maximum number of
bubble-down operations

Heaps of
height 4



For each new node joining two heaps:
mark path with maximum number of
bubble-down operations

Heaps of
height 4



For height-4 heap, leftmost branch not marked

Inductive argument: marking procedure will never mark all edges in heap, since the leftmost branch is never marked

- Note: leftmost branch in height- h heap: not marked
- When joining 2 heaps of height h to heap of height $h + 1$: new edges to be marked are
 - edge joining new node and right heap of height h , and
 - edges on left path in the right heap of height h
- We conclude: leftmost branch in height- $(h+1)$ heap is not marked

More efficient version (Kruskal's algorithm)

- union-find data structure

Kruskal's Algorithm

Algorithm Kruskal

Input: a weighted connected graph $G = (V, E)$

Output: an MST T for G

Data structure: Disjoint sets (lists or union-find) DS ;
sorted weights priority queue A ; and tree T

for each $v \in V$ **do** $C(v) \leftarrow DS.insert(v)$ **end** // one cluster per vertex

for each $(u, v) \in E$ **do** $A.insert((u, v))$ **end** // sort edges by weight

$T \leftarrow \emptyset$

while T has fewer than $n-1$ edges **do**

$(u, v) \leftarrow A.deleteMin()$ // edge with smallest weight

$C(u) \leftarrow DS.findCluster(u)$;

$C(v) \leftarrow DS.findCluster(v)$;

if $C(u) \neq C(v)$ **then**

 add edge (u, v) to T ;

$DS.insert(DS.union(C(v), C(u)))$; // merge two clusters

end

end

return T

Kruskal's Algorithm

Algorithm Kruskal

Time complexity analysis

Input: a weighted connected graph $G = (V, E)$

Output: an MST T for G

Data structure: Disjoint sets (lists or union-find) DS ;
sorted weights $A[]$; and tree T

for each vertex v in V of G **do** $C(v) \leftarrow DS.insert(v)$ **end**

$O(n)$

$A[] \leftarrow$ sort all edges by edge weight; // sorted array priority queue

$T \leftarrow \emptyset$; $k \leftarrow 0$;

$O(m \log m)$ or $O(m)$ with heap

while T has fewer than $n-1$ edges **do**

$(u, v) \leftarrow A[k]$; $k \leftarrow k + 1$; // next edge with smallest weight; deleteMin()

$C(v) \leftarrow DS.findCluster(v)$;

$C(u) \leftarrow DS.findCluster(u)$;

if $C(v) \neq C(u)$ **then**

add edge (v, u) to T ;

Total merging $O(m \log n)$

$DS.insert(DS.union(C(v), C(u)))$; // merge two clusters

end

end

return T

Total $O(m \log n)$

Union-find (or disjoint set) data structure

Disjoint Set Data Structure

- Given a set of elements, it is often useful to break them up or partition them into a number of separate, nonoverlapping sets
- A ***disjoint-set data structure*** is a data structure that keeps track of such a partitioning
- A ***union-find algorithm*** is an algorithm that performs two useful operations on such a data structure:
 - ***Find***: Determine which set a particular element is in. Also useful for determining if two elements are in the same set.
 - ***Union***: Combine or merge two sets into a single set.
 - ***MakeSet***: Creates a set containing only a given element

Disjoint Set Data Structure

- The universe consists of n elements, named $1, 2, \dots, n$
- The ADT is a collection of sets of elements
- Each element is in exactly one set
 - Sets are disjoint
 - To start, each set contains one element
- Each set has a name
 - Which is the name of one of its elements
 - Any name of one of its elements will do

Disjoint Set Operations

- **find(elementName)**
 - Returns the name of the unique set that contains the given element
- **union(setName1, setName2)**
 - Merges two sets and replaces them with one
- **Time complexity analysis**
 - Involves analyzing the *amortized* worst-case running time over a sequence of f find and u union operations

Disjoint Set Implementation I

- Create a linked list for each set and choose the element at the head of the list as the representative
- ***MakeSet*** creates a list of one element
- ***Union*** simply appends two lists, a constant-time operation
- ***Find*** requires linear time (i.e., may search entire list)
- A sequence of m union-find operations takes time $O(mn)$. The amortized time per operation is $O(n)$.

Disjoint Set Implementation II

- Two complementary techniques, *union by rank* and *path compression*, reduce the amortized time complexity per operation to $O(\alpha(n))$, where $\alpha(n)$ is the *inverse* of function $f(n) = A(n, n)$, and A is the extremely quickly-growing Ackermann function
- Since $\alpha(n)$ is its inverse it grows extremely slowly: it is less than 5 for all remotely practical values of n
- Thus, the amortized running time per operation is effectively a small constant
- Thus, n union and find operations require time $O(n\alpha(n))$

Union-find operations

- possible implementation
- Representing of sets in partition

Union-Find Operations:

$\text{union}(a, b)$

- Given a partition, let a and b be *canonical elements* of sets A and B , respectively
- If $a \neq b$:
 - Form a new set that is the union of the two sets
 - Destroy the two old sets
 - Select & return canonical element for new set

How should we represent the sets?

- Each set is rooted tree
- Each element of set corresponds to node in tree
- *Canonical element* (= name of set) is root of tree
- Each node e has reference $e.parent$ to its parent in tree, the root points to itself ($r.parent = r$)

Implementation I

Naive union-find algorithm

- $\text{makeset}(e)$: $e.\text{parent} \leftarrow e$
canonical element is root
- $\text{find}(e)$:
 - Follow parent pointers from e to root of tree that contains e
 - Return root
- $\text{union}(a,b)$
 - $a.\text{parent} \leftarrow b$
 - Return b as canonical element of new set

Union-Find Operations: $\text{makeSet}(e)$

Create singleton set containing single element e
[we do this if e is encountered but is previously in no set]


$e.\text{parent} \leftarrow e$

canonical element is root, i.e., e .

Union-Find Operations: $\text{makeSet}(e)$

Create singleton set containing single element e
[we do this if e is encountered but is previously in no set]

$e.\text{parent} \leftarrow e$




canonical element is root, i.e., e .

Union-Find Operations: $\text{makeSet}(e)$

Create singleton set containing single element e
[we do this if e is encountered but is previously in no set]

$e.\text{parent} \leftarrow e$



canonical element is root, i.e., e .

Union-Find Operations: $\text{find}(e)$


Return canonical element of the set containing e

- Follow parent pointers from e to root r of tree that contains e
- Return root r

Union-Find Operations: $\text{find}(e)$

Return canonical element of the set containing e


- Follow parent pointers from e to root r of tree that contains e
- Return root r



Union-Find Operations: $\text{find}(e)$

Return canonical element of the set containing e


- Follow parent pointers from e to root r of tree that contains e
- Return root r



Union-Find Operations: $\text{find}(e)$


Return canonical element of the set containing e

- Follow parent pointers from e to root r of tree that contains e
- Return root r




Union-Find Operations: $\text{union}(a, b)$

- Let a and b be canonical elements of sets A and B , respectively
- If $a \neq b$ then form a new set: union of two sets whose canonical elements are a and b
- Destroy two old sets
- Select and return canonical element for new set: $a.\text{parent} \leftarrow b$
- Return b as canonical element of new set




Union-Find Operations: $\text{union}(a, b)$

- Let a and b be canonical elements of sets A and B , respectively
- If $a \neq b$ then form a new set: union of two sets whose canonical elements are a and b
- Destroy two old sets
- Select and return canonical element for new set: $a.\text{parent} \leftarrow b$
- Return b as canonical element of new set




Union-Find Operations: $\text{union}(a, b)$

- Let a and b be canonical elements of sets A and B , respectively
- If $a \neq b$ then form a new set: union of two sets whose canonical elements are a and b
- Destroy two old sets
- Select and return canonical element for new set: $a.\text{parent} \leftarrow b$
- Return b as canonical element of new set



Union-Find Operations: $\text{union}(a, b)$

- Let a and b be canonical elements of sets A and B , respectively
- If $a \neq b$ then form a new set: union of two sets whose canonical elements are a and b
- Destroy two old sets
- Select and return canonical element for new set: $a.\text{parent} \leftarrow b$
- Return b as canonical element of new set



The naive union-find algorithm

- makeset(e): $e.\text{parent} \leftarrow e$ (canonical element is root)
- find(e):
 - Follow parent pointers from e to root of tree containing e
 - Return root
- union(a, b): $a.\text{parent} \leftarrow b$ (Return b as canonical element of new set)

The naive union-find algorithm

- makeset(e): $e.\text{parent} \leftarrow e$ (canonical element is root)
 $O(1)$
- find(e):
 - Follow parent pointers from e to root of tree containing e
 - Return root
- union(a, b): $a.\text{parent} \leftarrow b$ (Return b as canonical element of new set)

The naive union-find algorithm

- makeset(e): $e.\text{parent} \leftarrow e$ (canonical element is root)
 $O(1)$
- find(e):
 - Follow parent pointers from e to root of tree containing e
 - Return root
- union(a, b): $a.\text{parent} \leftarrow b$ (Return b as canonical element of new set)
 $O(1)$

The naive union-find algorithm

- makeset(e): $e.\text{parent} \leftarrow e$ (canonical element is root)
 $O(1)$
- find(e):
 - Follow parent pointers from e to root of tree containing e
 - Return root $O(n)$
- union(a, b): $a.\text{parent} \leftarrow b$ (Return b as canonical element of new set)
 $O(1)$

An *improved* union-find algorithm

- makeset(e): $e.\text{parent} \leftarrow e$, the canonical element is the root.
- find(e): Follow the parent pointers from e to the root of the tree containing e . Return the root.
- **union(a, b)**: $a.\text{parent} \leftarrow$ **canonical element of *larger* of two sets**

Why is this better?

Improved Union-Find: Union by Rank

Idea: Store with each node v the size of the subtree (*rank*) rooted at y . In a union operation, make the tree of the smaller set/rank a subtree of the other tree, and update the size field of the root of the resulting tree.

Union by Rank

- With each node x we store a nonnegative integer $x.\text{rank}$ that is an upper bound on the height of x .
- When carrying out $\text{makeSet}(x)$, we set $x.\text{rank} \leftarrow 0$.
- To carry out $\text{union}(x,y)$, we compare $x.\text{rank}$ and $y.\text{rank}$.
 - If $x.\text{rank} < y.\text{rank}$ then $x.\text{parent} \leftarrow y$.
 - If $y.\text{rank} < x.\text{rank}$ then $y.\text{parent} \leftarrow x$.
 - If $x.\text{rank} = y.\text{rank}$ then $x.\text{parent} \leftarrow y$ and $y.\text{rank} \leftarrow y.\text{rank} + 1$.

Amortized Time Complexity

- A series of n makeSet, (modified) union, and find operations starting from an initially empty partition take $O(n \log n)$ time.
- The *amortized* running time per operation is $O(\log n)$ time.

A series of n makeSet, (modified) union, and find operations starting from an initially empty partition take $O(n \log n)$ time

Lemma: Starting with sets of size 1, using the modified union algorithm, any tree of m nodes has height of at most $\log m$.