

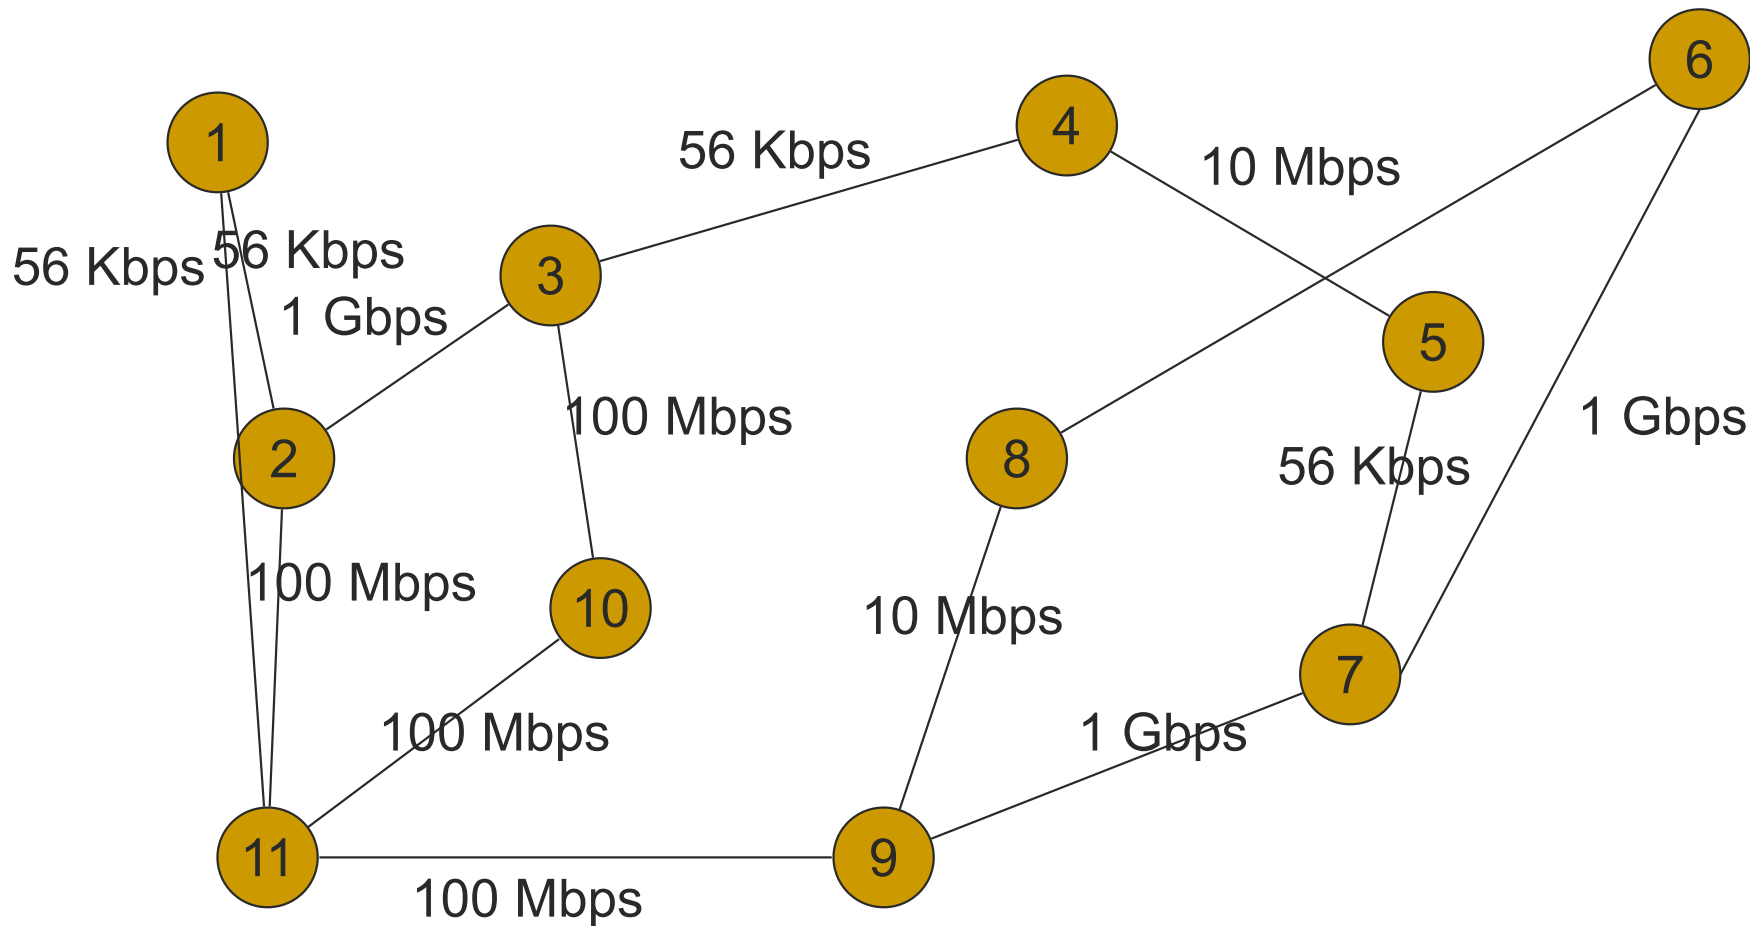


Shortest Paths

in edge-weighted graphs

Communication Speeds in a Computer Network

Find fastest way to route a data packet between two computers



[shortest paths in edge weighted graphs]

- Find the fastest way to travel across the country using a graph representing roads, with edge weights represented by
 - distances
 - travel times (to accommodate speed limits) between cities
 - Flight distances between airports

[Shortest Path problems]

- Single source shortest paths
- Find a shortest path between two given vertices
- All pair shortest paths

[Single Source Shortest Path problems]

- Undirected graphs with non-negative edge weights
- Directed graphs with non-negative edge weights
- Directed graphs with arbitrary weights


[Single Source Shortest Paths]

- If graph is not weighted (all edge-weights are unit-weight): BFS works
- Now assume: graph is edge-weighted
 - Every edge is associated a positive number
 - Possible weights: integers, real numbers, rational numbers
 - Edge-weights can represent: distance, connection cost, affinity


[Single Source Shortest Paths]

- **Input:** An edge-weighted undirected graph and a source node v with: for every edge e edge-weight $w(e) > 0$
- **Output:** All single-source shortest paths (and their weight) for v in G : for every node $w \neq v$ in G a shortest path from v to w .
 - Here, a *path* p from v to w consisting of edges e_1, e_2, \dots, e_n is shortest in G , if its length
$$w(p) = \sum_{i=1}^n w(e_i)$$
 is minimum (i.e, there is no path from v to w in G that is shorter).


Dijkstra's algorithm: a greedy algorithm




[Dijkstra's algorithm: Initializing]




[Dijkstra's algorithm: Initializing Cloud C (consisting of "solved" subgraph)]




[Dijkstra's algorithm: pull v into C]




Dijkstra's algorithm: update C 's neighborhood




[Dijkstra's algorithm: pick closest vertex u outside C]




]




]




[Dijkstra's algorithm: pick closest vertex u outside C]




[Dijkstra's algorithm: pull u into C]




Dijkstra's algorithm: update C 's neighborhood




[Dijkstra's algorithm: pick closest vertex u outside C]




[Dijkstra's algorithm: pull u into C]




Dijkstra's algorithm: update C 's neighborhood




[Dijkstra's algorithm: pick closest vertex u outside C]




]




Dijkstra's algorithm: update C 's neighborhood




[Dijkstra's algorithm: pick closest vertex u outside C]




[Dijkstra's algorithm: pull u into C]




Dijkstra's algorithm: update C 's neighborhood




[Dijkstra's algorithm: pick closest vertex u outside C]




[Dijkstra's algorithm: pull u into C]




Dijkstra's algorithm: update C 's neighborhood




[Dijkstra's algorithm: pick closest vertex u outside C]




[Dijkstra's algorithm: pull u into C]




Dijkstra's algorithm: update C 's neighborhood




[Dijkstra's algorithm: pick closest vertex u outside C]




[Dijkstra's algorithm: pull u into C]




Dijkstra's algorithm: update C 's neighborhood



[Dijkstra's algorithm: pick closest vertex u outside C]



[Dijkstra's algorithm: pull u into C]



[When pulling a neighbour u of C into C]

- The value associated with u denotes the length of a shortest path from v to u
- For any vertex x not in the cloud
 - the value associated with x denotes a shortest way from v to x without the use of other vertices outside of the cloud
 - $+\infty$ denotes that the vertex cannot be reached yet from x via cloud vertices only

[Algorithm

DijkstraShortestPaths(G, v)]

Input: A simple undirected graph G with nonnegative edge-weights, a distinguished vertex v in G

Output: A label $D[u]$ for each vertex u in G such that $D[u]$ is the distance from v to u in G .

[Algorithm DijkstraShortestPaths(G, v)]

$D[v] \leftarrow 0$

for each vertex $u \neq v$ of G **do**

$D[u] \leftarrow +\infty$

Let Q be a priority queue containing all vertices of G using $D[.]$ as keys

while Q is not empty **do**

$u \leftarrow Q.\text{removeMin}()$ // u is added to cloud

for each vertex $z \in N(u)$ with $z \in Q$ **do**

if $D[u] + w((u, z)) < D[z]$ **then**


Relaxation

$D[z] \leftarrow D[u] + w((u, z))$


update z 's key in Q to $D[z]$

return D


Dijkstra's algorithm using heaps




Dijkstra's algorithm using heaps




Dijkstra's algorithm using heaps



Dijkstra's algorithm using heaps



Dijkstra's algorithm using heaps



[Running time]

$D[v] \leftarrow 0$

for each vertex $u \neq v$ of G **do**

$D[u] \leftarrow +\infty$

Let Q be a priority queue containing all vertices of G using $D[.]$ as keys

while Q is not empty **do**

$u \leftarrow Q.\text{removeMin}()$ // u is added to cloud

for each vertex $z \in N(u)$ with $z \in Q$ **do**

if $D[u] + w((u, z)) < D[z]$ **then**

Relaxation

$D[z] \leftarrow D[u] + w((u, z))$

update z 's key in Q to $D[z]$

return D

Running time for $G=(V,E)$ with $|V|=n$ and $|E|=m$

- Insertion of vertices in priority queue Q
 - $O(n)$ when using heap
- While loop:
 - Per iteration:
 - Remove vertex from Q $O(\log n)$
 - Relaxation $O(\deg(u) \log(n))$
 - $\sum_{u \in G} (1 + \deg(u)) \log n$ is $O((n+m) \log n)$
- Overall running time: $O(m \log n)$

[In real life applications]

- Often the graphs are sparse
- Then $O(m \log n)$ may be $O(n \log n)$

Correctness

[Algorithm

DijkstraShortestPaths(G, v)

$D[v] \leftarrow 0$

for each vertex $u \neq v$ of G **do**

$D[u] \leftarrow +\infty$

Let Q be a priority queue containing all vertices of G using $D[.]$ as keys

while Q is not empty **do**

$u \leftarrow Q.\text{removeMin}()$ // u is added to cloud

for each vertex $z \in N(u)$ with $z \in Q$ **do**

if $D[u] + w((u, z)) < D[z]$ **then**

Relaxation

$D[z] \leftarrow D[u] + w((u, z))$

update z 's key in Q to $D[z]$

return D

[Correctness of Dijkstra's algorithm]

- **To show:** whenever u is pulled into cloud C , $D[u]$ stores the length from a shortest path from u to v
- **Definition:** For vertices u and v in G , we denote with $d(u, v)$ the length of a shortest path from u to v .

Whenever u is pulled into cloud C , $D[u]$ stores the length from a shortest path from u to v

Proof. Assume: claim is wrong. Then:


there exists a vertex t that is pulled into cloud C and $D[t] > d(v, t)$

We define:

- u the first such vertex (currently) pulled into C
- P a shortest path in G from source v to vertex u
- y the last vertex that lies on P and is pulled “correctly” into C
- z the vertex closest to y that lies on P and is not in C

[

]




[

]




- $D(u) > d(v, u)$
- $y \in C, D[y] = d(v, y)$
- $D(u) \leq D(z)$
- $D(z) = d(v, z)$




- $\underline{D(u)} \leq D(z) = d(v, z) \leq d(v, z) + d(z, u)$
 $\quad \quad \quad = \underline{d(v, u)}$

[Example (Input contains negative edges)]




[How far away is u from v ?



Does Dijkstra's algorithm work?

Directed graphs with positive edge weights




Does Dijkstra's algorithm work?

[Single source shortest paths for directed graphs with positive edge]

- Dijkstra's algorithm works without changes except here edges are directed, that is $(a,b) \neq (b,a)$
- The big-oh worst case running time remains the same


Shortest paths in graphs containing *negative* edges

- Not possible for undirected graphs
- What about directed graphs?



Shortest paths in directed graphs with negative edge weights


- Another example



Does Dijkstra's algorithm work?

Single Source shortest Paths in directed Graphs


- If G does not contain any negative-weight cycles: does Dijkstra's algorithm work?



Negative edges and negative-weight cycles

- If G is directed, compute single-source shortest path problem using **Bellman-Ford** shortest path algorithm
- Negative-weight cycles are discovered

]



[Algorithm Bellman-Ford(G, v)]

Input: A simple undirected graph G with nonnegative edge-weights, a distinguished vertex v in G

Output: A label $D[u]$ for each vertex u in G such that $D[u]$ is the distance from v to u in G .


[Algorithm Bellman-Ford(G, v)]




```
D[v] ← 0
for each vertex u ≠ v of G do
    D[u] ← +∞
for i ← 1 to n-1 do
    for each edge (u,z) in G do
        if D[u] + w((u,z)) < D[z] then
            D[z] ← D[u] + w((u,z))
if there are no edges left with potential
relaxation operations then
    return D
else
    return "G contains a negative cycle"
```

performs $n-1$
times a
relaxation of
every edge
in the graph


[Initialize ...]




[i=1]




[i=1]




[i=1]




[i=1]




[i=1]




[i=1]




[i=1]




[i=1]




[i=1]




[i=1]




[i=1]




[i=1]




[i=1]




[i=1]




[i=1]



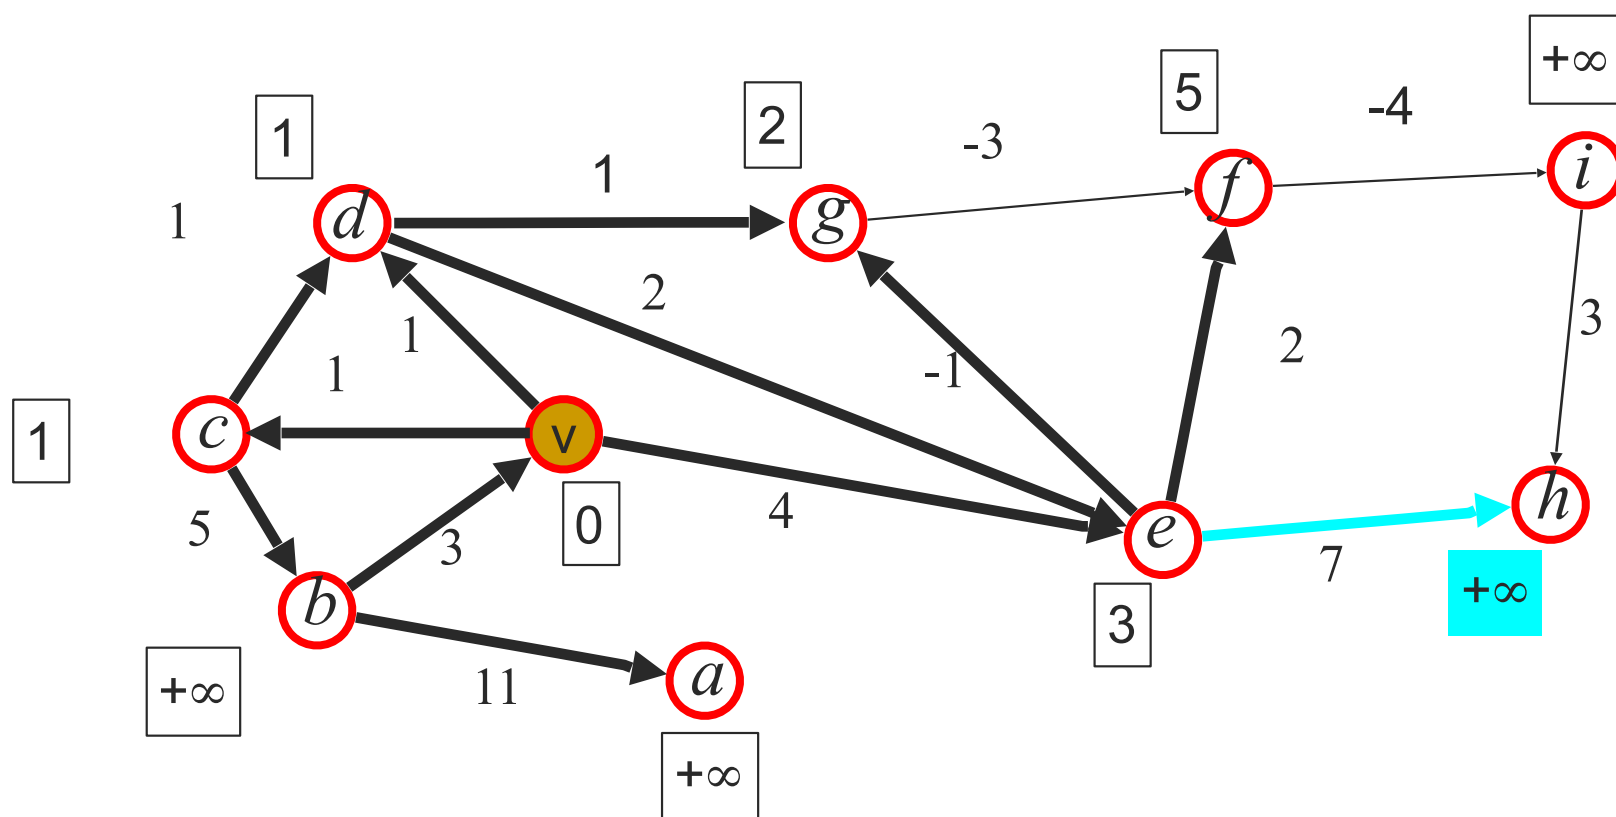
[i=1]



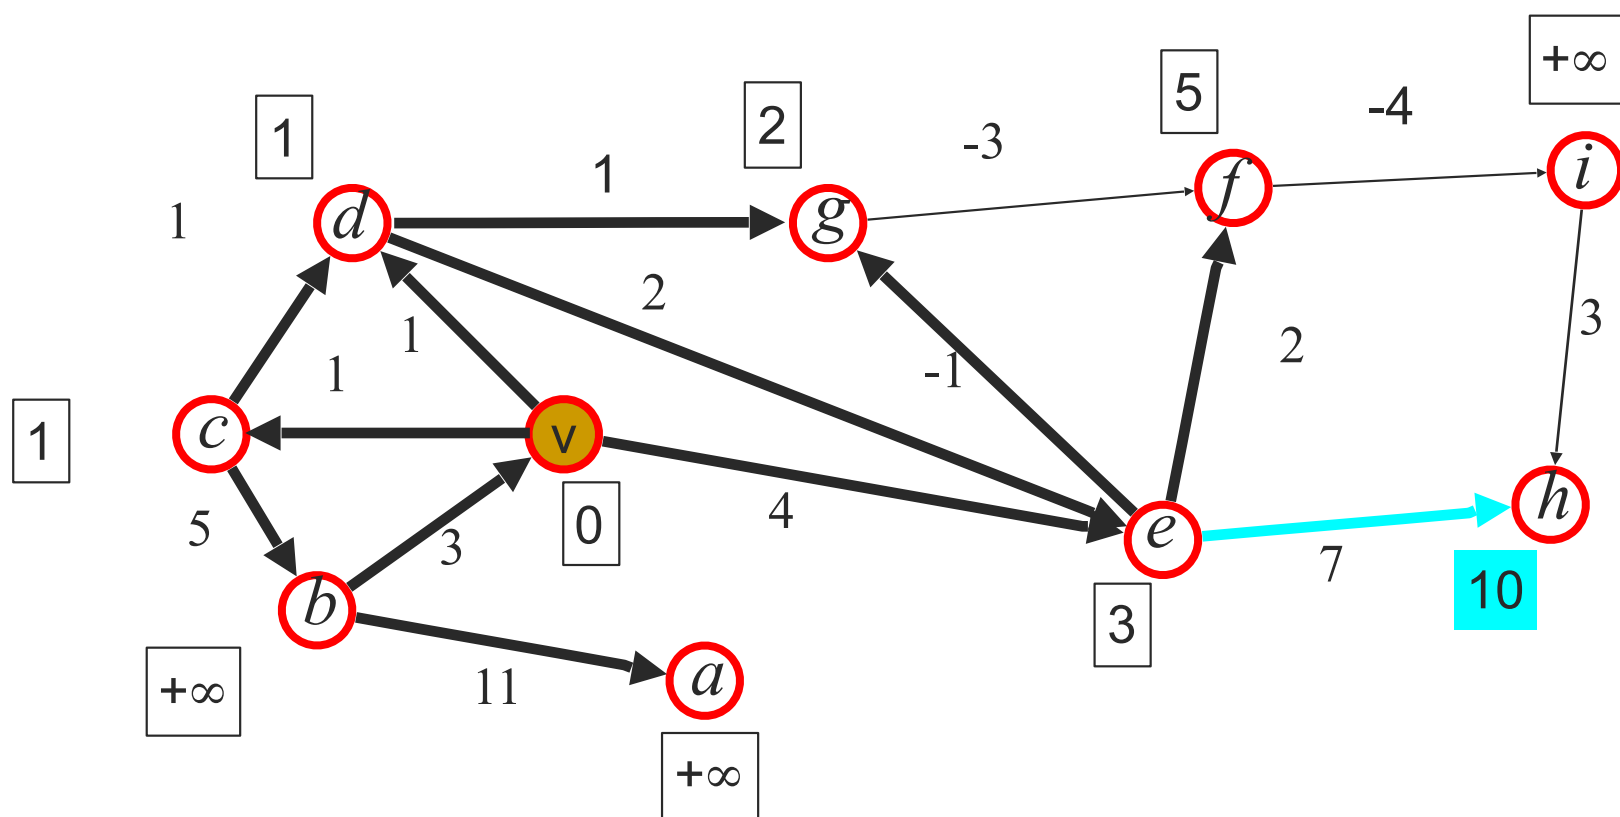
[i=1]




[i=1]




[i=1]



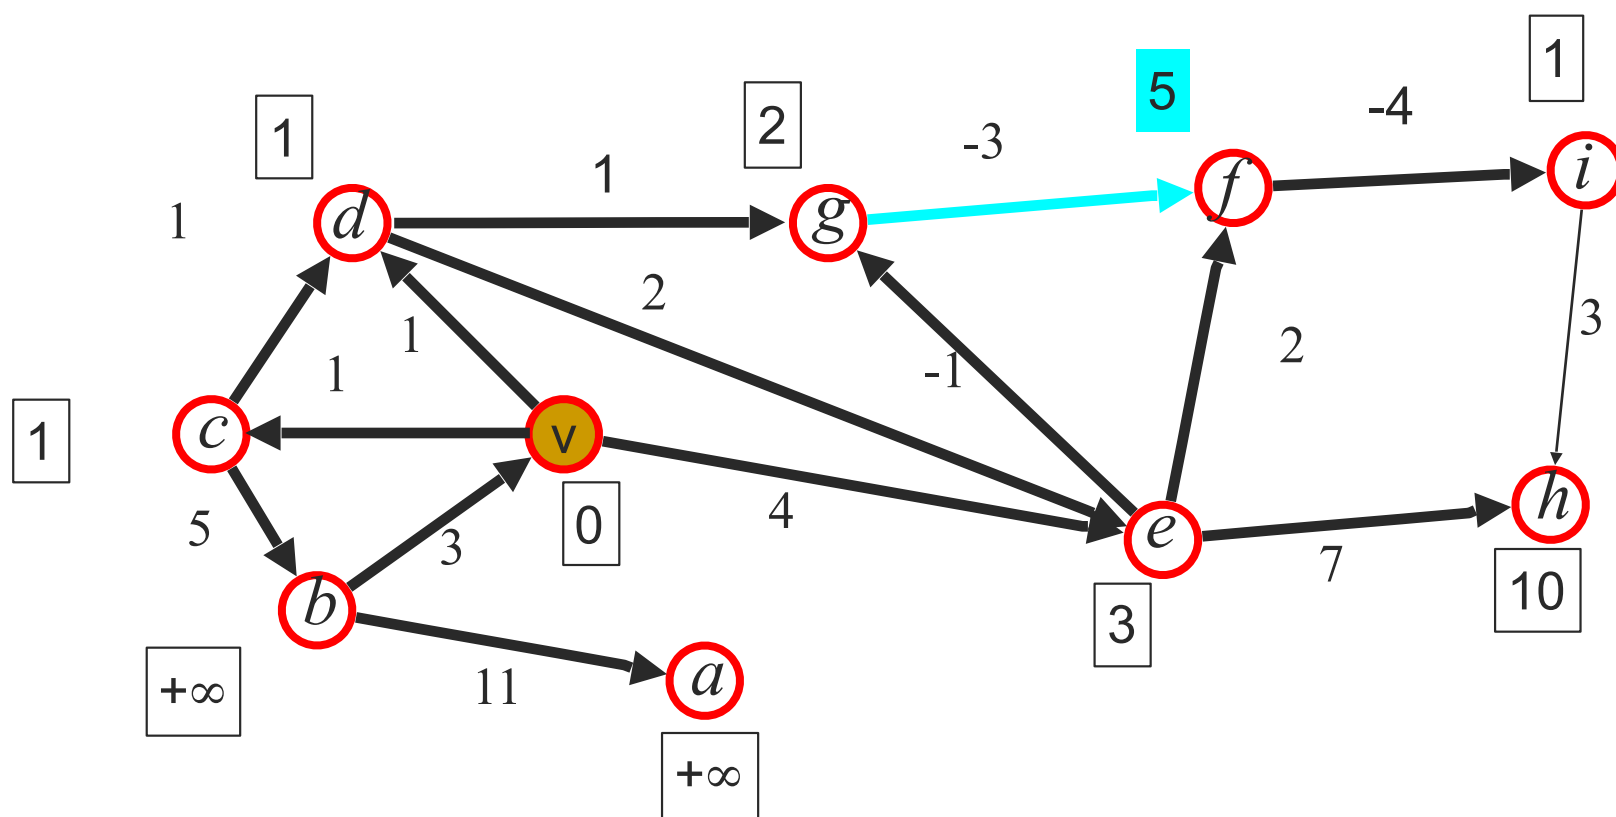
[i=1]




[i=1]



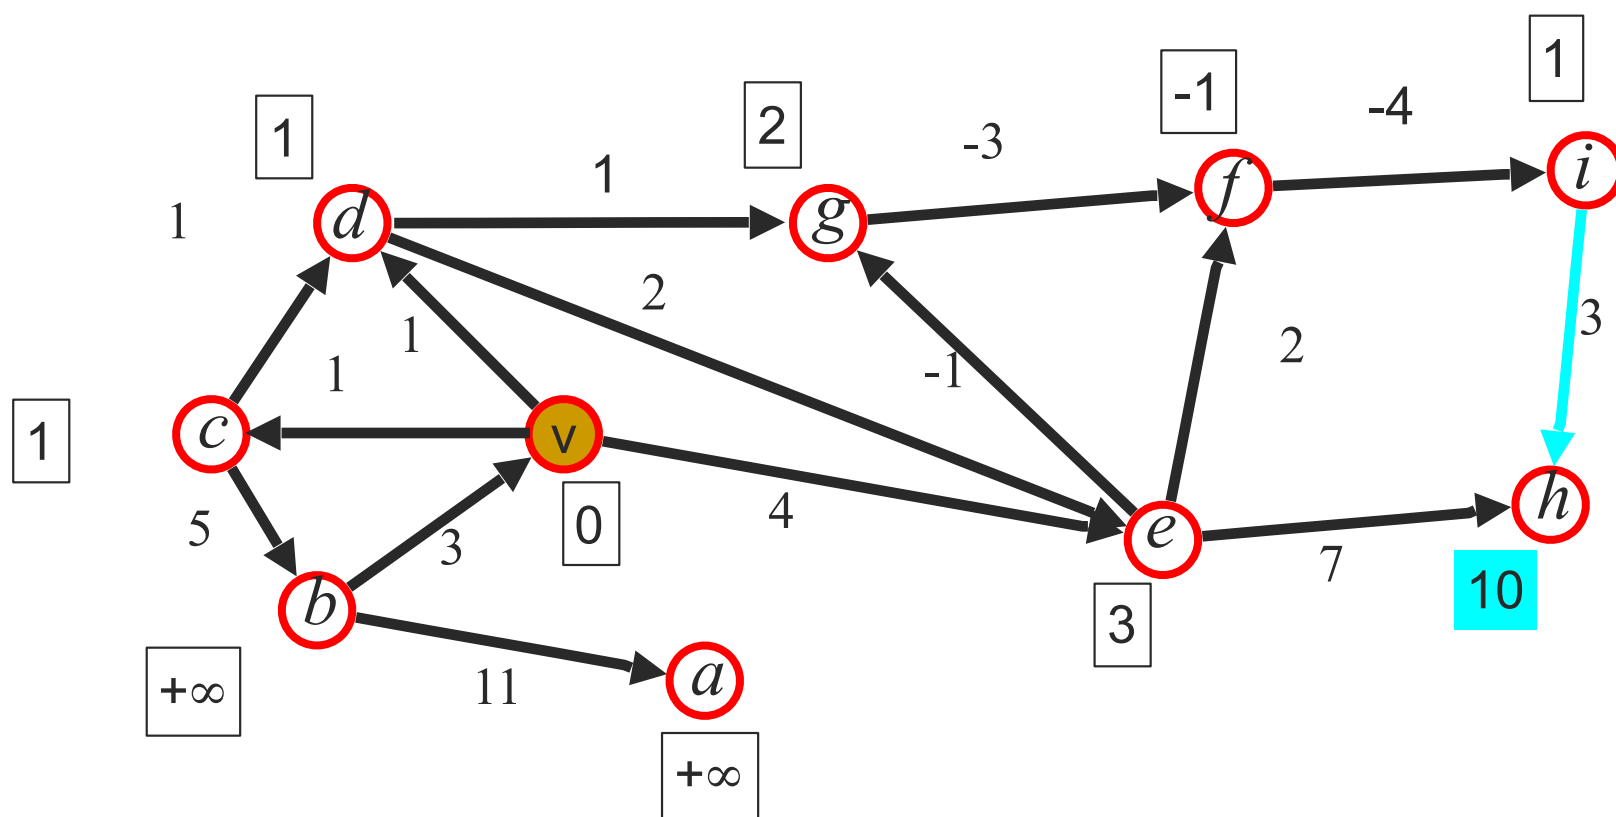
[i=1]



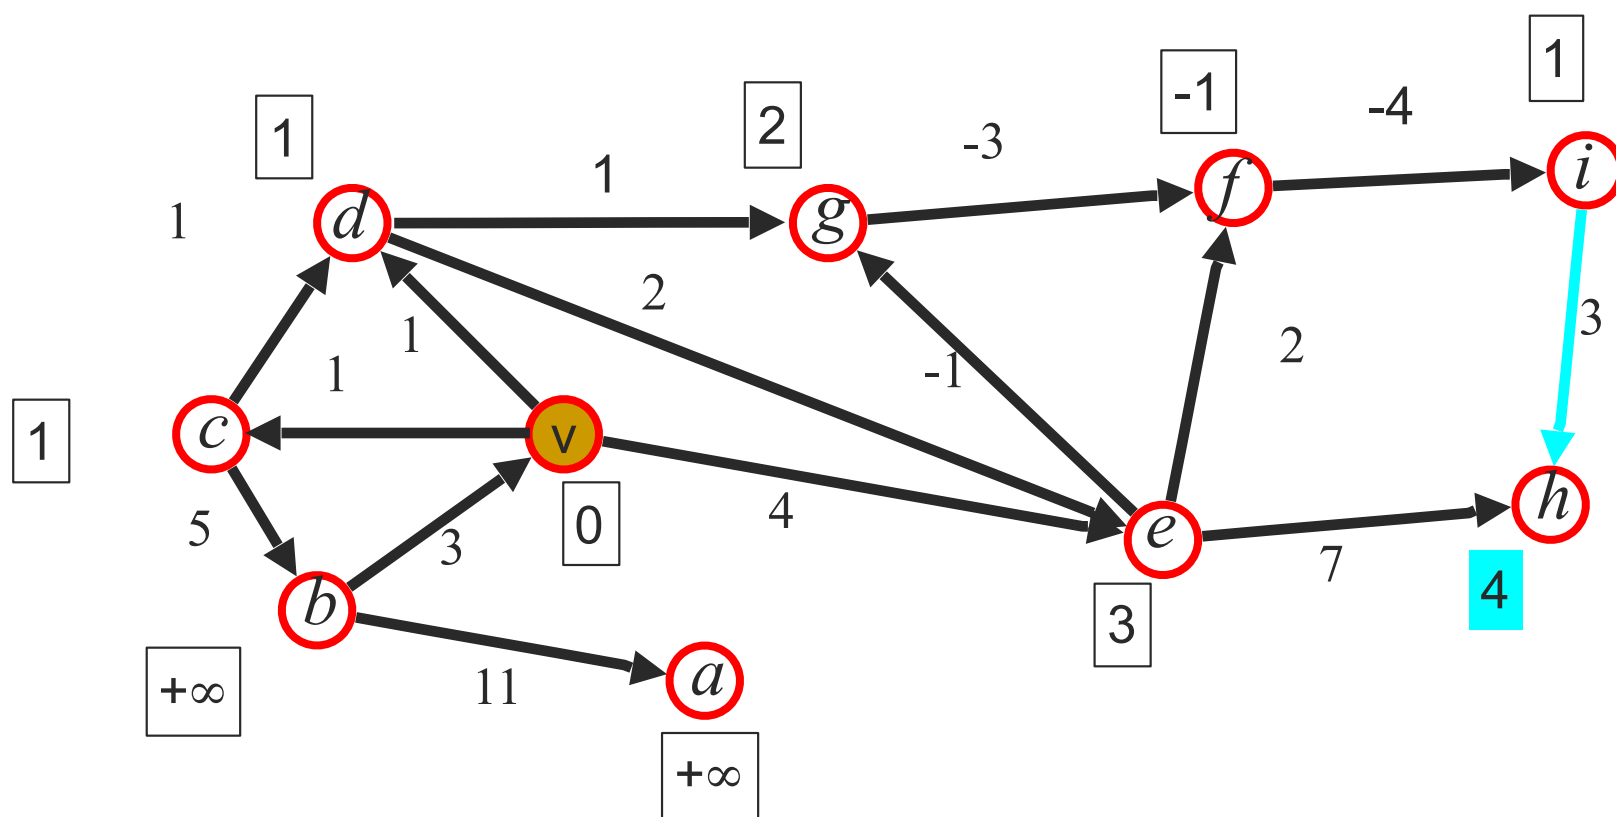
[i=1]



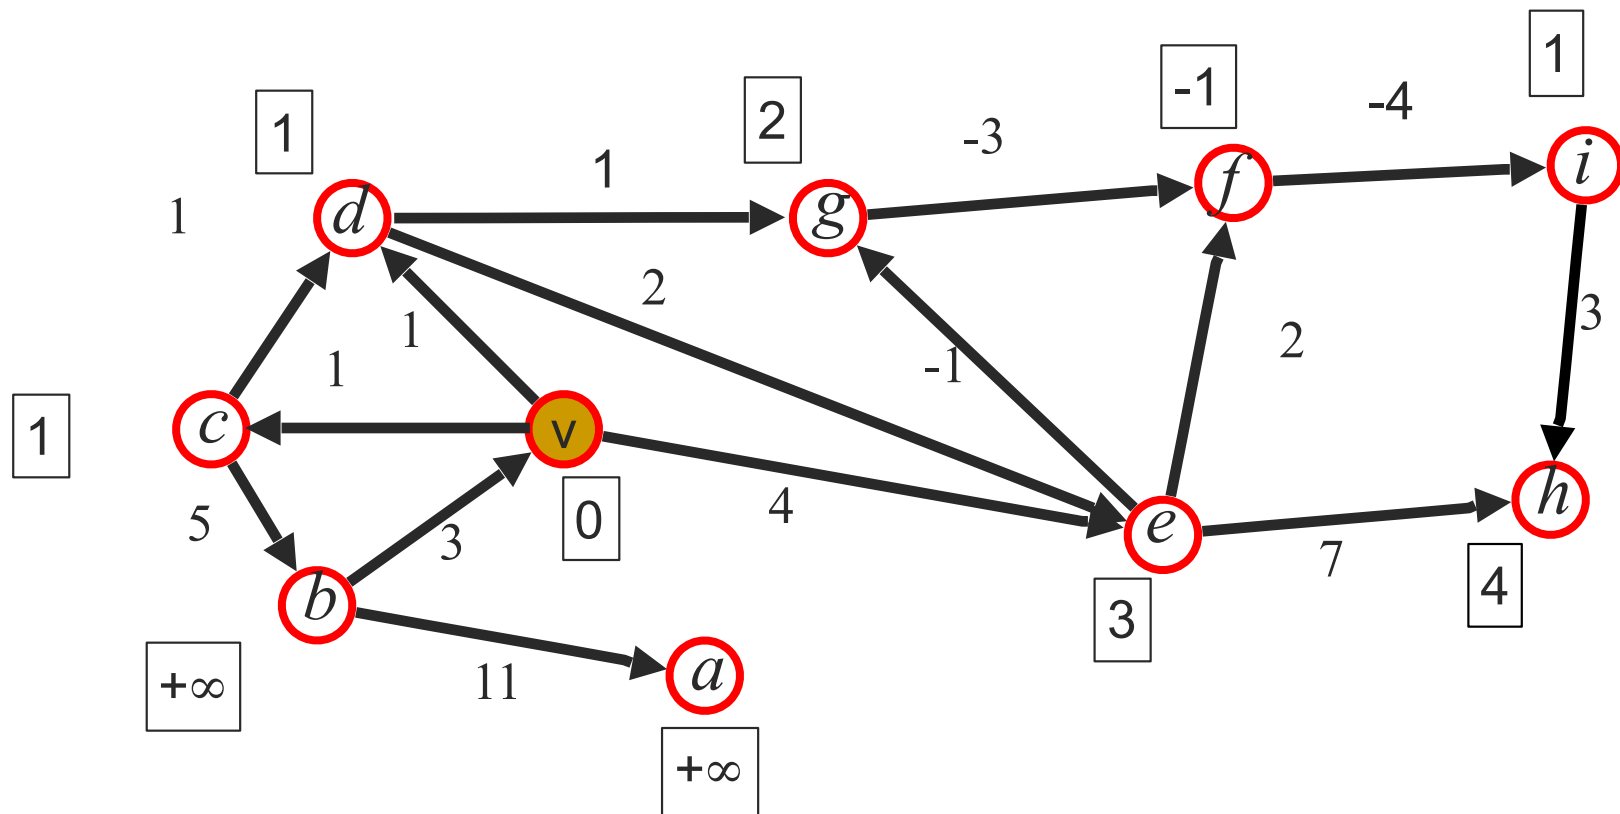
[i=1]




[i=1]




[i=1]




[i=1]




[i=2]




[i=2]




[i=2]




[i=2]




[i=2]




[i=2]




[i=2]




[i=2]




[i=2]




[i=2]




[i=2]



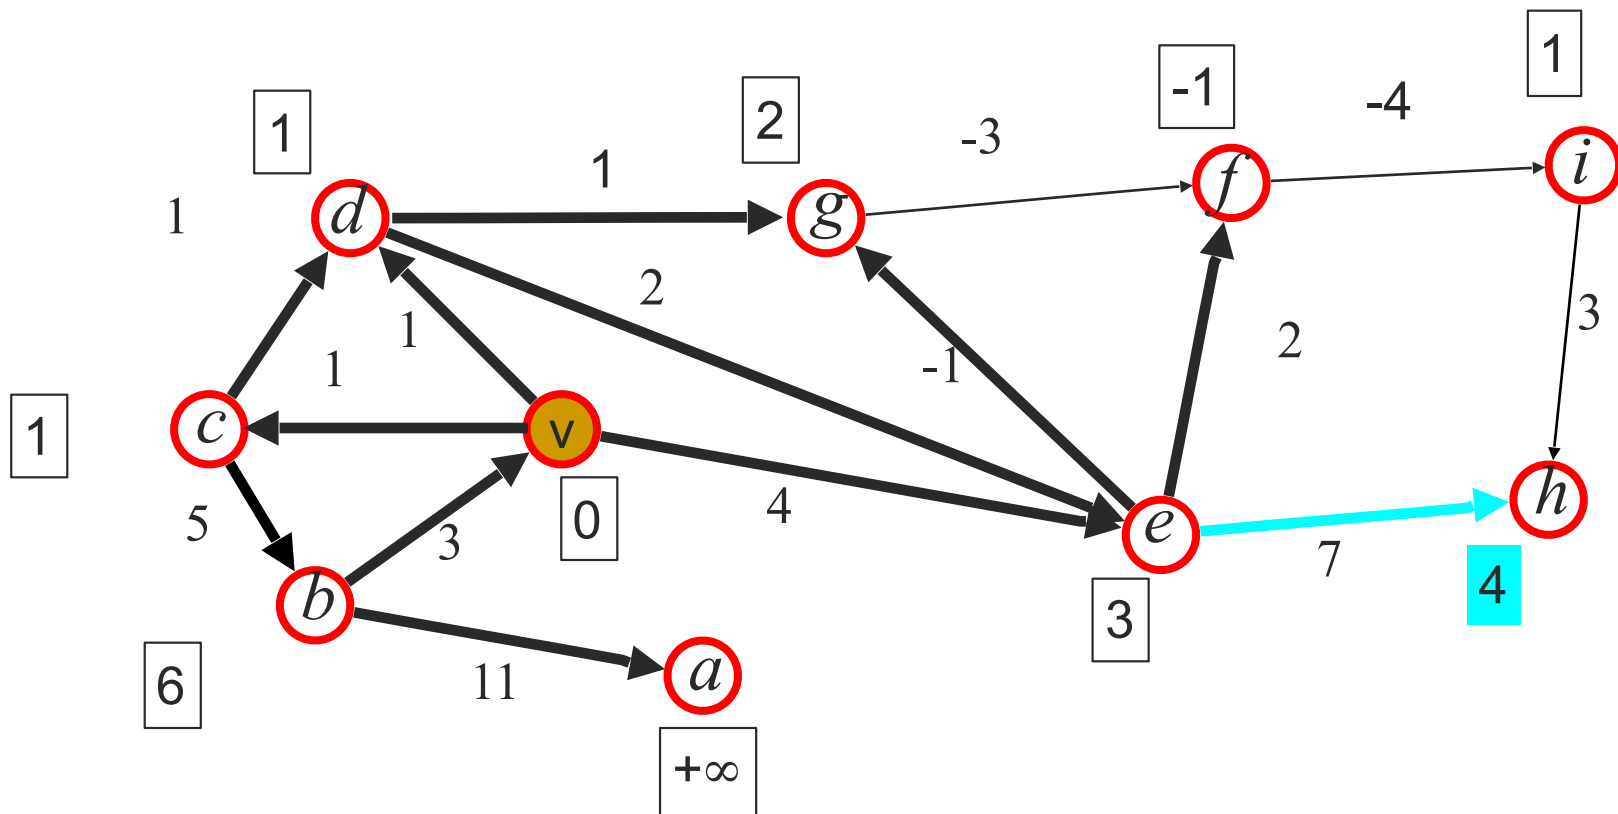
[i=2]




[i=2]



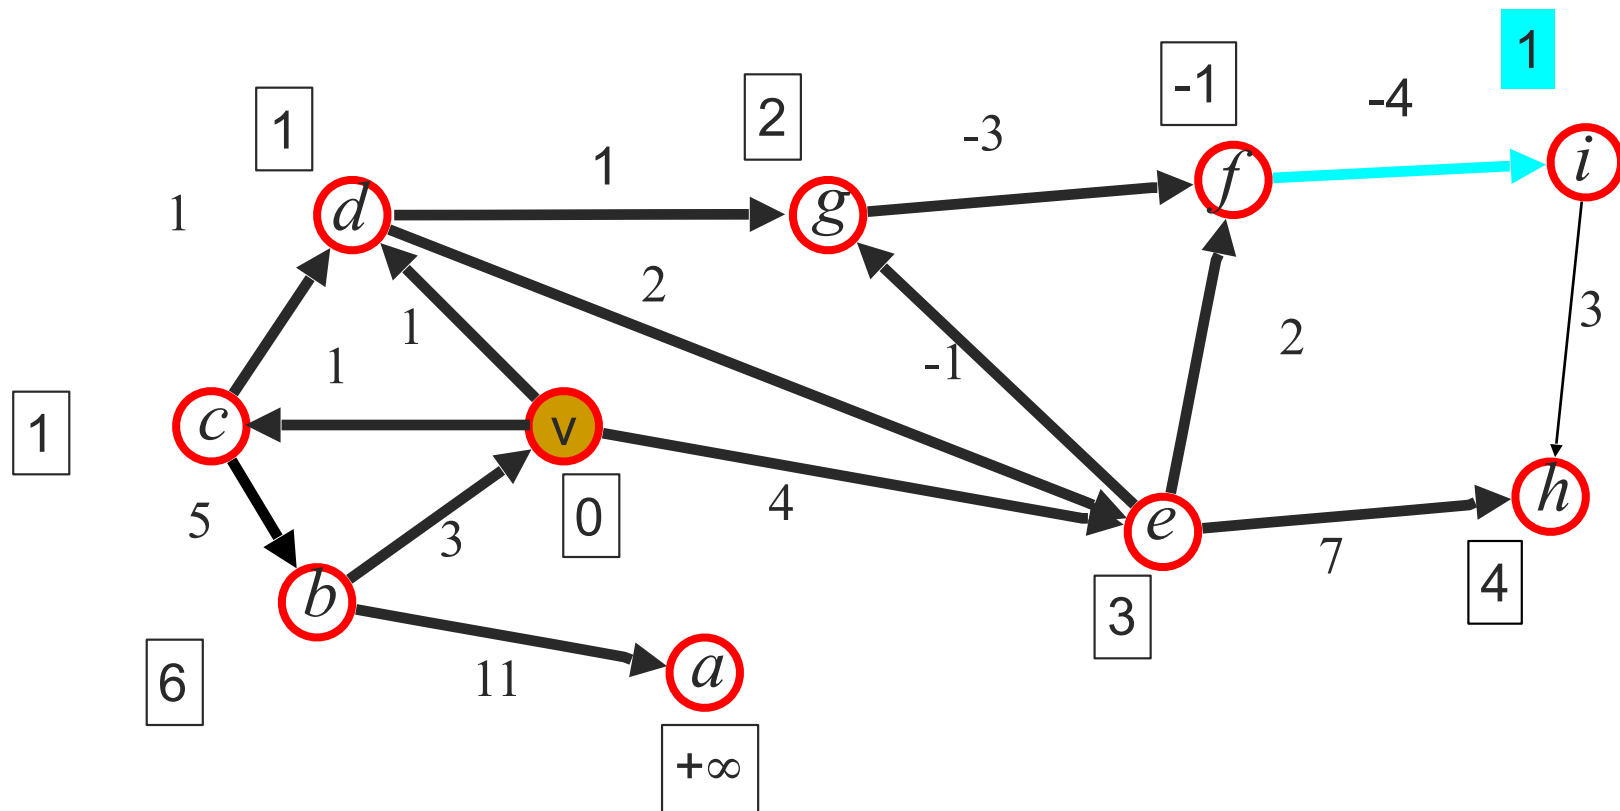
[i=2]



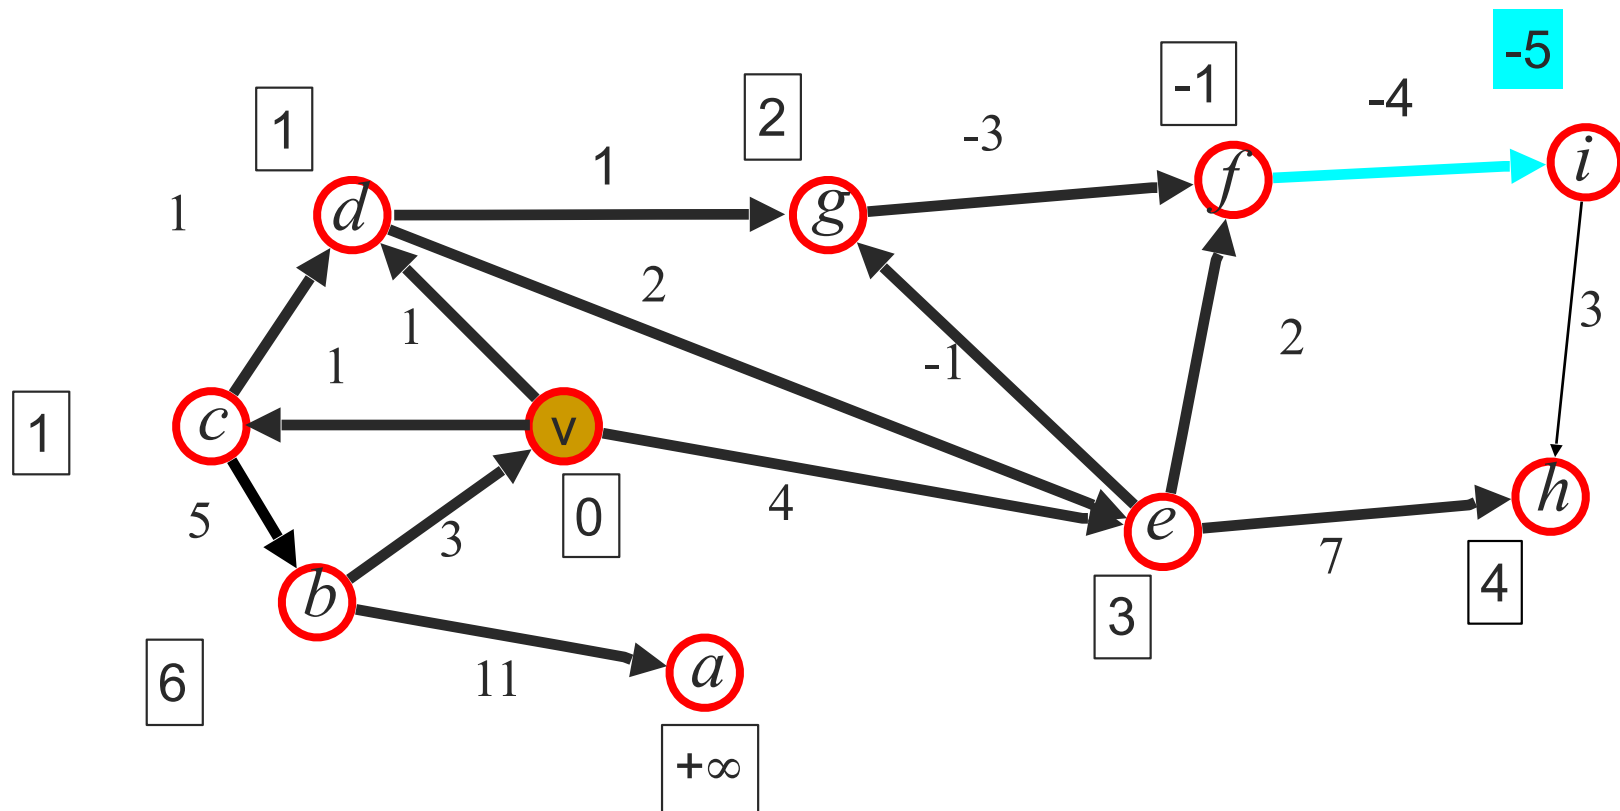
[i=2]



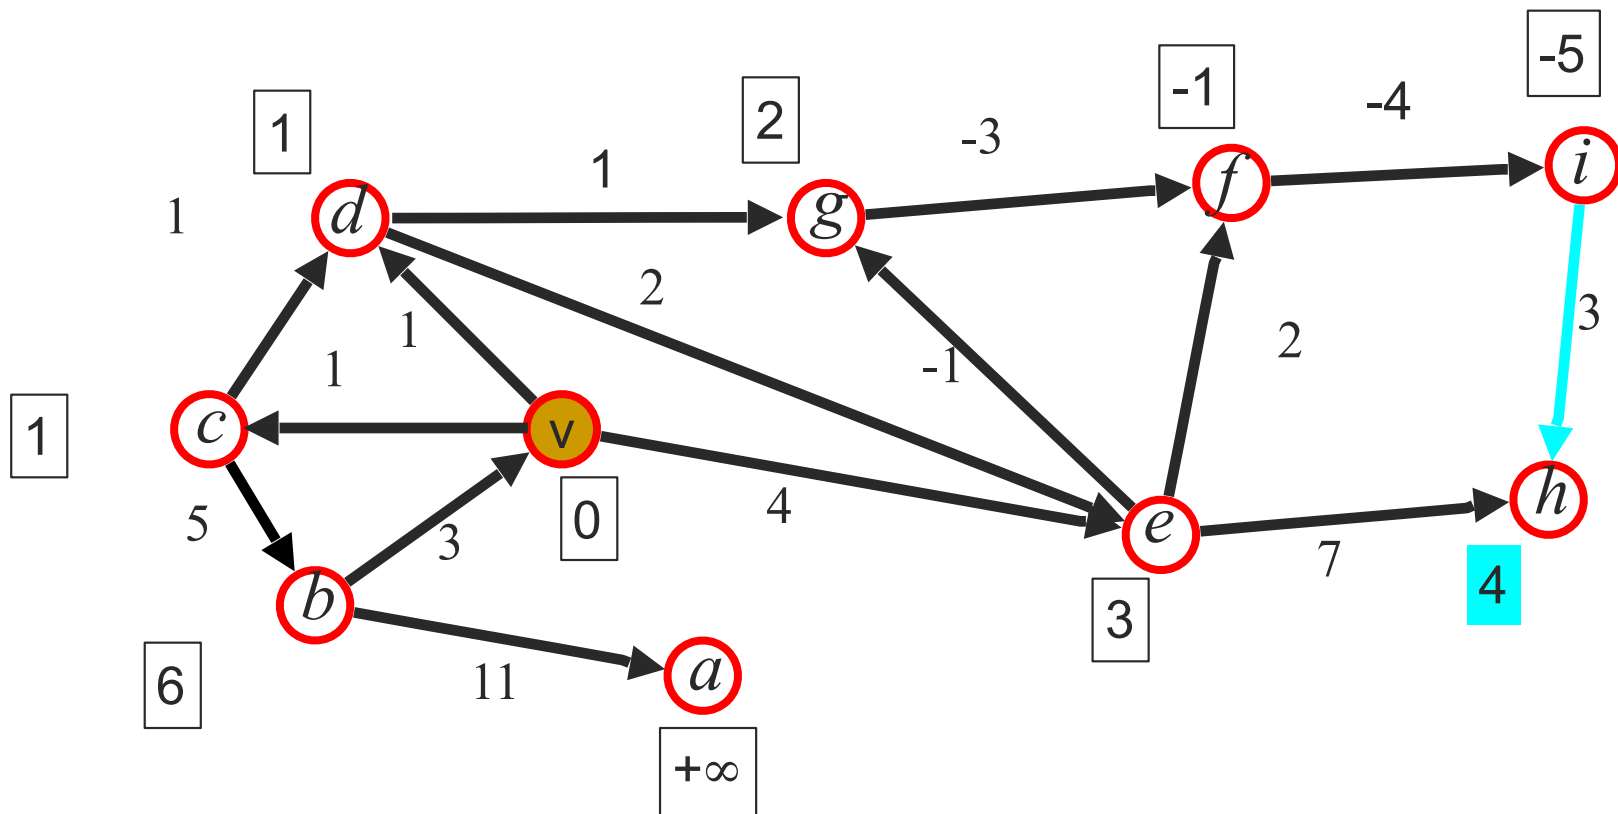
[i=2]



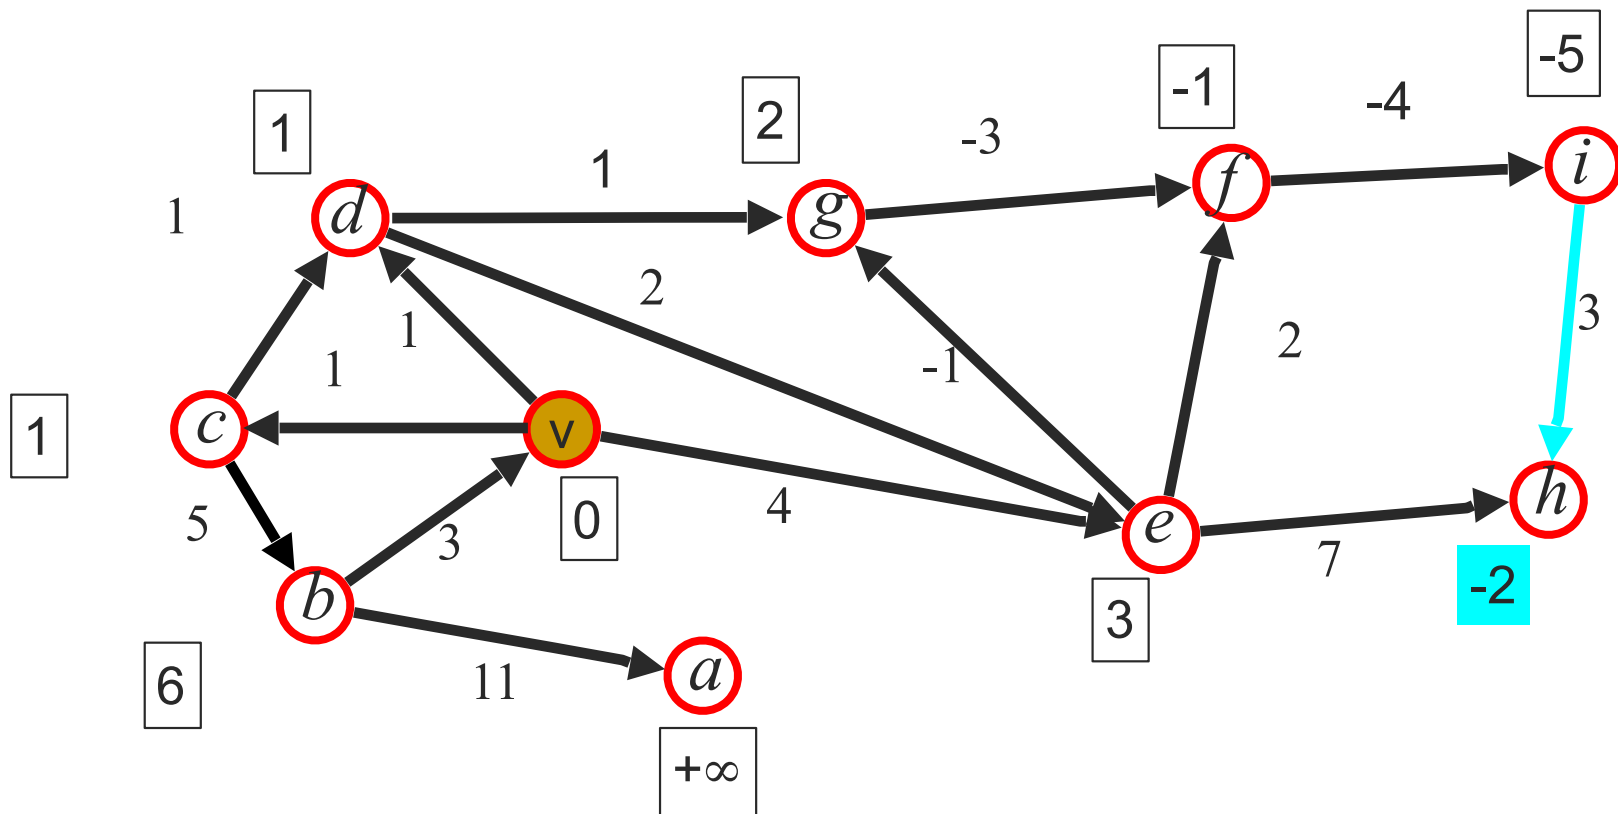
[i=2]




[i=2]




[i=2]




[i=2]




[i=3]




[i=3]




[i=3]




[i=3]




[i=3]




[i=3]




[i=3]




[i=3]




[i=3]



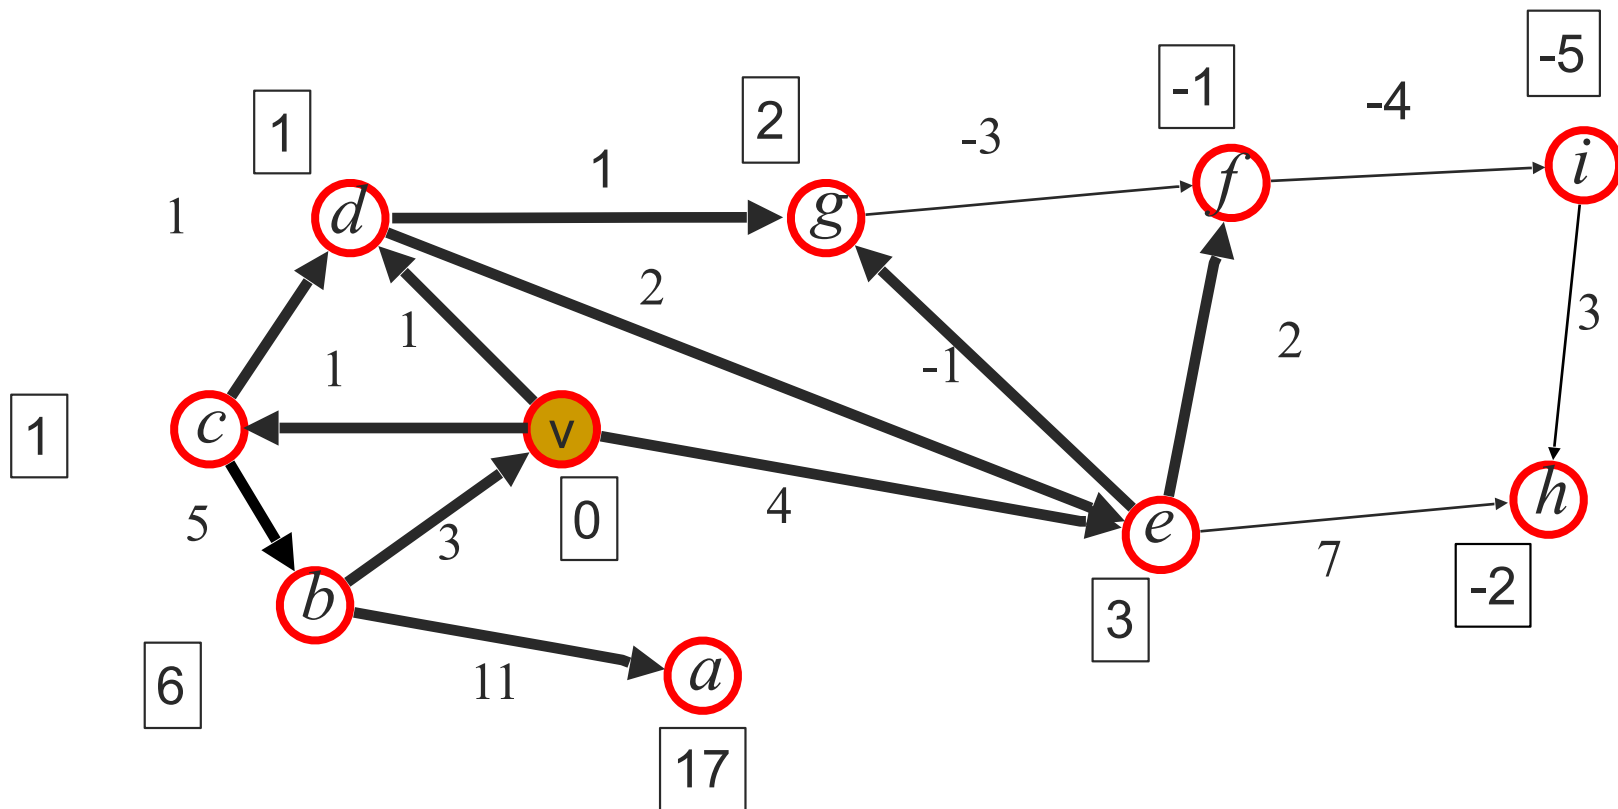
[i=3]



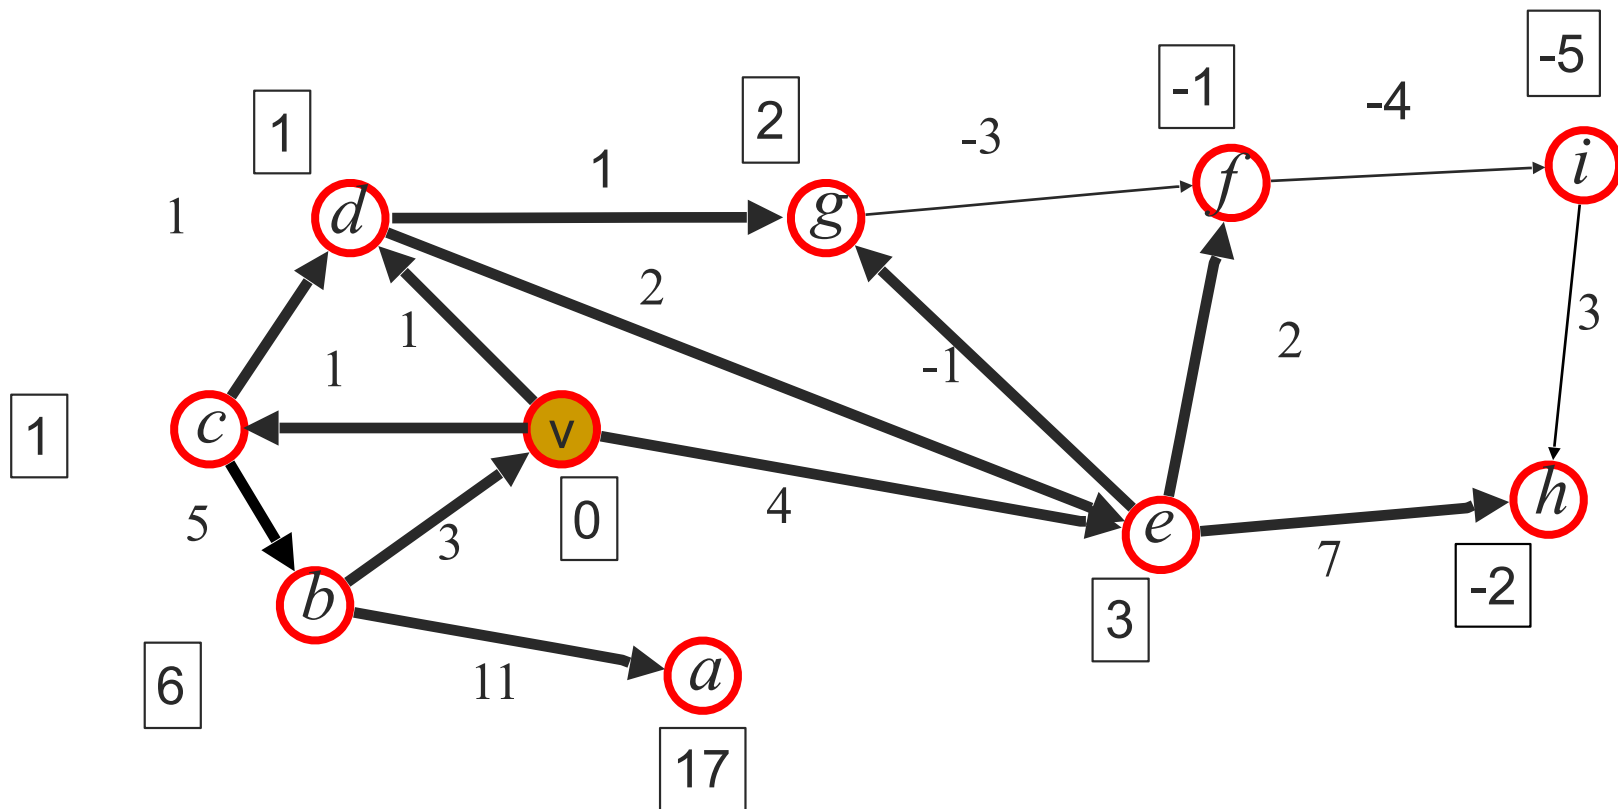
[i=3]



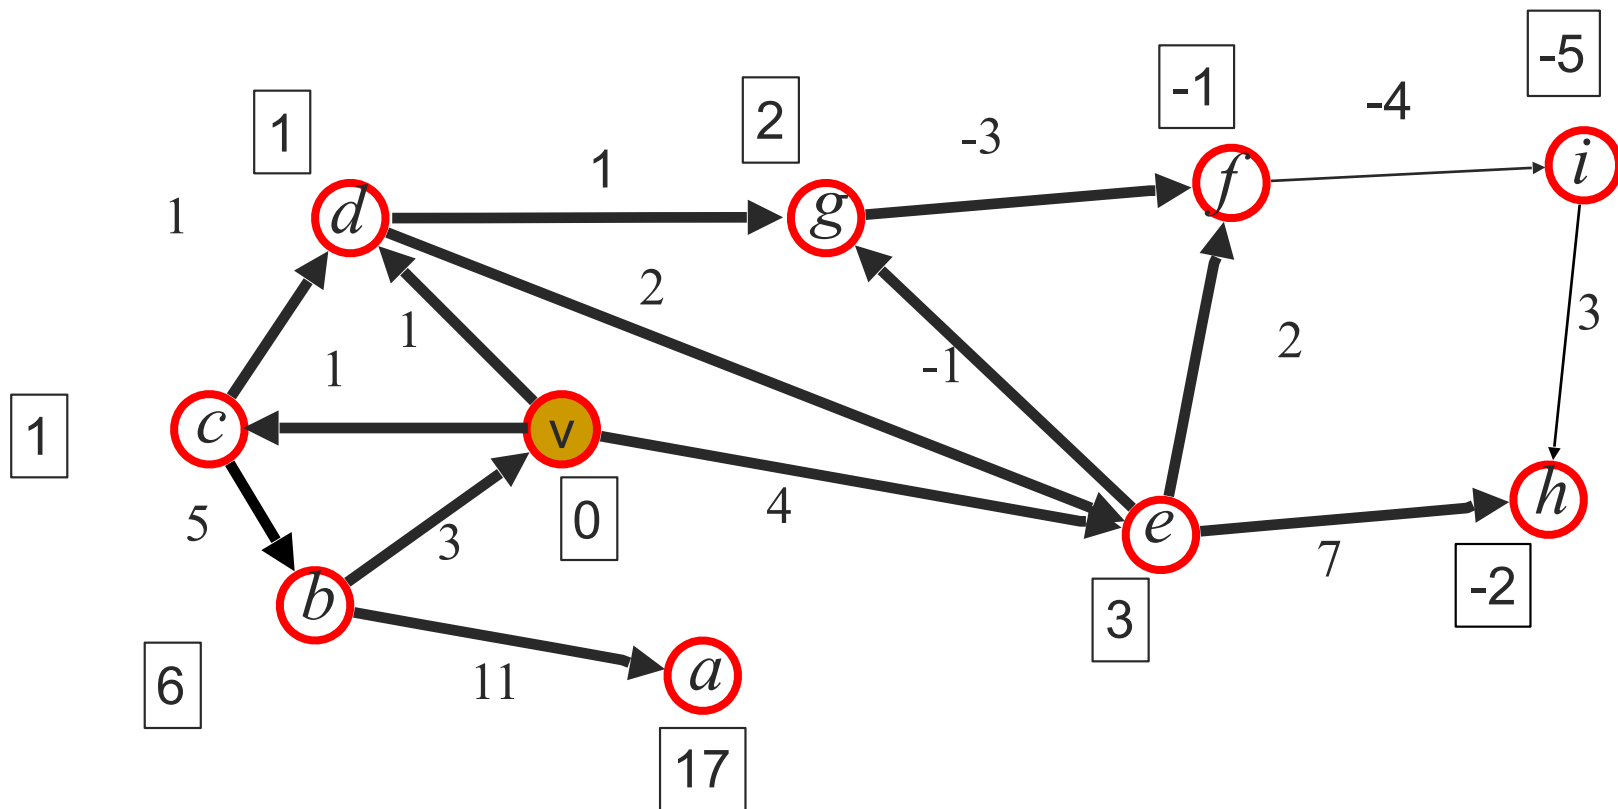
[i=3]



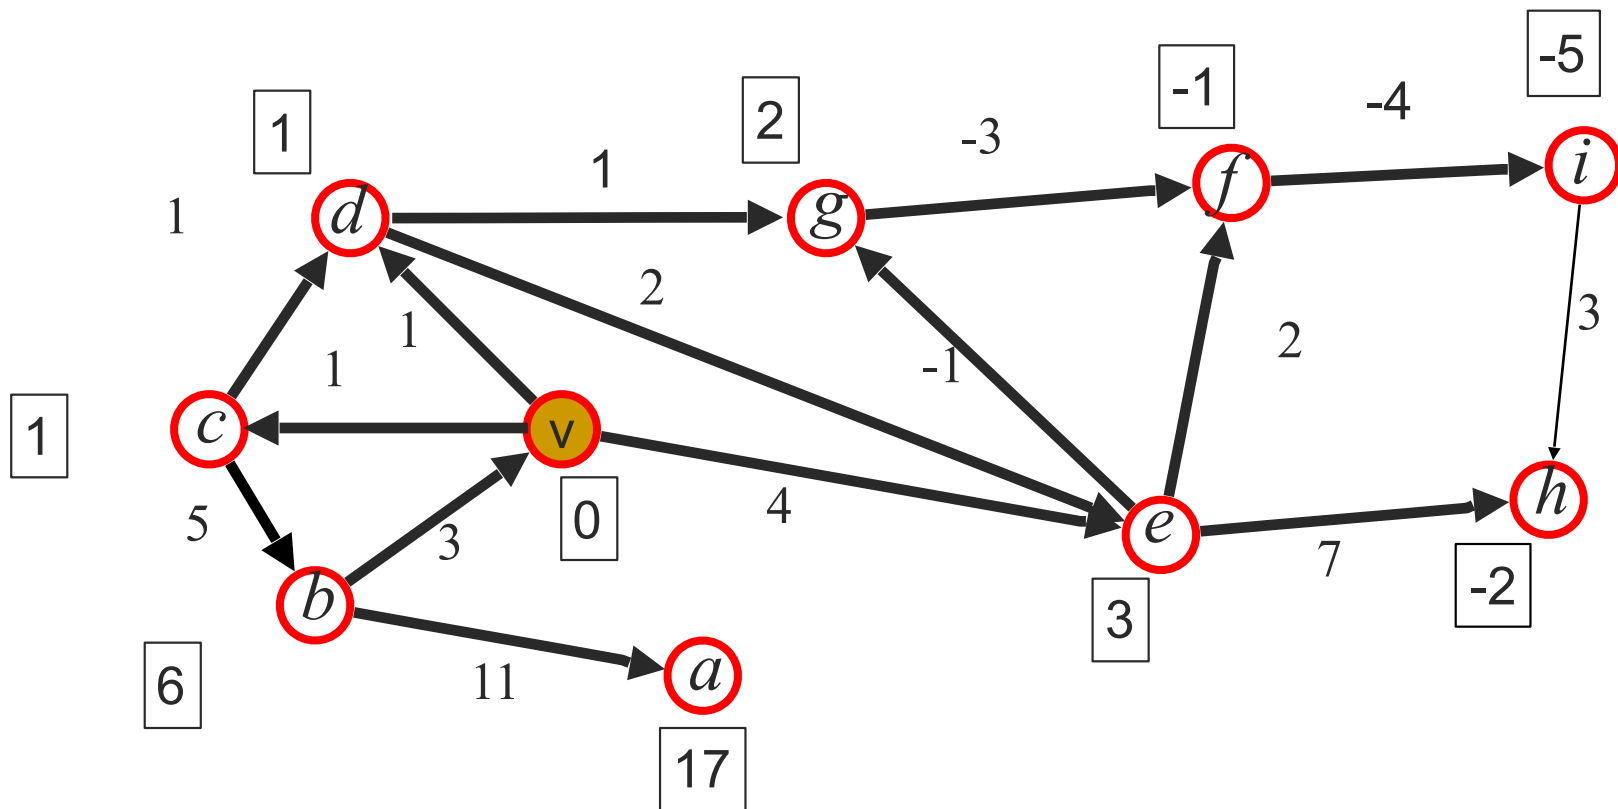
[i=3]



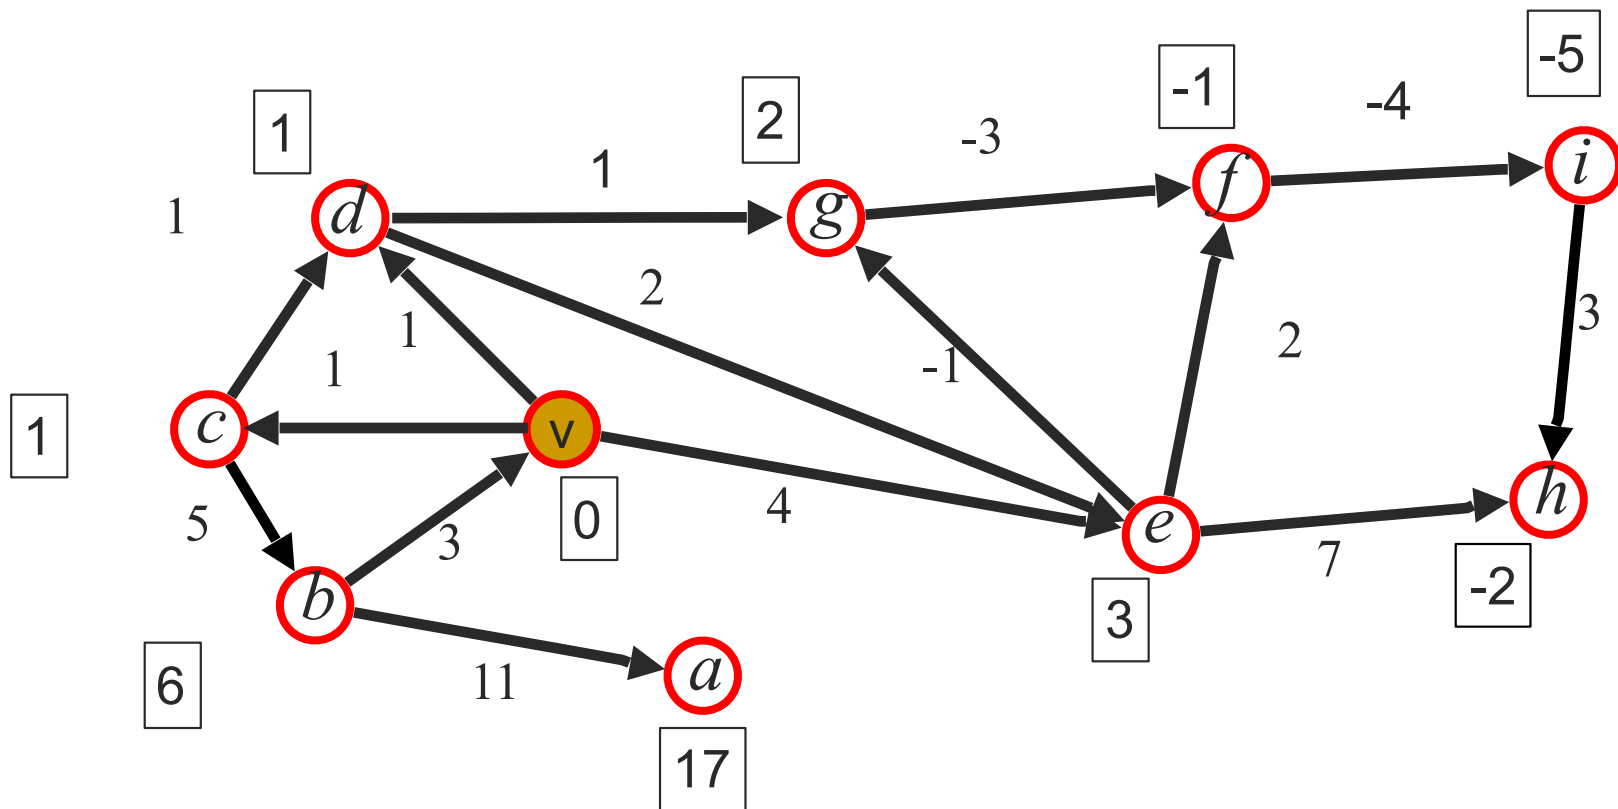
[i=3]



[i=3]



[i=3]



[Algorithm continues until $i=n-1$]

- In this example: $i = 9$, no more changes starting at $i = 4$

Running Time

[Algorithm Bellman-Ford(G, v)]

```
D[v] ← 0
for each vertex u ≠ v of G do
    D[u] ← +∞
for i ← 1 to n-1 do
    for each edge (u,z) in G do
        if D[u] + w((u,z)) < D[z] then
            D[z] ← D[u] + w((u,z))
if there are no edges left with potential
relaxation operations then
    return D
else
    return "G contains a negative cycle"
```

performs $n-1$
times a
relaxation of
every edge
in the graph

[Running time of Bellman-Ford algorithm]

- $O(nm)$

[Shortest Paths in directed acyclic graphs]

- Can we do faster than Bellman-Ford?
- Great final exam question 😊

[All-pairs shortest paths]

- For graphs with nonnegative edges
 - Run Dijkstra for each vertex (as a source).
 - n times $O(m \log n)$ is: $O(n m \log n)$
- For digraphs with negative edges
 - Run Bellman-Ford for each vertex (as a source).
 - n times $O(n m)$ is: $O(n^2 m)$
- Use programming Dynamic Programming

Dynamic Programming: the three steps

1. Characterize the structure of an optimal solution (**optimal substructure**)
2. Define the value of an optimal solution recursively in terms of the optimal solutions to subproblems (**overlapping subproblems**)
3. Construct an optimal solution with help of a look-up table

[Before computing All-Shortest Paths]

- Dynamic programming approach to solving the Longest Common Subsequence Problem

[Longest Common Subsequence]

Similarity of two strings (e.g. words, texts, biological sequences, ...)

Longest Common Subsequence

Input: Two strings s and t

Output: Find the longest common subsequence that appears in both s and t .