

# **16 Performance (1)**

## **CSC 230**



**Department of Computer Science**  
**University of Victoria**

**M&H: 6.5, 6.6**

**Stallings: 2.6 (skip Amdahl's law till later)**

# Performance issues: an introduction

how quickly can a program be executed?




**Review**

**Execution =  
OS time +  
program  
time**

# Unix "time" command

how long does a program take?



Answer is in 3 numbers:

- real
- user
- system

real = (t0 to t5)  
user = (t1 to t2) + (t3 to t4)  
system = (t0 to t1) + (t2 to t3) + (t4 to t5)

# Execution Time

## ❑ Elapsed Time (or Real Time)

- counts everything (*disk and memory accesses, I/O , etc.*)
- a useful number, but often not good for comparison purposes

## ❑ CPU time


- doesn't count I/O or time spent running other programs
- can be broken up into *system* time and *user* time

## ❑ Our focus: user CPU time

- time spent executing the lines of code that are "in" our program

## Which of these airplanes has the best performance?

<u>Airplane</u>	<u>Passengers</u>	<u>Range (mi)</u>	<u>Speed (mph)</u>
Boeing 737-100	118	1860	485
Boeing 747-400	416	7284	584
Boeing 787-8	210	8200	561
BAC/Sud Concorde	128	4500	1350
Douglas DC-8-50	146	8720	544
Airbus A320-200	148	3000	594
Airbus A340-200	261	8000	557
Airbus A380-800	525 to 853	8200	647



- ✓ Which plane delivers the greatest capacity on a route?  
(As measured in passenger-miles per hour)
- ✓ What factors does an airline consider when purchasing a fleet of planes?

# Computer Performance: TIME, TIME, TIME

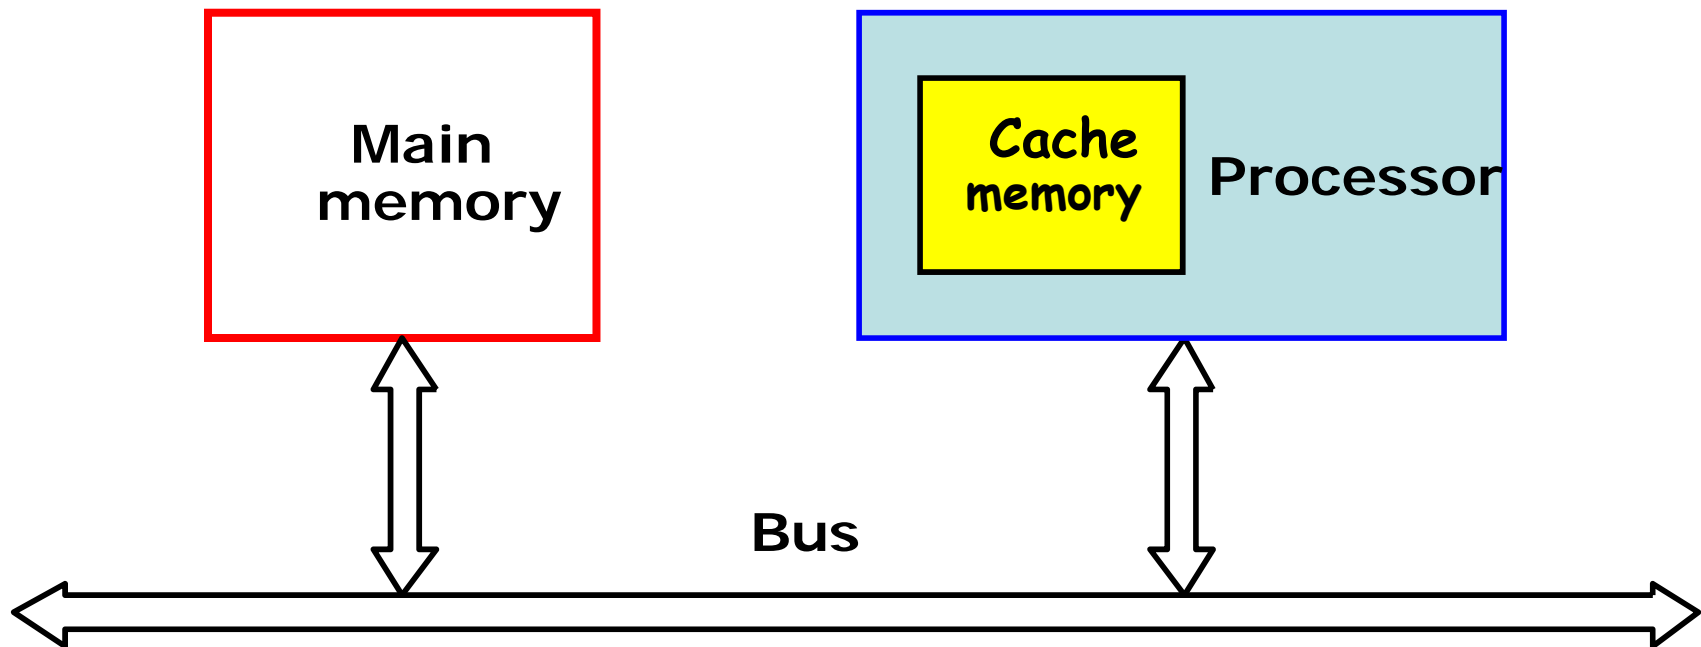
## □ Response Time (latency)

- How long does it take for a job to run?
- How long does it take to execute a job?
- How long must one wait for the database query?

## □ Throughput

- How many jobs can the machine run at once?
- What is the average execution rate?
- How much work is getting done?
- How many processes are executed in a unit of time?

- ✓ *If we upgrade a machine with a new processor what do we increase?*
- ✓ *If we add a new machine to the lab what do we increase?*




**Fetch-Decode-Execute Cycle**

**Processor speed versus memory access speed**

## Processor Clock

- ❑ Clock is a timing signal
- ❑ It defines intervals called clock cycles, each of length  $P$



*Clock rate  $R = 1 / P$*

- ❑ A machine instruction is divided into small steps
  - ➔ 1 clock cycle per step
- ❑ “Cycles per second” is a measure known as *Hertz (Hz)*

500 million cycles per second = 500 megahertz or 500 MHz

➔ 1 cycle every 2 ns

*Is the clock speed the most important factor?*



## Two small examples

### Reminders

$$1 \text{ ns} = 10^{-9} \text{ sec}$$

$$1 \text{ ms} = 10^{-3} \text{ sec}$$

$$1 \text{ sec} = 10^9 \text{ ns} = 10^3 \text{ ms}$$

$$1 \text{ MHz} = 10^6 \text{ hertz}$$

$$1 \text{ GHz} = 10^9 \text{ hertz}$$

$R = 1/P$  where  $R$  is  
clock rate,  $P$  time for  
1 clock cycle

If  $P = 25 \text{ ns/cycle}$ ,

$$\text{Then } R = \frac{1 \text{ cycles}}{25 \text{ ns}}$$

$$= 0.04 \text{ cycles/ns}$$

Translate to seconds  $\rightarrow$  multiply by  $10^9$

$\rightarrow 40,000,000 \text{ cycles/sec}$

$\rightarrow 40 \text{ MHz}$  is the clock rate  $R$

If  $R = 1 \text{ GHz}$  (clock rate) then

$$1 \text{ GHz} = 10^9 \text{ (hertz) cycles/sec} = 1,000,000,000 \text{ cycles/sec}$$

$\rightarrow$  Divide by  $10^9$  to get cycles/ns

$$\rightarrow (1,000,000,000 / 10^9) = 1 \text{ cycles/ns} = R$$

Since  $R = 1/P$

$\rightarrow$  invert to get ns/cycles

$$R = \frac{1 \text{ cycles}}{1 \text{ ns}} \Rightarrow P = \frac{1 \text{ ns}}{1 \text{ cycles}}$$

## Basic Performance Equation

$$T = \frac{N \times S}{R}$$

T = performance (time) for a program

N = actual number of executed instructions (not just number of instructions, i.e. loops are multiplied by number of iterations)

S = number of basic 1-clock steps needed for one instruction

R = clock rate (cycles per second)

**GOAL: reduce T → reduce (N or S or both) AND (increase R)**

*NOTE 1: these parameters are NOT independent*

*NOTE 2: architectural issues affect them*

## Performance Helpers

- ❑ Cache (done here)
- ❑ Pipelining (will do here an introduction)
- ❑ Superscalar architectures (will do here an introduction)
- ❑ Optimizing compilers (a bit)
- ❑ CISC (Complex Instruction Set Computers)
- ❑ RISC (Reduced Instruction Set Computers)

*Increase in performance is proportional to the square root of the complexity*  
*- 2 simpler components better than 1 large complex one?*

## SPEC Ratings

- ❑ SPEC (System Performance Evaluation Corporation)
- ❑ SPEC Benchmarks (SPEC CPU2000 benchmarks use a 300-MHz Sun Ultra5\_10 as the reference machine)

$$\text{SPEC rating} = \frac{\text{running time on reference computer}}{\text{running time on computer under test}}$$

For example, the Intel D925XECV2 motherboard (3.80 GHz, Pentium 4 processor 570J) has a SPECint2000 rating of 1671.

See <http://www.spec.org/cpu2006/results/>

*All these concepts will be reviewed and re-evaluated after the description of more architectural issues*

Higher Performance  Lower Execution Time

$$\text{Performance} = \frac{1}{\text{ExecutionTime}}$$


$X$  is  $n$  times faster than  $Y \rightarrow \frac{\text{Performance}_X}{\text{Performance}_Y} = n$

If  $X$  is  $n$  times faster than  $Y$ , then the execution time on  $Y$  is  $n$  times longer than on  $X$

$$\frac{\text{Performance}_X}{\text{Performance}_Y} = \frac{\text{Execution time}_Y}{\text{Execution time}_X} = n$$

## CPU Performance and clock

$$\text{CPU Execution time for a program} = \frac{\# \text{ CPU clock cycles for a program}}{\text{Clock rate}} \times \text{clock cycle time}$$



$$\text{CPU Execution time for a program} = \frac{\# \text{ CPU clock cycles for a program}}{\text{Clock rate}}$$

*Independent of the number of instructions in a program*

## **CPI : clock cycles per instruction**

Think of execution time as:

$$\text{CPU clock cycles} = \frac{\text{Instruction count for a program}}{\text{Average clock cycles per instruction (CPI)}}$$

OR

$$\text{CPI} = \frac{\text{CPU clock cycles}}{\text{Instruction count}}$$

## Small Example: same program on 2 machines

### CPU A:

CPI = 5 *(i.e. 5 clock cycles per instruction)*

Clock period  $P = 100 \text{ ns}$

### CPU B:

CPI = 3.5 *(i.e. 3.5 clock cycles per instruction)*

Clock period  $P = 120 \text{ ns}$

Test yourself: what is the  
clock rate  $R$  for each CPU?



## Small Example: same program on 2 machines

### CPU A:

CPI = 5 *(i.e. 5 clock cycles per instruction)*

Clock period P = 100 ns

### CPU B:

CPI = 3.5 *(i.e. 3.5 clock cycles per instruction)*

Clock period P = 120 ns

Test yourself: what is the  
clock rate R for each CPU?

*Clock A = 25%  
faster*

### CPU A:

Clock period = 100 ns/cycle

$$\Rightarrow R = \frac{1}{100} \text{ cycles / ns}$$

= 0.01 cycles/ns

*Multiply by  $10^9$  to get seconds*

= 10,000,000

= 10 MHz clock rate R

### CPU B:

Clock period = 120 ns/cycle

$$\Rightarrow R = \frac{1}{120} \text{ cycles / ns}$$

= 0.008 cycles/ns

*Multiply by  $10^9$  to get seconds*

= 8,000,000

= 8 MHz clock rate R

## Small Example: same program on 2 machines

### CPU A:

CPI = 5 *(i.e. 5 clock cycles per instruction)*

Clock period P = 100 ns

### CPU B:

CPI = 3.5 *(i.e. 3.5 clock cycles per instruction)*

Clock period P = 120 ns

Let I = number of instructions

CPU clock cycles [A] = I x 5

CPU clock cycles [B] = I x 3.5

CPU time = CPU clock cycles x clock cycle time

CPU time[A] = I x 5 x 100 ns = I x 500 ns

CPU time[B] = I x 3.5 x 120 ns = I x 420 ns

*How much faster is B?*

$$\frac{\text{Performance}_B}{\text{Performance}_A} = \frac{\text{Execution time}_A}{\text{Execution time}_B} = \frac{I \times 500 \text{ ns}}{I \times 420 \text{ ns}} = 1.19$$

→ B is 19% faster than A

## Classic CPU Performance Equation

CPU time = Instruction count x CPI x Clock cycle time

OR

CPU time =  $\frac{\text{Instruction count} \times \text{CPI}}{\text{Clock rate}}$

## Comparing Code Segments

A compiler designer needs to decide between two code sequences for a system. Here are the data from hardware for each class of instructions which can be used:

CPI for each instruction class			
	A	B	C
CPI	1	2	3

Here is the data for each code sequence:

Code Sequence	Instruction count for each instruction class		
	A	B	C
1	2	1	2
2	4	1	1

## Which Code Segment executes the most instructions?

Sequence 1 =  $2+1+2 = 5$  instructions

Sequence 2 =  $4+1+1 = 6$  instructions

Code Sequence	Instruction count for each instruction class		
	A	B	C
1	2	1	2
2	4	1	1

CPI for each instruction class

	A	B	C
CPI	1	2	3

## Which Code Segment is faster?

Need to compute total CPU clock cycles, using CPI and instruction count

CPU clock cycles = for each class (Instruction count x CPI)

CPU clock cycles[1] = (2x1) + (1x2) + (2x3) = 10 cycles

CPU clock cycles[2] = (4x1) + (1x2) + (1x3) = 9 cycles

*Sequence 2 is faster even if it executes more instructions*

Code Sequence	Instruction count for each instruction class		
	A	B	C
1	2	1	2
2	4	1	1

CPI for each instruction class

	A	B	C
CPI	1	2	3

## What is the CPI for each sequence?

Since Sequence 2 is faster even if it has more instructions, it should have a lower CPI overall

$$\text{CPI} = \frac{\text{CPU clock cycles}}{\text{Instruction count}}$$

$$\text{CPI}_1 = \frac{\text{CPU clock cycles}}{\text{Instruction count}} = \frac{10}{5} = 2.0$$

$$\text{CPI}_2 = \frac{\text{CPU clock cycles}}{\text{Instruction count}} = \frac{9}{6} = 1.5$$

Code Sequence	Instruction count for each instruction class		
	A	B	C
1	2	1	2
2	4	1	1

CPI for each instruction class

	A	B	C
CPI	1	2	3

## Instruction Frequency

Frequency of occurrence of instruction types for a variety of languages. The percentages do not sum to 100 due to roundoff. (Adapted from Knuth, D. E., *An Empirical Study of FORTRAN Programs, Software—Practice and Experience*, 1, 105-133, 1971.)

Statement	Average Percent of Time
Assignment	47
If	23
Call	15
Loop	6
Goto	3
Other	7



## Complexity of Assignments

Percentages showing complexity of assignments and procedure calls. (Adapted from Tanenbaum, A., *Structured Computer Organization*, 4/e, Prentice Hall, Upper Saddle River, New Jersey, 1999.)

	<b>Percentage of Number of Terms in Assignments</b>	<b>Percentage of Number of Locals in Procedures</b>	<b>Percentage of Number of Parameters in Procedure Calls</b>
0	—	22	41
1	80	17	19
2	15	20	15
3	3	14	9
4	2	8	7
$\geq 5$	0	20	8

## Speedup: another view

Speedup  $S$  is the ratio of the time needed to execute a program without improvement ( $T_{wo}$ ) to the time required with an improvement ( $T_w$ )

$$S = \frac{T_{wo}}{T_w}$$

$$S = \frac{T_{wo} - T_w}{T_w} \times 100$$

Time  $T$  is computed as: (the instruction count  $IC$ )  $\times$  (the number of cycles per instruction  $CPI$ )  $\times$  (the cycle time  $t$ )

$$T = IC \times CPI \times t$$

Substituting  $T$  into the speedup percentage calculation above yields:

$$S = \frac{(IC_{wo} \times CPI_{wo} \times t_{wo}) - (IC_w \times CPI_w \times t_w)}{(IC_w \times CPI_w \times t_w)} \times 100$$

## Speedup: another view

Speedup  $S$  is the ratio of the time needed to execute a program without improvement ( $T_{wo}$ ) to the time required with an improvement ( $T_w$ )

$$S = \frac{T_{wo}}{T_w} \quad S = \frac{T_{wo} - T_w}{T_w} \times 100$$

Time  $T$  is computed as: (the instruction count  $IC$ )  $\times$  (the number of cycles per instruction  $CPI$ )  $\times$  (the cycle time  $t$ )

$$T = IC \times CPI \times \tau$$

Substituting  $T$  into the speedup percentage calculation above yields:

$$S = \frac{IC_{wo} \times CPI_{wo} \times \tau_{wo} - IC_w \times CPI_w \times \tau_w}{IC_w \times CPI_w \times \tau_w} \times 100$$