

# **21 Cache Organization**

## **CSC 230**

**Department of Computer Science**  
**University of Victoria**

**M&H: 7.5, insert (p.284) on Intel vs Motorola**

**Stallings: 4.2, 4.3, 4.5**

## Cache storage : locality of reference

*Cache: a safe place for hiding or storing things (Webster's)*

90% of execution time spent on 10% of code

### *Temporal locality*

when the CPU accesses a piece of information, there is a high probability it will be accessed *again soon in time*

### *Spatial locality*

when the CPU accesses a piece of information, there is a high probability that the data *in nearby locations* will be accessed as well

## Cache : locality of reference

→ when data is accessed 1<sup>st</sup> time, put in cache, ready for next time



*Temporal locality*


→ when data is accessed, transfer an adjacent block to cache (and/or use a large block size)



*Spatial locality*


- ❑ Put ACTIVE segments of program in fast cache memory
- ❑ Cache operations are supported by hardware control

## Cache : placement and speed



**The cache consists of a number of cache lines  
{slots/blocks}**

**Each cache line consists of two parts**



**cache hit:** when data is found in the cache

**cache miss:** when data is not found in the cache  
(next go to memory)

**cache full:** need to decide how to create space  
→ replacement

## A few more definitions about cache

**Miss rate** – the fraction of memory accesses not found in a level of the memory hierarchy (e.g. L1 cache miss because data is in RAM)

**Hit time** – the time required to access a level of the memory hierarchy, including the time needed to determine whether the access is a hit or a miss

**Miss penalty** – the time required to fetch a block into a level of the memory hierarchy from the lower level, including:

- the time to access the block,
- transmit it from one level to the other,
- insert it in the level that experienced the miss

# How the Cache and the CPU operate together

1. CPU requests contents of memory location
2. Cache is checked first for this data (*hardware-supported search*)
3.
  - [3A] If present, get data from cache: HIT (fast)
  - [3B.1] If not present, MISS, issue a READ for the required block *from main memory to cache*
  - [3B.2] Then copy from cache to CPU
4. Cache includes tags to identify which block of main memory is in each cache slot

# Cache design issues

✓ Size

✓ Mapping Function

*In detail here*

✓ Replacement Algorithm

✓ Write Policy

✓ Block Size

✓ Number of Caches

*Definitions only in this course*



# Write Policy

## ❑ Write-through protocol

- + cache and memory are updated at the same time
- + simple
- + memory is consistent when there are two CPUs
- extra unnecessary writes

## ❑ Write-back protocol

- use Dirty/Modified bit to flag cache
- can still write too much as it will update whole block
- complications if two CPUs access same block
- + update memory later when cache block is removed

# Cache Replacement Policies

## Possible algorithms

- ☐ Least Recently Used (LRU)  
e.g. in 2-way set-associative which of the 2 blocks is LRU?
- ☐ First-in First-out (FIFO)  
replace block that has been in cache longest
- ☐ Least Frequently Used  
replace block which has had fewest hits
- ☐ Random
- ☐ Instructions available to programmers?

## Sizes

### Example:

Cache of 64K Byte

Cache block/line of 4 bytes

i.e. cache is 16K ( $2^{14}$ ) lines of 4 bytes

16 Mb main memory

24 bit address ( $2^{24}=16\text{M}$ )

### Example: a home laptop

Cache of 512K Byte

Cache blocks/lines of 16 bytes

768 Mb main memory

# Cache Mapping Strategies: preview summary

## Direct

- a memory block always goes to the same cache line
- memory block  $j$  goes to cache block  $(j \bmod \text{size})$
- needs tagging

## Associative


- a memory block can go to any cache line
- needs tagging

## Set-Associative

- cache is divided into a number of sets
- each set contains a number of lines
- block  $j$  goes to set  $(j \bmod \#sets)$   
then to *any* line in that set
- needs tagging

## [1]: DIRECT-MAPPED CACHE

- ❑ Each memory location is mapped to exactly one cache location using a mapping function
- ❑ *Mapping function:*  
*(Block address) modulo (Number of cache blocks)*
- ❑ There is only one choice of what to replace



# DIRECT-MAPPED CACHE


*Mapping:*

*(Block address) modulo (Number of cache blocks)*


*Example below: (address) MOD (8)*

*e.g. (01101) MOD (8)*

- If the number of entries in the cache is a power of two, then modulo is given by the low order bits of the address → where the number of bits =  $\log_2$  (cache size)




## Example



Cache

each block = 16 words  
total 2K words

Direct-mapped cache.




Main memory

16-bit addresses  
64K words  
(4K blocks each of 16 words)

*each block  
contains 16 words*


*Blocks 0 to 4095 =  
 $2^{12}$  blocks =  
 $2^2 \times 2^{10} = 4K$  blocks*



### Cache

each block = 16 words  
total 2K words

**Direct-mapped cache.**



### Main memory

16-bit addresses  
64K words  
(4K blocks each of 16 words)

**Block J → line (J mod 128)**

block number  
in memory


line number  
in cache

*e.g. 212 → 84*



# DIRECT-MAPPED CACHE

(from textbook) – read example




## [2]: ASSOCIATIVE CACHE

**Fully associative:** A block is placed in *any* location in cache  
No cache index is needed  
Reduces miss ratio

### Searching:

- To find a given block, *all the entries* in the cache must be searched
- Search is done *in parallel in hardware*
- Practical only for caches with small number of blocks (else hardware too expensive)

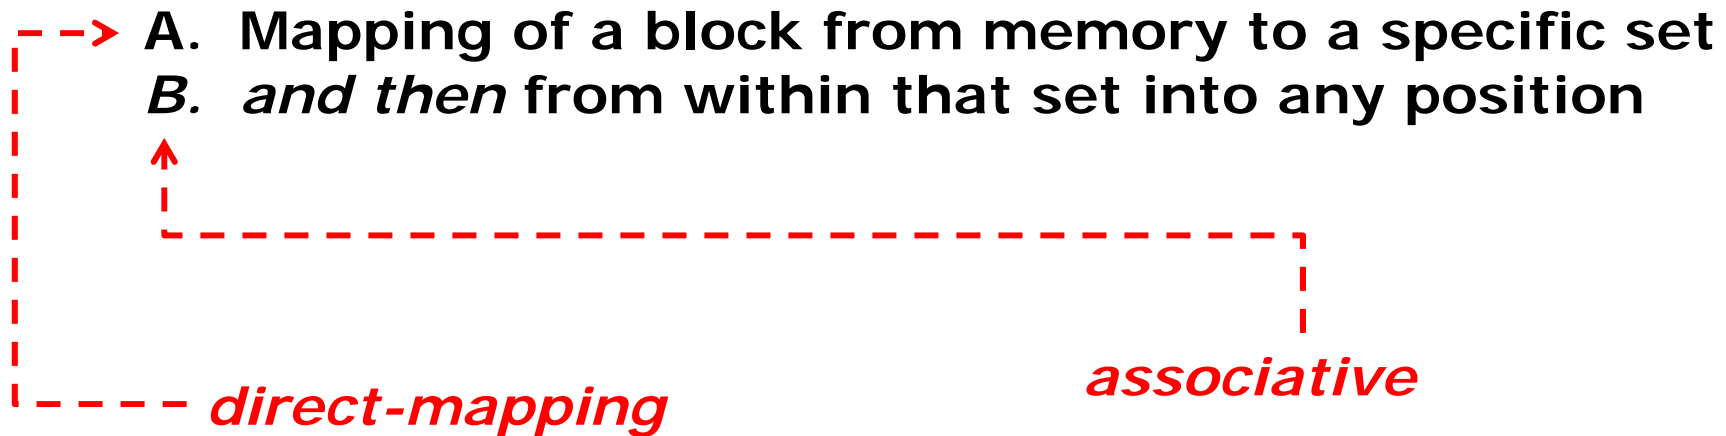


➔ *Most flexible and most expensive*

## [3]: SET-ASSOCIATIVE CACHE

*Combination of direct-mapped and fully associative*

□ Blocks of the cache are grouped into sets



Given a memory block **J**:

1. The set is given by:

(*block number J*) modulo (*Number of sets in the cache*)

2. Then any line number within that set

## SET-ASSOCIATIVE CACHE: notes

- *A Direct-mapped cache is simply a one-way set-associative cache*

In direct-mapped, each cache entry holds **1** line and forms **1** set with **1** element

- *A fully associative cache with  $m$  entries is simply an  $m$ -way set-associative cache*

It has **one** set with  **$m$**  lines

In *set-associative*, an entry can reside in any line within a set


**Example. Set-associative-mapped cache with 2 lines per set.**

**Main memory**  
(same as before)  
16-bit addresses  
64K words  
(4K lines each of 16 words)

**Cache**

tag	Line 0
tag	Line 1
tag	Line 2
tag	Line 3
...	
tag	Line 126
tag	Line 127

**Main memory**



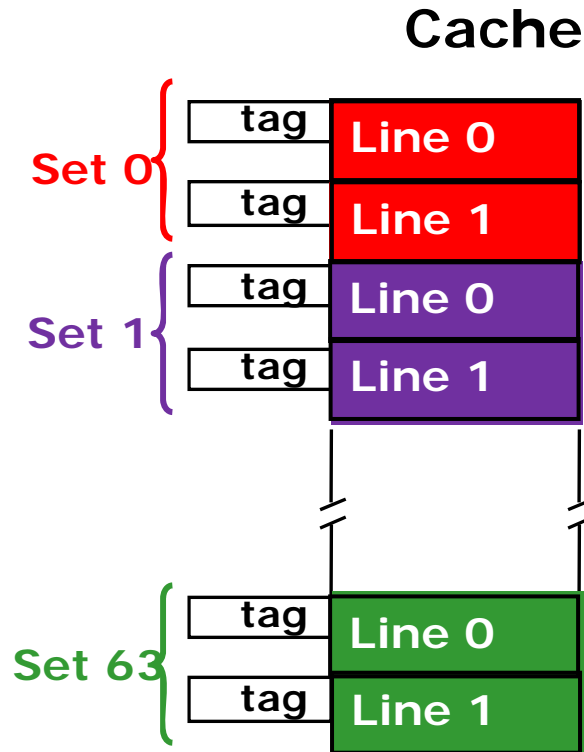
**Cache**

each line = 16 words  
128 blocks organized into 64  
sets of 2 lines each  
(total 2K words as before)



**Example. Set-associative-mapped cache with 2 lines per set.**

**Main memory**  
(same as before)  
16-bit addresses  
64K words  
(4K lines each of 16 words)

**Cache**  
each line = 16 words  
128 blocks organized into 64 sets of 2 lines each  
(total 2K words as before)




**Main memory**




**Example. Set-associative-mapped cache with 2 lines per set.**

**Main memory**  
(same as before)  
16-bit addresses  
64K words  
(4K lines each of 16 words)

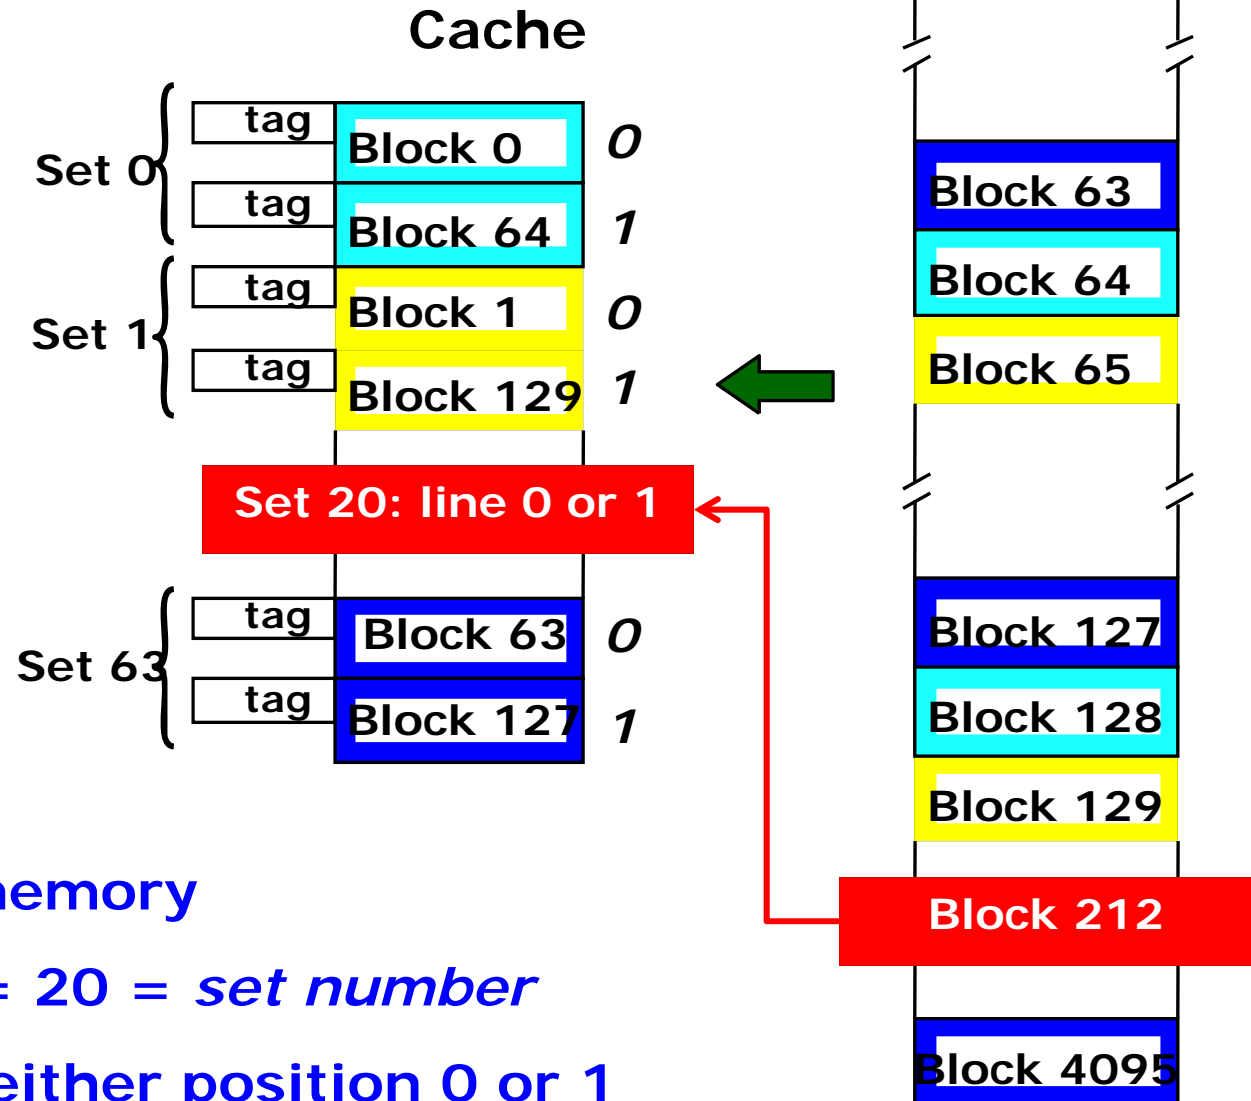


**Cache**  
each line = 16 words  
128 blocks organized into 64 sets of 2 lines each  
(total 2K words as before)

Main memory



**Example. Set-associative-mapped cache with 2 lines per set.**



**Example:**

**Block 212 from memory**

**→  $(212) \bmod 64 = 20 = \text{set number}$**

**→ within set 20, either position 0 or 1**



## To find a block in a set-associative cache

### Decompose the address: index + tag + offset

- Index is used to select the set
- Tag is used to choose the block by parallel comparison within the selected set
- Block offset is the address of the desired data within a block

### Advantages:

- Decreases the miss rate as it reduces the misses that compete for the same location

### Disadvantages:

- Requires *parallel* comparators
- Requires more tag bits per cache block
- Extra delay for data, increasing hit time

➔ *Do not need to become competent at computing tags/index/offset here*

## MULTI-LEVEL CACHE

Often primary cache is on the same chip as the processor

Add 2<sup>nd</sup> SRAMs as second cache (on chip or on board)

**First cache** closes the speed gap between the CPU and memory

**Multilevel** cache reduces the miss penalty

- Miss penalty goes down if data is in 2<sup>nd</sup> level cache
- Primary cache can be smaller and have a higher miss rate

If the 2<sup>nd</sup> level cache contains the desired data, the miss penalty = access time of the 2<sup>nd</sup> level cache

If the 2<sup>nd</sup> level cache does not contain the data, a main memory access is required, *a larger miss penalty* is incurred

➔ Primary cache is often 10 or more times faster than secondary

# Comparing single and multi-level cache

**Primary cache** of a multilevel cache:

- often smaller
- often uses a smaller block size
- tries to optimize the hit time for a shorter clock cycle
- many recent systems use *direct-mapped primary caches* because of their advantage in access time and simplicity

**Secondary cache:**

- usually larger
- access time less critical
- often uses larger block size
- Tries to optimize the miss rate to reduce the penalty of long memory access time

On-chip caches have associativity of two to four, while off-chip caches rarely have associativity greater than two

## Another Design Concept: Split cache

Two independent caches which operate in parallel

→ instruction cache (I cache)

→ data cache (D cache)

❑ To minimize collisions, separate fetches of data and instructions


- Then instructions only have collisions with instructions, data with data
- The two have different behaviour patterns
- Goal is to decrease the miss ratio from a combined cache

❑ The miss ratio does not necessarily improve!

*(do not confuse split and multi-level cache)*

## Instruction cache miss steps

1. Send original address to memory
2. Send "READ" command to memory and wait for access cycle
3. Write the cache entry, plus tag and V bit
4. *RESTART* instruction cycle from fetch  
→ it will cause a cache hit



**Miss rate goes UP if block size is too large relative to cache size**

## CACHE DESIGN: what the designer must consider

- ☐ Cache size
- ☐ Block size
- ☐ Associativity
- ☐ Replacement policy
- ☐ Write-hit policy (write-through, write-back)

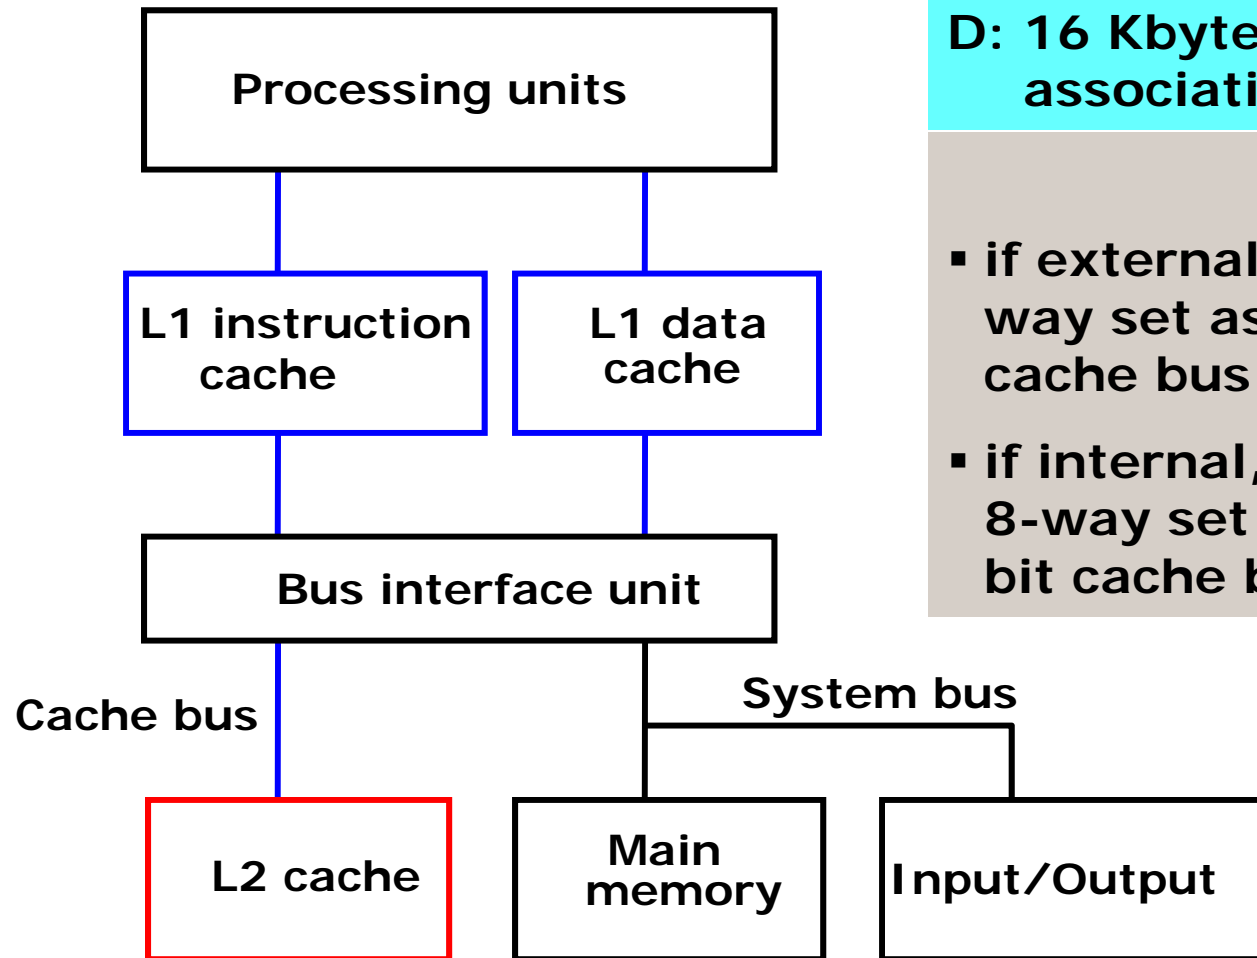
Design change	Effect on miss rate	Possible negative performance effect
Increase size	decrease misses	may increase access time
Increase associativity	decreases miss rate due to conflict misses	may increase access time
Increase block size	decreases miss rate for a wide range of block sizes	may increase miss penalty





## ARM Cache

- ❑ One cache only
- ❑ similar organization to 68040
- ❑ 4-way set associative
- ❑ 64 sets
- ❑ each set with 4 lines
- ❑ each block has 16 bytes = 4 32-bit words
- ❑ write-through protocol and random replacement
  - ➔ simplicity



## L1

**I: 16 Kbyte instruction 2-way set associative**

**D: 16 Kbyte data, 4-way set associative**

## L2

- if external, 512 Kbyte, 4-way set associative, 64 bit cache bus
- if internal, 256 Kbyte data, 8-way set associative, 256 bit cache bus

**Caches and external connections in Pentium III processor**

## Examples

Time to transfer a block of data from memory to cache after miss


- cache with 8-word blocks
- 1 clock cycle to send address to memory
- in DRAM, 1<sup>st</sup> word accessed in 8 cycles, following words in 4 cycles
- 1 clock cycle to send 1 word data back to cache

### EX 1: Single memory module


$$1 + 8 + (7 \times 4) + 1 = 38 \text{ cycles}$$

The diagram illustrates the components of the 38-cycle calculation:

- send address* points to the first **1**.
- Word 1* points to the **8**.
- Next 7 words* points to the **(7 x 4)**.
- Write to cache* points to the final **1**.




a. One-word wide memory organization



b. Wide memory organization

*Read only and be aware!*



c. Interleaved memory organization

## Examples

Time to transfer a block of data from memory to cache after miss

- cache with 8-word blocks
- 1 clock cycle to send address to memory
- in DRAM, 1<sup>st</sup> word accessed in 8 cycles, following words in 4 cycles
- 1 clock cycle to send 1 word data back to cache

EX 1: Single memory module

$$1 + 8 + (7 \times 4) + 1 = 38 \text{ cycles}$$

EX 2: Four interleaved memory modules

$$1 + 8 + 4 + 4 = 17 \text{ cycles}$$

*first 4 words accessed in parallel*

*first 4 words to cache in parallel*

*plus 4 more words accessed in parallel*

*last 4 words to cache in parallel*

[1] Consider memory access time with different memory widths: one-word wide

*Read only and be aware!*

If cache block = 4 words,

DRAM is 1 word wide,

→ then miss penalty:


$$1 + (4 \times 15) + (4 \times 1) = 65$$

memory bus clock cycles

Bytes transferred per bus clock cycle for a miss is:

$$(4 \times 4) / 65 = 0.25$$

→ effective bandwidth



a. One-word wide memory organization

**[2] Consider memory access time with different memory widths: 2-word wide**

**If cache block = 4 words,  
DRAM is 2 words wide,**

**→ then miss penalty:**


$$1 + (2 \times 15) + (2 \times 1) = 33$$

**memory bus clock cycles**

**Bytes transferred per bus clock cycle for a miss is:**

$$(4 \times 4) / 33 = 0.48$$


**→ effective bandwidth**



**b. Wide memory organization**

*Read only and  
be aware!*

**[3] Consider memory access time with different memory widths: 4-word wide**



**If DRAM is 4 words wide,  
→ then miss penalty:**

$$1 + (1 \times 15) + (1 \times 1) = 17$$

memory bus clock cycles

**b. Wide memory organization**

**Bytes transferred per bus clock cycle  
for a miss is:**

$$(4 \times 4) / 17 = 0.94$$

**→ effective bandwidth**

*Read only and  
be aware!*



In general, average access time from CPU

$$T(\text{average}) = h \times C + (1 - h) \times M$$

The diagram shows the formula  $T(\text{average}) = h \times C + (1 - h) \times M$  with several annotations. The variables  $h$ ,  $C$ ,  $(1 - h)$ , and  $M$  are each enclosed in a red dashed circle. An arrow points from the text *hit rate* to the circle containing  $h$ . Another arrow points from the text *cache access time* to the circle containing  $C$ . A third arrow points from the text *miss penalty* to the circle containing  $M$ . A bracket below the plus sign and the  $(1 - h)$  term is labeled *miss = no hit*.

*hit rate*

*cache access time*

*miss penalty*

*miss = no hit*

## Look at another example (follow up)

$$T(\text{ave}) = h \times C + (1 - h) \times M$$

- 10 cycles for DRAM access, no cache
- cache with 8-word blocks, and interleaved memory
- 30% of all instructions need a Read or Write
  - ➔ i.e. 130 memory accesses every 100 instructions executed (fetch!)

### GIVEN:

- hit rate in instruction cache = 0.95,
- hit rate in data cache = 0.90
- need 17 cycles to load a block in cache (see earlier example)

## Calculate the speedup rate

Calculate the ratio between  
a "no cache" and a "with  
cache" system

$$\frac{\text{no cache}}{\text{with cache}} =$$

$$\text{no cache} = 130 \times 10$$

Accesses to memory for  
100 instructions + 30 data

Given time for each  
Memory access

$$\text{with cache} = 100 (0.95 \times 1 + 0.05 \times 17)$$

Accesses to  
memory for  
100 instructions

$$+ 30 (0.9 \times 1 + 0.1 \times 17)$$

Accesses to  
memory for  
30 data

$$\frac{130 \times 10}{100 (0.95 \times 1 + 0.05 \times 17) + 30 (0.9 \times 1 + 0.1 \times 17)} = 5.04$$

$$T(\text{ave}) = h \times C + (1 - h) \times M$$

(1) 10 cycles for DRAM access, no cache

(2) cache with 8-word blocks, and interleaved memory

30% instructions need Read or Write

i.e. 130 memory accesses every 100 instructions executed (fetch!)

hit rate in instruction cache = 0.95, and for data = 0.90

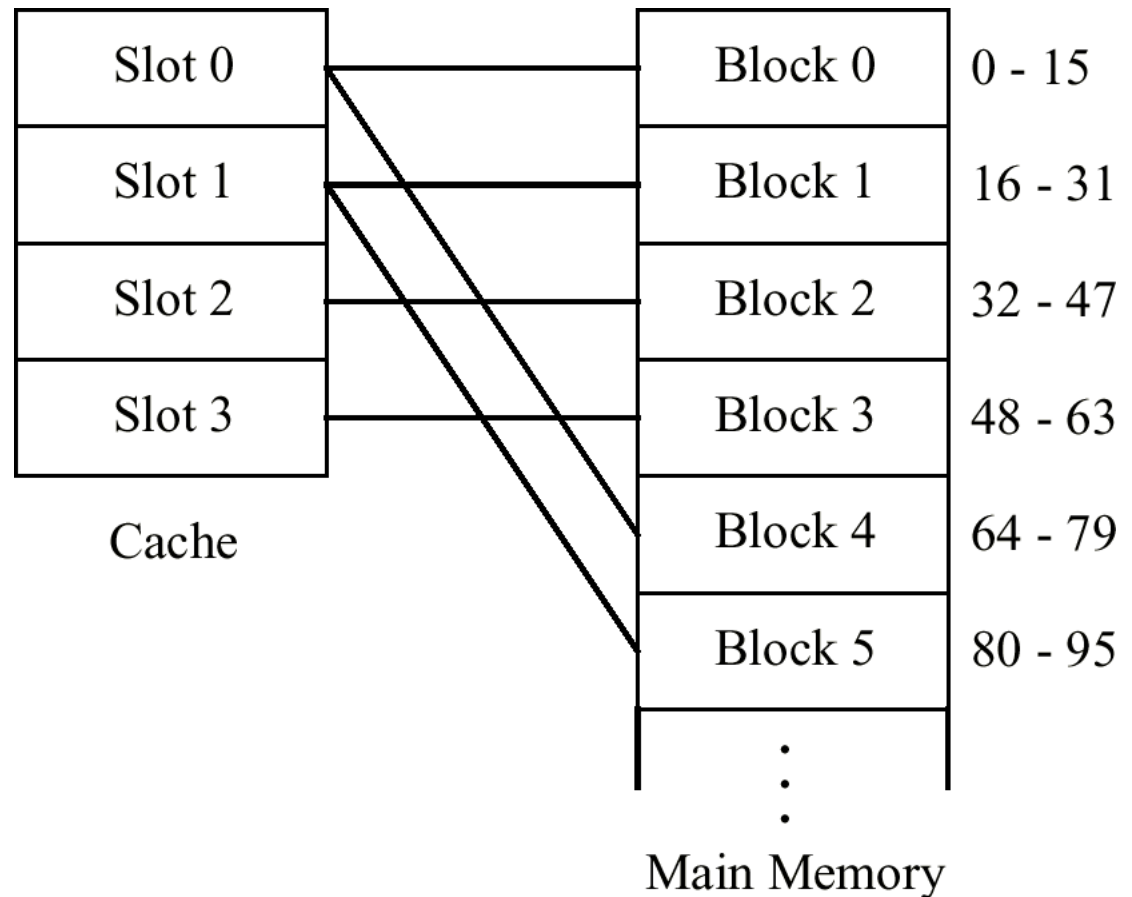
need 17 cycles to load block in cache (see earlier)

$$\frac{\text{no cache}}{\text{with cache}} =$$

$$\frac{130 \times 10}{100 (0.95 \times 1 + 0.05 \times 17) + 30 (0.9 \times 1 + 0.1 \times 17)} = 5.04$$

## Direct Mapped Cache Example

- Compute hit ratio and effective access time for a program that executes from memory locations 48 to 95, and then loops 10 times from 15 to 31.
- The direct mapped cache has four 16-word slots, a hit time of 80 ns, and a miss time of 2500 ns. Load-through is used. The cache is initially empty.



## Table of Events for Example Program

Event	Location	Time	Comment
1 miss	48	2500ns	Memory block 3 to cache slot 3
15 hits	49-63	$80\text{ns} \times 15 = 1200\text{ns}$	
1 miss	64	2500ns	Memory block 4 to cache slot 0
15 hits	65-79	$80\text{ns} \times 15 = 1200\text{ns}$	
1 miss	80	2500ns	Memory block 5 to cache slot 1
15 hits	81-95	$80\text{ns} \times 15 = 1200\text{ns}$	
1 miss	15	2500ns	Memory block 0 to cache slot 0
1 miss	16	2500ns	Memory block 1 to cache slot 1
15 hits	17-31	$80\text{ns} \times 15 = 1200\text{ns}$	
9 hits	15	$80\text{ns} \times 9 = 720\text{ns}$	Last nine iterations of loop
144 hits	16-31	$80\text{ns} \times 144 = 11,520\text{ns}$	Last nine iterations of loop
Total hits = 213    Total misses = 5			

## Calculation of Hit Ratio and Effective Access Time for Example Program

$$\text{Hit ratio} = \frac{213}{218} = 97.7\%$$

$$\text{EffectiveAccessTime} = \frac{(213)(80ns) + (5)(2500ns)}{218} = 136ns$$

## Can the hit rate be improved to almost optimal?

- Larger cache → increased costs
  - Increase block size, same total cache size → ok only up to certain relative ratio between sizes
  - BEST: block sizes not too small, not too large  
usually 16 to 128 bytes
- 
- Speed is improved if cache is on processor chip  
→ however space on processor is at premium
  - Split caches → instruction and data
  - Multi-level → L1, L2 or even L3
  - Prefetching
  - Lockup free cache



## Multi-level Cache: Example

As an example, consider a two-level cache in which the L1 hit time is 5 ns, the L2 hit time is 20 ns, and the L2 miss time is 100 ns. There are 10,000 memory references of which 10 cause L2 misses and 90 cause L1 misses. Compute the hit ratios of the L1 and L2 caches and the overall effective access time.

$H_1$  is the ratio of the number of times the accessed word is in the L1 cache to the total number of memory accesses. There are a total of 85 (L1) and 15 (L2) misses, and so:

$$H_1 = \frac{(10,000 - 10 - 90) \text{ hits}}{10,000 \text{ accesses}} = 99\%$$

(Continued on next slide.)

## Multi-level Cache: Example (cont.)

$H_2$  is the ratio of the number of times the accessed word is in the L2 cache to the number of times the L2 cache is accessed, and so:

$$H_2 = \frac{(100 - 10) \text{ hits}}{100 \text{ accesses}} = 90\%$$

The effective access time is then:

$$\begin{aligned} T_{EFF} &= (10,000 - 10 - 90 \text{ L1 hits})(5 \text{ ns per L1 hit}) + \\ &\quad (100 - 10 \text{ L2 hits})(20 \text{ ns per L2 hit}) + \\ &\quad (10 \text{ L2 cache misses})(100 \text{ ns per L2 cache miss}) \\ &\quad / 10,000 \text{ accesses} \\ &= 5.23 \text{ ns per access} \end{aligned}$$

## Designing memory to support cache more efficiently

- ❑ DRAM for main memory designed for density rather than access time
- ❑ Increase bandwidth from memory to cache to reduce miss penalty?
- ❑ Processor connected to memory over bus → the bus is at least 10 times slower than CPU clock cycle

## Consider memory access time (example):

- ✓ 1 memory bus clock cycle to send the address
- ✓ 15 memory bus clock cycles for each DRAM access
- ✓ 1 memory bus clock cycle to send a word of data

If cache block = 4 words, and DRAM is 1 word wide, then miss penalty:

$$1 + (4 \times 15) + (4 \times 1) = 65 \text{ memory bus clock cycles}$$

Bytes transferred per bus clock cycle for a miss is:

$$(4 \times 4) / 65 = 0.25 \rightarrow \text{bandwidth}$$