

23. Virtual Memory

**Department of Computer Science
University of Victoria**

MH: 7.6

HVZ: 6


Stallings: 8.3, 8.5

(Some of the textbook subsections are more OS-oriented, but it will be helpful to you to read through them)

PROGRAM RELOCATABILITY

In multiple user systems, each user program 'appears' to use a contiguous block of memory beginning at location 0

Hardware maps user address space to the physical address space using registers that can only be set by the operating system



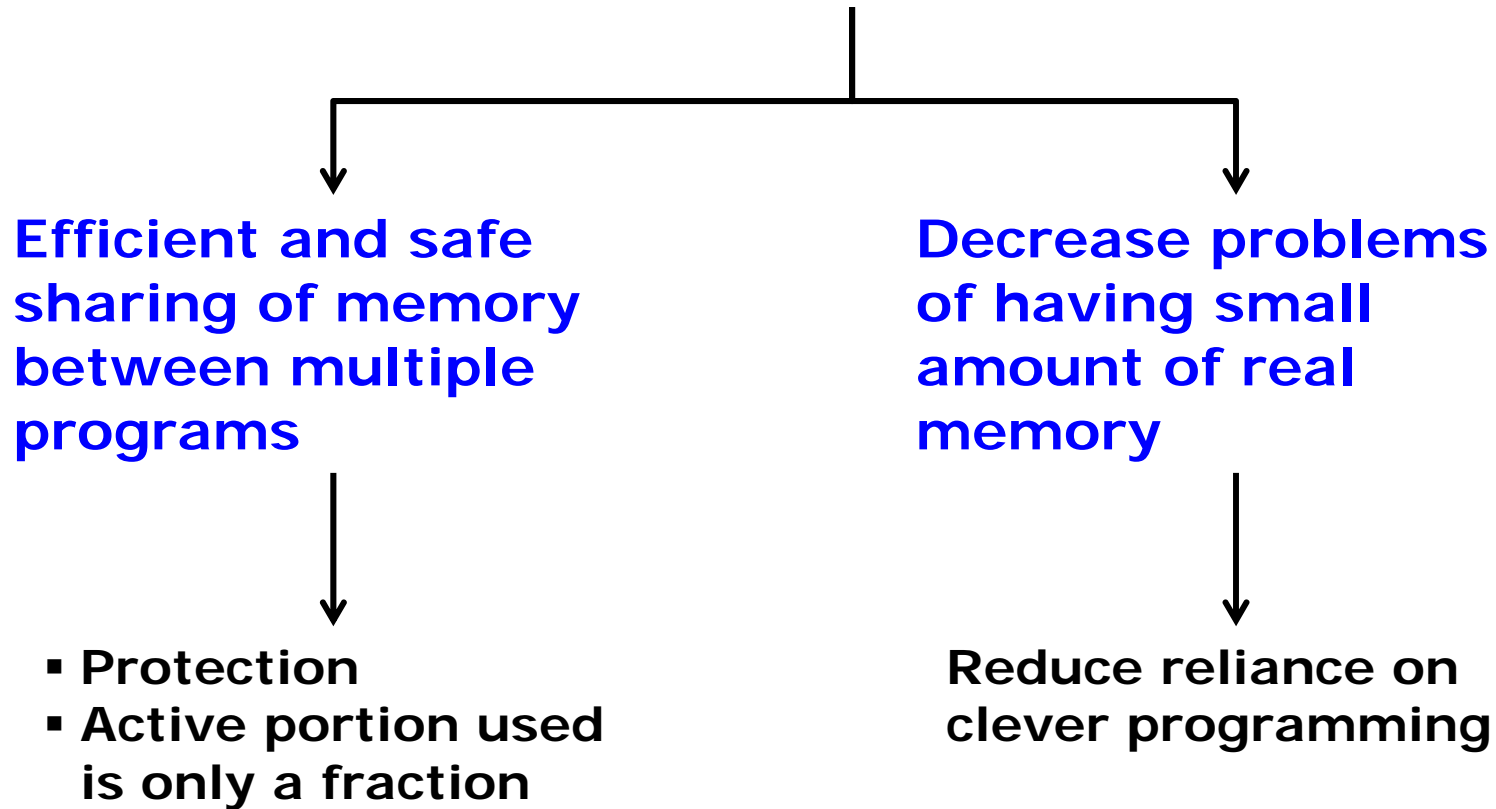
➔ This scheme provides easy program relocation

What is Virtual Memory?

Virtual memory is a memory management technique which virtualizes all storage system and makes them appear as one large memory device

The side effect is that the main (physical) memory can be seen to act as a cache for the secondary storage

Why Virtual Memory?



What does Virtual Memory do?

- Give an **illusion** of an essentially **unbounded** amount of memory
- Allow efficient and safe **sharing** of memory among multiple programs
- Remove the programming burdens of a small, limited amount of main memory
- Provide **relocation**, which simplifies loading the program for execution (programs can be loaded into any location)


How does Virtual Memory do it?

1. Each program is compiled to its address space
2. Virtual memory implements the translation of a program's address space to physical addresses

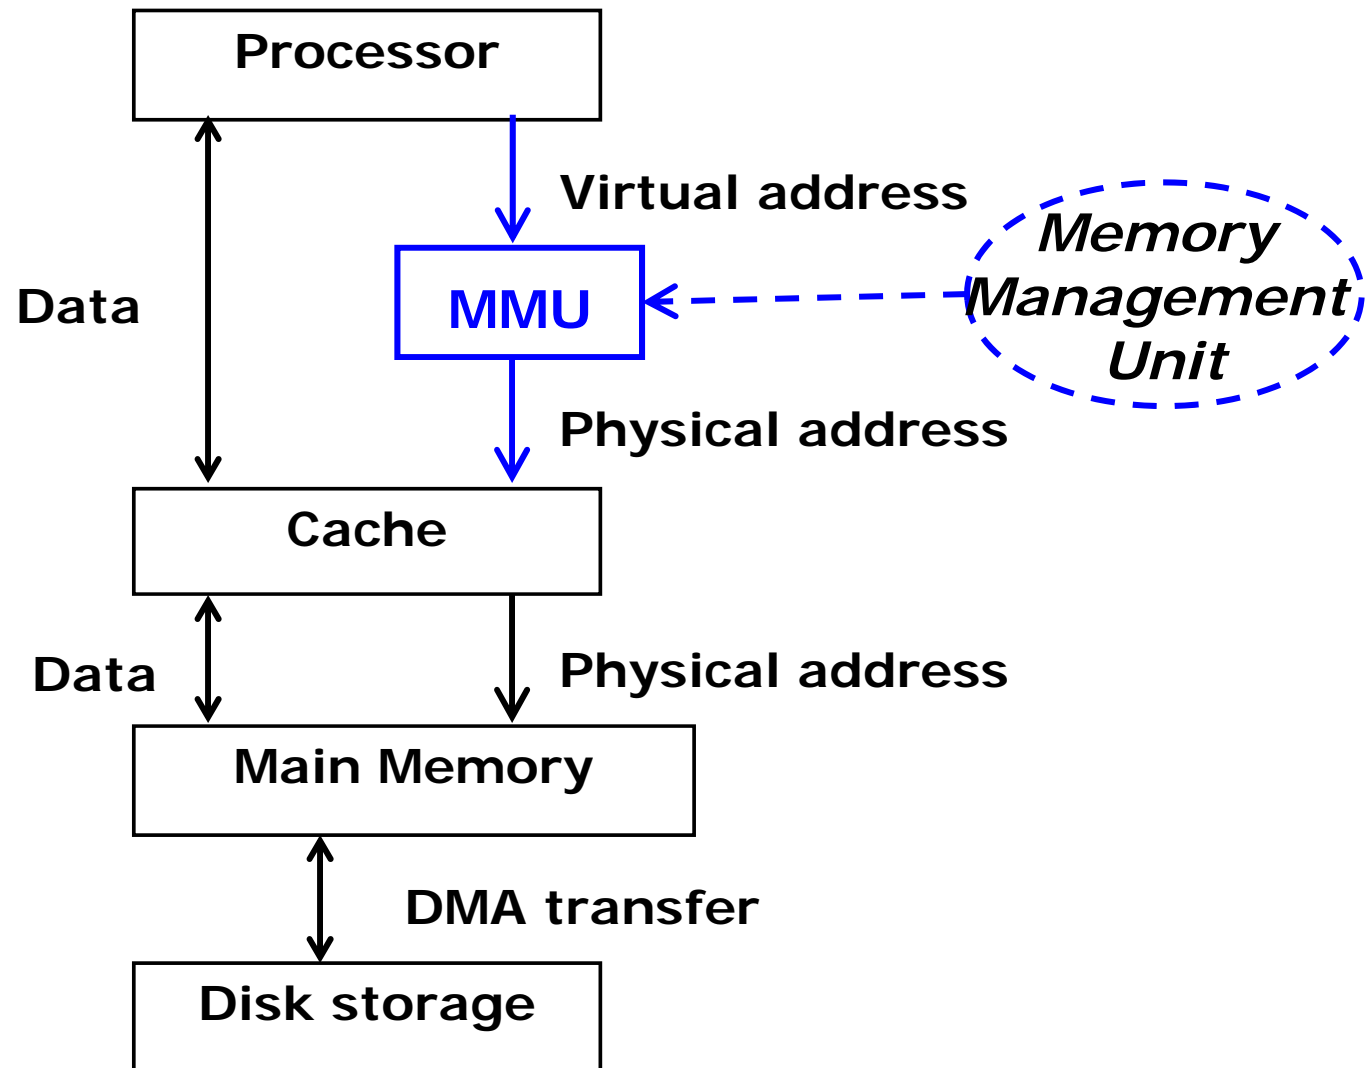
Memory mapping/address translation:



Overview



Virtual memory organization



How can programs make effective and efficient use of such large address space?

- ❑ Not long ago, main memory (RAM) was measured in MB or 1-2 Gb
- ❑ But theoretical address space is *much* larger
 - ➔ 32-bit addresses can represent 4GB of byte-addressable space
 - ➔ 33-bit addresses can represent 8GB of byte-addressable space

Observation:

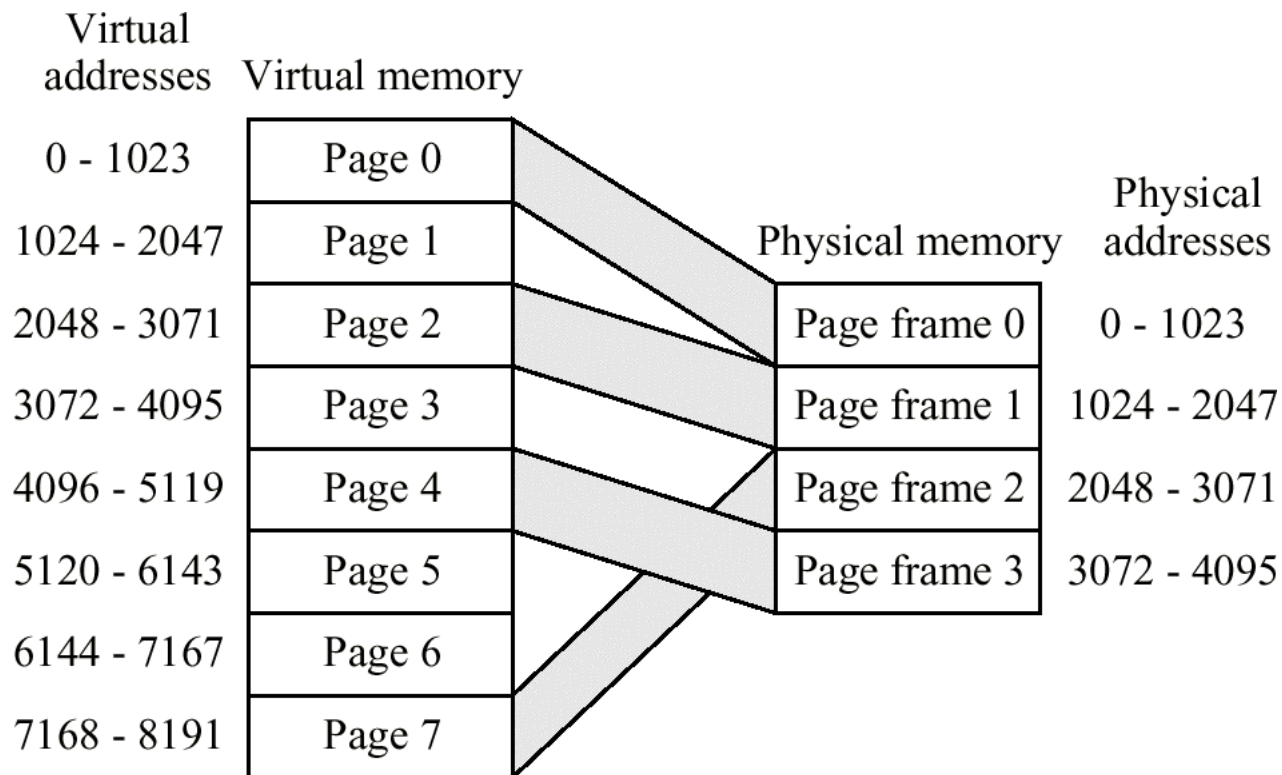
Few programs need all their code and data at once. In fact, most code and data are accessed rarely: the 90/10 or 80/20 “rules”.

➔ Solution:

- Only use RAM for “active” code and data.
- Store the rest on disk.

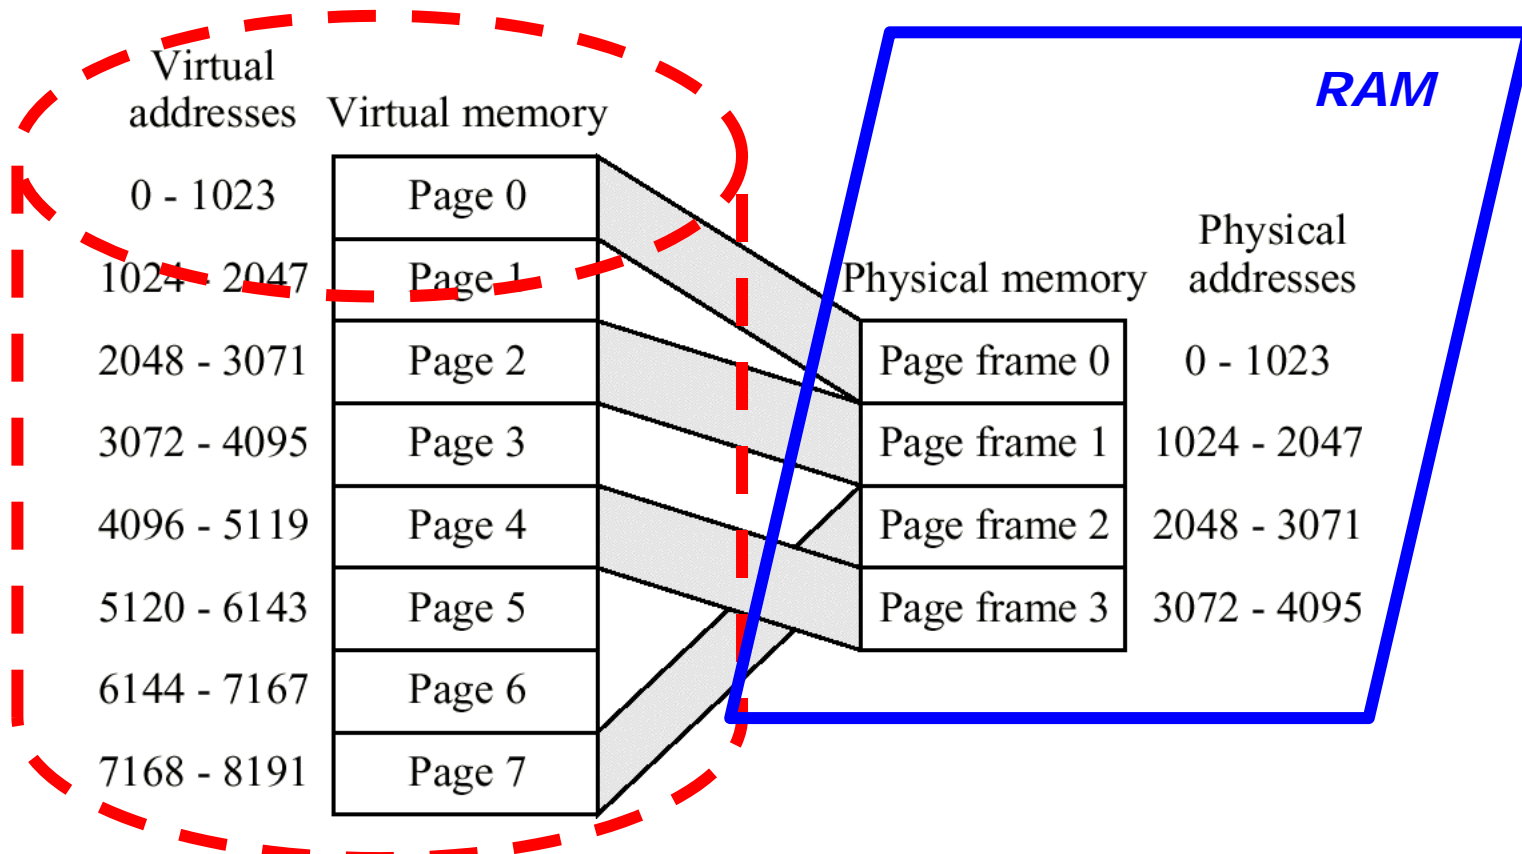
Virtual Memory: swap space

- ❑ Virtual memory is stored in a specialized hard disk image (often called a “swap” partition, or just “swap” space).
- ❑ The physical memory holds a small number of virtual *pages* in physical page *frames*.
- ❑ A mapping between a virtual and a physical memory:



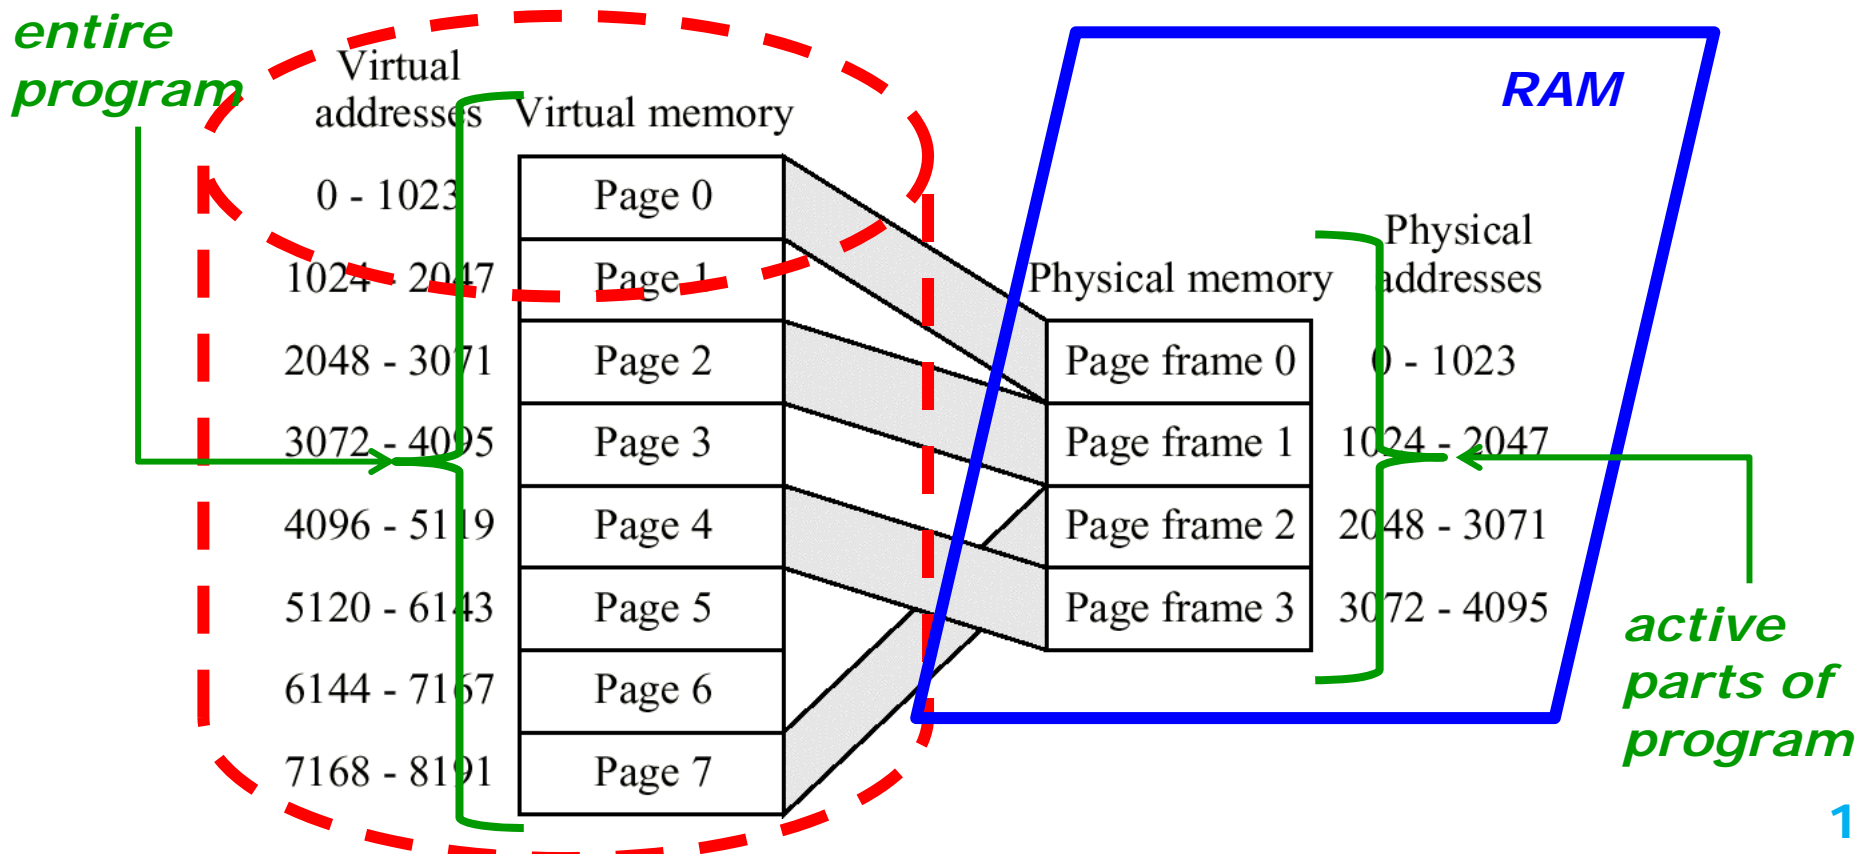
Virtual Memory: swap space

- ❑ Virtual memory is stored in a specialized hard disk image (often called a “swap” partition, or just “swap” space).
- ❑ The physical memory holds a small number of virtual *pages* in physical *page frames*.
- ❑ A mapping between a virtual and a physical memory:



Virtual Memory: swap space

- ❑ Virtual memory is stored in a specialized hard disk image (often called a “swap” partition, or just “swap” space).
- ❑ The physical memory holds a small number of virtual *pages* in physical page *frames*.
- ❑ A mapping between a virtual and a physical memory:



Questions on Virtual Memory

- **How much total virtual memory should be represented?**
 - How much disk space is available?
 - What is the ratio of allocated vs. active memory used by typical programs?
- **What granularity of pages to use?**
 - Smaller pages = more precision, less waste (later)
 - Smaller pages = more bookkeeping and swapping
- **How to replace (“swap”) pages?**
 - Similar to caching, need a replacement policy: FIFO, LIFO, LRU, ...?

What is the Page Table?

The page table (kept by the OS) records the mapping of virtual pages to physical frames

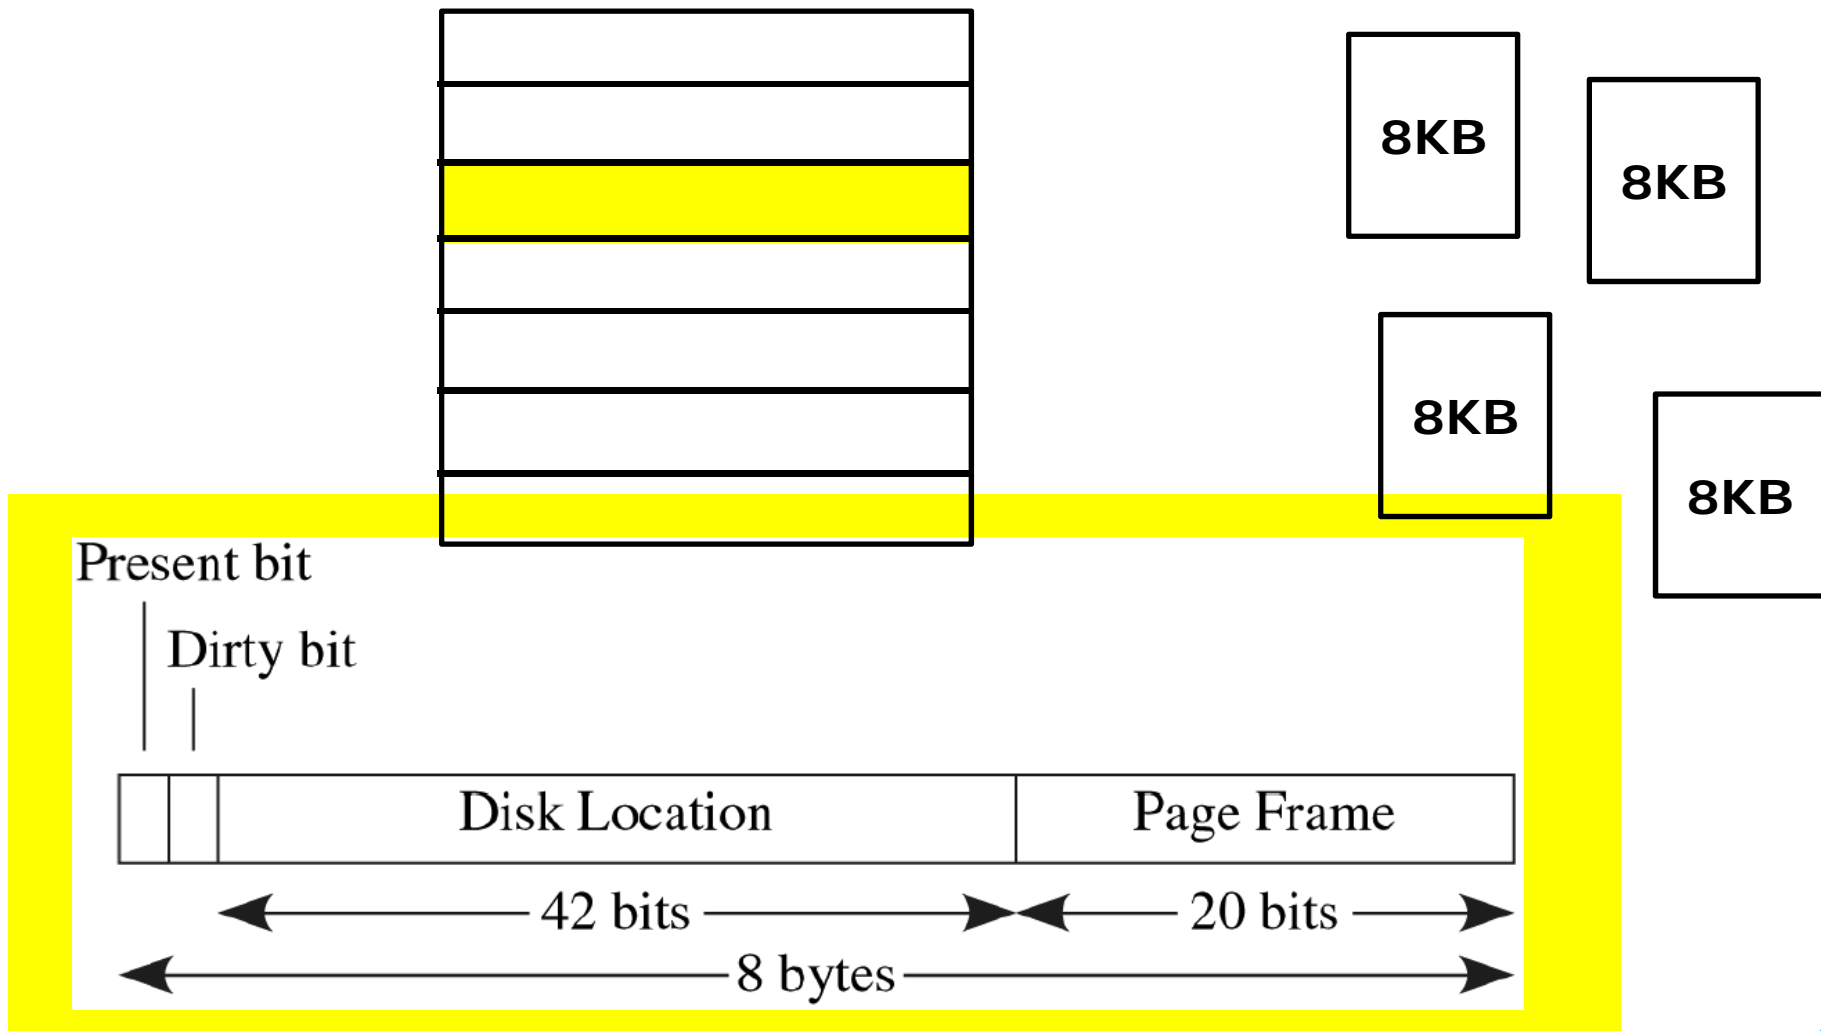
Where is the Page Table?
→ in RAM main memory

Present bit:
0: Page is not in
physical memory
1: Page is in physical
memory

Pag	Present bit	Disk address	Page frame
0	1	01001011100	00
1	0	11101110010	xx
2	1	10110010111	01
3	0	00001001111	xx
4	1	01011100101	11
5	0	10100111001	xx
6	0	00110101100	xx
7	1	01010001011	10

Page Table – Example.1

Consider a paging scheme that uses pages of size 8KB and a table with 64 bit entries




Page Table – Example.2

Consider a paging scheme that uses pages of size 8KB and a table with 64 bit entries

1. How much RAM can we have?

- 20 bits for the frame number (see below) imply
 - ➔ the maximum number of entries (pages) = 2^{20}
 - ➔ then $2^{20} \times 8\text{KB}$ (size of each page) = 2^{33}B = 8GB of RAM

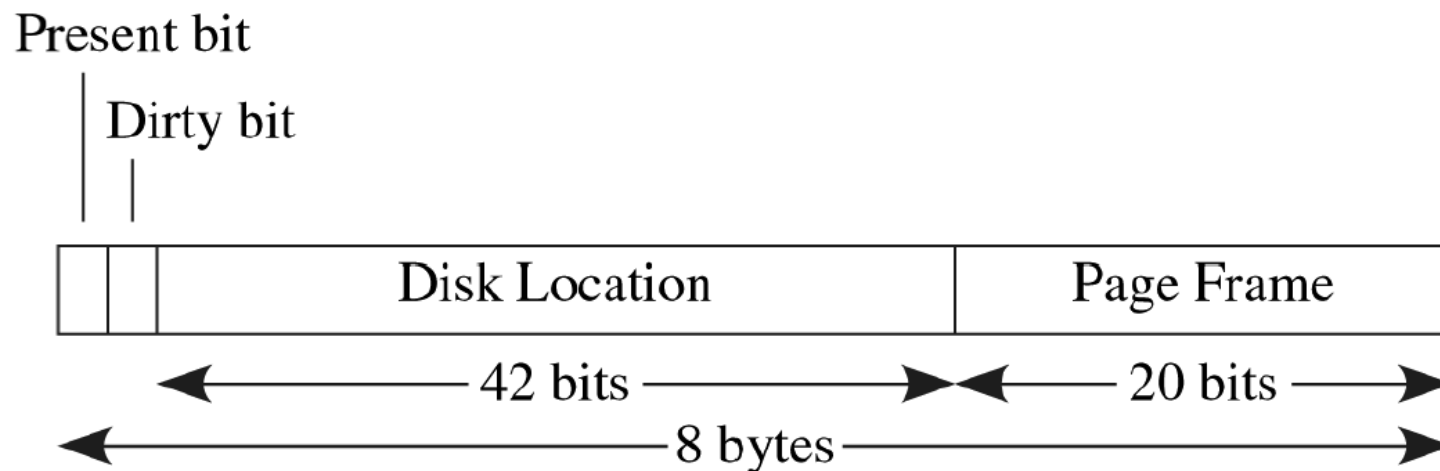


Page Table – Example.3

Consider a paging scheme that uses pages of size 8KB and a table with 64 bit entries

2. How much (maximum) virtual memory?

- 42 bits for page number on disk imply
 ➔ $2^{42} \times 8\text{KB} = 2^{55}\text{B} = 2^{15}\text{TB!}$ (maximum)



Virtual page number



Disk storage

Page Tables – How they work

Page Table is a table to index memory to locate pages

- ❑ It may itself be in memory or in MMU
- ❑ Contains the corresponding physical page number
- ❑ May contain entries for pages not present in memory
- ❑ Indexed with the page number from the virtual address
- ❑ Each program has its own page table


Page Tables – How they work

Hardware in the system includes a register that points to the start of the page table

- ❑ A valid bit is used in each page table entry
 - a) if bit is *off*, the page is not in main memory
→ page fault occurs → must go to disk
 - b) if bit is *on*, the page is valid
→ entry + offset = physical page number in memory
- ❑ No tags are required:
page table contains a mapping for every possible page

Virtual Address from Processor

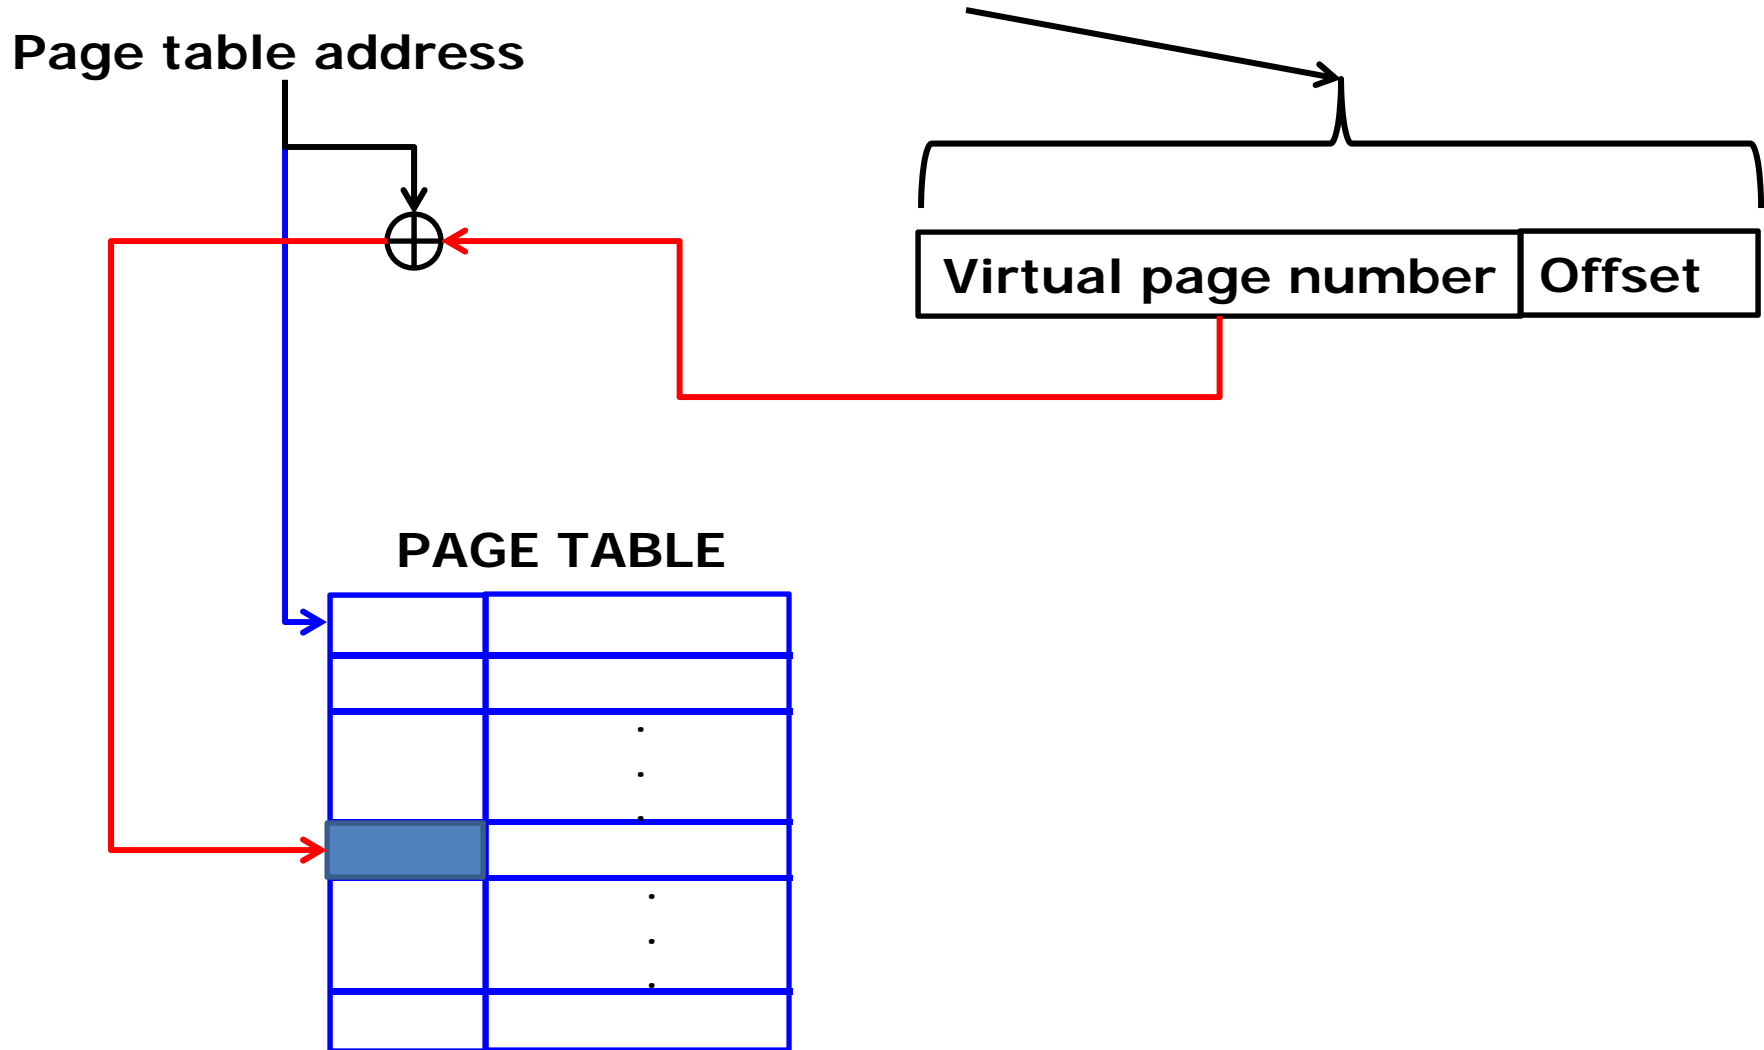
(Register)
Page table address



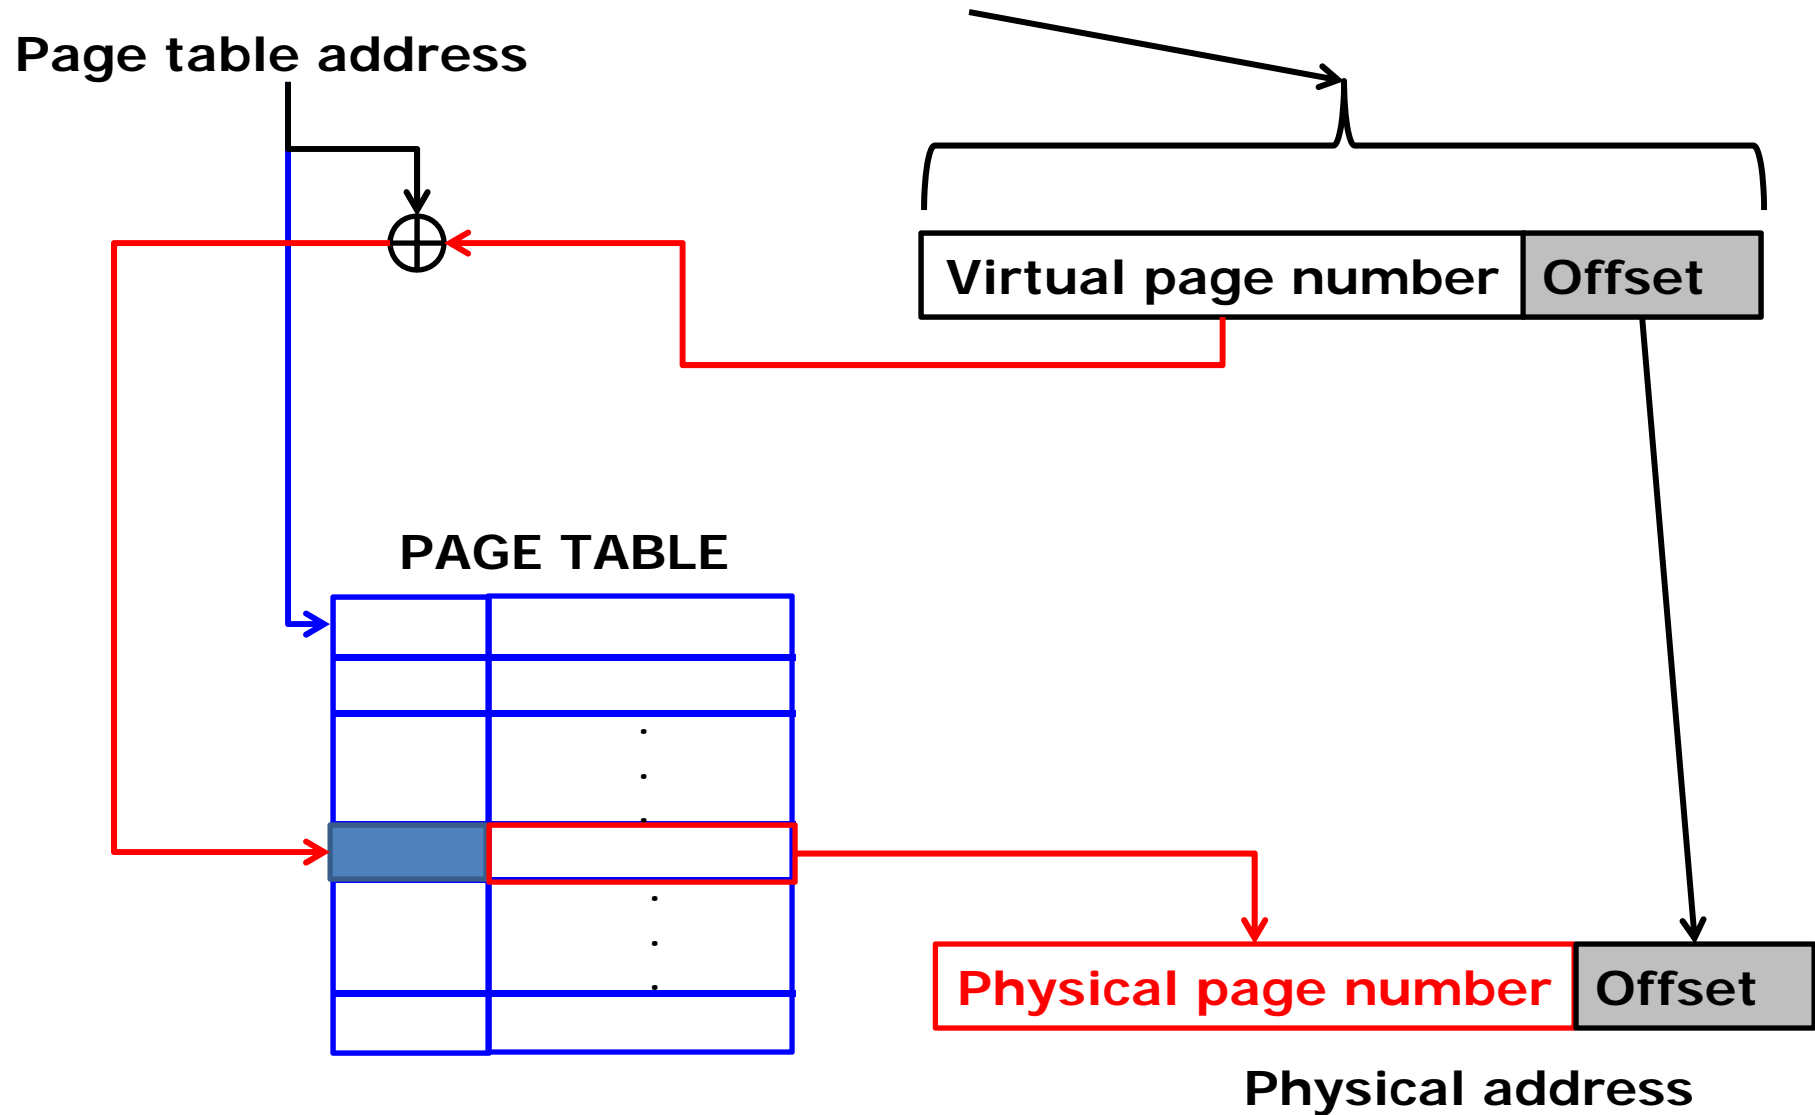
PAGE TABLE

	⋮
	⋮
	⋮
	⋮
	⋮

Virtual Address from Processor



Virtual Address from Processor



Example from Spring 2009 final exam (1)

Consider a computer which has a virtual memory system with the following characteristics. The computer has a 32-bit address space with byte addressable memory; this implies a virtual address space which is 2^{32} bytes or 4 GB in size. The main memory (RAM memory) is 1 GB in size. Pages are 4 KB in size.

How many entries can the page table have?

Each page is 4 KB = $2^2 \times 2^{10} = 2^{12}$

Total address space = 2^{32} bytes

$$\rightarrow 2^{32} / 2^{12} = 2^{20} \text{ pages}$$

→ it will need 20 bits to give a number entry to each page (5 hex digits)

Example from Spring 2009 final exam (2)

Suppose that the CPU generates the sequence of memory addresses shown below (written in hexadecimal).



Example from Spring 2009 final exam (3)

Suppose that the page table contains the entries shown on the left. Invalid (unmapped) entries in the page table are shown as empty.

Page Table

0:	0x00443
1:	0x015C2
2:	-
3:	0x28AE1
4:	0x016DE4
5:	-
6:	0x3EF04

.

Example from Spring 2009 final exam (4)


Show the corresponding sequence of addresses in main memory which the virtual addresses are mapped to.

If any address cannot be mapped, show your answer as the words 'page fault'.

Example from Spring 2009 final exam (5)

Page Table

0:	0x00443
1:	0x015C2
2:	-
3:	0x28AE1
4:	0x016DE4
5:	-
6:	0x3EF04



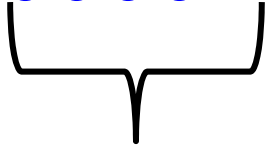
Virtual Addresses
from processor

Example from Spring 2009 final exam (6)

Consider the first address = 0x00001020

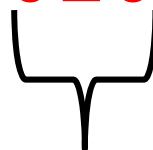
It was previously decided that there are 2^{20} pages and we will need 20 bits to give a number entry to each page

0x00001020



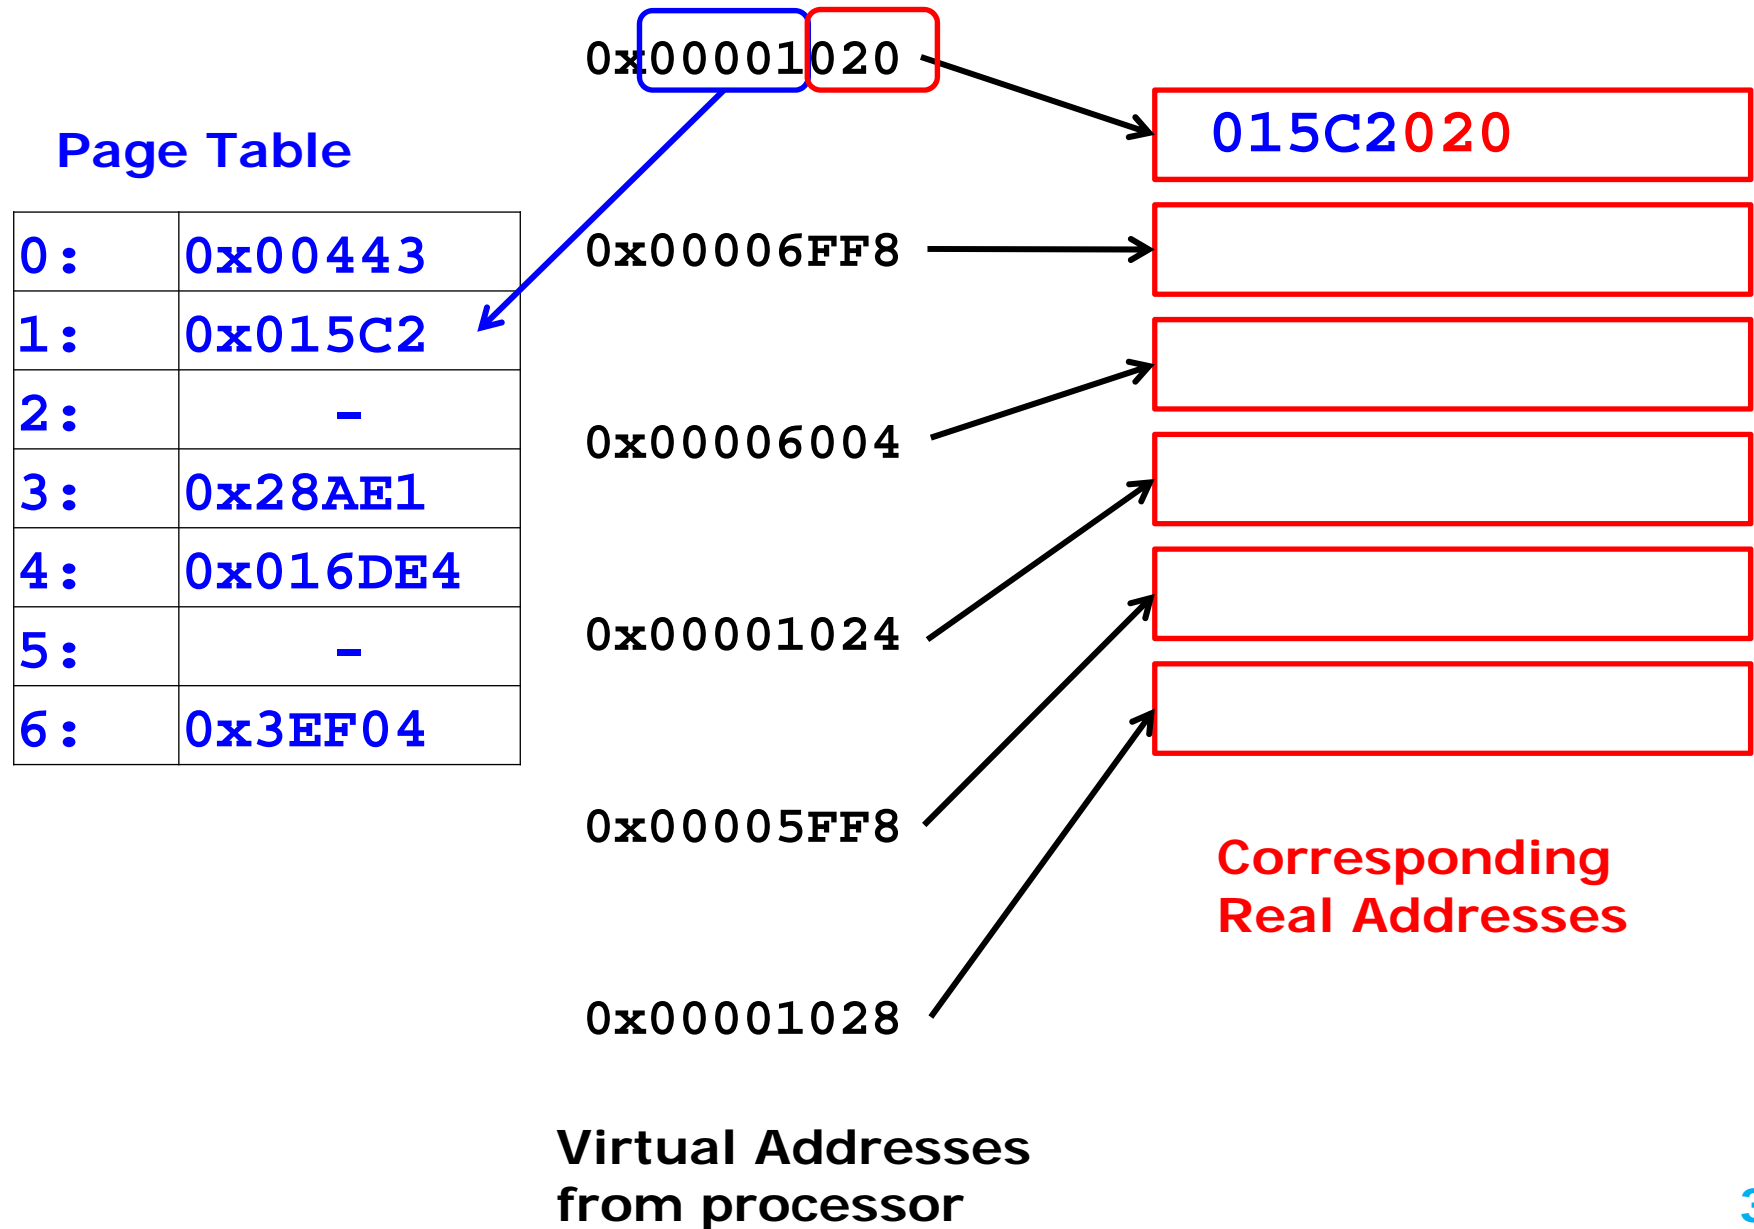
5 Hex digits = 20 bits = 0001 → page entry number

0x00001020

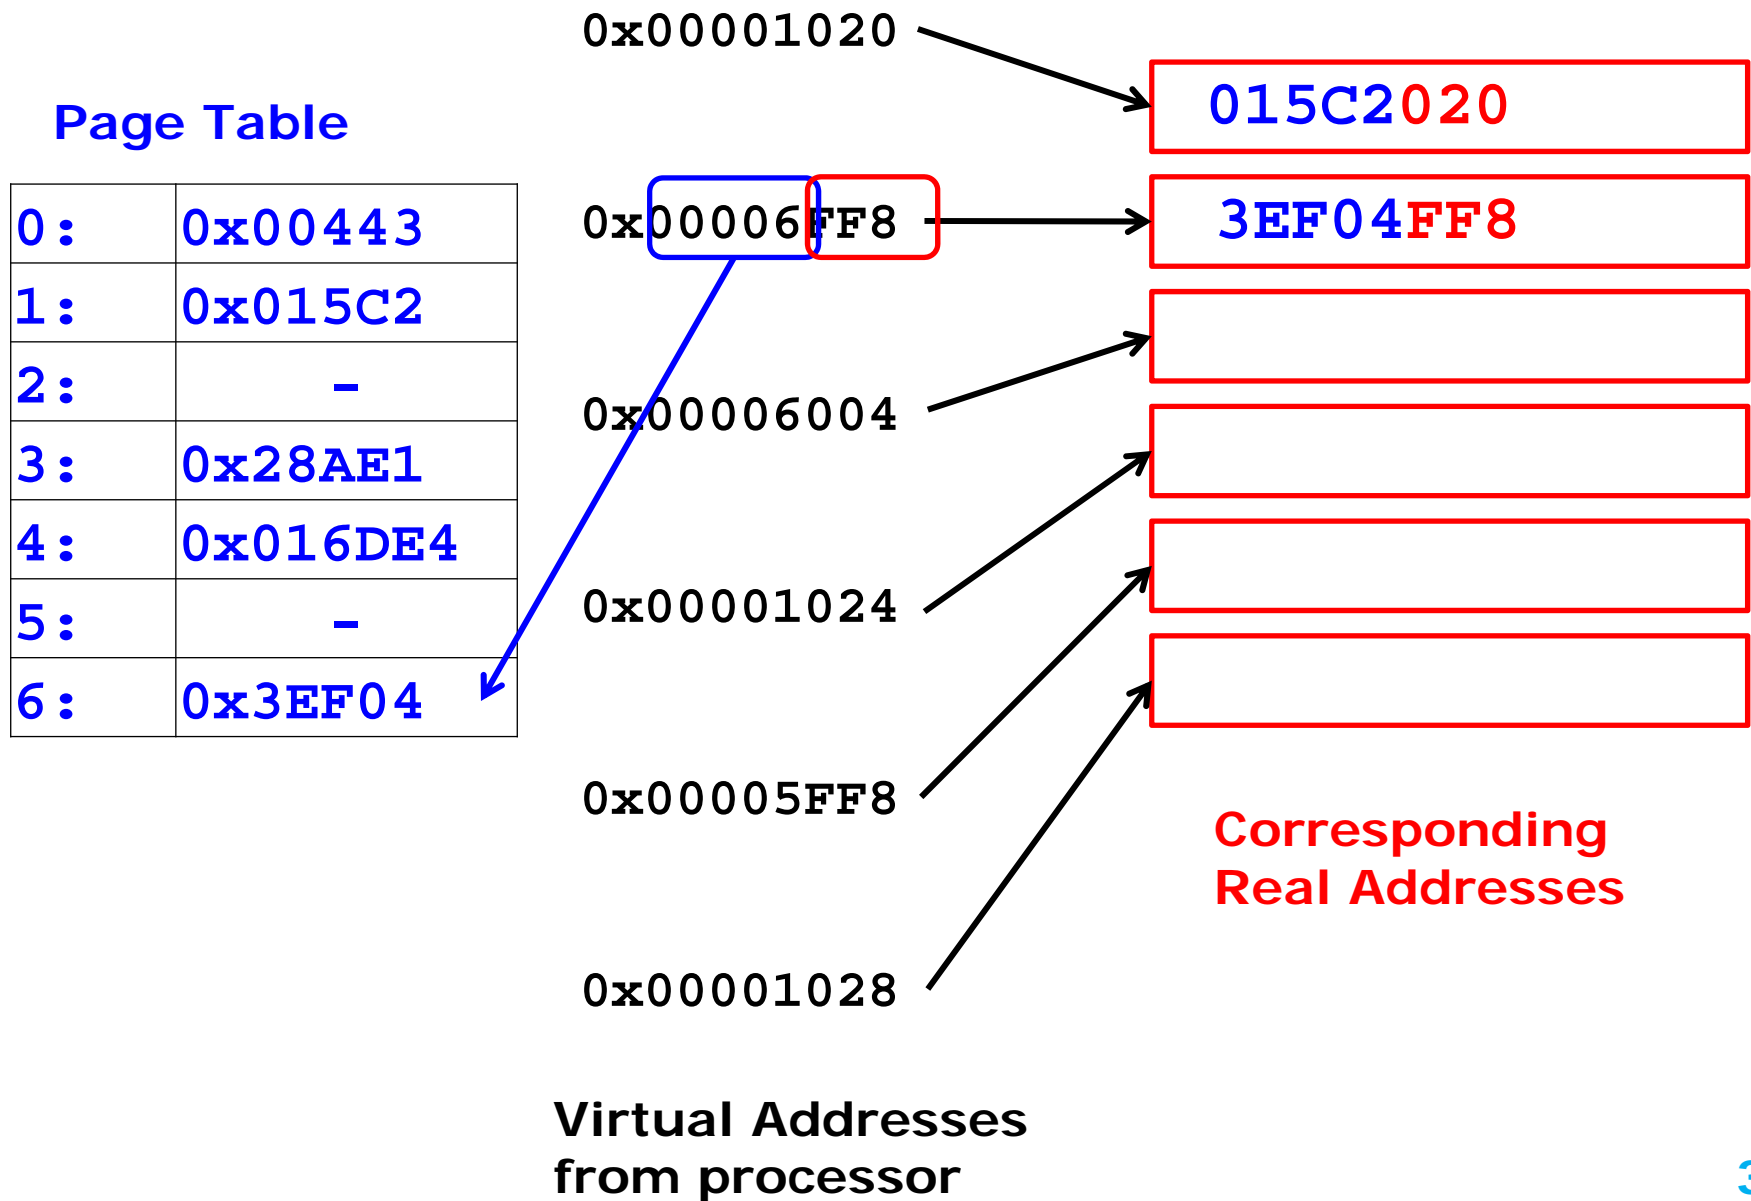


Remaining 3 Hex digits = 12 bits = 020 → offset

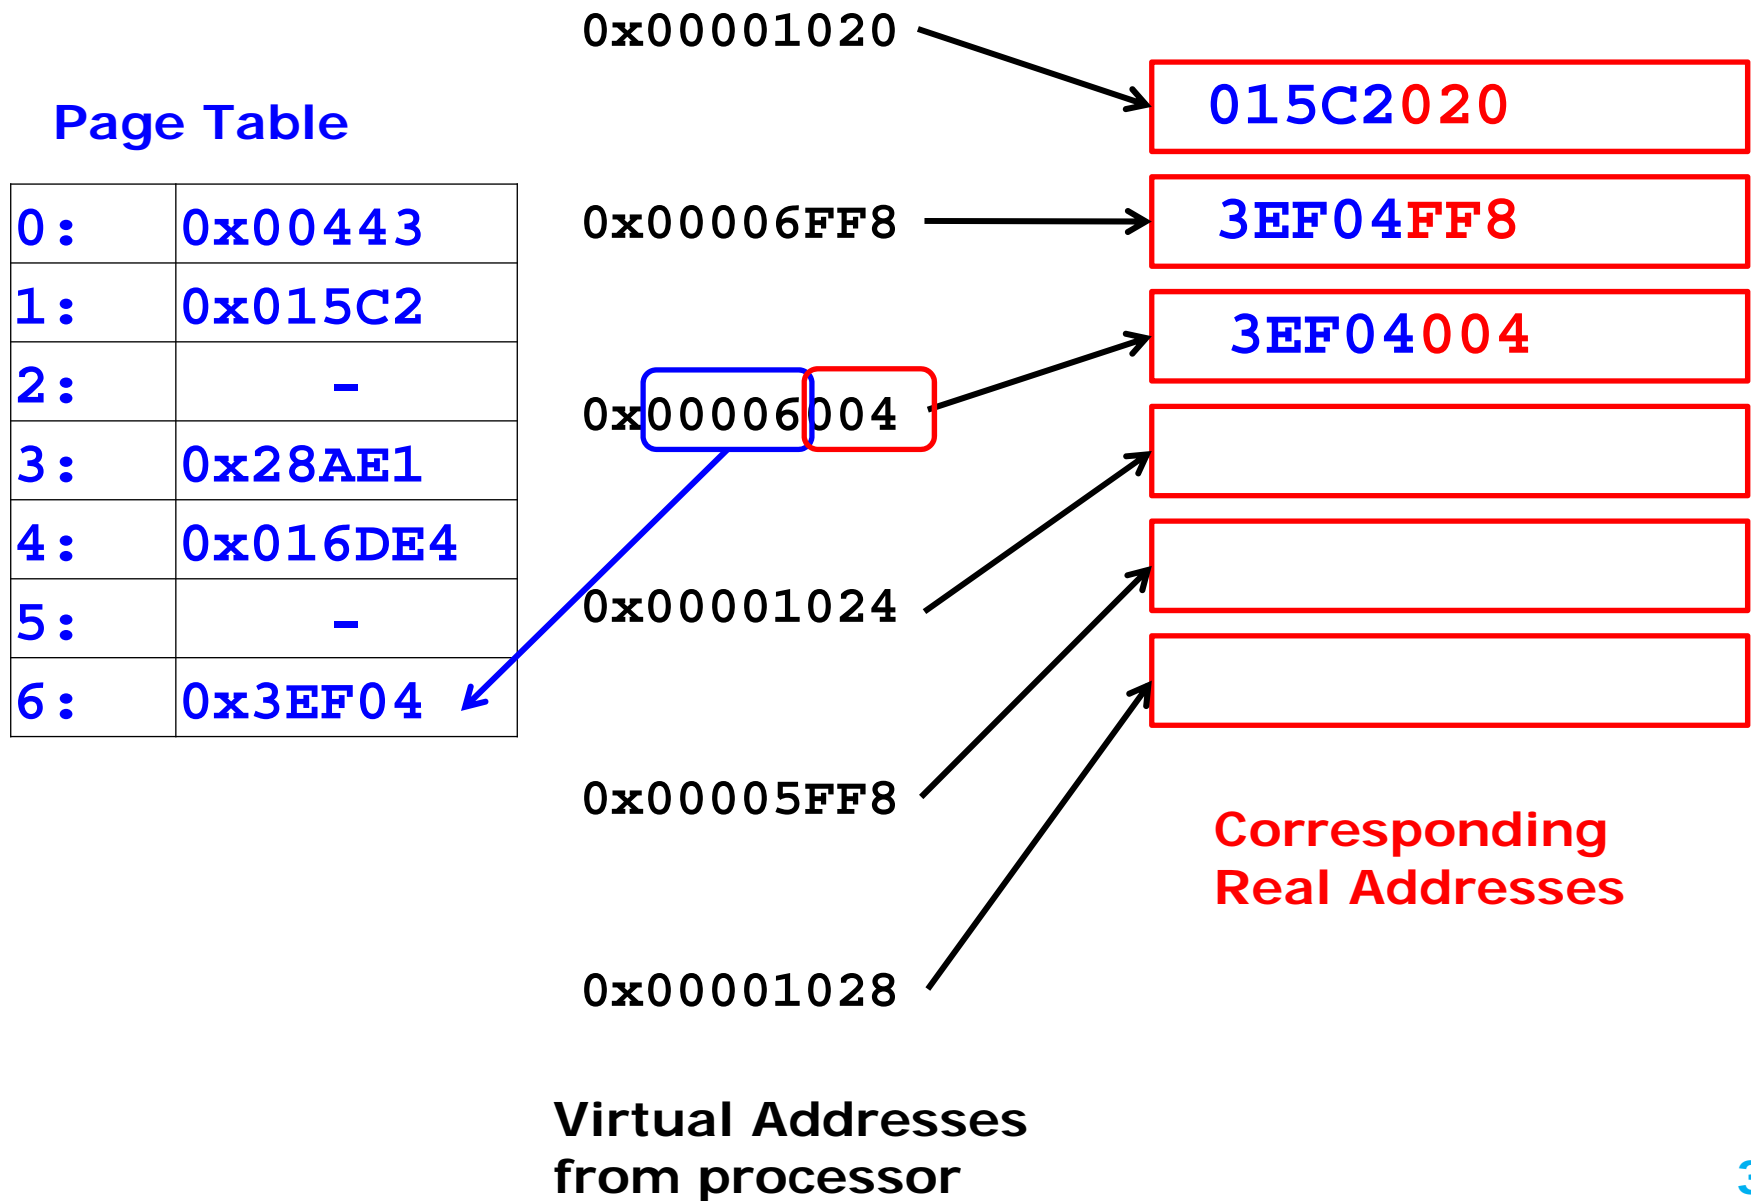
Example from Spring 2009 final exam (7)



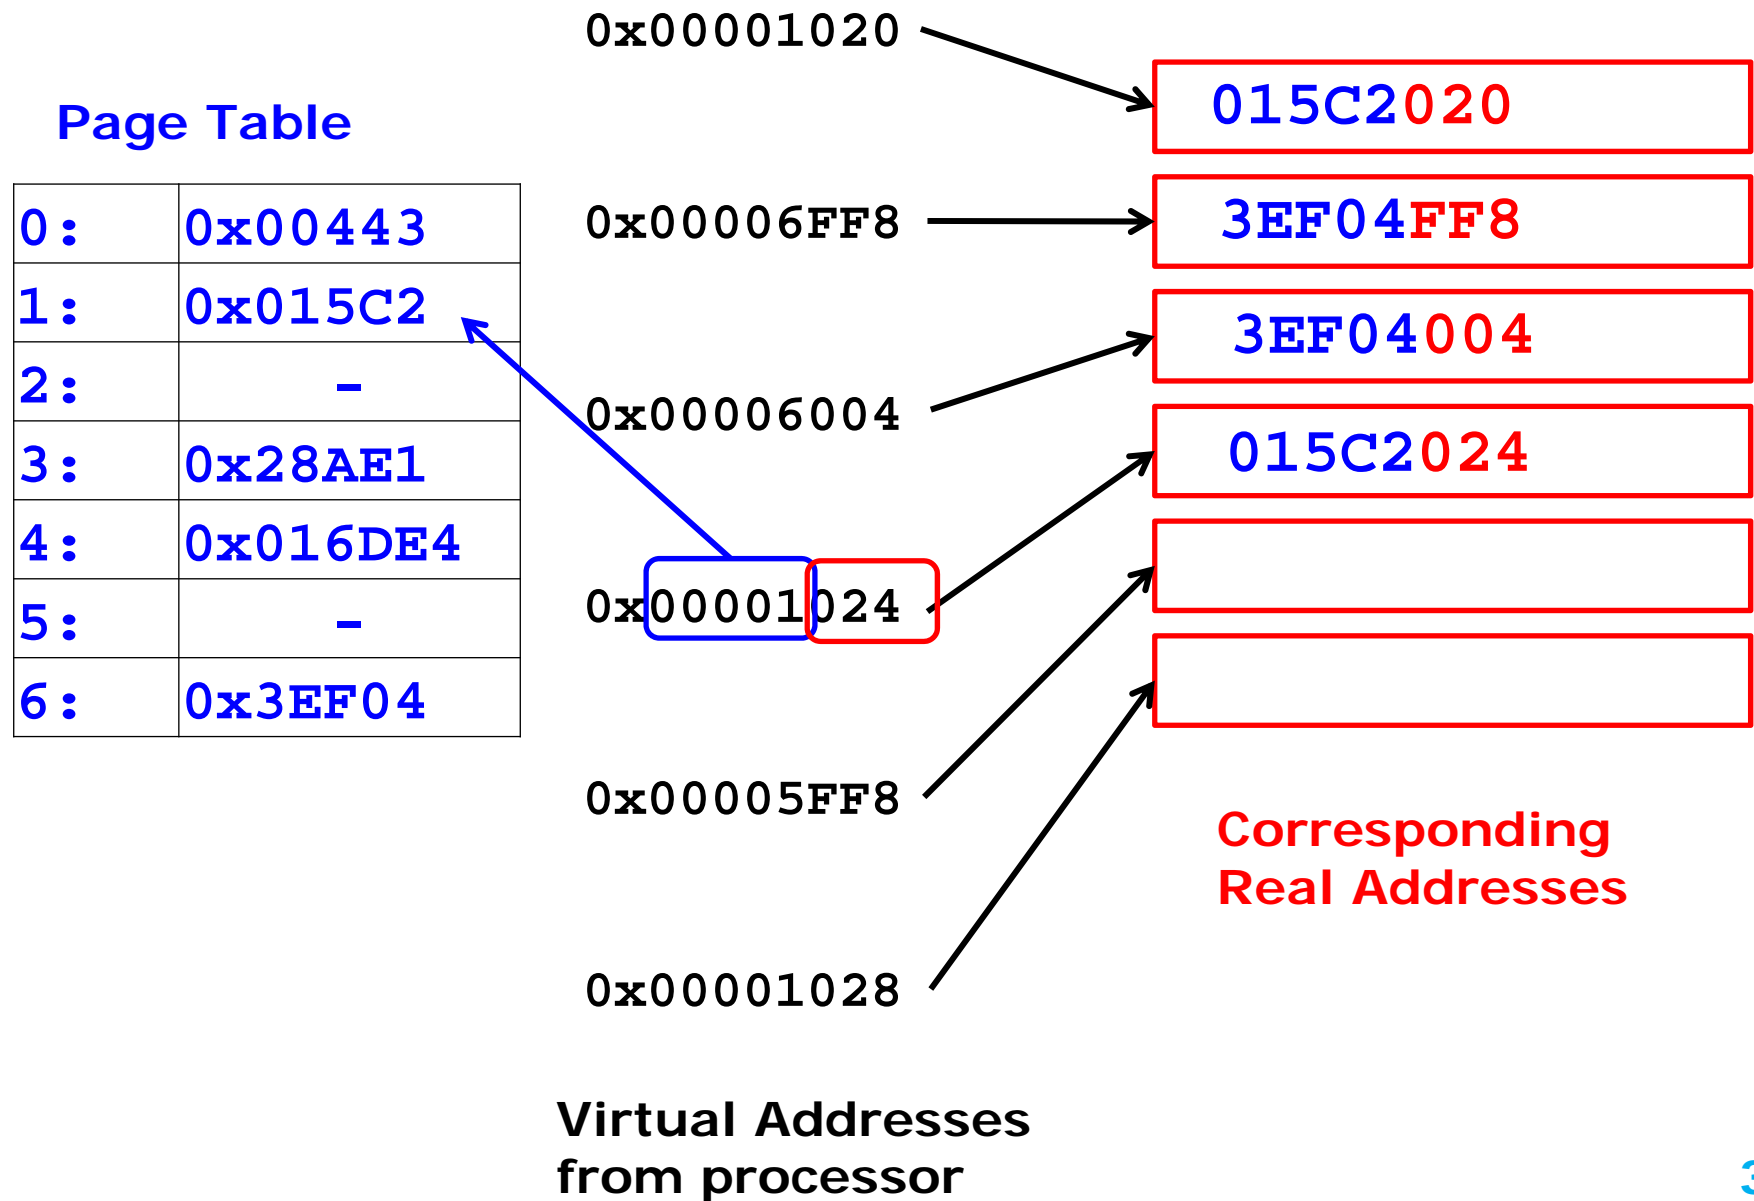
Example from Spring 2009 final exam (8)



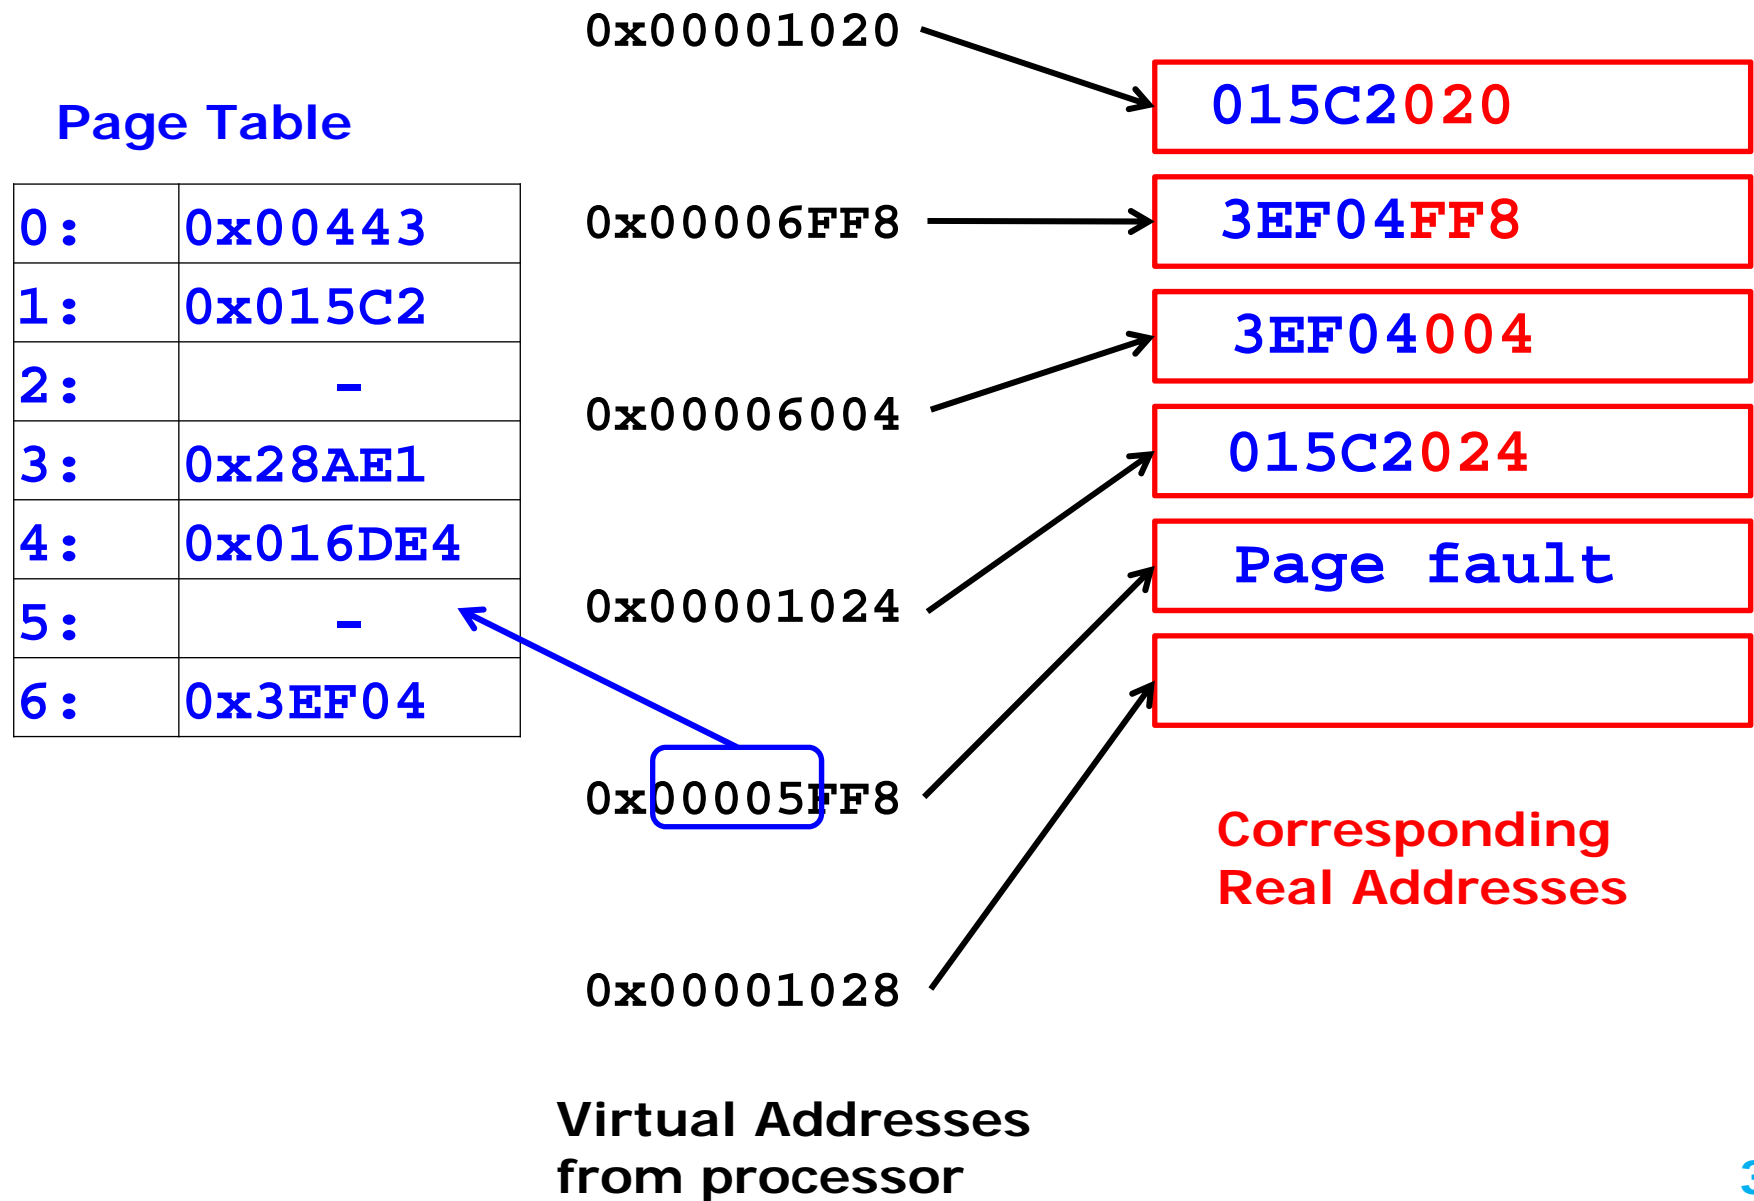
Example from Spring 2009 final exam (9)



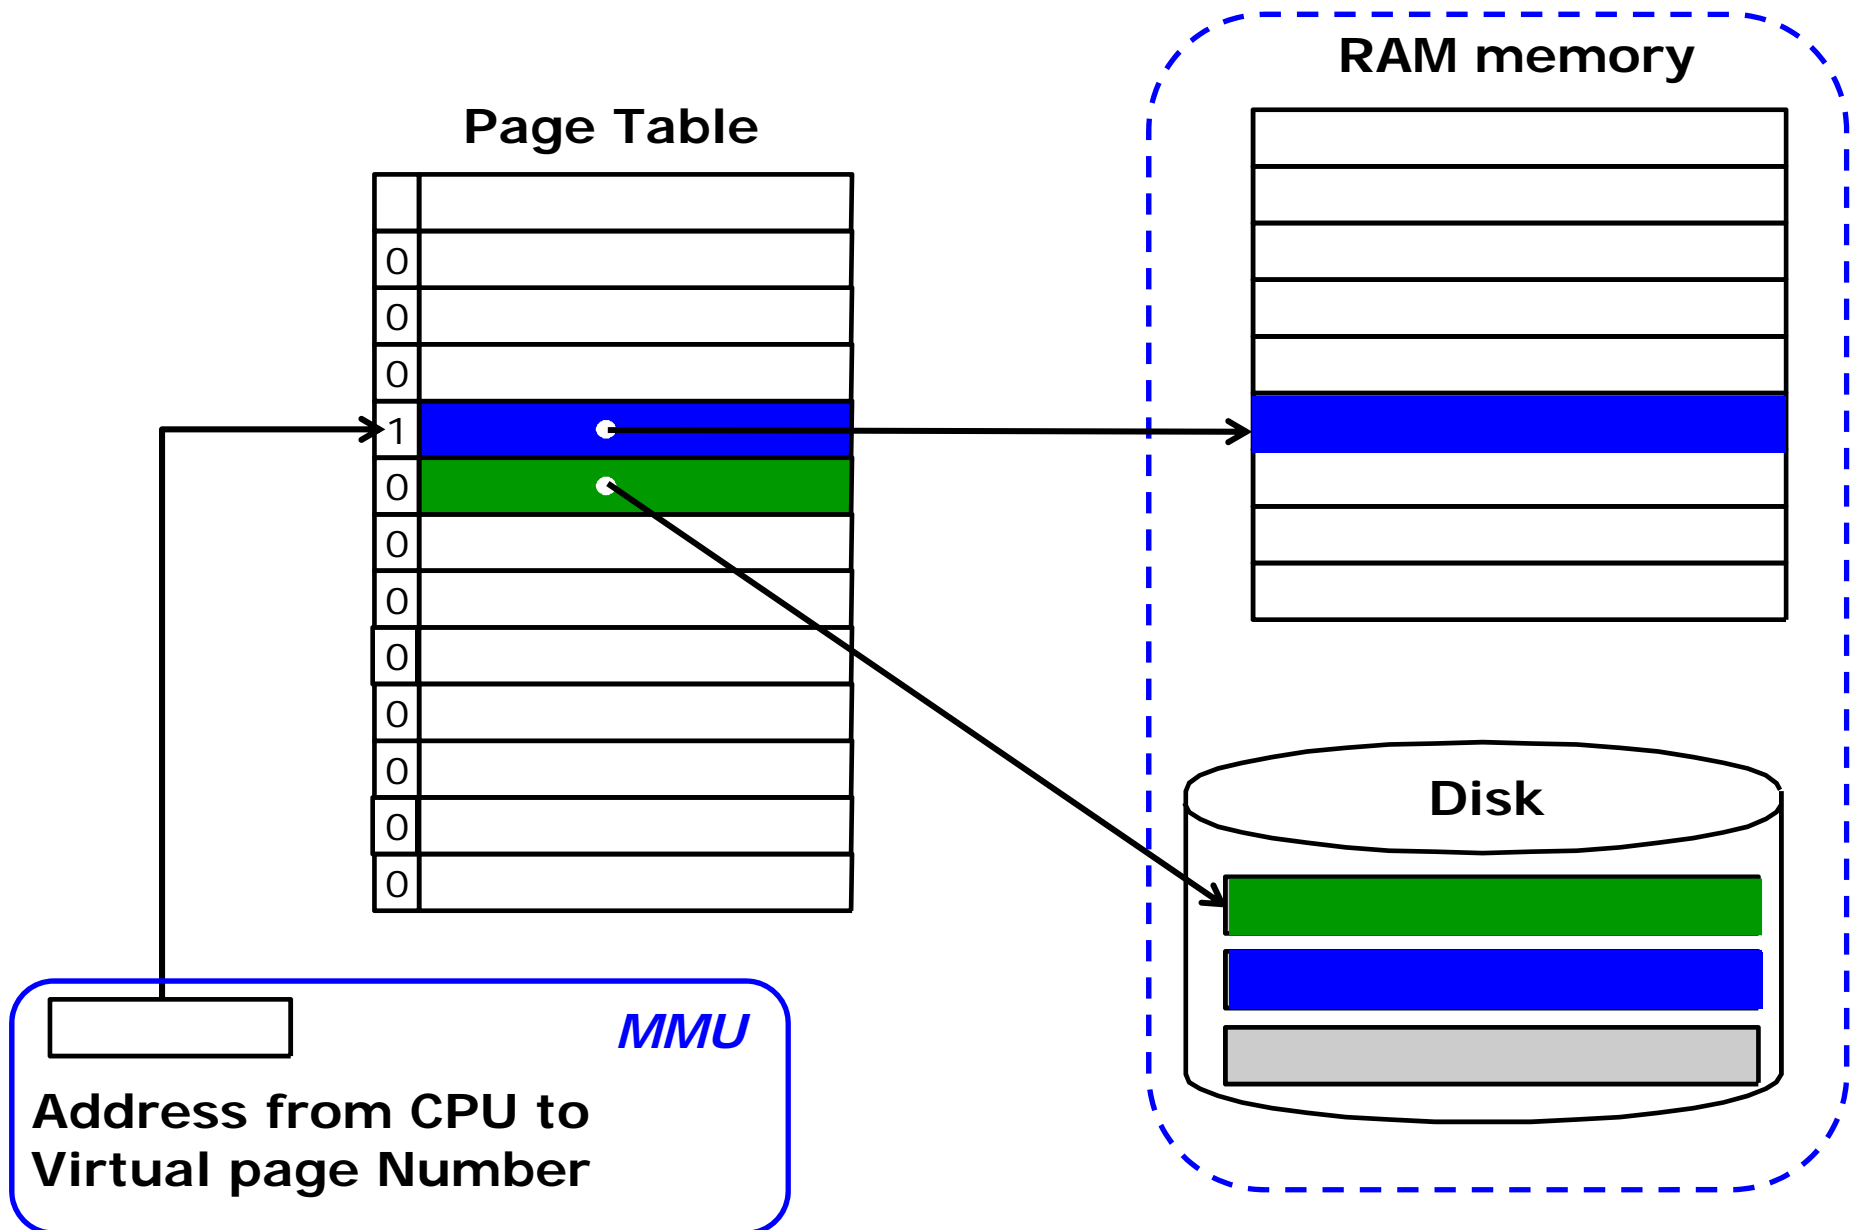
Example from Spring 2009 final exam (10)



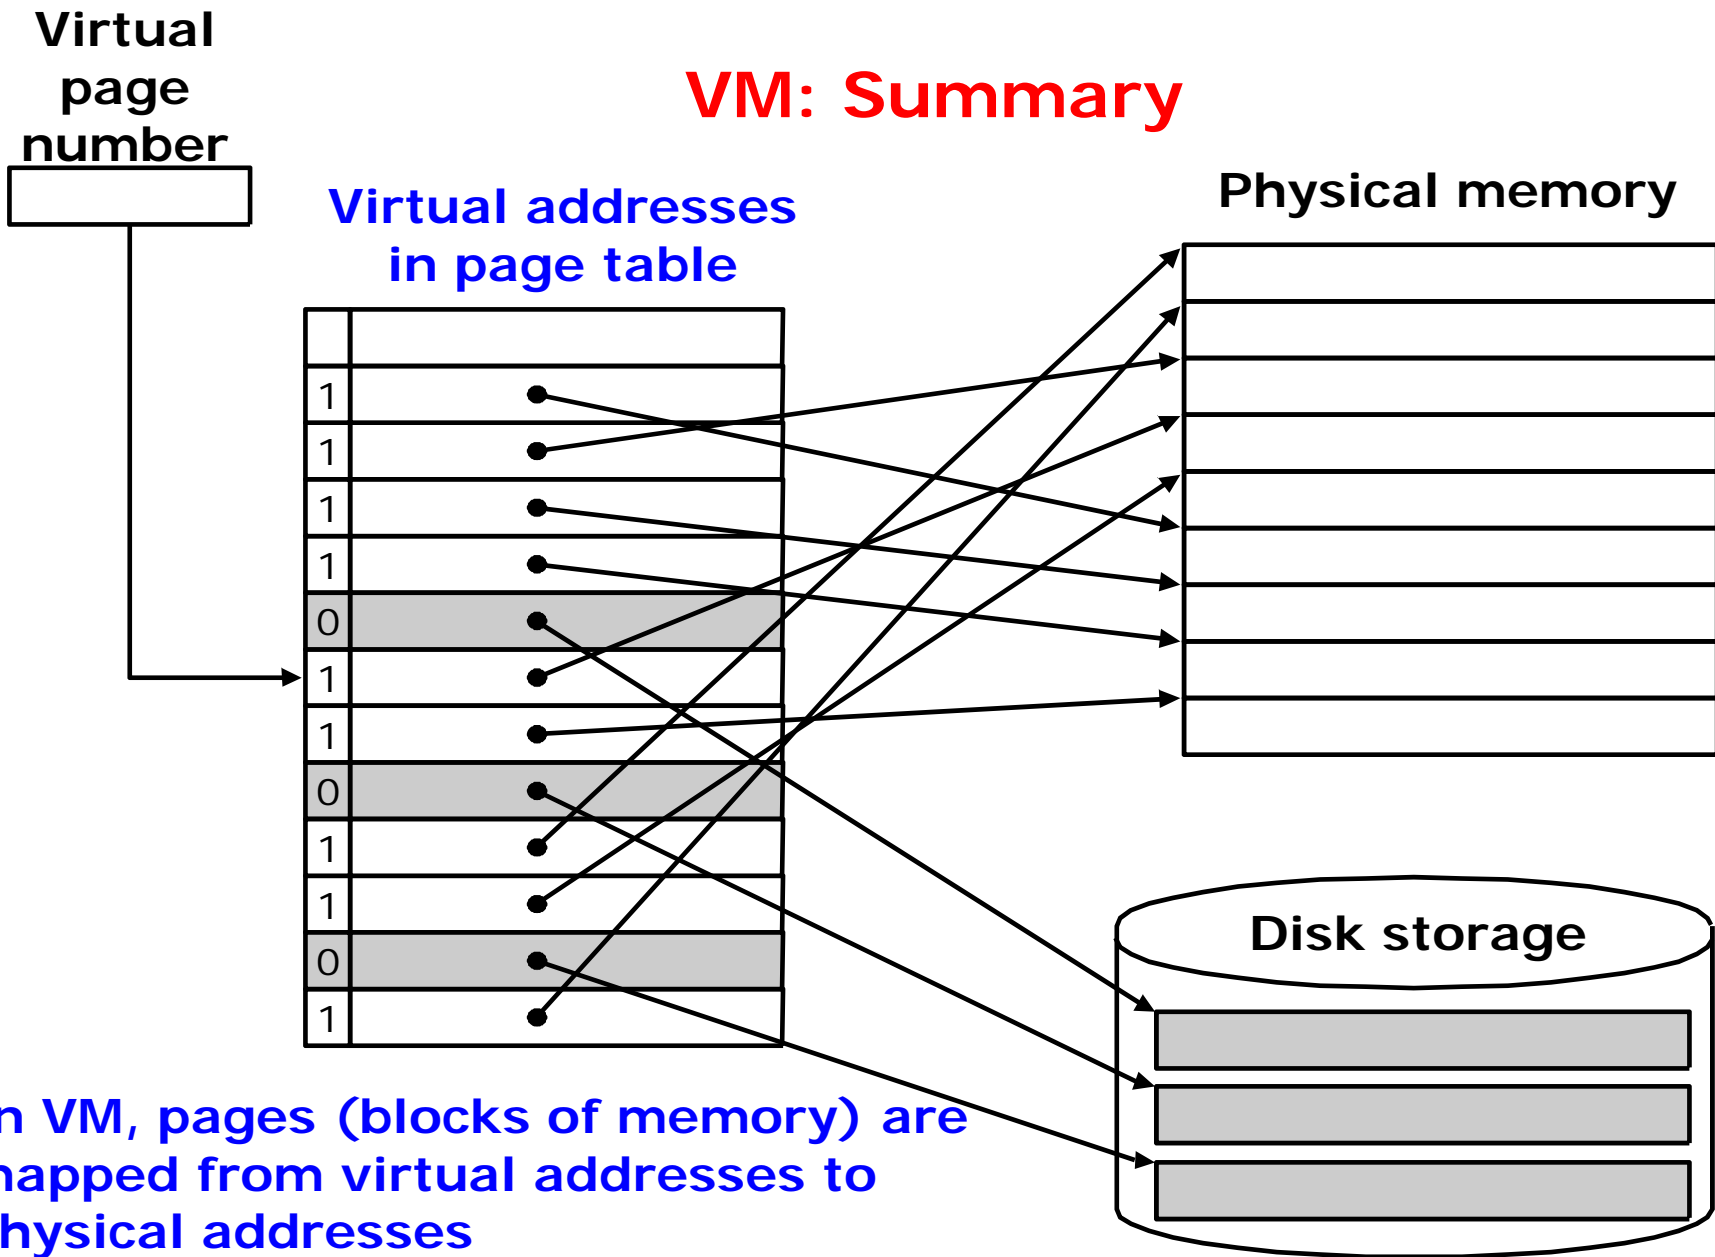
Example from Spring 2009 final exam (11)



VM: Summary



VM: Summary



In VM, pages (blocks of memory) are mapped from virtual addresses to physical addresses

NOTE: pages could map to disk

Page Faults: : the data is not in memory, retrieve it from disk

The OS finds the page in the next level of the hierarchy and decides where to place it in the main memory

Design challenges:

- *Huge miss penalty*
 - pages should be fairly large (e.g. 32 KB)
 - optimize the page placement
- *Reducing page fault rate is important*
 - allow fully associative placement of pages
- *Reduce the overhead*
 - handle the faults in software or hardware?
 - allow using clever algorithms?
- *using write-through is too expensive since it takes too long*
 - use write back
 - a dirty bit is used to decide whether to write back or not

Making Address Translation Faster: The TLB

Every memory access by a program can have two parts:

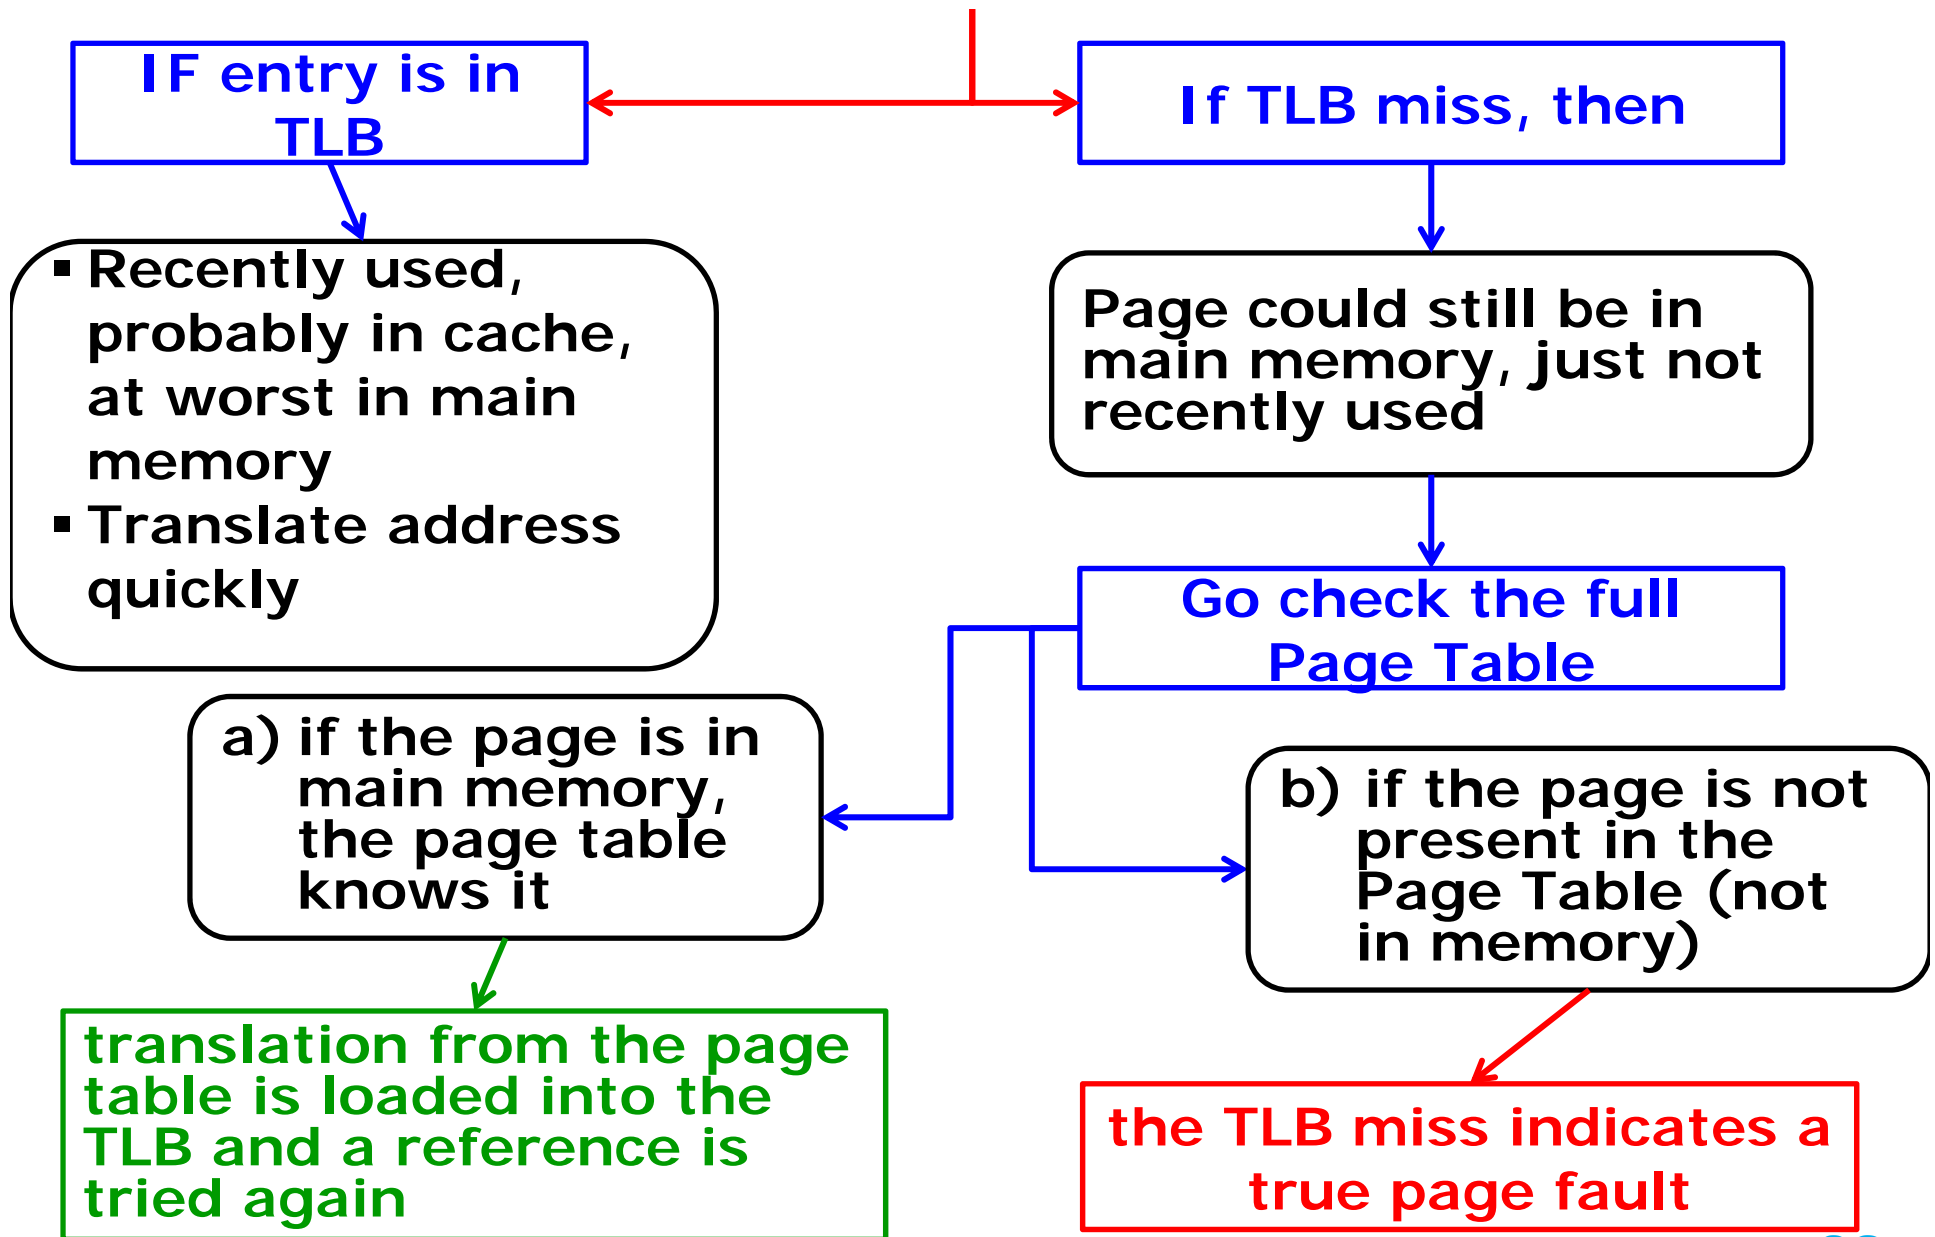
1. one memory access to obtain the physical address
2. a second access to get the data

To speed up step 1 (translation) use a **TLB = Translation-lookaside buffer (TLB)**:

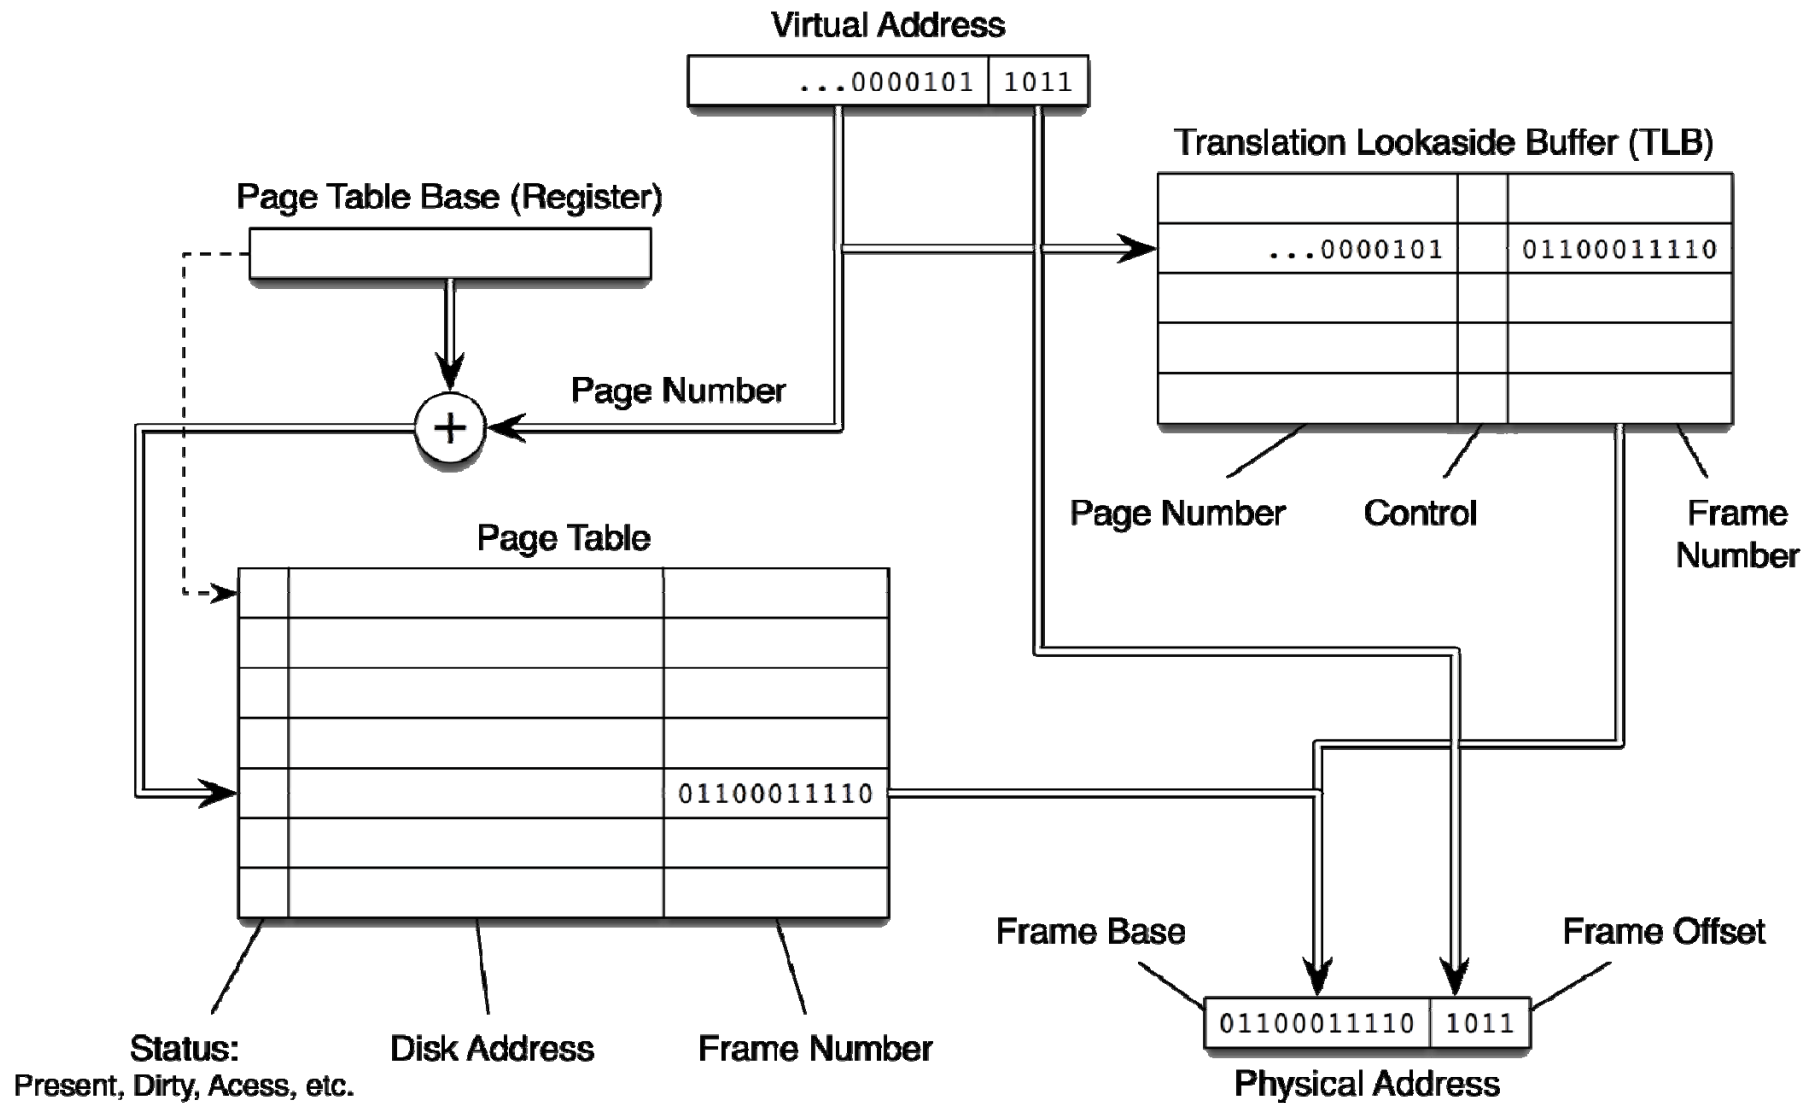
a cache for recent address translations
(recent page entries)

- TLB can be: fully associative, set associative, or direct mapped
- TLBs are usually small, typically not more than 128 - 256 entries
- TLB access time should be comparable to cache access time

Making Address Translation Faster: The TLB



Address Translation



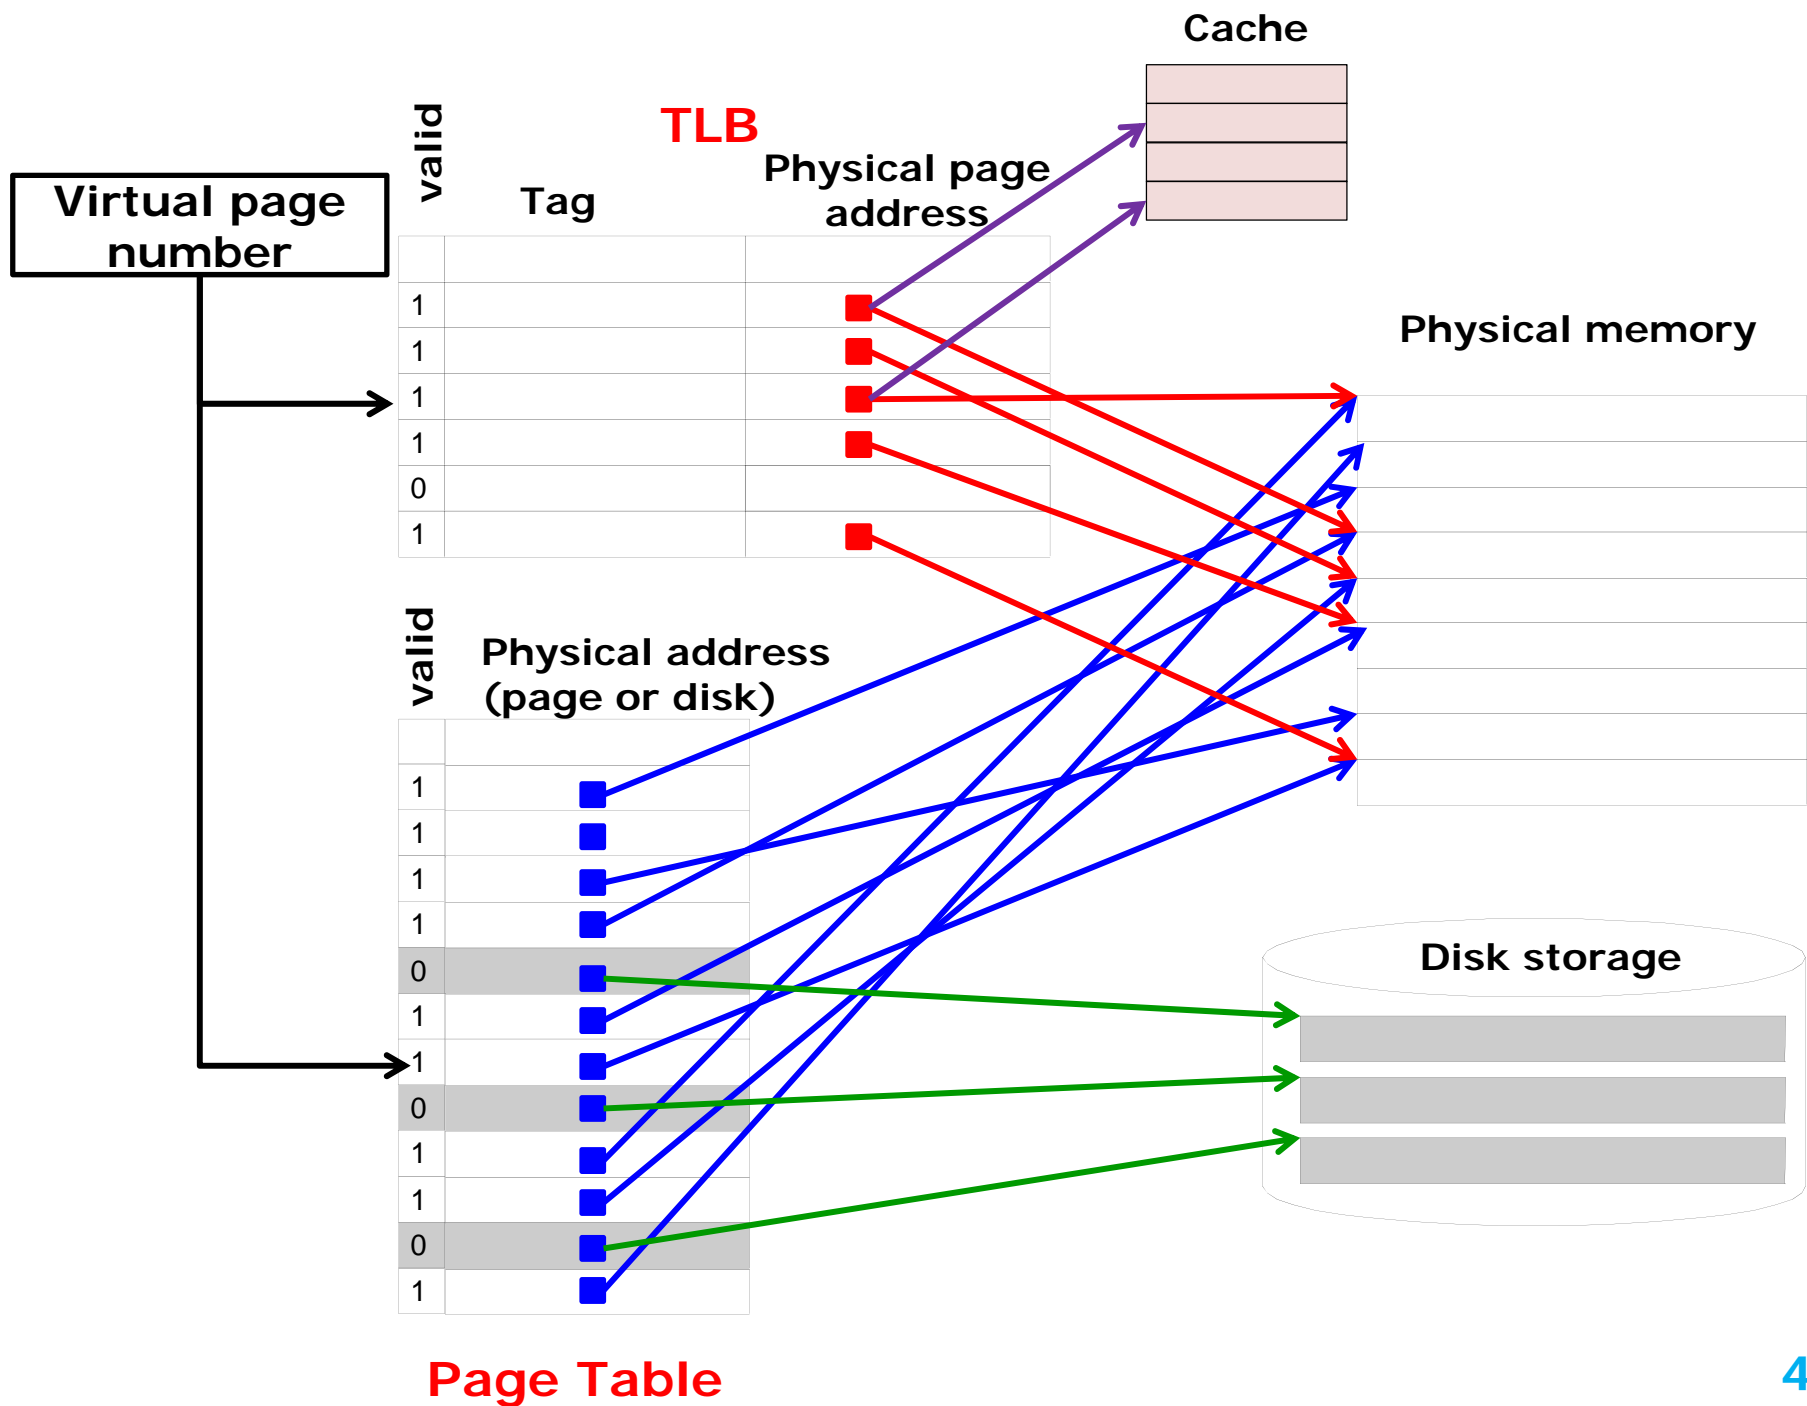
Translation Lookaside Buffer (TLB)

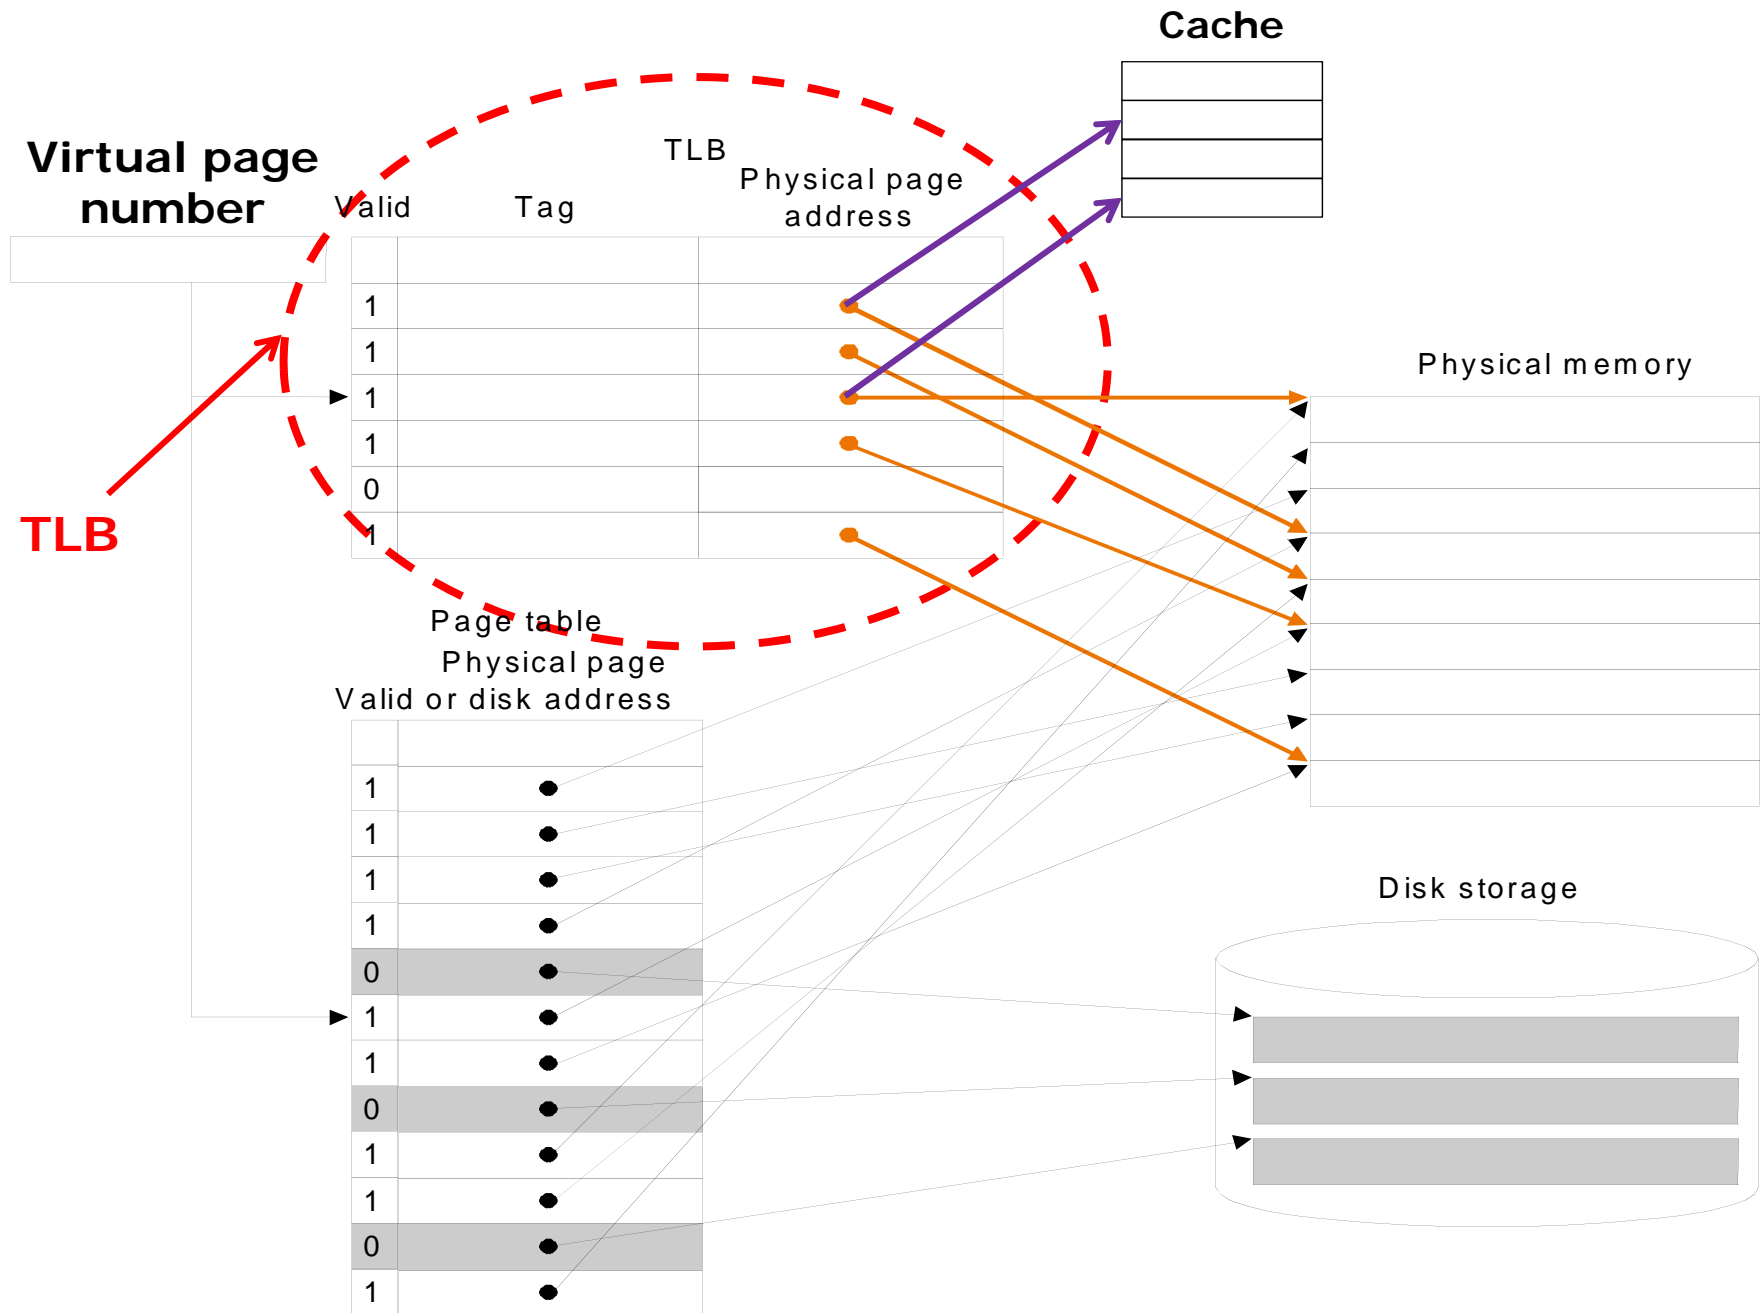
Used to speed virtual address translation, a “cache” of page table entries: the TLB (implemented in hardware)

- TLB holds recent page to frame mappings
- Normally resides in MMU

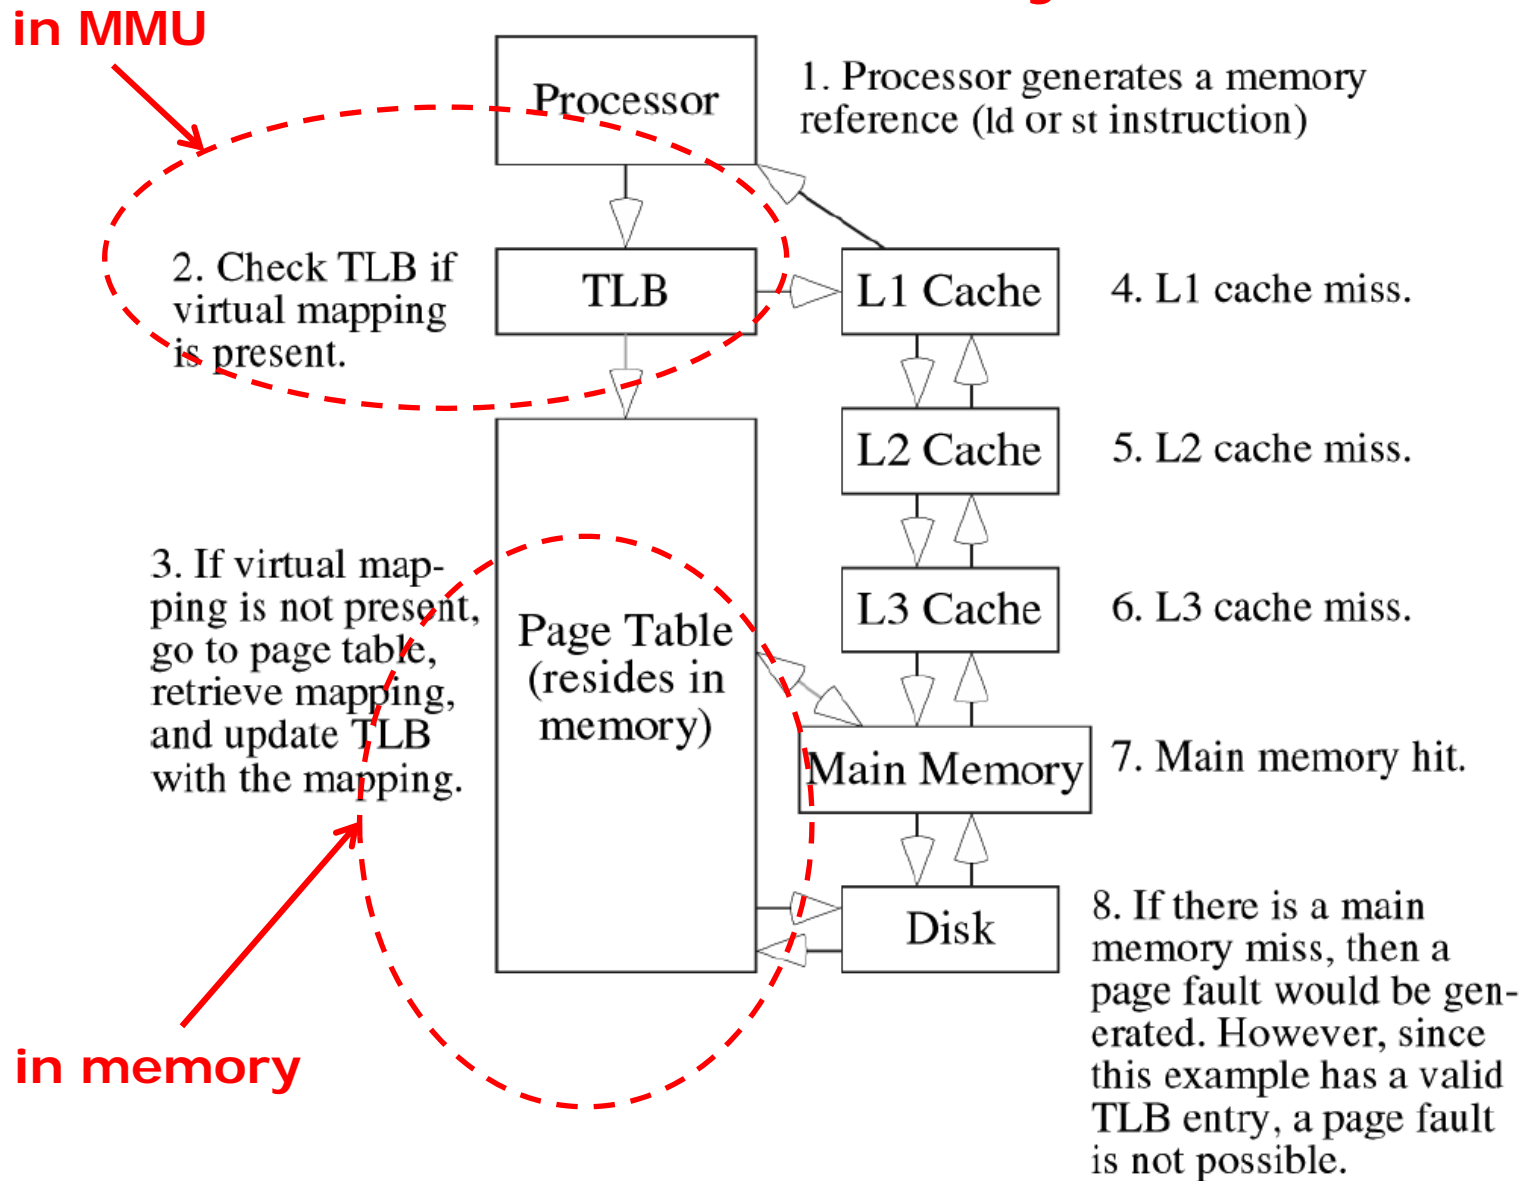
Entries may be invalid when pages get ejected (swapped) out

Valid	Virtual page number	Physical page number
1	0 1 0 0 1	1 1 0 0
1	1 0 1 1 1	1 0 0 1
0	- - - - -	- - - -
0	- - - - -	- - - -
1	0 1 1 1 0	0 0 0 0
0	- - - - -	- - - -
1	0 0 1 1 0	0 1 1 1
0	- - - - -	- - - -






Putting it All Together: example when data is in memory



Putting it All Together: more details about TLB and cache



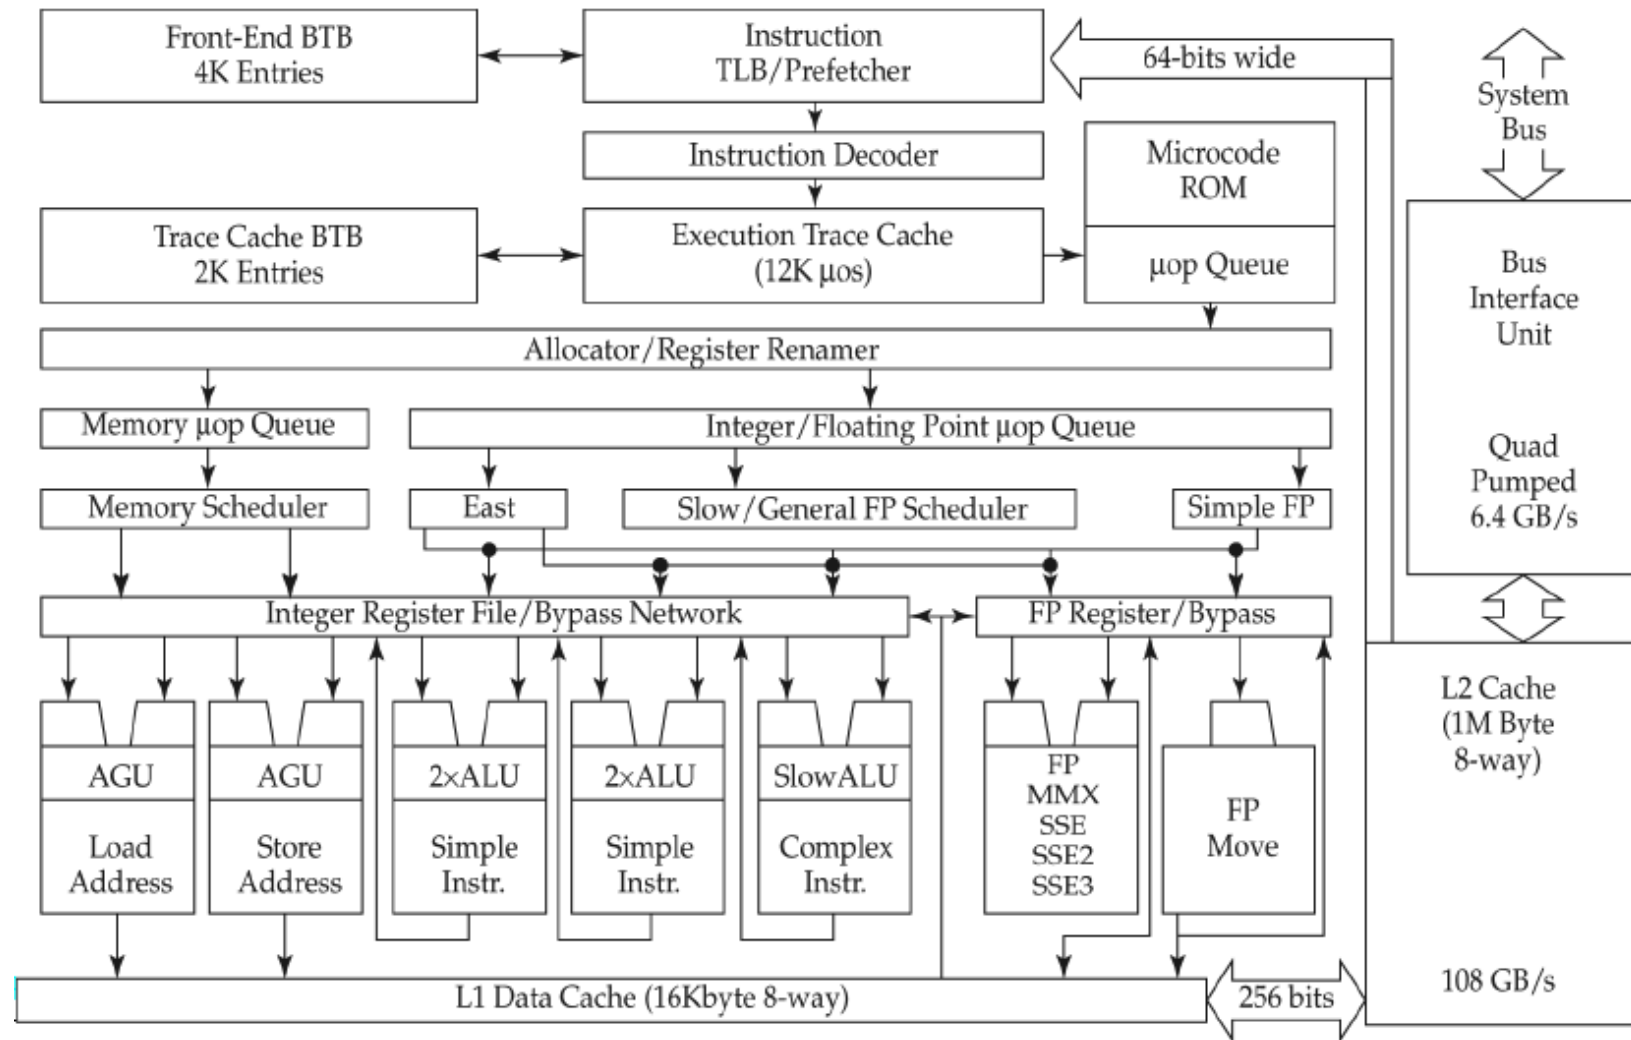
1. Initial address sent in parallel to TLB and L1 cache
2. If in L1
→ cache hit, done
3. Else parallel search continues in TLB and L2 and L3
→ if quick answer with hit comes from TLB or L2 or L3
→ done
4. Else possibly
→ TLB hit to memory
5. Etc. etc.

- All supported by hardware until Main memory hit
- OS software takes over for page Fault and access to disk

Typical values for a TLB

Size:	16-512 entries
Block size:	1-2 page table entries of typically 4-8 bytes each
Hit time:	0.5 – 1 clock cycles
Miss penalty:	10-100 clock cycles
Miss rate:	0.01% - 1%

Putting it All Together





ARM's Virtual Memory System Architecture

MMU within processor

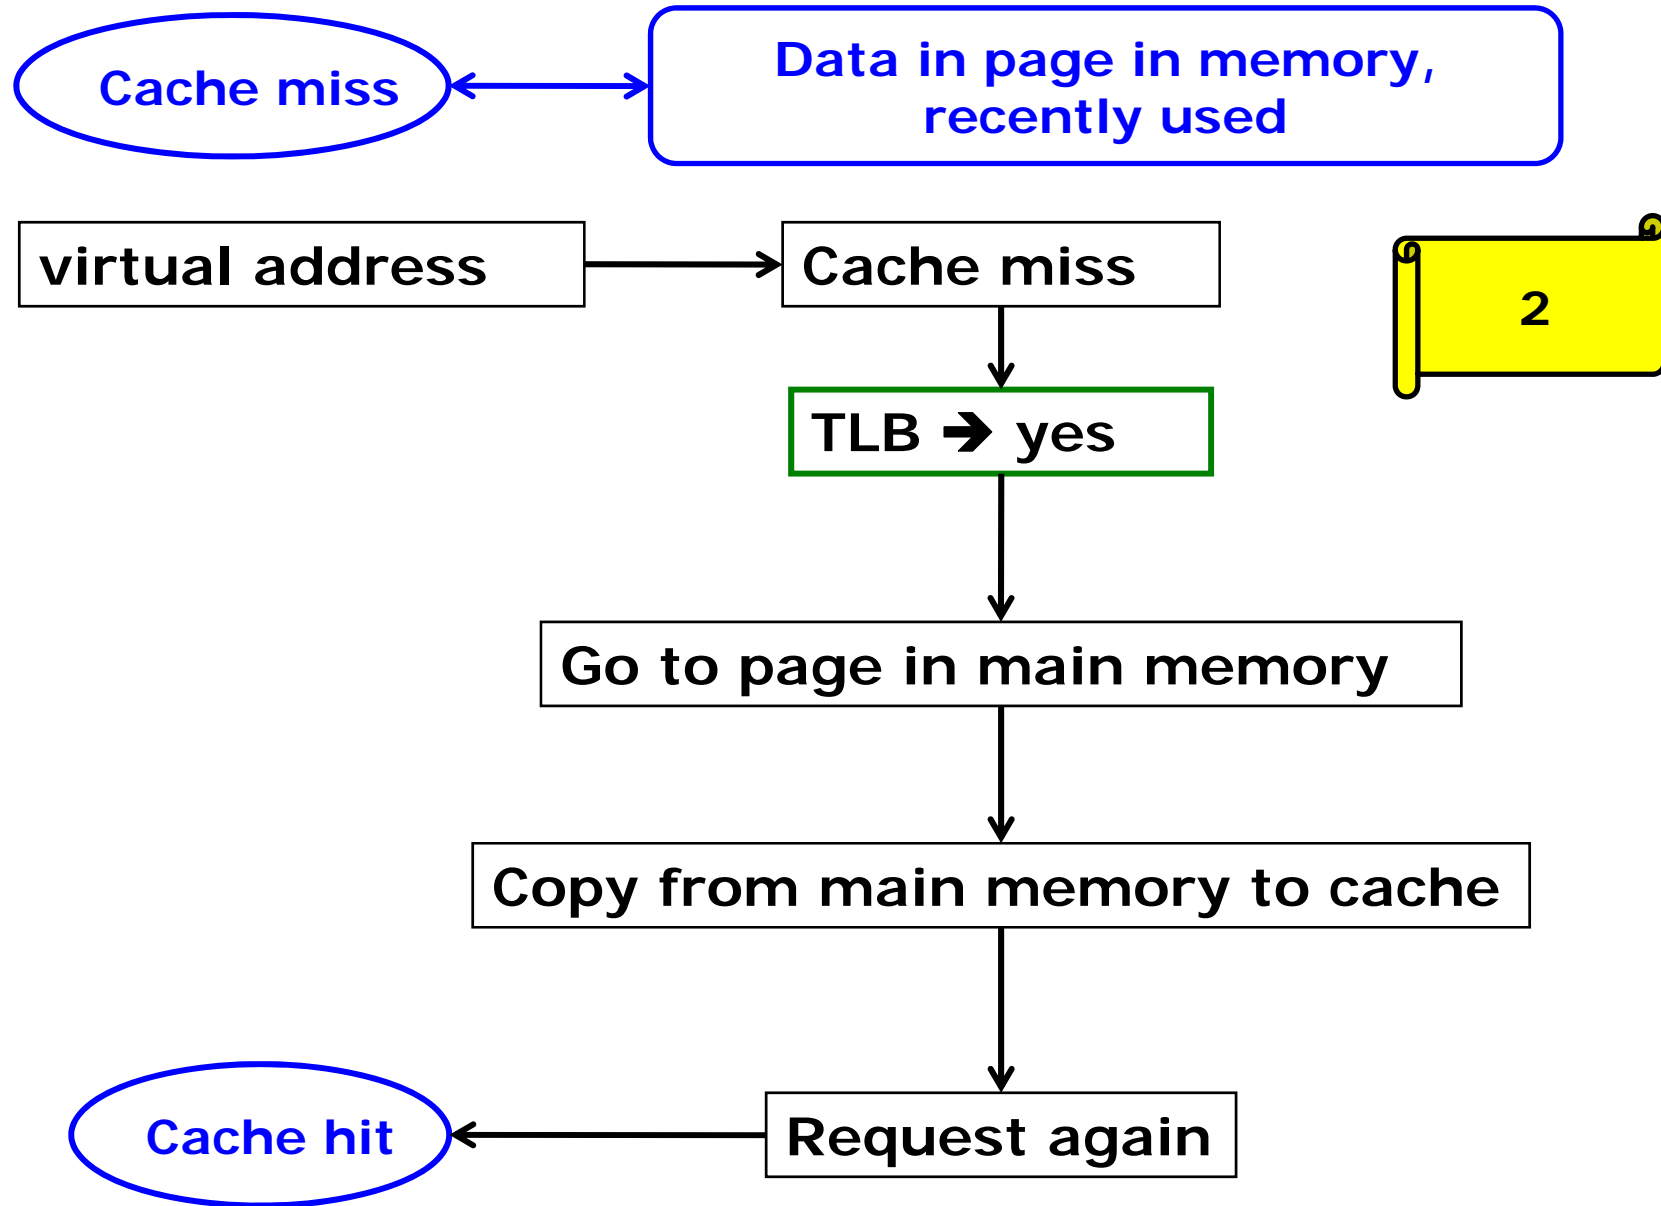
Defines 1Kb, 4KB and 64Kb pages

TLB all hardware

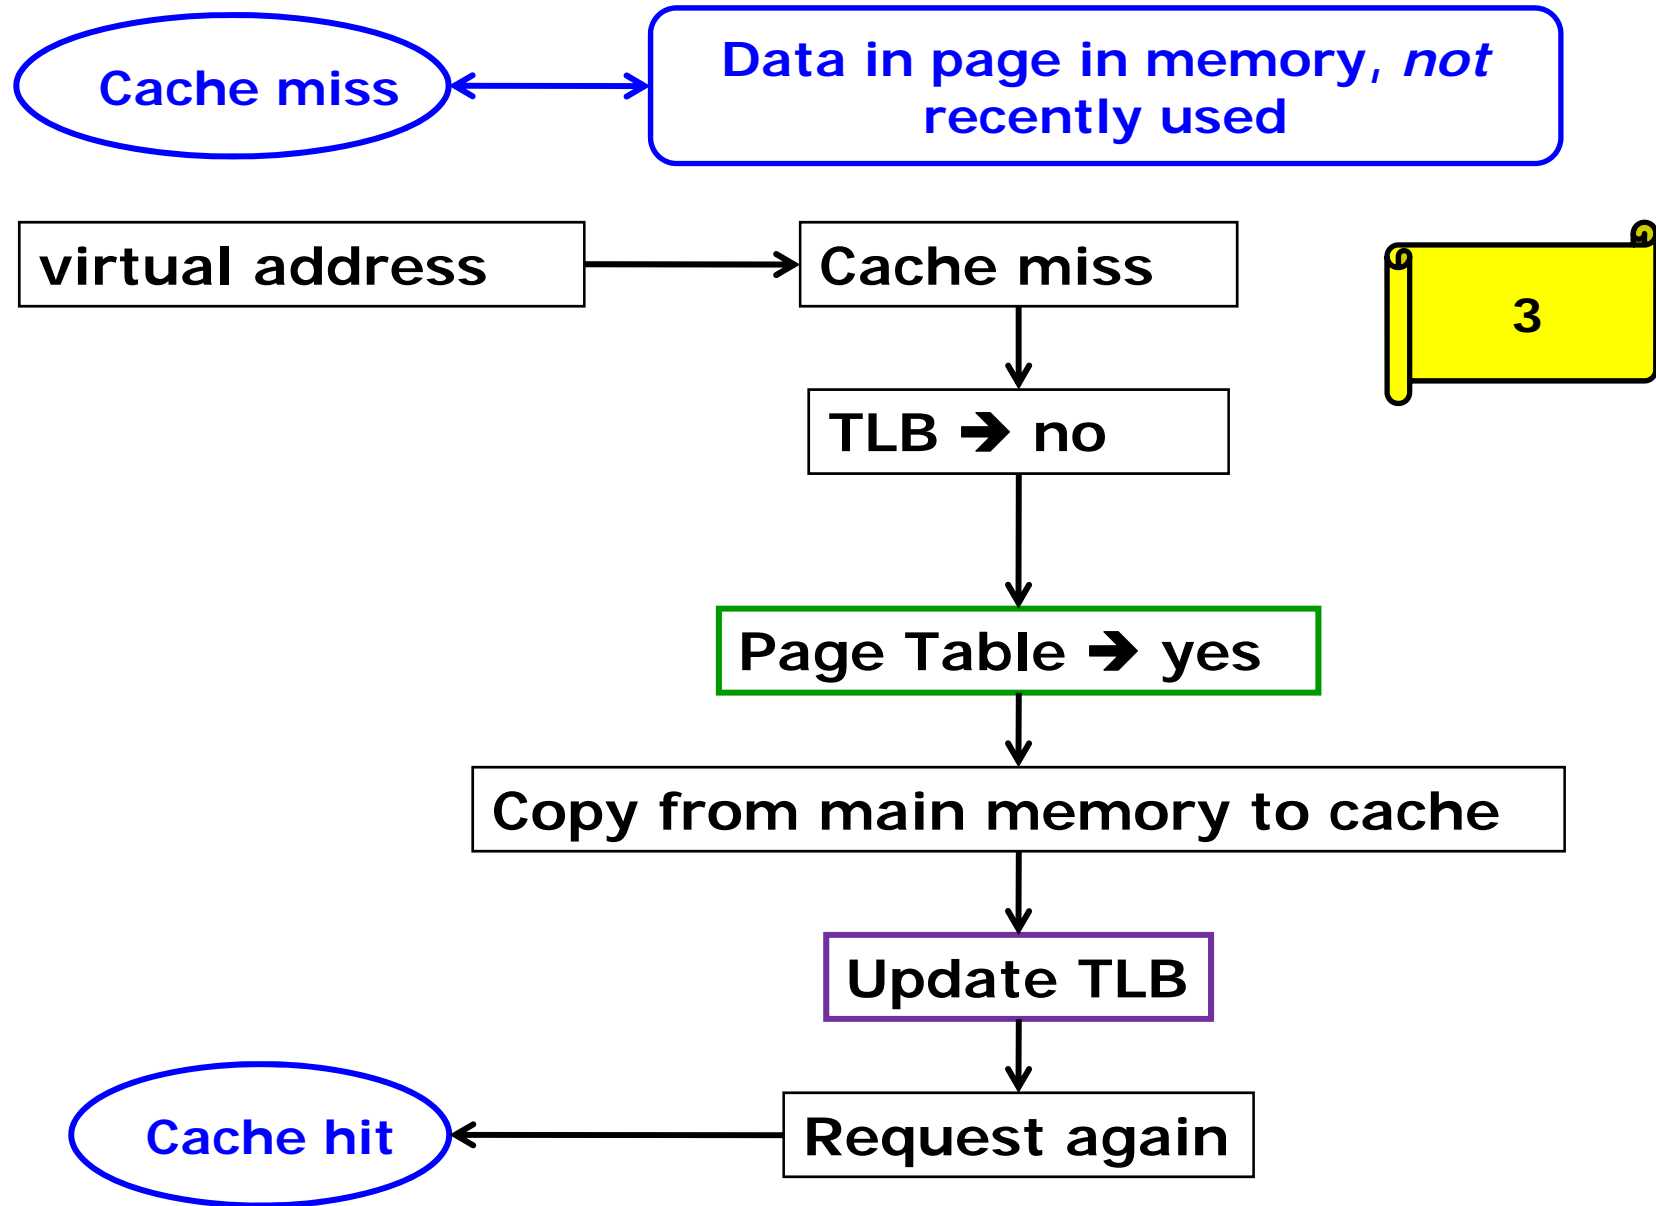
Integrating Virtual Memory, TLBs, and Caches



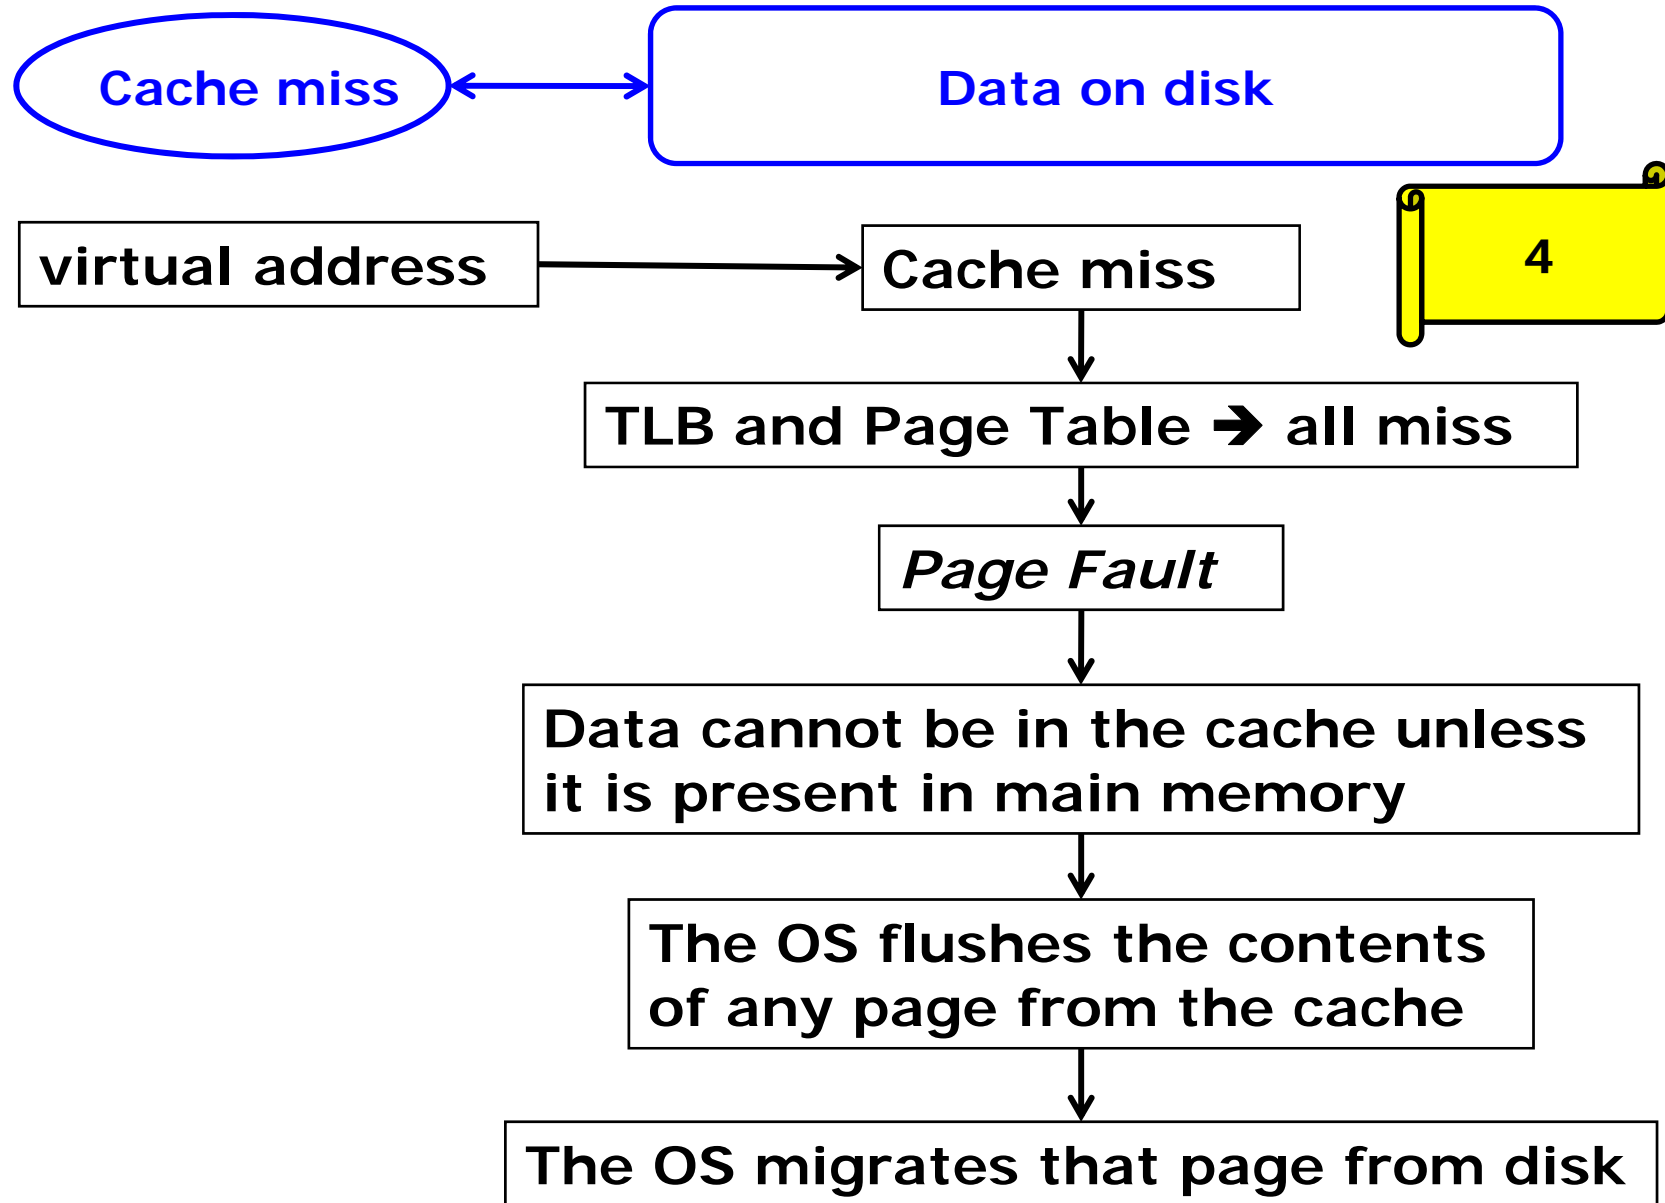
Integrating Virtual Memory, TLBs, and Caches



Integrating Virtual Memory, TLBs, and Caches



Integrating Virtual Memory, TLBs, and Caches



While the notion of cache and virtual memory seem similar there are differences:

	Cache	Virtual Memory
definition	cache is a subset of main memory	memory is a subset of what is on disk
purpose	speed	expand memory (provide relocatability)
data unit	line/block	page
implemented	hardware	hardware / software

Integrating Virtual Memory, TLBs, and Caches

Cache	TLB	Virtual memory	Possible? If so, under what circumstance?
miss	hit	hit	Possible, although the page table is never really checked if TLB hits.
hit	miss	hit	TLB misses, but entry found in page table; after retry, data is found in cache.
miss	miss	hit	TLB misses, but entry found in page table; after retry, data misses in cache.
miss	miss	miss	TLB misses and is followed by a page fault; after retry, data must miss in cache.
miss	hit	miss	Impossible; cannot have a translation in TLB if page is not present in memory.
hit	hit	miss	Impossible; cannot have a translation in TLB if page is not present in memory.
hit	miss	miss	Impossible; data cannot be allowed in cache if the page is not in memory.

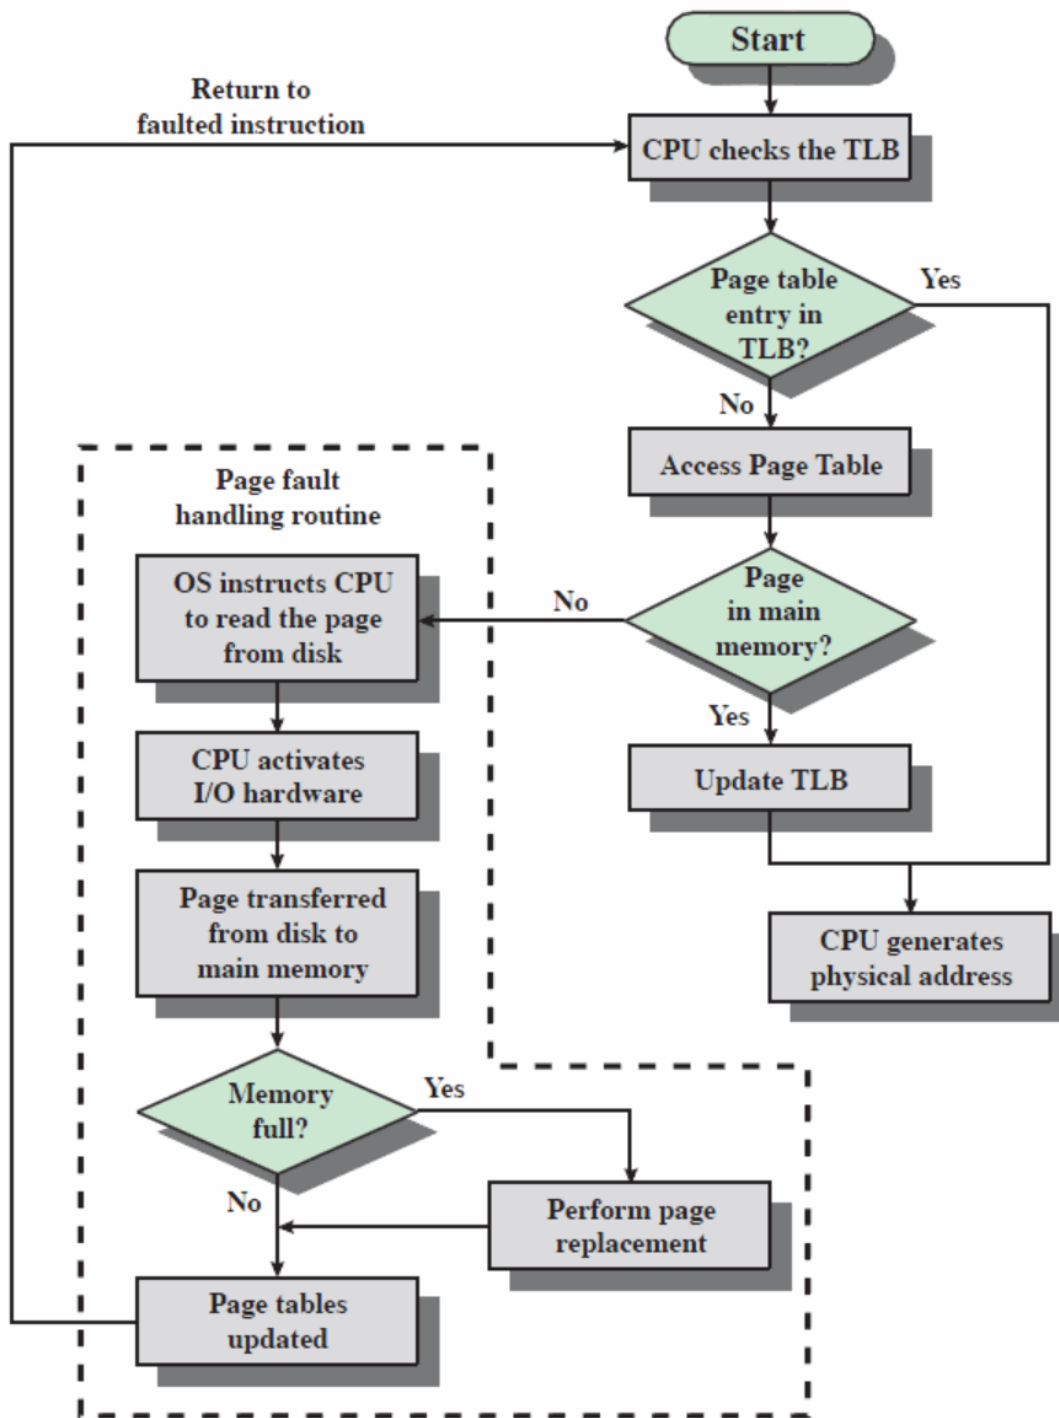


Figure 8.18

To figure 8.19

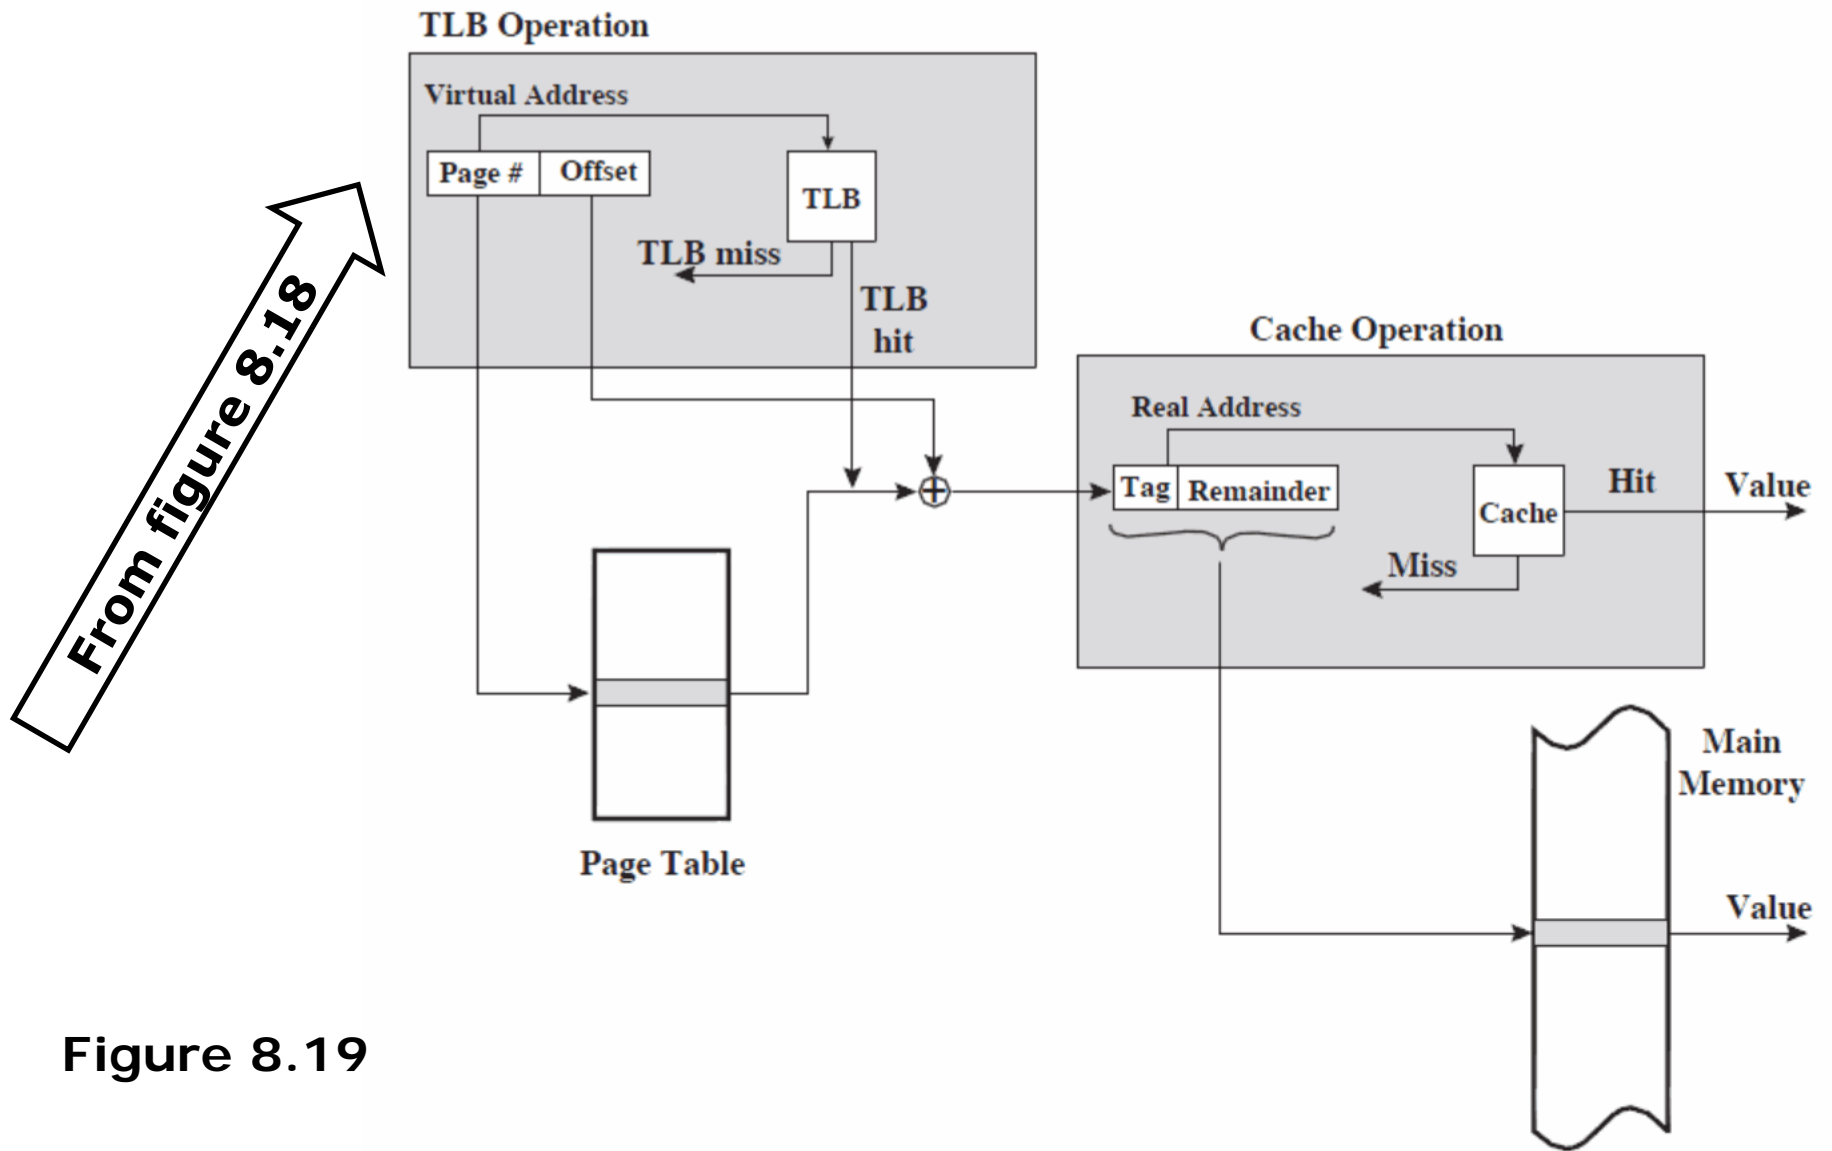



Figure 8.19

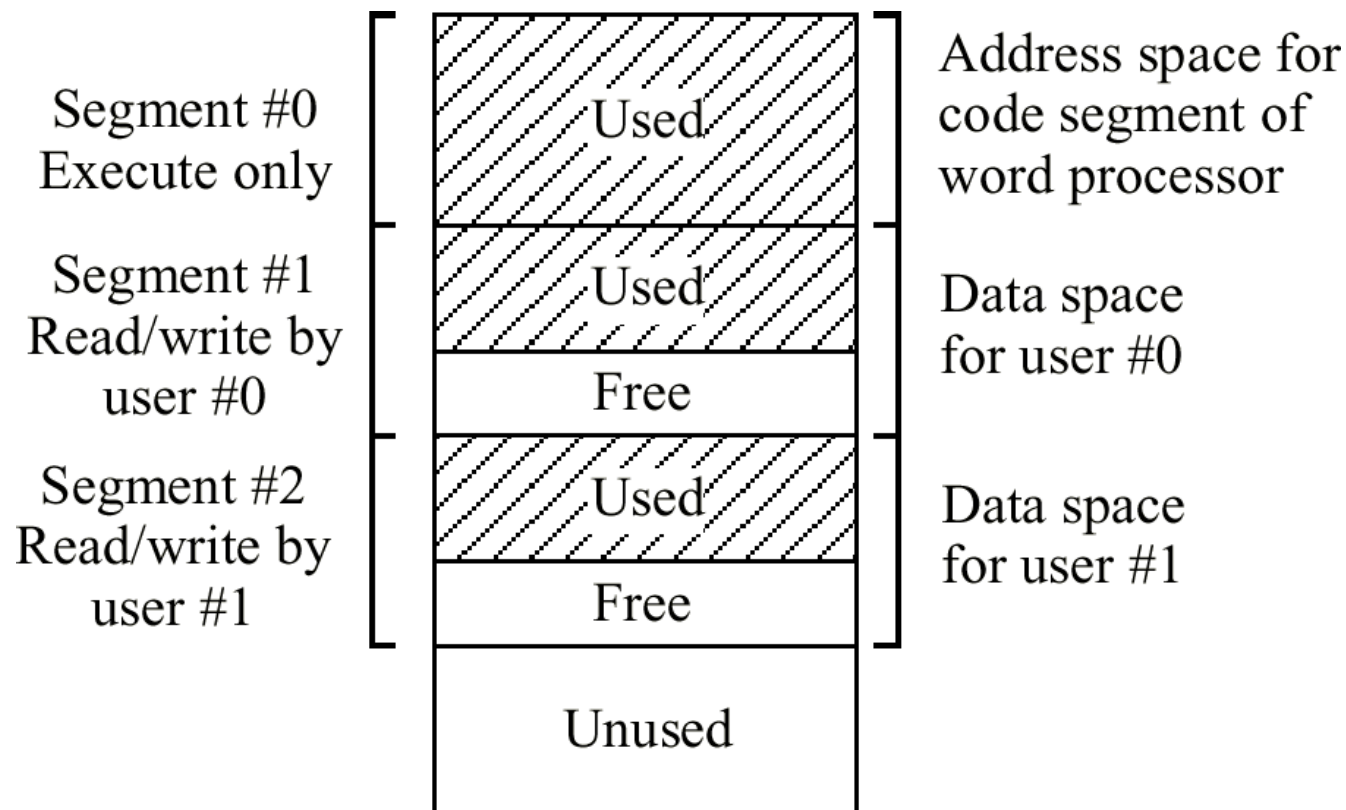
Useful utility: CPUZ



<http://www.cpuid.com/softwares/cpu-z.html>

Segmentation

A segmented memory allows two users to share the same word processor code, with different data spaces:




Fragmentation

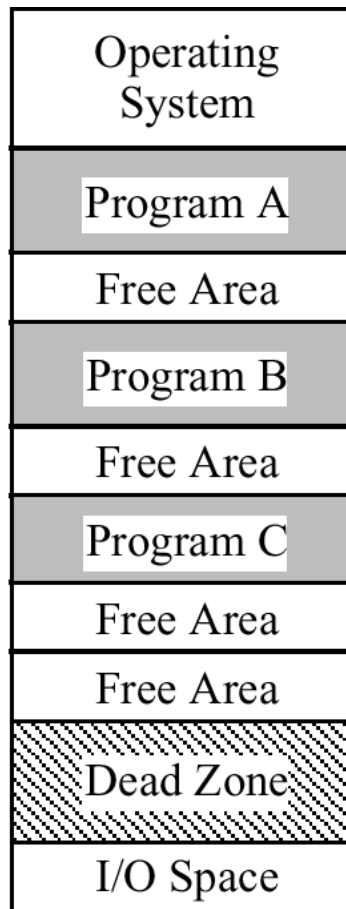
(a) Free area of memory after initialization;

(b) after fragmentation;

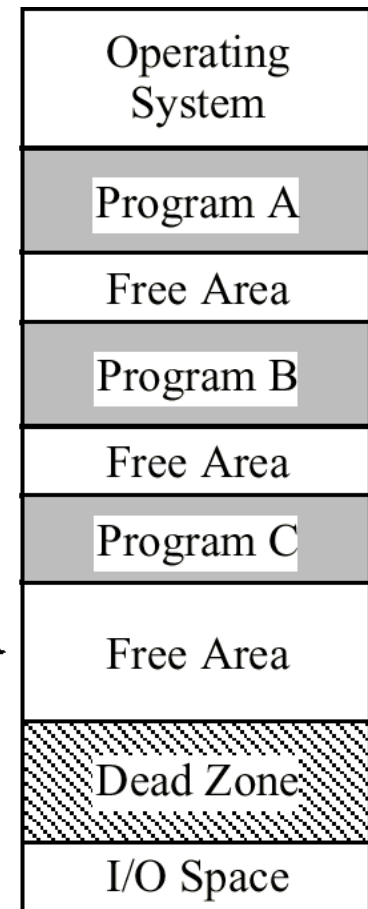

(c) after coalescing.



(a)



(b)



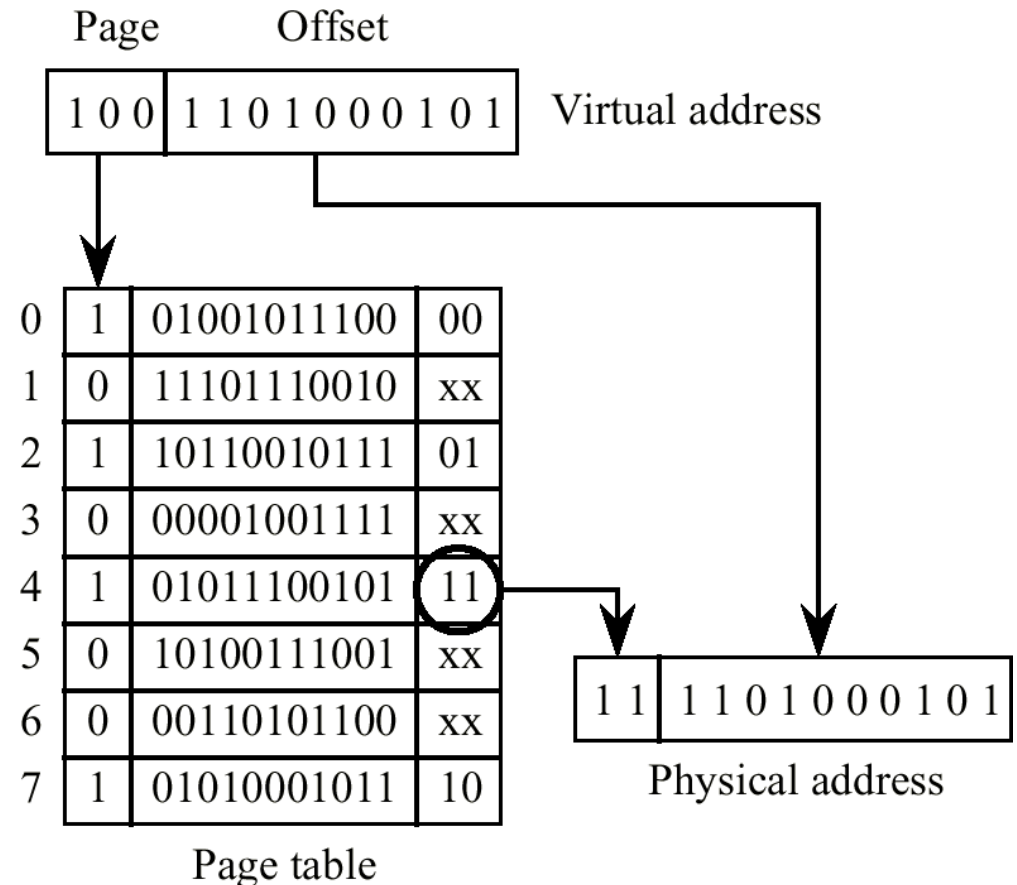
(c)

Using the Page Table (read from MH)

Virtual addresses are translated to physical addresses by breaking them into two parts:

1. page number
2. offset

- The page number identifies a location in the page table which contains the current frame for the page
- The frame number and offset are combined to form a physical address.



Using the Page Table - Example

Assume sequential access to pages 1, 2, 3, 4, and 5 (with 4 frames).

- Table is initially empty
- When page is not in table, trigger a "page fault" and load the needed page
- Replace old pages when out of frames

0	0	01001011100	xx
1	1	11101110010	00
2	0	10110010111	xx
3	0	00001001111	xx
4	0	01011100101	xx
5	0	10100111001	xx
6	0	00110101100	xx
7	0	01010001011	xx

After
fault on
page #1

0	0	01001011100	xx
1	1	11101110010	00
2	1	10110010111	01
3	0	00001001111	xx
4	0	01011100101	xx
5	0	10100111001	xx
6	0	00110101100	xx
7	0	01010001011	xx

0	0	01001011100	xx
1	1	11101110010	00
2	1	10110010111	01
3	1	00001001111	10
4	0	01011100101	xx
5	0	10100111001	xx
6	0	00110101100	xx
7	0	01010001011	xx

After
fault on
page #3

0	0	01001011100	xx
1	0	11101110010	xx
2	1	10110010111	01
3	1	00001001111	10
4	1	01011100101	11
5	1	10100111001	00
6	0	00110101100	xx
7	0	01010001011	xx