

Reductions

Basic idea: we don't have to solve problems from scratch. Use existing problems to solve new problems.

Can be useful in practice (e.g. SAT solvers) but is also a way to show that new problems are e.g., undecidable

Say, e.g., we have a language L and we want to know whether it is decidable.

Suppose we can show that *if* we *could* decide L , *then* we *could* decide A_{TM}

We conclude that we *can't* decide L

How do we show *if* . . . *then*? Reductions!

$$HALT_{TM}$$

$HALT_{TM} = \{\langle M, w \rangle \mid M \text{ is a TM and } M \text{ halts on input } w\}.$

Theorem: $HALT_{TM}$ is undecidable.

IDEA: Use machine for $HALT_{TM}$ to solve A_{TM} .

Proof: Assume there is a TM R which decides $HALT_{TM}$. Construct the TM S to decide A_{TM} as follows:

1. S runs TM R on input $\langle M, w \rangle$
2. If R rejects, reject.
3. If R accepts, simulate M on w until it halts.
4. If M accepts, accept. Else if M rejects, reject.

We have shown that if there is such an R then A_{TM} is decidable, but that is false, so our assumption is false and $HALT_{TM}$ is undecidable.

Test for emptiness

$E_{TM} = \{\langle M \rangle \mid M \text{ is a TM and } L(M) = \emptyset\}$.

Theorem: E_{TM} is undecidable.

Proof: For any M, w we can let M_1 be the TM which takes as input string x :

1. If $x \neq w$ M_1 rejects.
2. If $x = w$ M_1 runs M on input w and accepts if M does.

Now we construct TM S to decide A_{TM} . Let R be a hypothetical TM which decides E_{TM} :

S has input $\langle M, w \rangle$

1. Use $\langle M \rangle, w$ to construct $\langle M_1 \rangle$ for these M, w .
2. Run R on $\langle M_1 \rangle$.
3. If R accepts, reject; if R rejects, accept.

If R decided the emptiness of $L(M_1)$, then S decides A_{TM} . Therefore R can't exist and E_{TM} is undecidable.

Whether a language recognized by a TM is regular

$REGULAR_{TM} = \{\langle M \rangle \mid M \text{ is a TM and } L(M) \text{ is a regular language}\}.$

Theorem: $REGULAR_{TM}$ is undecidable.

Proof idea: We construct a machine M_2 which recognizes a non-regular language if M does not accept w and recognizes all binary strings if M accepts w .

Proof that $REGULAR_{TM}$ is undecidable

Proof: Let R be a TM that decides $REGULAR_{TM}$. We construct S to decide A_{TM} as follows. S has input $\langle M, w \rangle$:

1. Construct M_2 such that on input x , M_2 does:
 - (a) If x has the form $1^n 0^n$ accept.
 - (b) If x has any other form, run M on w and accept x if M accepts w .
2. S runs R on $\langle M_2 \rangle$.
3. If R accepts, accept. If R rejects, reject.

If M accepts w then M_2 accepts all strings, which implies R accepts, which implies S accepts. If M doesn't accept w , then for all strings not in $0^n 1^n$, M_2 does not accept, which implies $L(M_2)$ is not regular, which implies R rejects, which implies S rejects. So S decides A_{TM} which gives a contradiction. Hence $REGULAR_{TM}$ is undecidable.

NOTE: testing any property of languages is undecidable. RICE's THEOREM—Prob. 5.28.

Equivalence of TM's

$$EQ_{TM} = \{\langle M_1, M_2 \rangle \mid M_1 \text{ and } M_2 \text{ are TM's and } L(M_1) = L(M_2)\}.$$

Theorem: EQ_{TM} is undecidable.

IDEA: Can use such a TM to test if a language is empty.

Proof: Let R be the (hypothetical) TM which decides EQ_{TM} . Construct S to decide E_{TM} as follows: S is given $\langle M \rangle$ as input.

1. Run R on $\langle M, M' \rangle$, where M' is the TM which rejects all strings.
2. If R accepts, S accepts. If R rejects, S rejects.

If $L(M)$ is empty, then R accepts and S accepts. If $L(M)$ is not empty then R rejects and S rejects. So S decides E_{TM} . But E_{TM} is undecidable so R cannot exist and EQ_{TM} is undecidable.

Mapping reducibility

Formalizing what we've done:

A function $f : \Sigma^* \rightarrow \Sigma^*$ is a *computable function* if some TM on every input w halts with just $f(w)$ on its tape.

Examples: adding $m + n$, modifying a description of a TM $\langle M \rangle$ in a simple way.

Formal definition of reduction

A language A is *mapping reducible* to a language B written $A \leq_m B$ if there is a computable function $f : \Sigma^* \rightarrow \Sigma^*$ where for every w

$$w \in A \text{ iff } f(w) \in B$$

Observation: $A \leq_m B$ iff $\overline{A} \leq_m \overline{B}$.

Proving problems are decidable

Theorem: If $A \leq_m B$ and B is decidable then A is decidable.

Proof: Let M be a TM which decides B . Construct a TM N which decides A using M as a subroutine:

1. Compute $f(w)$
2. Run M on $f(w)$

Corollary: If $A \leq_m B$ and A is undecidable then B is undecidable.

Examples from before.

Reductions also prove Turing-recognizability

Theorem: If $A \leq_m B$ and B is Turing-recognizable then A is Turing-recognizable.

Corollary: If $A \leq_m B$ and A is not Turing-recognizable then B is not Turing-recognizable.

Recall that $\overline{A_{TM}}$ is not Turing recognizable. Why not? We use this to show that EQ_{TM} is neither Turing recognizable nor co-Turing recognizable.

We can't recognize whether two TM's are equivalent

Theorem: EQ_{TM} is not Turing recognizable Proof: To show EQ_{TM} is not Turing recognizable, we show that $\overline{A_{TM}} \leq_m EQ_{TM}$ or alternatively $A_{TM} \leq_m \overline{EQ_{TM}}$.

Let the reducing function f be defined as follows:

$f(\langle M, w \rangle)$ outputs two machine descriptions $\langle M_1, M_2 \rangle$ such that

1. M_1 rejects all inputs;
2. M_2 on any input disregards its input and runs M on w and accepts if M accepts.

We show f is a reducing function. If $\langle M, w \rangle \in A_{TM}$ then M accepts w . But then M_1 rejects w and M_2 accepts w . Hence $f(\langle M, w \rangle) = \langle M_1, M_2 \rangle \in \overline{EQ_{TM}}$.

If $\langle M, w \rangle \notin A_{TM}$ then both M_1 and M_2 reject all strings and $\langle M_1, M_2 \rangle \notin \overline{EQ_{TM}}$.

We can't recognize whether two TM's are not equivalent

Theorem: EQ_{TM} is not co-Turing recognizable

Proof: To show $\overline{EQ_{TM}}$ is not Turing recognizable, we show a reduction from A_{TM} to EQ_{TM} . Let the reducing function g be defined as follows: $f(\langle M, w \rangle)$ outputs two machine descriptions $\langle M_1, M_2 \rangle$ such that

1. M_1 accepts all inputs;
2. M_2 on any input disregards its input and runs M on w and accepts if M accepts.

Similar argument as before except M accepts w iff both M_1 and M_2 accept.

Examples

Does M halt on the empty tape?

Is there any string which M halts on?

Given a TM and a state q , does the TM ever enter state q ?

(Remaining slides are optional)

Computation history

A *computation history* for a TM on an input is the sequence of configurations that the machine goes through as it processes the input.

An *accepting computation history* for M on w is the sequence of configurations C_1, C_2, \dots, C_t where C_1 is a start configuration and each C_i follows according to the rules of M and C_t is an accepting configuration. A *rejecting computation history* is similar except C_t is a rejecting configuration.

Suppose M is a deterministic TM. Then $\langle M, w \rangle \in A_{TM}$ iff there is *exactly one* accepting computation history for M on w . We will try to use this to get some more undecidability results.

Linear bounded automaton

A *linear-bounded automaton* is a type of TM where the tape head isn't permitted to move off the portion of the tape containing the input. If the machine tries to move off, it stays put. The tape alphabet may be larger than the input alphabet, giving in effect a constant factor increase in tape size.

Theorem: $A_{LBA} = \{\langle M, w \rangle \mid M \text{ is an LBA that accepts string } w\}$ is decidable.

Lemma: Let M be an LBA with q states and g symbols in the tape alphabet. There are exactly qng^n distinct configurations of M for a tape of length n .

Proof of lemma: A configuration is determined by the state, tape head position, and tape contents. Number of states = q . Number of tape head positions = n . Number of possible tape contents = g^n .

Decidability of LBA acceptance

Proof of theorem:

IDEA: If LBA M doesn't halt within qng^n steps then it repeats forever.

A TM S to decide A_{LBA} simulates M on w for qng^n steps or until it halts.

1. If M accepts, S accepts
2. If M rejects, S rejects
3. If M doesn't halt, S rejects.

Emptiness for LBA's is undecidable

Theorem: $E_{LBA} = \{\langle M \rangle \mid M \text{ is an LBA where } L(M) = \emptyset\}$ is undecidable.

Proof Idea: We will show that $A_{TM} \leq_m \overline{E_{LBA}}$. Define a mapping reduction f such that $f(\langle M, w \rangle) = \langle B \rangle$, where B is a LBA that recognizes accepting computation histories of M on w .

Clearly, if M does not accept w , then $L(B) = \emptyset$. Otherwise, $L(B)$ contains one element (what is it?)

So $\langle M, w \rangle \in A_{TM}$ iff $\langle B \rangle \in \overline{E_{LBA}}$.

Defining B

B operates as follows:

It checks its input string to see if it is a sequence of configurations, with the start configuration equal to $q_0 w_1 w_2 \dots w_n$ (w is hard-coded into B and the final configuration an accepting configuration (state is q_{accept} .)

It then checks back and forth (using marking) to see if each configuration follows from the next according to the transitions of M . That is, it checks that the tapes are the same except where the head is, the state change is according to δ of M . B accepts if the configuration history is according to δ and ends in an accepting state (δ is hard-coded into B .)

Since $\overline{E_{LBA}}$ is not decidable, neither is E_{LBA} .

Does a CFG generate all possible strings?

Theorem: $ALL_{CFG} = \{\langle G \rangle \mid G \text{ is a CFG and } L(G) = \Sigma^*\}$ is undecidable.

Proof Idea: Show that $A_{TM} \leq_m \overline{ALL_{CFG}}$ by designing a mapping reduction $f(\langle M, w \rangle) = \langle D \rangle$, where D is a PDA accepting all strings *except* the accepting computation history of M on w .

So M accepts w iff $L(D) \neq \Sigma^*$.

D guesses where its input fails to be an accepting computation history of M on w , uses stack to verify this.

See textbook for details.

An undecidable problem that doesn't involve automata

Post Correspondence Problem (PCP)

Input – a set of *tiles* of the form $\begin{bmatrix} u \\ v \end{bmatrix}$, where u, v are strings over Σ^*

Question: can we form a sequence $\begin{bmatrix} u_1 \\ v_1 \end{bmatrix}, \dots, \begin{bmatrix} u_k \\ v_k \end{bmatrix}$ of tiles from the set (with possible repetitions), such that $u_1 u_2 \dots u_k = v_1 v_2 \dots v_k$?

E.g., for the following set of tiles

$$\left\{ \begin{bmatrix} b \\ ca \end{bmatrix}, \begin{bmatrix} a \\ ab \end{bmatrix}, \begin{bmatrix} ca \\ a \end{bmatrix}, \begin{bmatrix} abc \\ c \end{bmatrix} \right\}$$

we have the following sequence.

$$\begin{bmatrix} a \\ ab \end{bmatrix} \begin{bmatrix} b \\ ca \end{bmatrix} \begin{bmatrix} ca \\ a \end{bmatrix} \begin{bmatrix} a \\ ab \end{bmatrix} \begin{bmatrix} abc \\ c \end{bmatrix}$$

PCP is undecidable

Show that $A_{TM} \leq_m PCP$ – given $\langle M, w \rangle$, tiles are used to encode configurations such that there is a match iff M accepts w .