

Advanced Computer Networks

P2P Swarming

Jianping Pan
Spring 2017

VOLUNTEERS NEEDED – Please email
Kyle Price (ksprice@uvic.ca) if you can help!

A poster for the Hi-Tech Fair. The background is orange with a pattern of white wireframe cubes in the bottom left. The text is in white and black. At the top, it says 'CO-OP + CAREER' in a small font. Below that, 'HI-TECH FAIR' is written in large, stylized, outlined letters. Underneath the title, it says 'FEBRUARY 7 - 8 | 10 A.M. TO 3 P.M. | ENGINEERING LAB WING LOBBY'. To the right, there is a list of bullet points under the heading 'Looking to connect with hi-tech employers?'. At the bottom, there is a dark blue banner with the website 'uvic.ca/coopandcareer/hitechfair', the hashtag '#uvichitech', and the University of Victoria logo.

CO-OP + CAREER

HI-TECH FAIR

FEBRUARY 7 - 8 | 10 A.M. TO 3 P.M. | ENGINEERING LAB WING LOBBY

Looking to connect with hi-tech employers?

- + Ask questions about working in industry
- + Learn what competencies employers look for
- + Discover co-op and career opportunities
- + Network with professionals

uvic.ca/coopandcareer/hitechfair #uvichitech

 University of Victoria

Review: going P2P

- Client-server
 - ↳ server is well-known and serves all client requests
 - ↳ scalability issue
- Peer-to-peer
 - ↳ structured or unstructured
 - ↳ every peer is a (potential) server
 - search is a challenge
 - ↳ one request is still served by one peer
 - until the peer fails, then try to use another peer

Napster and Gnutella

- Napster

- ↳ centralized directory server

- list uploading and query handling

- ↳ peer-to-peer file download

- Gnutella

- ↳ fully distributed


- scoped flooding search

- ↳ peer-to-peer file download

- Improving Gnutella

- ↳ node hierarchy

- ↳ non-flooding search



More design choices

- If more than one peer can serve, why do they not serve the same request together?
- Benefit
 - ↳ more resilient to node dynamic
 - does not rely on any particular peer
 - ↳ fit better with the asymmetric access link
 - higher download bandwidth than upload
- Overhead
 - ↳ how to get served from multiple peers
 - work together constructively

The BitTorrent approach

- Chop a file into small, fixed-size pieces
 - ↳ e.g., pieces (usually 256 KB each)
 - ↳ and then into blocks (usually 16 KB each)
- .torrent
 - ↳ meta information about the file
 - ↳ out-of-band retrieval
- Tracker
 - ↳ return a list of peers may have some pieces
- Seed and leecher/downloader
 - ↳ peers have the complete/incomplete file

.torrent

- Tracker URL
- File info
 - ↳ name, length
- Piece info
 - ↳ length, hash
- Other info
 - ↳ date, comment, etc
- Bencoding
 - ↳ strings, integers, lists, directories
 - ↳ e.g., `4:spam`, `i3e`, `l4:spam4:eggse`, `d4:spaml1:a1:bee`

Tracker protocol

- HTTP GET request
 - info_hash: to identify the file
 - peer_id: of the requesting peer
 - client address and port: to respond to incoming requests
 - bytes uploaded, downloaded, left, etc
 - numwant: the number of peers in the response list
- Tracker response
 - failure reason, if any
 - contact interval
 - peer list and stat (seed and leecher, etc)
- Tracker-less mode (on Kademlia DHT)

Tit-for-tat

- Download while upload: tit-for-tat
 - ↳ upload to whom from which download: trading pieces
 - ↳ prevent free-riding
 - fairness?
- Choking/unchoking
 - ↳ a limited number of uploads
 - default: 4+1
 - ↳ evaluate peers based on their recent download speed
 - 20-second average
 - ↳ upload to the peers with the fastest download speed
 - adjust every 10 seconds

Optimistic unchoking

- Stuck with poor peers?
- Optimistic unchoking
 - ↳ upload to other peers as well
 - rotate every 30 seconds
 - ↳ hope to get better download
 - ↳ also help bootstrap other peers
- Seed's unchoking
 - ↳ seed does not download from other peers
 - ↳ try to equally distribute its upload to leechers
 - ↳ or upload to the one downloads fastest

Peer wire protocol

- Messages over TCP

- handshake

- keep-alive

- choke/unchoke

- interested/not-interested

- a block is downloaded if the client is interested and unchoked
 - a block is uploaded if the peer is interested and unchoked

- have

- advertise new pieces

- request/piece

- request blocks in a piece

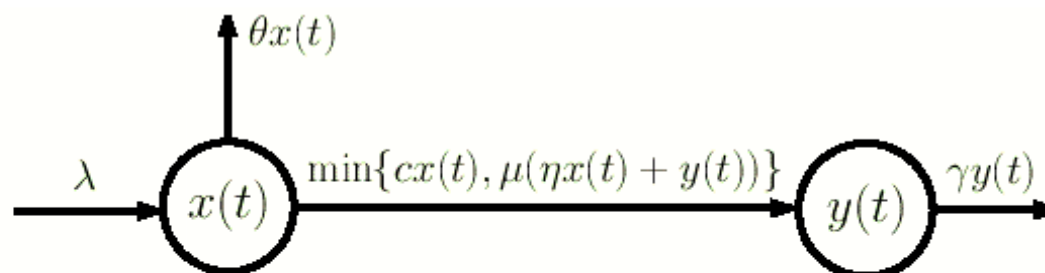
Piece selection

- Initially, a few random pieces
 - ↳ anything is better than nothing
 - ↳ easy to find at the beginning
- Then, rarest-first in neighborhood
 - ↳ become less dependent on seed
 - ↳ more interested by peers
- Finally, “end game” mode
 - ↳ look for missing pieces aggressively
 - ↳ send requests to all peers
 - ↳ cancel requests after last pieces are collected

BitTorrent performance

- Modeling and analysis
 - [QS04] Dongyu Qiu, R. Srikant. Modeling and Performance Analysis of Bit Torrent-Like Peer-to-Peer Networks. SIGCOMM 2004 [BitTorrent]

Fluid model



- ▶ $x(t)$: number of downloaders, $y(t)$: number of seeds
- ▶ λ : arrival rate of new requests
- ▶ θ : the rate at which a downloader aborts the download
- ▶ μ : uploading bandwidth of a peer
- ▶ η : effectiveness parameter (Yang and de Veciana)
- ▶ c : downloading bandwidth of a peer
- ▶ γ : seed departure rate

$$\frac{dx}{dt} = \lambda - \theta x(t) - \min\{cx(t), \mu(\eta x(t) + y(t))\}$$

$$\frac{dy}{dt} = \min\{cx(t), \mu(\eta x(t) + y(t))\} - \gamma y(t)$$

Steady-state performance

- $dx(t)/dt = dy(t)/dt = 0$

- ▶ If $\frac{1}{c} \geq \frac{1}{\eta}(\frac{1}{\mu} - \frac{1}{\gamma})$, the downloading bandwidth is the constraint:

$$\bar{x} = \frac{\lambda}{c(1 + \frac{\theta}{c})}, \quad \bar{y} = \frac{\lambda}{\gamma(1 + \frac{\theta}{c})}$$

- ▶ If $\frac{1}{c} \leq \frac{1}{\eta}(\frac{1}{\mu} - \frac{1}{\gamma})$, the uploading bandwidth is the constraint:

$$\bar{x} = \frac{\lambda}{\nu(1 + \frac{\theta}{\nu})}, \quad \bar{y} = \frac{\lambda}{\gamma(1 + \frac{\theta}{\nu})},$$

where $\frac{1}{\nu} = \frac{1}{\eta}(\frac{1}{\mu} - \frac{1}{\gamma})$.

Analytical insights

- Intrinsic scalability


- ▶ Little's law: average downloading time

$$T = \frac{1}{\theta + \beta}, \text{ where } \frac{1}{\beta} = \max \left\{ \frac{1}{c}, \frac{1}{\eta} \left(\frac{1}{\mu} - \frac{1}{\gamma} \right) \right\}$$

- ▶ Scalability: T is not a function of λ , the request arrival rate
- ▶ When the seed departure rate γ increases, T increases
- ▶ Even if $c \gg \mu$, the downloading bandwidth c may still be the bottleneck (e.g. if $\gamma < \mu$)
- ▶ Prior work assumes $c = \infty$ (motivated by the asymmetry in cable modem and DSL rates): doesn't capture the above effect

Evaluation results

- Fluid model vs real trace
 - the number of seeds in the system



More discussion

This lecture

- BitTorrent
 - ↳ P2P swarming
 - ↳ protocol overview
 - ↳ performance analysis
- Explore further
 - ↳ measurement-based modeling
 - ↳ measurement-based performance analysis
 - ↳ BitTorrent extensions
 - <http://wiki.theory.org/BitTorrentSpecification>

Next lecture

- Next: Skype

- [BS06] Salman A. Baset and Henning Schulzrinne, "An Analysis of the Skype Peer-to-Peer Internet Telephony Protocol", IEEE Infocom 2006. [Skype]

- Notice

- reading list and schedule are crosscourse

- check the list and submit reading summary on time

Reading summaries

- Challenge the paper presented by the other group
 - ↳ The problem(s)
 - ↳ Main Idea(s)
 - ↳ Major (at least three) strengths
 - ↳ ***Major (at least three) Weaknesses, then and now***
 - ↳ Possible improvement
- Template but submit in PDF only
 - ↳ <http://www.cs.uvic.ca/~pan/csc466/rs.txt> (Text)
 - ↳ <http://www.cs.uvic.ca/~pan/csc466/rs.tex> (LaTeX)